

744. Find Smallest Letter Greater Than Target

Easy Array Binary Search

Problem Description

The problem presents an array `letters` of sorted characters in non-decreasing order and a `target` character. The task is to find the smallest character in the array that is lexicographically greater than the target character. If no such character exists in the array, the requirement is to return the first character in `letters`. Lexicographically greater means the character should come after the target character in the alphabet. It is guaranteed that the array contains at least two different characters.

Intuition

The intuitive approach to solving this problem is to use a [binary search](#) because the characters are sorted, which makes it possible to quickly eliminate half of the search space at each step. We search for the point where we can find the first character in `letters` that is greater than `target`. Here's a step-by-step intuition:

1. Initialize two pointers, `left` and `right`, at the beginning and at the end (plus one) of the array respectively.
2. While the `left` pointer is less than the `right` pointer, we repeatedly divide the search space in half. We calculate a middle index `mid` by averaging `left` and `right`.
3. We compare the character at the middle index with the `target` using their ASCII values (through the `ord` function).
4. If the character at the middle index is greater than the `target` character, we have to keep searching in the left part of the array to find the smallest character larger than `target`, so we move the `right` pointer to `mid`.
5. If the character in the middle is less than or equal to `target`, we have to search in the right part of the array, so we increment the `left` pointer to `mid + 1`.
6. When the `left` and `right` pointers meet, they point at the smallest character greater than the target if such a character exists in the array. If such character does not exist, `left` will be equal to `len(letters)`, meaning we need to wrap around to the start of the array.
7. To address the wrapping, we use modulus operation `left % len(letters)`. This ensures that if `left` is beyond the last index, we return the first character in the array.

This approach efficiently narrows down the search space using the properties of sorted arrays and direct character comparisons, allowing us to find the answer in logarithmic time complexity, which is much faster than linearly scanning through the array.

Solution Approach

The implementation uses a [binary search](#) algorithm, which is efficient for sorted arrays. The binary search algorithm repeatedly divides the search space in half, gradually narrowing down the section of the array where the answer could be, thereby reducing the number of comparisons that need to be made. Here's how the provided code works:

1. **Initialization:** `left` is set to 0 (the start of the array), and `right` is set to `len(letters)`, which is one past the end of the array.
2. **Binary Search Loop:**
 - The condition `while left < right` ensures that the loop runs until `left` and `right` meet.
 - `mid = (left + right) >> 1` finds the middle index. The `>> 1` is a bit shift to the right by one position, effectively dividing the sum by 2, but faster.
3. **Character Comparison:**
 - `if ord(letters[mid]) > ord(target)`: If the character at the `mid` index in `letters` is lexicographically greater than `target`, search to the left by updating `right = mid`.
 - `else`: statement means the mid character is not greater than `target`, so we search to the right by updating `left = mid + 1`.
4. **Finding the Answer:**
 - After the loop exits, `left` is the index of the smallest character that is lexicographically greater than the target. If such a character doesn't exist, `left` will be equal to `len(letters)`, indicating that we have searched the entire array without finding a character greater than `target`.
5. **Return Statement:**
 - `return letters[left % len(letters)]` ensures we return the correct character:
 - If `left` is less than `len(letters)`, it means we've found a character that is greater than `target`, and we return that character.
 - If `left` is equal to `len(letters)`, the modulus operation causes it to wrap around to 0, returning the first character of the array, which is the required behavior when a greater character isn't found.

In terms of data structures, only the input array `letters` is used, and pointers (indices) are manipulated to traverse the array. No additional data structures are required for this algorithm, which keeps the space complexity low.

By implementing the [binary search](#) pattern, the time complexity is $O(\log n)$, where `n` is the length of the input array. This is significantly more efficient than a linear search, which would have a time complexity of $O(n)$.

Example Walkthrough

Let's apply the solution approach to a small example:

Assume we have the sorted array `letters = ["c", "f", "j"]` and the target character `target = "a"`.

Now, let's walk through the binary search algorithm step by step:

1. **Initialization:** We set `left` to 0 and `right` to 3 (the length of `letters` is 3).
2. **Binary Search Loop:**
 - Begin the loop since `left < right` ($0 < 3$).
 - Calculate `mid` which is $(0 + 3) >> 1$, yielding 1.
3. **Character Comparison:**
 - At index 1, the character is `"f"`. We compare `"f"` with `target` which is `"a"`.
 - Since `ord("f")` (102) is greater than `ord("a")` (97), we set `right` to `mid`, which is now 1.
4. **Continue Binary Search Loop:**
 - The loop condition still holds as `left < right` ($0 < 1$).
 - Calculate `mid` again, which is $(0 + 1) >> 1$, yielding 0.
 - Compare character at index 0, `"c"`, with `target`.
 - Since `ord("c")` (99) is greater than `ord("a")` (97), we set `right` to `mid`, which is now 0.
5. **Loop Ends:** The loop now terminates since `left` is not less than `right` (both are 0).
6. **Finding the Answer:** The `left` is 0, and `left` is less than `len(letters)`.
7. **Return Statement:** We return `letters[left % len(letters)]`, which is `letters[0]`. The character at index 0 is `"c"`, so `"c"` is returned as the smallest lexicographically greater character than `target`.

In this example, the desired character is found within two iterations of the binary search loop, demonstrating the efficiency of the binary search algorithm in quickly locating the answer. The result, `"c"`, is indeed the smallest character in the array that is lexicographically greater than the target `"a"`. If the `target` had been `"k"`, the search would conclude with `left` being 3, and the modulus operation would wrap around to return the first element in the array, `"c"`.

Solution Implementation

Python

```
class Solution:
    def nextGreatestLetter(self, letters: List[str], target: str) -> str:
        # Initialize the left and right pointers to the start and end of the list respectively
        left, right = 0, len(letters)

        # Use binary search to find the position of the next greatest letter
        while left < right:
            # Find the middle index
            mid = (left + right) // 2 # the '>> 1' is a bitwise operation equivalent to integer division by 2

            # If the middle letter is greater than the target, look to the left half
            if ord(letters[mid]) > ord(target):
                right = mid
            else:
                # Otherwise, look to the right half
                left = mid + 1

        # The modulo operation ensures wrapping around if the target letter is greater than any letter in the list
        return letters[left % len(letters)]
```

Java

```
class Solution {
    public char nextGreatestLetter(char[] letters, char target) {
        // Initialize the start and end pointers for binary search
        int start = 0;
        int end = letters.length;

        // Perform binary search to find the smallest letter greater than target
        while (start < end) {
            // Calculate the mid point to split the search into halves
            int mid = (start + end) >> 1; // Using unsigned shift for safe mid calculation

            // If the middle letter is greater than the target
            if (letters[mid] > target) {
                // We have a new possible candidate for next greatest letter (inclusive)
                // and we need to search to the left of mid (exclusive)
                end = mid;
            } else {
                // If mid letter is less than or equal to the target,
                // we need to search to the right of mid (exclusive)
                start = mid + 1;
            }
        }

        // After the search, start is the least index where letters[index] > target.
        // Since the array is circular, we use modulo operator to wrap around the index
        return letters[start % letters.length];
    }
}
```

C++

```
#include <vector>

class Solution {
public:
    char nextGreatestLetter(vector<char>& letters, char target) {
        // Initialize the pointers for the binary search.
        int left = 0;
        int right = letters.size();

        // Perform binary search.
        while (left < right) {
            // Find the middle index.
            int mid = left + (right - left) / 2; // Prevent potential overflow.

            // If the middle letter is strictly greater than the target,
            // move the right pointer to mid, as we want the smallest letter
            // that is greater than the target.
            if (letters[mid] > target) {
                right = mid;
            } else {
                // If the middle letter is less than or equal to the target,
                // move the left pointer past mid.
                left = mid + 1;
            }
        }

        // Since the list is circular, if we go past the end,
        // we return the first element (modulo operation).
        return letters[left % letters.size()];
    }
};
```

TypeScript

```
function nextGreatestLetter(letters: string[], target: string): string {
    // Initialize the number of elements in the 'letters' array.
    const numLetters = letters.length;
    // Set the initial search interval between the start and end of the array.
    let leftIndex = 0;
    let rightIndex = numLetters;

    // Perform a binary search to find the smallest letter greater than the target.
    while (leftIndex < rightIndex) {
        // Calculate the middle index.
        let middleIndex = leftIndex + ((rightIndex - leftIndex) >> 1); // Same as Math.floor((left + right) / 2)

        // If the letter at the middle index is greater than the target,
        // it could be a potential answer, so move the right index to the middle.
        // Otherwise, move the left index to one position after the middle.
        if (letters[middleIndex] > target) {
            rightIndex = middleIndex;
        } else {
            leftIndex = middleIndex + 1;
        }
    }

    // Since the letters wrap around, we use the modulo operator (%) with the number of letters.
    // This ensures we get a valid index if the 'leftIndex' goes beyond the array bounds.
    return letters[leftIndex % numLetters];
}

class Solution:
    def nextGreatestLetter(self, letters: List[str], target: str) -> str:
        # Initialize the left and right pointers to the start and end of the list respectively
        left, right = 0, len(letters)

        # Use binary search to find the position of the next greatest letter
        while left < right:
            # Find the middle index
            mid = (left + right) // 2 # the '>> 1' is a bitwise operation equivalent to integer division by 2

            # If the middle letter is greater than the target, look to the left half
            if ord(letters[mid]) > ord(target):
                right = mid
            else:
                # Otherwise, look to the right half
                left = mid + 1

        # The modulo operation ensures wrapping around if the target letter is greater than any letter in the list
        return letters[left % len(letters)]
```

Time and Space Complexity

Time Complexity

The provided code performs a binary search on the sorted list of letters to find the smallest letter in the list that is larger than the `target` character. The time complexity of binary search is $O(\log n)$, where `n` is the number of elements in the input list. This is because, with each comparison, it effectively halves the size of the search space.

Initially, the search space is the entire `letters` list. With each iteration of the loop, either the `left` or `right` index is adjusted to narrow down the search space. Since the search space is divided by two in every step, the maximum number of steps is proportional to $\log_2(n)$. Therefore, the time complexity is $O(\log n)$.

Space Complexity

The space complexity of the code is $O(1)$. This is because the space required does not scale with the size of the input list. The code uses a constant amount of extra space for variables such as `left`, `right`, and `mid`. There are no additional data structures or recursive calls that would increase the space complexity. Regardless of the list size, the amount of memory required remains the same.