

340. Longest Substring with At Most K Distinct Characters

MediumHash TableStringSliding Window

Problem Description

In this problem, we are given a string s and an integer k . Our task is to find the length of the longest substring (which is a contiguous block of characters in the string) within s that contains at most k distinct characters.

A distinct character means that no matter how many times the character appears in the substring, it is counted only once. For example, if $k = 2$, the substring "aabbc" has 3 distinct characters ('a', 'b', and 'c'), thus, it does not meet the requirement.

The goal here is to achieve this while maximizing the length of the substring. A substring could range from containing a single character up to the length of the entire string, if k is sufficiently large to cover all distinct characters in the string s .

Intuition

The core intuition behind the solution is to use the two-pointer technique, or more specifically, the [sliding window](#) approach. Here is the thinking process that leads us to this approach:

- We need to examine various substrings of s efficiently without having to re-scan the string repeatedly.
- We can start with a window (or two pointers) at the beginning of the string and expand the window to the right until we exceed k distinct characters within the window.
- Upon exceeding k distinct characters, we need to move the left side of the window to the right to potentially drop one of the distinct characters out of the count.
- We keep track of the count of distinct characters in the current window using a data structure, such as a counter (dictionary), that is updated when the window is adjusted.
- As we expand and contract the window, we keep a record of the maximum window size (i.e., substring length) that has appeared so far that contains at most k distinct characters. This is the number we want to return.

Now, the implementation uses a counter and two indices, i and j , where i is the end of the [sliding window](#), and j is the start. We iterate over the string with i :

- We include the character at position i in the current window by incrementing its count in the counter.
- If adding the current character has led to more than k distinct characters in the window, we increment j , effectively reducing the size of the window from the left, until the number of distinct characters drops to k or lower.
- At each step, we calculate the length of the current window ($i - j + 1$) and update the answer if it's the largest such length we have seen so far.
- This process continues until i has reached the end of the string.

The [sliding window](#) is moved across the string efficiently to identify the longest substring that satisfies the condition of having at most k distinct characters, thus leading us to the solution.

Solution Approach

The solution code implements the [sliding window](#) pattern using two pointers to keep track of the current window within the string s . This pattern is completed using the Python [Counter](#) from the collections module as the key data structure to keep counts of each character in the current window.

Here is how it is done step by step:

- We initialize a [Counter](#) object, which is a dictionary that will keep track of the counts of individual characters in the current [sliding window](#).
- Two pointers, i and j , are created. i is used to traverse the string s character by character from the start to the end, while j is used to keep track of the start of the [sliding window](#). j starts at 0 and moves rightward to narrow the window whenever necessary.
- We initialize a variable ans to keep track of the maximum length of the substring found that meets the criteria.
- We start iterating over the characters of the string s , with i acting as the end boundary of the window. For each character c at index i , we:
 - Increment its count in the [Counter](#) by 1 ($cnt[c] += 1$).
 - Check if the number of distinct characters in the [Counter](#) has exceeded k . This is done by evaluating $len(cnt) > k$. As long as this condition is true, meaning we have more than k distinct characters, we need to shrink the window from the left side by increasing j .
 - In the inner [while](#) loop, decrease the count of the character at the beginning of the window ($cnt[s[j]] -= 1$).
 - If the count of that character drops to 0, we remove it from the [Counter](#) ($cnt.pop(s[j])$), as it's no longer part of the window.
 - We increment j to shrink the window from the left.
- After adjusting the size of the window (if it was needed), we calculate the length of the current window ($i - j + 1$) and update ans to be the maximum of its current value and this newly calculated length.
- After iterating through all characters of s , the process concludes, and the final value of ans contains the length of the longest substring with at most k distinct characters. This value is returned as the result.

Through this method, we avoid unnecessary re-scans of the string and implement an $O(n)$ time complexity algorithm, where n is the length of the string s . By dynamically adjusting the [sliding window](#) and using the [Counter](#) to track distinct characters, we achieve an efficient approach to solving the problem as outlined.

Example Walkthrough

Let's assume we have a string $s = \text{"aabcbabb"}$, and the integer $k = 2$. We want to find the length of the longest substring with at most k distinct characters.

Here is a step-by-step walkthrough of the solution approach applied to this example:

- Initialize the Counter and Pointers:** Create a [Counter](#) object, cnt , to count characters in the current window. Set i, j to 0 to represent the start and end of the sliding window, and ans to 0 as the longest substring length found so far.

- Start Iterating with Pointer i :** Move i from the start to the end of the string s .

$i = 0$: $cnt = \{\text{'a'}: 1\}$, substring is a, $ans = 1$.

$i = 1$: $cnt = \{\text{'a'}: 2\}$, substring is aa, $ans = 2$.

$i = 2$: $cnt = \{\text{'a'}: 2, \text{'b'}: 1\}$, substring is aab, $ans = 3$.

- Exceeding k Distinct Characters:** At the next character, we check if the number of distinct characters will exceed k .

$i = 3$: We attempt to insert c into cnt . Doing so would increase distinct character count to 3 ($cnt = \{\text{'a'}: 2, \text{'b'}: 1, \text{'c'}: 1\}$), which is greater than k . We need to move j right until we get at most k distinct characters.

- Shrinking the Window:**

Start incrementing j , decreasing the count of the character $s[j]$ in cnt , and remove the character from cnt if its count reaches 0.

$j = 1$: We decrease the count of a and update $cnt = \{\text{'a'}: 1, \text{'b'}: 1, \text{'c'}: 1\}$. We still have more than k distinct characters.

$j = 2$: We decrease the count of a and remove it from cnt ($cnt = \{\text{'b'}: 1, \text{'c'}: 1\}$). Now we have k distinct characters.

After shrinking, the substring is now bc ($s[2:4]$), and ans remains 3.

- Continue Process:**

$i = 4$: Add a to cnt , $cnt = \{\text{'b'}: 1, \text{'c'}: 1, \text{'a'}: 1\}$, substring is bca, ans remains 3.

$i = 5$: Add b to cnt , $cnt = \{\text{'b'}: 2, \text{'c'}: 1, \text{'a'}: 1\}$. We have exceeded k again, so start moving j .

$j = 3$: Reduce count of b, $cnt = \{\text{'b'}: 1, \text{'c'}: 1, \text{'a'}: 1\}$. Still too many distinct characters.

$j = 4$: Remove c, $cnt = \{\text{'b'}: 1, \text{'a'}: 1\}$, and now the substring is ab ($s[4:6]$), update ans to 2.

$i = 6$: Add b, $cnt = \{\text{'b'}: 2, \text{'a'}: 1\}$, and substring is abb ($s[4:7]$). Now, update ans to 3, as this is the length of the current window.

- Finish Iteration:**

After the final iteration, ans contains the length of the longest satisfactory substring, which is 3 in this case, representing the substring aab or abb.

By using the sliding window technique, the algorithm efficiently finds the longest substring that satisfies the given constraint with a time complexity of $O(n)$, where n is the length of the string s .

Python Solution

```
1 from collections import Counter
2
3 class Solution:
4     def lengthOfLongestSubstringKDistinct(self, s: str, k: int) -> int:
5         # Initialize a counter to keep track of the frequency of each character
6         char_count = Counter()
7
8         # Length of the input string
9         n = len(s)
10
11         # Initialize the answer and the start index (j) of the current window
12         max_length = start_index = 0
13
14         # Enumerate over the characters of the string with index (i) and character (char)
15         for i, char in enumerate(s):
16             # Increment the count of the current character
17             char_count[char] += 1
18
19             # If the number of distinct characters exceeds k, shrink the window
20             while len(char_count) > k:
21                 # Decrement the count of the character at the start index
22                 char_count[s[start_index]] -= 1
23
24                 # If the count goes to zero, remove the character from the counter
25                 if char_count[s[start_index]] == 0:
26                     del char_count[s[start_index]]
27
28                 # Move the start index forward
29                 start_index += 1
30
31             # Update the maximum length found so far
32             max_length = max(max_length, i - start_index + 1)
33
34         # Return the maximum length of substring with at most k distinct characters
35         return max_length
36
37 # The code can now be run with an instance of the Solution class
38 # For example: Solution().lengthOfLongestSubstringKDistinct("eceba", 2) should return 3
39
```

Java Solution

```
1 import java.util.HashMap;
2 import java.util.Map;
3
4 class Solution {
5
6     public int lengthOfLongestSubstringKDistinct(String s, int k) {
7         // Map to store the frequency of each character in the current window
8         Map<Character, Integer> charCountMap = new HashMap<>();
9         int n = s.length();
10         int longestSubStringLength = 0; // variable to store the length of the longest substring
11         int left = 0; // left pointer for the sliding window
12
13         // Iterate through the string using the right pointer of the sliding window
14         for (int right = 0; right < n; ++right) {
15             // Step 1: Update the count of the current character
16             char currentChar = s.charAt(right);
17             charCountMap.put(currentChar, charCountMap.getOrDefault(currentChar, 0) + 1);
18
19             // Step 2: Shrink the window from the left if count map has more than 'k' distinct characters
20             while (charCountMap.size() > k) {
21                 char leftChar = s.charAt(left);
22                 charCountMap.put(leftChar, charCountMap.get(leftChar) - 1);
23                 // Remove the character from map when count becomes zero
24                 if (charCountMap.get(leftChar) == 0) {
25                     charCountMap.remove(leftChar);
26                 }
27                 left++; // shrink the window from the left
28             }
29
30             // Step 3: Update the longest substring if the current window is larger
31             longestSubStringLength = Math.max(longestSubStringLength, right - left + 1);
32         }
33
34         return longestSubStringLength; // Return the length of the longest substring found
35     }
36 }
37
```

C++ Solution

```
1 #include <string>
2 #include <unordered_map>
3 #include <algorithm>
4
5 class Solution {
6 public:
7     // Finds the length of the longest substring with at most k distinct characters
8     int lengthOfLongestSubstringKDistinct(std::string s, int k) {
9         std::unordered_map<char, int> charCountMap; // Map to store character counts
10
11         int maxSubStringLength = 0; // Maximum length of substring found
12         int leftPointer = 0; // Left pointer for sliding window
13         int n = s.size(); // Length of the input string
14
15         for (int rightPointer = 0; rightPointer < n; ++rightPointer) {
16             // Increase char count for the current position
17             charCountMap[s[rightPointer]]++;
18
19             // If we have more than k distinct chars, contract the window from the left
20             while (charCountMap.size() > k) {
21                 // Decrease the count of the char at the left pointer
22                 charCountMap[s[leftPointer]]--;
23
24                 // If the count drops to 0, remove it from the map
25                 if (charCountMap[s[leftPointer]] == 0) {
26                     charCountMap.erase(s[leftPointer]);
27                 }
28
29                 // Move the left pointer to the right
30                 ++leftPointer;
31             }
32
33             // Update maxSubStringLength if we've found a larger window
34             maxSubStringLength = std::max(maxSubStringLength, rightPointer - leftPointer + 1);
35         }
36
37         return maxSubStringLength; // Return the max length found
38     }
39 };
40
```

Typescript Solution

```
1 // Imports a generic map class to emulate the unordered_map feature in C++
2 import { Map } from "typescript-collections";
3
4 // Finds the length of the longest substring with at most k distinct characters
5 function lengthOfLongestSubStringKDistinct(s: string, k: number): number {
6     let charCountMap = new Map<char, number>(); // Map to store character counts
7
8     let maxSubStringLength = 0; // Maximum length of substring found
9     let leftPointer = 0; // Left pointer for sliding window
10    let n = s.length; // Length of the input string
11
12    for (let rightPointer = 0; rightPointer < n; rightPointer++) {
13        // Increase char count for the current character
14        let count = charCountMap.getValue(s[rightPointer]) || 0;
15        charCountMap.setValue(s[rightPointer], count + 1);
16
17        // If we have more than k distinct characters, contract the window from the left
18        while (charCountMap.size() > k) {
19            // Decrease the count of the character at the left pointer
20            let leftCharCount = charCountMap.getValue(s[leftPointer]) || 0;
21            charCountMap.setValue(s[leftPointer], leftCharCount - 1);
22
23            // If the count drops to 0, remove it from the map
24            if (charCountMap.getValue(s[leftPointer]) === 0) {
25                charCountMap.remove(s[leftPointer]);
26            }
27
28            // Move the left pointer to the right
29            leftPointer++;
30        }
31
32        // Update maxSubStringLength if we've found a larger window
33        maxSubStringLength = Math.max(maxSubStringLength, rightPointer - leftPointer + 1);
34    }
35
36    return maxSubStringLength; // Return the max length found
37 }
38
```

Time and Space Complexity

The given code snippet defines a function that determines the length of the longest substring with no more than k distinct characters in a given string s .

Time Complexity

The time complexity of the function can be analyzed as follows:

- There are two pointers (i and j) that traverse the string s only once. The outer loop runs for all characters from position 0 to $n - 1$ where n is the length of the string s , hence contributing $O(n)$ to the time complexity.
- Inside the loop, there is a [while](#) loop that shrinks the sliding window from the left when the number of distinct characters exceeds k . This [while](#) loop does not run for each element in s multiple times. Each character is removed from the window only once, accounting for another $O(n)$ run for the whole element in the algorithm.

Therefore, the total time complexity of the algorithm is $O(n)$ where n is the number of characters in the input string s .

Space Complexity

The space complexity of the function can be analyzed as follows:

- A [Counter](#) object is used to keep track of the frequency of each character in the current window. In the worst case, the counter could store a distinct count for every character in the string s . However, since the number of distinct characters is limited by k , the [Counter](#) would hold at most k elements.

Therefore, the space complexity of the algorithm is $O(k)$ where k is the maximum number of distinct characters that the substring can have.