56. Merge Intervals

Sorting

Problem Description

Medium Array

The given problem requires us to merge all the overlapping intervals in a list. An interval is represented as a list with two elements where the first element is the start and the second is the end of the interval ([start, end]). An "overlap" occurs when one interval's start is less than or equal to the end of another interval. The goal is to simplify the list of intervals to a list where no intervals overlap, ensuring that the new list collectively spans the same range as the original intervals.

After merging all the overlapping intervals, we get: [[1,6],[8,10],[15,18]] In the merged intervals, there is no pair of intervals such that one overlaps with another.

equal to the current interval's start time.

Original list of intervals: [[1,3],[2,6],[8,10],[15,18]]

Here's an example for clarification:

ntuition

1. If we sort the intervals based on their start times, any overlapping intervals will be placed next to each other in the list. 2. To merge intervals, we only need to track the end time since the sorted order ensures that the next interval's start time is always greater than or

The approach is as follows:

2. Initialize a new list to hold the merged intervals and add the first interval to it.

The intuition behind the solution comes from two realizations:

1. Sort the list of intervals based on their start times.

3. Iterate through the rest of the intervals, and for each one, compare its start time with the end time of the last interval in the merged list. 4. If the start time is greater than the end time of the last interval in the merged list, then there is no overlap, and we can add the current interval

as a new entry to the merged list.

- 6. The process continues until we have gone through all the intervals.
- 5. If there is overlap (the start time is less than or equal to the end time), we update the end time of the last interval in the merged list to be the maximum of the end times of the last interval and the current interval.
- 7. We return the merged list of intervals as the answer. This solution guarantees that we merge all overlapping intervals and result in a list of intervals with no overlaps.
- Solution Approach
- The solution to the problem involves sorting the intervals and then iterating through the sorted list to merge any overlapping intervals.

the intervals based on their first element (the start times).

Here's a step-by-step breakdown of the implementation:

First, we sort the given list of intervals. This is done in-place using the native sort function provided by Python, which sorts

By sorting the intervals, we are able to take advantage of the fact that any intervals that might need merging will appear next

intervals.sort()

to each other in the list.

ans.append([s, e])

overlapping intervals.

Before sort: [[5,7],[1,3],[3,4],[2,6]]

After sort: [[1,3],[2,6],[3,4],[5,7]]

else:

ans = [intervals[0]]

We then proceed to iterate over the rest of the intervals, starting from the second interval onward, to check for overlapping

We then initialize a new list called ans, which will store our merged intervals, and we start by adding the first interval to it.

with the currently last interval in our ans list. for s, e in intervals[1:]:

Here, (s, e) represents the start and end times of the current interval we are looking at.

This acts as a comparison base for merging subsequent intervals.

and we can simply add this interval to ans. if ans[-1][1] < s:

If the start time s of the current interval is greater than the end time of the last interval in ans, it means there is no overlap

However, if an overlap exists, we need to merge the current interval with the last interval in ans. To do this, we update the

ans $[-1][1] = \max(ans[-1][1], e)$ This ensures that the intervals are merged, covering the overlapping time spans without duplicating intervals.

Once we finish iterating through all intervals, the ans list contains the merged intervals. We return ans as the final set of non-

The algorithm uses the sort-merge pattern, which is common for interval problems. By sorting and then merging, we bring the

overall run-time complexity down to O(N log N) where N is the number of intervals, with the sort contributing to the log N factor

and the merge process being linear in nature. Regarding data structures, the solution leverages lists and the use of tuple

end time of the last interval in ans with the maximum end time between the two overlapping intervals.

unpacking for readability. **Example Walkthrough** Let's take a small example to illustrate the solution approach with the provided intervals: [[5,7],[1,3],[3,4],[2,6]].

We used intervals.sort() to achieve this. **Initializing and Iterating for Merging** We then initialize the ans list with the first sorted interval, treating it as the base for our merged intervals: ans = [[1,3]]

We look at [2,6] and compare it to the last element of ans, which is [1,3]. Since the start 2 is within [1,3] (as 3 is greater

Proceeding to [3,4], we compare it to the last element [1,6]. Again, it overlaps because 4 is not greater than 6. No need to

However, our merging logic must have the current start to be greater than the last end to avoid overlap. Therefore, we should

• [5,7] is checked again and it actually overlaps with [1,6] (since 5 is less than or equal to 6). We merge them by updating the end of the last

Having completed the iteration over the sorted intervals, we have a list of merged intervals where no two intervals overlap. The

than 2), they overlap. We merge them by updating the end of the last interval in ans to the max end of both intervals, now

adjust the last step:

Final Merged List

final ans is:

Python

class Solution:

Solution Implementation

intervals.sort()

merged_intervals = [intervals[0]]

for start, end in intervals[1:]:

Return the merged intervals

return merged_intervals

if merged intervals[-1][1] < start:</pre>

// List that holds the merged intervals.

mergedIntervals.add(intervals[0]);

int start = intervals[i][0];

int end = intervals[i][1];

} else {

TypeScript

} else {

intervals.sort()

else:

Time Complexity

For a list of n intervals:

merged intervals = [intervals[0]]

for start, end in intervals[1:]:

if merged intervals[-1][1] < start:</pre>

merged_intervals.append([start, end])

// Return the merged intervals

function merge(intervals: number[][]): number[][] {

for (let i = 1: i < intervals.length: ++i) {</pre>

if (lastInterval[1] < intervals[i][0]) {</pre>

mergedIntervals.push(intervals[i]);

const mergedIntervals: number[][] = [intervals[0]];

intervals.sort((a, b) => a[0] - b[0]);

// First, we sort the intervals array based on the start times

// Initialize the merged intervals array with the first interval

// Iterate through the intervals starting from the second element

const lastInterval = mergedIntervals[mergedIntervals.length - 1];

lastInterval[1] = Math.max(lastInterval[1], intervals[i][1]);

// Get the last interval in the mergedIntervals array

def merge(self, intervals: List[List[int]]) -> List[List[int]]:

Sort the interval list based on the start times of intervals

Initialize the merged intervals list with the first interval

Iterate over the intervals, starting from the second interval

return mergedIntervals;

List<int[]> mergedIntervals = new ArrayList<>();

for (int i = 1; i < intervals.length; ++i) {</pre>

// Get the last interval in the merged list.

// Add the first interval to the list as starting interval for merging.

int[] lastMergedInterval = mergedIntervals.get(mergedIntervals.size() - 1);

// Check if there is an overlap with the last interval in the merged list.

// If there is an overlap, merge the current interval with the last interval

mergedIntervals.back()[1] = std::max(mergedIntervals.back()[1], intervals[i][1]);

// in the result by updating the end time to the maximum end time seen

// If the current interval does not overlap with the last interval in mergedIntervals

// If there is an overlap, merge the current interval with the last interval

// by updating the end time of the last interval to the maximum end time

// Add the current interval to the mergedIntervals array as it cannot be merged

Check if the current interval does not overlap with the last interval in merged_intervals

updating the end time of the last interval to the maximum end time seen so far

If it does not overlap, add the current interval to merged_intervals

If it does overlap, merge the current interval with the last one by

 $merged_intervals[-1][1] = max(merged_intervals[-1][1], end)$

// Loop through all the intervals starting from the second one.

// Get the start and end times of the current interval.

ans becomes [[1,6]].

Sorting the Intervals

Finally, we look at [5,7]. This does not overlap with [1,6] as 5 is not greater than 6. Since 7 is greater than 6, we add [5,7] as a new interval to ans. After the addition, ans becomes [[1,6],[5,7]].

Next, we start iterating from the second element of the sorted intervals:

change the end time since 6 is already the maximum end.

def merge(self, intervals: List[List[int]]) -> List[List[int]]:

Sort the interval list based on the start times of intervals

Initialize the merged intervals list with the first interval

Iterate over the intervals, starting from the second interval

interval in ans to 7, and ans now becomes [[1,7]].

First, we need to sort the intervals by their starting points to align any intervals that might overlap:

[[1,7]]This means we successfully merged all intervals to cover the same range without having any overlapping intervals. The solution approach, by sorting and then merging, streamlined the process and ensures an efficient way to obtain the merged intervals.

merged_intervals.append([start, end]) else: # If it does overlap, merge the current interval with the last one by # updating the end time of the last interval to the maximum end time seen so far $merged_intervals[-1][1] = max(merged_intervals[-1][1], end)$

Check if the current interval does not overlap with the last interval in merged_intervals

If it does not overlap, add the current interval to merged_intervals

// Method to merge overlapping intervals. public int[][] merge(int[][] intervals) { // Sort the intervals by their starting times. Arrays.sort(intervals, (a, b) -> a[0] - b[0]);

class Solution {

import java.util.Arrays;

import java.util.ArrayList;

import iava.util.List:

Java

if (lastMergedInterval[1] < start) {</pre> // No overlap, so we can add the current interval as it is. mergedIntervals.add(intervals[i]); } else { // Overlap exists, so we extend the last interval's end time. lastMergedInterval[1] = Math.max(lastMergedInterval[1], end); // Convert the merged intervals list to a 2D array and return it. return mergedIntervals.toArray(new int[mergedIntervals.size()][]); C++ #include <vector> #include <algorithm> class Solution { public: // Function to merge overlapping intervals std::vector<std::vector<int>> merge(std::vector<std::vector<int>>& intervals) { // First, sort the intervals based on the starting times std::sort(intervals.begin(), intervals.end()); // This will be the result vector for merged intervals std::vector<std::vector<int>> mergedIntervals; // Initialize the result vector with the first interval mergedIntervals.push_back(intervals[0]); // Iterate through all the intervals starting from the second one for (int i = 1; i < intervals.size(); ++i) {</pre> // If the current interval does not overlap with the last interval in the result, // then simply add the current interval to the result if (mergedIntervals.back()[1] < intervals[i][0]) {</pre> mergedIntervals.push_back(intervals[i]);

// Return the array containing all the merged intervals return mergedIntervals; class Solution:

The given code has two main operations: 1. Sorting the intervals list. 2. Iterating through the sorted list and merging overlapping intervals.

Return the merged intervals

return merged_intervals

Time and Space Complexity

insertion sort) for sorting lists. • The iteration over the list has a complexity of O(n), because we go through the intervals only once.

Hence, the total time complexity is the sum of these two operations, which is 0(n log n) + 0(n). Since 0(n log n) is the higher order term, it dominates the total time complexity, which simplifies to $O(n \log n)$.

Space Complexity The space complexity consists of the additional space used by the algorithm apart from the input:

• The ans list which contains the merged intervals is the main additional data structure used, and in the worst case, if no intervals overlap, it will contain n intervals. • Since the ans list reuses the intervals from the original input list, and the input list size itself is not included in the additional space used for computing space complexity, the space used to store ans can be considered 0(1) (constant space) in this context.

• The sort operation typically has a complexity of O(n log n), since Python uses TimSort (a hybrid sorting algorithm derived from merge sort and

Thus, the space complexity is 0(1).