253. Meeting Rooms II

Medium | Greedy | Array | Two Pointers | Prefix Sum | Sorting | Heap (Priority Queue)

Problem Description

The problem presents a scenario where we have an array of meeting time intervals, each represented by a pair of numbers `[start_i, end_i]`. These pairs indicate when a meeting starts and ends. The goal is to find the minimum number of conference rooms required to accommodate all these meetings without any overlap. In other words, we want to allocate space such that no two meetings occur in the same room simultaneously.

Intuition

The core idea behind the solution is to track the changes in room occupancy over time, which is akin to tracking the number of trains at a station at any given time. We can visualize the timeline from the start of the first meeting to the end of the last meeting, and keep a counter that increments when a meeting starts and decrements when a meeting ends. This approach is similar to the sweep line algorithm, often used in computational geometry to keep track of changes over time or another dimension.

By iterating through all the meetings, we apply these increments/decrements at the respective start and end times. The maximum value reached by this counter at any point in time represents the peak occupancy, thus indicating the minimum number of conference rooms needed. To implement this:

1. We initialize an array `delta` that is large enough to span all potential meeting times. We use a fixed size in this solution, which assumes the meeting times fall within a predefined range (0 to 1000009 in this case).
2. Iterate through the intervals list, and for each meeting interval `[start, end]`, increment the value at index `start` in the `delta` array, and decrement the value at index `end`. This effectively marks the start of a meeting with +1 (indicating a room is now occupied) and the end of a meeting with -1 (a room has been vacated).
3. Accumulate the changes in the `delta` array using the `accumulate` function, which applies a running sum over the array elements. The maximum number reached in this accumulated array is our answer, as it represents the highest number of simultaneous meetings, i.e., the minimum number of conference rooms required.

This solution is efficient because it avoids the need to sort the meetings by their start or end times, and it provides a direct way to calculate the running sum of room occupancy over the entire timeline.

Solution Approach

The solution uses a simple array and the concept of the prefix sum (running sum) to keep track of room occupancy over time—an approach that is both space-efficient and does not require complex data structures.

Here's a step-by-step breakdown of the implementation:

1. **Initialization**: A large array `delta` is created with all elements initialized to 0. The size of the array is chosen to be large enough to handle all potential meeting times (1 more than the largest possible time to account for the last meeting's end time). In this case, `1000010` is used.
2. **Updating the `delta` Array**: For each meeting interval, say `[start, end]`, we treat the start time as the point where a new room is needed (increment counter) and the end time as the point where a room is freed (decrement counter).

   ```
   for start, end in intervals:
       delta[start] += 1
       delta[end] -= 1
   ```

   This creates a timeline indicating when rooms are occupied and vacated.
3. **Calculating the Prefix Sum**: We use the `accumulate` function from the `itertools` module of Python to create a running sum (also known as a prefix sum) over the `delta` array. The result is a new array indicating the number of rooms occupied at each time.

   ```
   occupied_rooms_over_time = accumulate(delta)
   ```

4. **Finding the Maximum Occupancy**: The peak of the `occupied_rooms_over_time` array represents the maximum number of rooms simultaneously occupied, hence the minimum number of rooms we need.

   ```
   min_rooms_required = max(occupied_rooms_over_time)
   ```

   The `max` function is used to find this peak value, which completes our solution.

The beauty of this approach is in its simplicity and efficiency. Instead of worrying about sorting meetings by starts or ends or using complex data structures like priority queues, we leverage the fact that when we are only interested in the max count, the order of increments and decrements on the timeline does not matter. As long as we correctly increment at the start times and decrement at the end times, the `accumulate` function ensures we get a correct count at each time point.

In conclusion, this method provides an elegant solution to the problem using basic array manipulation and the concept of prefix sums.

Example Walkthrough

Let's consider a small set of meeting intervals to illustrate the solution approach:

```
Meeting intervals: [[1, 4], [2, 5], [7, 9]]
```

Here we have three meetings. The first meeting starts at time 1 and ends at time 4, the second meeting starts at time 2 and ends at time 5, and the third meeting starts at time 7 and ends at time 9.

Following the solution steps:

1. **Initialization:** We create an array `delta` of size 1000010, which is a bit overkill for this small example, but let's go with the provided approach. Initially, all elements in `delta` are set to 0.
2. **Updating the `delta` Array:** We iterate through the meeting intervals and update the `delta` array accordingly.

   ```
   delta[1] += 1  # Meeting 1 starts, need a room
   delta[4] -= 1  # Meeting 1 ends, free a room
   delta[2] += 1  # Meeting 2 starts, need a room
   delta[5] -= 1  # Meeting 2 ends, free a room
   delta[7] += 1  # Meeting 3 starts, need a room
   delta[9] -= 1  # Meeting 3 ends, free a room
   ```

   After the updates, the `delta` array will reflect changes in room occupancy at the start and end times of the meetings.
3. **Calculating the Prefix Sum:** Using an `accumulate` operation (similar to a running sum), we calculate the number of rooms occupied at each point in time. For simplicity, we will perform the cumulation manually:

   ```
   time       1  2  3  4  5  6  7  8  9
   delta     +1 +1  0 -1 -1  0 +1  0 -1
   occupied   1  2  2  1  0  0  1  1  0    (summing up `delta` changes over time)
   ```

   The maximum number during this running sum is 2, which occurs at times 2 and 3.
4. **Finding the Maximum Occupancy:** We can see that the highest value in the occupancy timeline is 2, therefore we conclude that at least two conference rooms are needed to accommodate all meetings without overlap.

```
The minimum number of conference rooms required is 2.
```

Solution Implementation

Python

```python
from typing import List
from itertools import accumulate

class Solution:
    def minMeetingRooms(self, intervals: List[List[int]]) -> int:
        # Initialize a list to keep track of the number of meetings starting or ending at any time
        # The range is chosen such that it covers all possible meeting times
        meeting_delta = [0] * 1000010

        # Go through each interval in the provided list of intervals
        for start, end in intervals:
            meeting_delta[start] += 1  # Increment for a meeting starting
            meeting_delta[end] -= 1    # Decrement for a meeting ending

        # The `accumulate` function is used to compute the running total of active meetings at each time
        # `max` is then used to find the maximum number of concurrent meetings, which is the minimum number of rooms required
        return max(accumulate(meeting_delta))

# Example usage:
# sol = Solution()
# print(sol.minMeetingRooms([[0, 30], [5, 10], [15, 20]]))  # Output: 2
```

Java

```java
class Solution {

    // Function to find the minimum number of meeting rooms required
    public int minMeetingRooms(int[][] intervals) {
        // Define the size for time slots (with assumed maximum time as 10^6+10)
        int n = 1000010;
        int[] delta = new int[n]; // Array to hold the changes in ongoing meetings

        // Iterate through all intervals
        for (int[] interval : intervals) {
            // Increment the start time to indicate a new meeting starts
            ++delta[interval[0]];
            // Decrement the end time to indicate a meeting ends
            --delta[interval[1]];
        }

        // Initialize res to the first time slot to handle the case if only one meeting
        int res = delta[0];

        // Traverse over the delta array to find maximum number of ongoing meetings at any time
        for (int i = 1; i < n; ++i) {
            // Cumulate the changes to find active meetings at time i
            delta[i] += delta[i - 1];
            // Update res if the current time slot has more meetings than previously recorded
            res = Math.max(res, delta[i]);
        }

        // Return the maximum value found in delta, which is the minimum number of rooms required
        return res;
    }
}
```

C++

```cpp
#include <vector>
#include <algorithm> // Include necessary headers

class Solution {
public:
    // Function to find the minimum number of meeting rooms required
    int minMeetingRooms(vector<vector<int>>& intervals) {
        int n = 1000010; // Define the maximum possible time point
        vector<int> countDelta(n); // Create a vector to keep track of the count changes

        // Iterate over each interval
        for (auto &interval : intervals) {
            ++countDelta[interval[0]]; // Increment count at the start time of the meeting
            --countDelta[interval[1]]; // Decrement count at the end time of the meeting
        }

        // Prefix sum - accumulate the count changes to find the count at each time
        for (int i = 0; i < n - 1; ++i) {
            countDelta[i + 1] += countDelta[i];
        }

        // Find and return the maximum count at any time, which is the minimum number of rooms required
        return *max_element(countDelta.begin(), countDelta.end());
    }
};
```

TypeScript

```typescript
// Import necessary module for max element search
import { max } from "lodash";

// Function to find the minimum number of meeting rooms required
function minMeetingRooms(intervals: number[][]): number {
    const maximumTimePoint: number = 1000010; // Define the maximum possible time point
    const countChange: number[] = new Array(maximumTimePoint).fill(0); // Create an array to track the count changes, initialized to

    // Iterate over each interval
    intervals.forEach(interval => {
        countChange[interval[0]]++; // Increment count at the start time of the meeting
        countChange[interval[1]]--; // Decrement count at the end time of the meeting
    });

    // Prefix sum calculation - accumulate the count changes to find the count at each time point
    for (let i = 0; i < maximumTimePoint - 1; ++i) {
        countChange[i + 1] += countChange[i];
    }

    // Find and return the maximum count at any time point, which is the minimum number of rooms required
    return max(countChange);
}

// Example usage:
// const intervals: number[][] = [[0, 30], [5, 10], [15, 20]];
// const numberOfRooms: number = minMeetingRooms(intervals);
// console.log(numberOfRooms); // Should output the minimum number of meeting rooms required
```


```python
from typing import List
from itertools import accumulate

class Solution:
    def minMeetingRooms(self, intervals: List[List[int]]) -> int:
        # Initialize a list to keep track of the number of meetings starting or ending at any time
        # The range is chosen such that it covers all possible meeting times
        meeting_delta = [0] * 1000010

        # Go through each interval in the provided list of intervals
        for start, end in intervals:
            meeting_delta[start] += 1  # Increment for a meeting starting
            meeting_delta[end] -= 1    # Decrement for a meeting ending

        # The `accumulate` function is used to compute the running total of active meetings at each time
        # `max` is then used to find the maximum number of concurrent meetings, which is the minimum number of rooms required
        return max(accumulate(meeting_delta))

# Example usage:
# sol = Solution()
# print(sol.minMeetingRooms([[0, 30], [5, 10], [15, 20]]))  # Output: 2
```


Time and Space Complexity

The provided Python code is meant to determine the minimum number of meeting rooms required for a set of meetings, represented by their start and end times. The code employs a difference array to keep track of the changes in the room occupancy count caused by the start and end of meetings, then uses `accumulate` to find the peak occupancy which gives the required number of rooms.

Time Complexity

The time complexity of the code is determined by three steps:

1. Initializing the `delta` list, which has a fixed length of `1000010`. This step is in `O(1)` since it does not depend on the number of intervals but on a constant size.
2. The first `for` loop that populates the `delta` array with `+1` and `-1` for the start and end times of meetings. It runs `O(n)` times, where `n` is the number of intervals (meetings).
3. The use of `accumulate` which calculates the prefix sum of the `delta` array. In the worst case, this operation is `O(m)`, where `m` is the size of the `delta` array, which is a constant of `1000010`.

The final time complexity is dominated by the larger of the two variables, which in this case is the constant time for using `accumulate`. Hence, the time complexity is `O(m)`, which translates to `O(1)` since `m` is a constant.

Space Complexity

The space complexity of the code is determined primarily by the `delta` array.

1. The `delta` list, which has a fixed length of `1000010`. This is a constant space allocation and does not depend on the input size.
2. No additional data structures that grow with the size of the input are used.

Therefore, the space complexity is `O(1)` due to the constant size of the `delta` array.