347. Top K Frequent Elements

Medium Array Hash Table Divide and Conquer Bucket Sort Counting Quickselect Sorting Heap (Priority Queue)

Leetcode Link

Problem Description

The LeetCode problem provides us with an integer array `nums` and an integer `k`. Our task is to find the `k` most frequent elements in the array. The "frequency" of an element is the number of times it occurs in the array. The problem specifies that we can return the result in any order, which means the sequence of the results does not matter.

Intuition

To solve this problem, we need to count the occurrences of each element and then find the `k` elements with the highest frequencies. The natural approach is to use a hash map (or dictionary in Python) to achieve the frequency count efficiently.

Once we have the frequency of each element, we want to retrieve the `k` elements with the highest frequency. A common data structure to maintain the `k` largest or smallest elements is a heap. In Python, we use a min-heap by default, which ensures that the smallest element is always at the top.

The intuition behind the solution is:

1. Count the frequency of each element using a hash map.
2. Iterate over the frequency map, adding each element along with its frequency as a tuple to a min-heap.
3. If the heap exceeds size `k`, we remove the smallest item, which is automatically done because of the heap's properties. This ensures we only keep the `k` most frequent elements in the heap.
4. After processing all elements, we're left with a heap containing `k` elements with the highest frequency.
5. We convert this heap into a list containing just the elements (not the frequencies) to return as our final answer.

This approach is highly efficient as it allows us to keep only the `k` most frequent elements at all times without having to sort the entire frequency map, which could be much larger than `k`.

Solution Approach

The implementation of the solution uses Python's `Counter` class from the `collections` module to calculate the frequency of each element in the `nums` array. `Counter` is essentially a hash map or a dictionary that maps each element to its frequency.

Here's a step-by-step walkthrough of the implementation:

1. First, we use `Counter(nums)` to create a frequency map that holds the count of each number in the `nums` array.

2. Next, we initialize an empty min-heap `hp` as a list to store tuples of the form `(frequency, num)`, where `frequency` is the frequency of the `num` number in the array.

3. We iterate over each item in the frequency map and add a tuple `(freq, num)` to the heap using the `heappush` function.

4. While we add elements to the heap, we maintain the size of the heap to not exceed `k`. If adding an element causes the heap size to become greater than `k`, we pop the smallest item from the heap using `heappop`. This is done to keep only the `k` most frequent elements in the heap.

5. After we finish processing all elements, the heap contains `k` tuples representing the `k` most frequent elements. The least frequent element is on the top of the min-heap, while the k-th most frequent element is the last one in the heap's binary tree representation.

6. Finally, we build the result list by extracting the `num` from each tuple `(freq, num)` in the heap using a list comprehension: `[v[1] for v in hp]`.

The `Counter` efficiently calculates the frequencies of each element in $O(n)$ time complexity, where n is the length of the input array. The heap operations (insertion and removal) work in $O(\log k)$ time, and since we perform these operations at most n times, the total time complexity of the heap operations is $O(n \log k)$. Thus, the overall time complexity of the solution is $O(n \log k)$, with $O(n)$ coming from the frequency map creation and $O(n \log k)$ from the heap operations. The space complexity of the solution is $O(n)$ to store the frequency map and the heap.

This efficient implementation ensures we're not doing unnecessary work by keeping only the top `k` frequencies in the heap, and it avoids having to sort large sets of data.

Example Walkthrough

Let's use a small example to illustrate the solution approach. Consider the array `nums = [1,2,3,2,1,2]` and `k = 2`. Our goal is to find the 2 most frequent elements in `nums`.

1. We first use `Counter(nums)` to create a frequency map. This gives us `{1: 2, 2: 3, 3: 1}` where the key is the number from `nums` and the value is its frequency.

2. We initialize an empty min-heap `hp`. It's going to store tuples like `(frequency, num)`.

3. We iterate over the frequency map and add each `num` and its frequency to `hp`. For example, `(2, 1)` for the number `1` with a frequency of `2`. We use `heappush` to add the tuples to `hp`, so after this step `hp` might have `[(1, 3), (2, 1)]`.

4. The heap should not exceed the size `k`. In our case, `k` is `2`, which means after we add the third element `(3, 2)`, we need to pop the smallest frequency. So we end up with `hp` as `[(2, 1), (3, 2)]` after all the operations since `(1, 3)` would be the popped element because it had the lowest frequency.

5. The heap now contains the tuples for the 2 most frequent elements. The tuple with the smallest frequency is at the top, ensuring that less frequent elements have been popped off when the size limit was exceeded.

6. Finally, to build our result list, we extract the number from each tuple in the heap. Using list comprehension `[v[1] for v in hp]` we get `[1, 2]`, which are the elements with the highest frequency. This is our final result and we can return it.

Following this approach, we implemented an efficient solution to the problem that avoids sorting the entire frequency map directly and instead maintains a heap of size `k` to track the `k` most frequent elements.

Python Solution

```python
1  from collections import Counter
2  from heapq import heappush, heappop
3
4  class Solution:
5      def topKFrequent(self, nums, k):
6          # Count the frequency of each number in nums using Counter.
7          num_frequencies = Counter(nums)
8
9          # Initialize a min heap to keep track of top k elements.
10         min_heap = []
11
12         # Iterate over the number-frequency pairs.
13         for num, freq in num_frequencies.items():
14             # Push a tuple of (frequency, number) onto the heap.
15             # Python's heapq module creates a min-heap by default.
16             heappush(min_heap, (freq, num))
17
18             # If the heap size exceeds k, remove the smallest frequency element.
19             if len(min_heap) > k:
20                 heappop(min_heap)
21
22         # Extract the top k frequent numbers by taking the second element of each tuple.
23         top_k_frequent = [pair[1] for pair in min_heap]
24
25         return top_k_frequent
```

Java Solution

```java
1  import java.util.*;
2  import java.util.function.Function;
3  import java.util.stream.Collectors;
4
5  class Solution {
6      public int[] topKFrequent(int[] nums, int k) {
7          // Create a Map to store the frequency of each number
8          Map<Integer, Long> frequencyMap = Arrays.stream(nums)
9                  .boxed() // box the ints to Integers
10                 .collect(Collectors.groupingBy(Function.identity(), // group by the number itself
11                         Collectors.counting())); // count the frequency
12
13         // Initialize a min-heap based on the frequency values
14         Queue<Map.Entry<Integer, Long>> minHeap = new PriorityQueue<>(Comparator.comparingLong(Map.Entry::getValue));
15
16         // Iterate over the frequency map
17         for (Map.Entry<Integer, Long> entry : frequencyMap.entrySet()) {
18             // Insert the current entry into the min-heap
19             minHeap.offer(entry);
20
21             // If the heap size exceeds 'k', remove the smallest frequency element
22             if (minHeap.size() > k) {
23                 minHeap.poll();
24             }
25         }
26
27         // Extract the top 'k' frequent numbers from the min-heap into an array
28         return minHeap.stream()
29                 .mapToInt(Map.Entry::getKey)
30                 .toArray();
31     }
32 }
33
```

C++ Solution

```cpp
1  #include <vector>
2  #include <queue>
3  #include <unordered_map>
4  using namespace std;
5
6  // Definition for a pair of integers
7  using IntPair = pair<int, int>;
8
9  class Solution {
10 public:
11     // Function that returns the k most frequent elements from 'nums'.
12     vector<int> topKFrequent(vector<int>& nums, int k) {
13         // HashMap to store the frequency of each number in 'nums'.
14         unordered_map<int, int> frequencyMap;
15         // Increment the frequency count for each number in 'nums'.
16         for (int value : nums) {
17             frequencyMap[value]++;
18         }
19
20         // Min-heap to keep track of the top k frequent numbers.
21         // It stores pairs of (frequency, number) and orders by smallest frequency first.
22         priority_queue<IntPair, vector<IntPair>, greater<IntPair>> minHeap;
23
24         // Iterate over the frequency map.
25         for (const auto& element : frequencyMap) {
26             int number = element.first;
27             int frequency = element.second;
28             // Push the current number and its frequency to the min-heap.
29             minHeap.push({frequency, number});
30
31             // If the heap size exceeds k, remove the least frequent element.
32             // This ensures that the heap always contains the top k frequent elements.
33             if (minHeap.size() > k) {
34                 minHeap.pop();
35             }
36         }
37
38         // Prepare a vector to store the result.
39         vector<int> topKFrequentElements(k);
40         // Retrieve the most frequent elements from the min-heap.
41         for (int i = 0; i < k; ++i) {
42             topKFrequentElements[i] = minHeap.top().second;  // Store the number, not the frequency.
43             minHeap.pop();
44         }
45
46         // Return the top k frequent elements in ascending frequency order.
47         return topKFrequentElements;
48     }
49 };
50
```

Typescript Solution

```typescript
1  function topKFrequent(nums: number[], k: number): number[] {
2      // Initialize a Map to hold the frequency of each number
3      const frequencyMap = new Map<number, number>();
4
5      // Iterate over the array of numbers and populate the frequencyMap
6      for (const num of nums) {
7          frequencyMap.set(num, (frequencyMap.get(num) || 0) + 1);
8      }
9
10     // Convert the Map into an array of key-value pairs
11     const frequencyArray = Array.from(frequencyMap);
12
13     // Sort the array based on frequency in descending order
14     frequencyArray.sort((a, b) => b[1] - a[1]);
15
16     // Initialize an array to hold the top k frequent elements
17     const topKElements: number[] = [];
18
19     // Iterate k times and push the most frequent elements onto the topKElements array
20     for (let i = 0; i < k; ++i) {
21         topKElements.push(frequencyArray[i][0]);
22     }
23
24     // Return the top k frequent elements
25     return topKElements;
26 }
```

Time and Space Complexity

Time Complexity

The time complexity of the function is determined by several factors:

- **Counting Elements**: The function begins with `cnt = Counter(nums)` which counts the frequency of each element in the array. Constructing this frequency count takes $O(N)$ time, where N is the number of elements in `nums`.

- **Heap Operations**: The function then involves a for loop, iterating over the frequency counter's items, and performing heap operations. For each unique element (up to N unique elements), a heap push is performed, which has a time complexity of $O(\log K)$, as the size of the heap is maintained at k. In the worst case, there are N heap push and pop operations, each taking $O(\log K)$ time. Therefore, the complexity due to heap operations is $O(N \cdot \log K)$.

The resulting overall time complexity is $O(N + N \cdot \log K)$. However, since $N \cdot \log K$ is the dominant term, we consider the overall time complexity to be $O(N \cdot \log K)$.

Space Complexity

The space complexity of the function is the additional space required by the data structures used:

- **Frequency Counter**: The `Counter` will at most store N key-value pairs if all elements in `nums` are unique, which takes $O(N)$ space.

- **Heap**: The heap size is maintained at k, so the space required for the heap is $O(k)$.

Therefore, the overall space complexity is $O(N + k)$. However, in most scenarios, we expect k to be much smaller than N, therefore the space complexity is often expressed as $O(N)$.