

528. Random Pick with Weight

Medium

Array

Math

Binary Search

Prefix Sum

Randomized

Leetcode Link

Problem Description

In this problem, we are given an array `w` of positive integers where each integer represents the weight of the corresponding index. The objective is to implement a function called `pickIndex()` that randomly selects an index from 0 to `w.length - 1`. However, this isn't just any random selection; the index must be chosen such that the probability of selecting any index `i` is proportional to its weight `w[i]` relative to the sum of all weights in array `w`. The probability of picking index `i` is calculated by dividing `w[i]` by the sum of all elements in the array `w` (`sum(w)`). The task is to select an index randomly, in a weighted manner, according to these probabilities.

Intuition

The intuition behind the solution is to use a prefix sum array to convert the weights into a range of cumulative sums. A prefix sum array is an array where each element at index `i` stores the sum of all elements of the original array from 0 to `i`. This way, we can represent the weight of each index as a range in the cumulative sum.

Once we have the prefix sum array `self.s`, the idea is to generate a random number `x` between 1 and the sum of the weights (`self.s[-1]`). This random number effectively chooses a "position" within the total weight. Our goal now is to find the index in the original weights array `w` that corresponds to this position if weights were laid out on a number line according to their weight sizes.

We do this by performing a binary search on the prefix sum array to find the smallest prefix sum that is equal to or greater than the randomly picked number `x`. The binary search narrows down the range of possible positions until it finds the correct index whose weight range contains `x`.

This method ensures the index is chosen randomly, and with a probability proportional to its weight, fulfilling the requirement of the problem.

Solution Approach

The solution is implemented in two parts: the constructor `__init__(self, w: List[int])` and the method `pickIndex(self) -> int`. Here's how each part contributes to the overall solution:

- Constructor (`__init__`):** Initialize the Solution class with the given weights array `w`. We calculate a prefix sums array which is stored in `self.s`. This array is built by starting with a 0 and then cumulatively adding the weights from the `w` array. The `self.s` array is one element longer than `w`, where `self.s[i]` represents the sum of weights from `w[0]` through `w[i-1]`.
 - If we have `w = [1, 3, 2]`, the resulting prefix sums array will be `self.s = [0, 1, 4, 6]`. Note how each element in `self.s` represents the cumulative weight up to but not including the current index in `w`.
- `pickIndex` method:** This method is where the random selection takes place, using the prefix sums array `self.s`.
 - First, we pick a random number `x` in the range from 1 to the cumulative weight of all elements (`self.s[-1]`).
 - Then, we perform a binary search to find the first element in the prefix sums array that is greater than or equal to this randomly chosen number `x`. The purpose is to find the segment where this random weight `x` would fall. We do this by maintaining two pointers `left` and `right` and repeatedly narrowing down the search space by adjusting these pointers based on the current middle element (`mid`), until `left` is just less than `right`.
 - The binary search condition `if self.s[mid] >= x` checks if the cumulative weight at `mid` is at least `x`. If so, we search to the left (adjust `right` to `mid`) as we may still find a smaller prefix sum that is still greater than or equal to `x`. Otherwise, we search to the right (set `left` to `mid + 1`) as we need a larger prefix sum to be greater than or equal to `x`.
 - Once the binary search completes, `left` will point to the first prefix sum that is greater than or equal to `x`, and hence `left - 1` will be the index of the weight in array `w` that corresponds to the random number `x`.

This solution efficiently simulates picking an index according to the weights' distribution. It uses the prefix sum to map a uniform distribution to the desired weighted distribution and employs binary search for fast index retrieval.

Example Walkthrough

Let's walkthrough a small example to illustrate the solution approach using the following array `w = [2, 5, 3]`:

Step 1: Initialize the Solution class

- Pass the array `w` to the constructor `__init__()`.
- Create a prefix sums array `self.s`. Starting with `[0]` and then adding each element from `w` cumulatively. For `w = [2, 5, 3]`, the prefix sums array will become `self.s = [0, 2, 7, 10]`.
 - `0` is the starting point.
 - `2` is the sum of weights up to but not including index `1` (`w[0]`).
 - `7` is the sum of weights up to but not including index `2` (`w[0] + w[1]`).
 - `10` is the sum of weights for all indexes (`w[0] + w[1] + w[2]`).

Step 2: Pick a random index with `pickIndex()` method

- Generate a random number `x` between `1` and `10` (the total weight).
- Suppose our random number `x` is `6`.
- Perform binary search to find the smallest element in `self.s` that is greater than or equal to `6`.
 - Initially, `left = 0` and `right = len(self.s) - 1`, which is `3`.
 - Our mid value in the first iteration will be `(0 + 3) // 2 = 1`.
 - Since `self.s[1]` is `2` and it's less than `6`, we make `left = mid + 1`, which is `2`.
 - On the next iteration, `left = 2`, `right = 3`, and so mid will be `(2 + 3) // 2 = 2`.
 - Now `self.s[2]` is `7` which is greater than `6`, set `right` to `mid`, now `right` becomes `2`.
 - The loop terminates when `left` is no longer less than `right`, so `left` will now point to `2`.

Since `left` is now `2`, index `1` (`left - 1`) in the original `w` array will be returned from `pickIndex()`. This is because `6` falls into the cumulative range `(2, 7]` (exclusive of 2 and inclusive of 7), corresponding to the second element (`5`) in the original weights array `w`. The choice of index `1` reflects the higher likelihood due to the weight of `5` in `w`.

This simple example demonstrates the weighted random selection using the prefix sums array and binary search.

Python Solution

```
1 import random
2 from typing import List
3
4 class Solution:
5     def __init__(self, weights: List[int]):
6         # Initialize an empty list to store cumulative weights
7         self.cumulative_weights = [0]
8         # Build up the cumulative weight list for later binary search
9         for weight in weights:
10             self.cumulative_weights.append(self.cumulative_weights[-1] + weight)
11
12     def pickIndex(self) -> int:
13         # Generate a random number between 1 and the total sum of weights
14         target = random.randint(1, self.cumulative_weights[-1])
15         # Perform a binary search to find the target within the cumulative weights
16         left, right = 1, len(self.cumulative_weights) - 1
17         while left < right:
18             # Calculate the middle index
19             mid = (left + right) // 2
20             # Since we want to find the first element that is not less than the target,
21             # move the right pointer to mid if the middle cumulative weight is >= target
22             if self.cumulative_weights[mid] >= target:
23                 right = mid
24             # Otherwise, move the left pointer to one after the current middle
25             else:
26                 left = mid + 1
27         # The final index will be right - 1, since the cumulative_weights includes
28         # an extra 0 at the beginning that we added during initialization
29         return right - 1
30
31 # How to use this class:
32 # Create a Solution object with a given list of weights
33 # obj = Solution(weights)
34 # Pick an index based on the weight distribution
35 # index = obj.pickIndex()
36
```

Java Solution

```
1 import java.util.Random;
2
3 class Solution {
4     private int[] prefixSums; // stores the prefix sums of the weights
5     private Random random = new Random(); // random number generator
6
7     public Solution(int[] weights) {
8         int n = weights.length;
9         prefixSums = new int[n + 1];
10        // Generate prefix sums array where each element represents the sum of weights up to that index.
11        for (int i = 0; i < n; ++i) {
12            prefixSums[i + 1] = prefixSums[i] + weights[i];
13        }
14    }
15
16    public int pickIndex() {
17        // Generate a random number between 1 and the total sum of weights.
18        int x = 1 + random.nextInt(prefixSums[prefixSums.length - 1]);
19        int left = 1, right = prefixSums.length - 1;
20
21        // Perform binary search to find the index for which prefixSums[index] is greater than or equal to x.
22        while (left < right) {
23            int mid = (left + right) >> 1; // Use unsigned right shift to avoid potential overflow
24            if (prefixSums[mid] >= x) {
25                // If the mid-index satisfies the condition, we search the left subarray.
26                right = mid;
27            } else {
28                // Otherwise, we search the right subarray.
29                left = mid + 1;
30            }
31        }
32        // Since we have shifted our prefixSums array by one, we subtract one to get the original index.
33        return left - 1;
34    }
35 }
36
37 /**
38  * The main class where instances of the Solution class can be created and the pickIndex() method can be called.
39  */
40 public class Main {
41     public static void main(String[] args) {
42         int[] weights = {1, 3, 4, 6}; // for example
43         Solution solution = new Solution(weights);
44         int index = solution.pickIndex();
45         System.out.println(index); // The picked index based on the weight
46     }
47 }
48
```

C++ Solution

```
1 #include <vector>
2 #include <cstdlib> // For rand()
3
4 class Solution {
5 public:
6     // Prefix sums array where each element at index i contains the sum of weights up to index i-1
7     std::vector<int> prefixSums;
8
9     // Constructor that initializes the Solution with a vector of weights
10    Solution(std::vector<int>& weights) {
11        int numWeights = weights.size();
12        prefixSums.resize(numWeights + 1);
13        // Build the prefix sums array
14        for (int i = 0; i < numWeights; ++i) {
15            prefixSums[i + 1] = prefixSums[i] + weights[i];
16        }
17    }
18
19    // Function to pick an index based on the weights (the weight at each index indicates the probability of picking that index)
20    int pickIndex() {
21        int numElements = prefixSums.size();
22        // Generate a random number between 1 and the sum of all weights
23        int randomNumber = 1 + rand() % prefixSums[numElements - 1];
24
25        // Perform binary search to find the right index
26        int left = 1, right = numElements - 1;
27        while (left < right) {
28            int mid = left + (right - left) / 2;
29            // If the prefix sum at mid is at least as large as the random number, search to the left
30            if (prefixSums[mid] >= randomNumber)
31                right = mid;
32            else
33                // Else, search to the right
34                left = mid + 1;
35        }
36        // The index in the original array is left-1 because of the extra element at the beginning of prefixSums
37        return left - 1;
38    }
39 };
40
41 /**
42  * Your Solution object will be instantiated and called as such:
43  * Solution* obj = new Solution(weights);
44  * int index = obj->pickIndex();
45  */
46
```

Typescript Solution

```
1 // Define the prefix sum array as a global variable.
2 let prefixSums: number[] = [];
3
4 /**
5  * Initializes the prefix sums array using the input weights.
6  * @param {number[]} weights - The list of weights, which corresponds to probabilities indirectly.
7  */
8 function initialize(weights: number[]): void {
9     const n = weights.length;
10    prefixSums = new Array(n + 1).fill(0);
11    for (let i = 0; i < n; ++i) {
12        prefixSums[i + 1] = prefixSums[i] + weights[i];
13    }
14 }
15
16 /**
17  * Picks an index randomly based on the weights initialized.
18  * The random pick is done using a binary search to find the interval
19  * that the random number falls into considering the prefix sums as intervals.
20  * @return {number} The picked index corresponding to the original weights' distribution.
21  */
22 function pickIndex(): number {
23     const n = prefixSums.length;
24     const randomNum = 1 + Math.floor(Math.random() * prefixSums[n - 1]);
25     let left = 1;
26     let right = n - 1;
27
28     // Binary search to find the smallest index such that prefixSums[index] >= randomNum
29     while (left < right) {
30         const mid = Math.floor((left + right) / 2);
31         if (prefixSums[mid] >= randomNum) {
32             right = mid;
33         } else {
34             left = mid + 1;
35         }
36     }
37     // left - 1 because the prefixSums array starts from 1 to n and we need to return 0 to n-1
38     return left - 1;
39 }
40
41 // Example usage:
42 // initialize([10, 20, 15]); // Initializes the weights
43 // console.log(pickIndex()); // Logs an index, where the probability correlates with weight.
44
```

Time and Space Complexity

Time Complexity

For the `__init__` method:

- The time complexity is $O(n)$, where `n` is the length of the input list `w`. This is because we iterate through the list `w` once to compute the prefix sum array `self.s`.

For the `pickIndex` method:

- The time complexity is $O(\log n)$ because we use binary search to find the index in the prefix sum array. The binary search divides the search space in half during each iteration, which leads to a logarithmic time complexity.

Space Complexity

For both methods:

- The space complexity is $O(n)$ due to the storage required for the prefix sum array `self.s`, which has one more element than the original input list `w`.