582. Kill Process

**Depth-First Search Breadth-First Search** 

# **Problem Description**

Medium <u>Tree</u>

a unique ID. The relationships between processes are given in a tree-like structure, where every process is a node in the tree, and has one parent node except for the root process, which has no parent and is identified by ppid[i] = 0. You are provided with two arrays: pid, which contains the ID of each process, and ppid, which contains the corresponding

In this problem, you are given the task of simulating the killing of processes in an operating system. Each process is identified by

Hash Table

<u>Array</u>

parent process ID for each process. The root of the tree is the process with no parent, and its ID is the one for which ppid[i] = 0.

Your task is to implement a function that, when given the ID of a process to kill, returns a list of all the process IDs that will be terminated. This includes the process itself and all of its descendant processes - in other words, its children, its children's children, and so on. The function should return the list of IDs in any order.

The problem statement implies that when a process is killed, it's not just that single process that stops - all processes that are 'below' it in the tree (its children) are also killed, recursively.

Intuition

To solve this problem, one intuitive approach is to map out the relationships between all processes in a way that allows you to

## easily identify a process's children. This can be effectively done using a graph data structure, where each process is a node and

new key.

the edges represent the parent-child relationships. First, construct a graph from the given pid and ppid arrays. In this graph, the key is the process ID of the parent, and the value is a list of IDs of its child processes. We represent this graph with a dictionary or a hashmap, where each entry corresponds to a process and its children. In Python, a defaultdict(list) is ideal for this because it automatically initializes a new list for each

Once you have the graph, perform a Depth-First Search (DFS) starting from the process ID that needs to be killed. DFS is a suitable algorithm here because it allows us to explore all descendants of a node by going as deep as possible into the tree's branches before backtracking, thus ensuring that we find all processes that need to be terminated. As you perform DFS, each time you visit a node (process), add its ID to the answer list. Once you exhaust all the child nodes of

the process to be killed, the recursive DFS function will naturally backtrack, and the process is complete. The DFS ensures all

child processes are visited and their IDs are added to the list, effectively giving us the full list of processes that will be killed.

The solution code implements this intuition, with the dfs function handling the actual traversal and the killProcess function preparing the graph and initiating the DFS call. Solution Approach

The solution to the problem is implemented in Python using a <u>Depth-First Search</u> (DFS) on the graph representation of the processes. Here's an explanation of the steps taken in the given solution:

The graph is created as a hashmap (specifically, a defaultdict(list) in Python) where each key-value pair corresponds to a

**Graph Construction** 

**Data Structure: Graph** 

Using the zip function, we iterate through both pid and ppid arrays in tandem. For each pair (i, p) from pid and ppid, we add process i to the graph under its parent p. The graph will store all the relations as adjacency lists.

The dfs function performs a classical depth-first search starting from the kill process ID. It takes an integer i as its argument,

## which is the process ID from where the search begins. The function works as follows:

**DFS Function** 

Initiating the Algorithm

After the DFS completes, 'ans' contains all the process IDs that were visited during the recursive search. Since DFS was executed

The solution code encapsulates this approach neatly and efficiently, making sure that by the end of the call to dfs(kill), the list

2. The function iterates over all the children of process i (if any) by looking up graph g[i]. For each child process j, the function calls itself

starting with the process ID of kill, this means it includes the kill process and all its descendant processes, recursively.

1. The current node (process ID) i is added to the list ans of processes to be killed.

recursively (dfs(j)), which leads to the child process and all its descendants being added to ans.

We initialize an empty list ans which will eventually contain all process IDs to be killed.

We then call dfs(kill), which begins the depth-first search starting from the process we want to kill.

ans includes every process that needs to be terminated as specified by the problem statement.

parent process ID (the key) and a list of its child process IDs (the value).

Imagine we have a system with the following processes and parent-child relationships: Processes: pid = [1, 3, 10, 5] Parent Processes: ppid = [0, 1, 3, 3]

ppid[1]=1 ppid[2]=3, ppid[3]=3

Here, process 1 is the root of the tree, and it has one child, process 3. Process 3, in turn, has two children, process 10 and

Step 1: Graph Construction First, we build the graph from pid and ppid arrays using a hashmap. Our graph would look like this:

This results in the graph representation being: {1: [3], 3: [10, 5]} with other processes linking to an empty list, since they

Step 2: Depth-First Search (DFS) Function We want to kill process 3. So, we perform a depth-first search starting from process

## Let's walk through the solution if we want to kill process 3.

don't have children.

Solution Implementation

from collections import defaultdict

return killed\_processes

import java.util.ArrayList;

import iava.util.HashMap;

import java.util.List;

import java.util.Map;

class Solution {

import java.util.Collections;

from typing import List

class Solution:

3.

**Python** 

Java

process 5.

**Example Walkthrough** 

• Key 1 will have a value [3], indicating that process 1's child is process 3. • Key 3 will have a value [10, 5], indicating that process 3's children are processes 10 and 5.

We add process 10 to the ans list. This process has no children, so the DFS on process 10 ends.

• We add process 5 to the ans list. This process has no children, so the DFS on process 5 ends.

def killProcess(self, pid: List[int], ppid: List[int], kill: int) -> List[int]:

# Create a graph using a dictionary where each key is a parent process ID

# Add the current process to the list of killed processes.

# Kill all the subprocesses of the current process.

for subprocess id in process graph[process\_id]:

kill\_all\_subprocesses(subprocess\_id)

# and the value is a list of its child process IDs.

# Return the list of all processes that were killed.

private Map<Integer, List<Integer>> processGraph = new HashMap<>();

private List<Integer> terminatedProcesses = new ArrayList<>();

// It returns the list of process ids that will be terminated.

int numberOfProcesses = pid.size();

// is a list of direct child process ids.

// all processes that will be terminated.

terminatedProcesses.add(processId);

depthFirstSearch(childProcessId);

// and the value is an array of child process ids

for (let i = 0;  $i < pid.length; ++i) {$ 

const processesToKill: number[] = [];

processesToKill.push(currentId);

depthFirstSearch(childId);

// Return the list of all processes to be killed

def kill all subprocesses(process id: int):

killed processes.append(process id)

**}**;

if (!processTree.has(ppid[i])) {

processTree.set(ppid[i], []);

// List to record all processes to be killed

processTree.get(ppid[i])?.push(pid[i]);

// to find all child processes that must be killed

// Add the current process to the kill list

const depthFirstSearch = (currentId: number) => {

const processTree: Map<number, number[]> = new Map();

// Populate the process tree map with child processes

// Helper function to perform depth-first search on the process tree

// Iterate over all child processes of the current process

for (const childId of processTree.get(currentId) ?? []) {

// Start the depth-first search with the process to be killed

// Recursively apply the depth-first search to the children

def killProcess(self, pid: List[int], ppid: List[int], kill: int) -> List[int]:

# Add the current process to the list of killed processes.

# Kill all the subprocesses of the current process.

for subprocess id in process graph[process\_id]:

kill all subprocesses(subprocess id)

process\_graph[parent\_pid].append(child\_pid)

# Return the list of all processes that were killed.

# Start the killing process from the process with ID kill.

# Initialize the list of killed processes.

# Recursive depth-first search (DFS) function to kill a process and its subprocesses.

depthFirstSearch(kill);

return terminatedProcesses;

for (int i = 0; i < numberOfProcesses; ++i) {</pre>

# Recursive depth-first search (DFS) function to kill a process and its subprocesses.

// This method takes in the process id list, parent process id list, and a process id to kill.

processGraph.computeIfAbsent(ppid.get(i), k -> new ArrayList<>()).add(pid.get(i));

for (int childProcessId : processGraph.getOrDefault(processId, Collections.emptyList())) {

public List<Integer> killProcess(List<Integer> pid, List<Integer> ppid, int kill) {

// Build a process graph where the key is the parent process id and the value

// Perform a depth-first search starting from the kill process id to find

// This helper method performs a depth-first search on the process graph.

// Add the current process id to the list of terminated processes.

// Recursively apply depth-first search on all child processes.

• The DFS retraces its steps back to process 3 and starts a DFS for process 5.

filled with the processes that were killed, which are [3, 10, 5].

def kill all subprocesses(process id: int):

killed processes append(process id)

The processes create a tree structure with their parent-child relationships like this:

ppid[0]=0 (root process)

• We add process 3 to the ans list. • We see that process 3 has two children: 10 and 5. We start a DFS for process 10.

problem by identifying the process to kill along with all of its descendants using a DFS approach on the process tree.

Step 3: Return the Result The list ans is returned from the function, which contains all the processes that were killed when

process 3 was terminated. In this case, the final output would be [3, 10, 5]. And so, the algorithm successfully solves the

The DFS has now visited all children and descendants of process 3. The process recursion completes and we have the list ans

process graph = defaultdict(list) for child pid, parent pid in zip(pid, ppid): process\_graph[parent\_pid].append(child pid) # Initialize the list of killed processes. killed processes = [] # Start the killing process from the process with ID kill. kill all subprocesses(kill)

```
private void depthFirstSearch(int processId) {
```

```
C++
#include <vector>
#include <unordered map>
#include <functional> // For std::function
using namespace std;
class Solution {
public:
    // Function to get all the processes that will be killed if process "kill" is terminated.
    vector<int> killProcess(vector<int>& pid, vector<int>& ppid, int kill) {
        // Create a graph that represents the parent-child relationship between processes.
        unordered map<int. vector<int>> processGraph;
        int numOfProcesses = pid.size();
        for (int i = 0; i < numOfProcesses; ++i) {</pre>
            processGraph[ppid[i]].push_back(pid[i]);
        // List that will hold all the processes to kill.
        vector<int> processesToKill;
        // Depth-First Search (DFS) function to traverse the graph and add process IDs to the list.
        @param currentProcess The process ID of the current process in DFS.
        function<void(int)> dfs = [&](int currentProcess) {
            // Add current process to the list of processes to kill.
            processesToKill.push back(currentProcess);
            // Recur for all the processes that the current process is a parent of.
            for (int childProcess : processGraph[currentProcess]) {
                dfs(childProcess);
        };
        // Kick-start DFS from the process we want to kill.
        dfs(kill);
        // Return the final list of processes to be killed.
        return processesToKill;
};
TypeScript
function killProcess(pid: number[], ppid: number[], kill: number): number[] {
    // Create a map to represent the process tree, where the key is parent process id,
```

#### # Create a graph using a dictionary where each key is a parent process ID # and the value is a list of its child process IDs. process graph = defaultdict(list) for child pid, parent pid in zip(pid, ppid):

killed processes = []

kill all subprocesses(kill)

return killed\_processes

Time and Space Complexity

depthFirstSearch(kill);

return processesToKill;

from typing import List

class Solution:

from collections import defaultdict

The input lists pid and ppid are traversed once to build the graph g, contributing O(n) to the time complexity, where n is the number of processes. The depth-first search (DFS) dfs is called once for each node in the graph within its own subtree. Since each node will be visited exactly once, the total time for all DFS calls is also O(n). Thus, the total time complexity of the algorithm is O(n).

The space complexity is O(n) as well, primarily due to two factors: the space taken by the graph g, which can have at most n edges, and the space taken by the ans list. There is also the implicit stack space used by the DFS calls, which in the worst case (when the graph is a straight line), could be <code>0(n)</code> . However, since this stack space does not exceed the size of the input, the overall space complexity remains O(n).