The given LeetCode problem involves creating a log system that can store log entries with a unique ID and timestamp, and then

Problem Description

retrieve logs that fall within a given time range based on varying levels of granularity. A timestamp in this problem is given in the string format "Year:Month:Day:Hour:Minute:Second", where each part of the date-time is zero-padded. The system needs to support two operations: • put(int id, string timestamp): This function takes a log entry's unique ID and its timestamp, then stores it in the system.

- retrieve(string start, string end, string granularity): This method returns the list of log IDs that have timestamps within the inclusive range specified by start and end. The granularity parameter determines how precise the time range should be, ranging
- from the year down to the second. The granularity truncates the timestamp to the specified level of detail, and only that level of detail is considered when filtering logs. For example, if the granularity is "Day", then the logs are filtered based on their year, month, and day, without considering hours,

Intuition

minutes, or seconds. The intuition behind the solution involves storing the log entries efficiently so that they can be easily retrieved based on time range

queries with varying granularities. To retrieve the correct logs, we need to compare only the parts of the timestamps that are

indices are used to truncate the timestamp strings to the required precision. The retrieval function takes advantage of Python's string comparison, which can be used directly on the truncated timestamp strings. Since the timestamps are fixed-length and zero-padded, lexicographical comparison of the strings works equivalently to

relevant to the specified granularity. The solution defines a dictionary d that contains the cutoff indices for each granularity. These

comparing the numerical values. Here's the thought process for arriving at the solution approach:

2. Define a dictionary that maps each granularity to the index up to which the timestamp should be considered. For instance, if the granularity is "Year", we only consider the first four characters of the timestamp. 3. When retrieving logs within a specific range, determine the cutoff index for the timestamps based on granularity. Use slicing to

4. Iterate through all stored logs, compare the truncated timestamp of each log to the truncated start and end timestamps, and

1. Store all logs in a list as they are received. There's no need to sort them since we can filter them during retrieval.

filter accordingly. 5. Return the list of IDs of the logs that satisfy the range and granularity criteria.

def put(self, id: int, timestamp: str) -> None:

granularity to corresponding substring indices.

self.logs.append((id, timestamp))

truncate the start, end, and current log timestamp to the relevant precision.

- By following this approach, we ensure that we can efficiently retrieve the correct log entries without the need for complex date-time parsing or conversion.
- Solution Approach
- The implementation of the LogSystem class involves two key methods: put and retrieve. The put method is straightforward. Each call to put adds a tuple consisting of the id and timestamp to the logs list:

timestamp. The retrieve method does the main work of filtering the log entries according to the provided time range and granularity:

The use of dictionary d to store the indices for each granularity level allows us to easily retrieve the correct substring length to

compare. For example, for granularity Year, d['Year'] would be 4, so start[:4] would give us the year component of the start

1. Upon initialization (__init__), an empty list logs is created to store the log entries. A dictionary d is also initialized to map

2. The put method simply appends the provided log entry to the logs list without any additional processing.

index i. This way, we only compare the parts of the timestamps that matter for the specified granularity.

Here we don't need to worry about sorting or the position of the log entry because the retrieval will handle the search based on the

timestamp.

A comprehensive breakdown of the algorithm used:

consistent format, this works out as effectively as numerical comparison.

1 results = logSystem.retrieve("2017:01:01:00:00:00", "2017:01:31:23:59:59", "Month")

character for a "Month" granularity, which corresponds to the year and month components.

First, we add a few logs to the system with varying timestamps:

logSystem.put(3, "2017:02:01:23:59:59")

In our logs data structure, we now have:

6. The method finally returns a list of IDs whose timestamps fall within the specified range.

i = self.d[granularity] # Find the index up to which we should compare timestamps

1 def retrieve(self, start: str, end: str, granularity: str) -> List[int]:

return [id for id, ts in self.logs if start[:i] <= ts[:i] <= end[:i]]

list comprehension to iterate through all the stored logs and includes the ones that match the range condition. 4. For each log, the comparison is done by slicing the timestamps of the logs and the provided start and end parameters up to

5. We rely on Python's string comparison to compare the sliced timestamps. Since timestamps are zero-padded and have a

3. The retrieve method calculates the index i up to which the timestamps should be considered, based on the granularity. It uses

Example Walkthrough

Let's say we instantiate a LogSystem and perform a series of put and retrieve operations as per the solution approach described.

This implementation is efficient as it separates the storage and querying of log entries. The dictionary d effectively eliminates

complex parsing or time conversion by using string slicing based on predefined indices corresponding to each granularity.

- 1 logSystem = LogSystem() 2 logSystem.put(1, "2017:01:01:23:59:59") logSystem.put(2, "2017:01:02:23:59:59")
- 1 [(1, "2017:01:01:23:59:59"), (2, "2017:01:02:23:59:59"), (3, "2017:02:01:23:59:59")] No sorting is needed; we just store the logs as they come.

The d dictionary in the LogSystem has an entry "Month": 7, indicating that we should compare the timestamps up to the seventh

Suppose we want to retrieve logs with timestamps in January 2017; we can call the retrieve method accordingly:

1. Determine the index i for month granularity, which is 7. Thus, we will compare timestamps up to "2017:01".

range and granularity.

9

11

12

13

14

15

16

24

25

26

27

32

8

9

10

11

12

13

14

15

16

17

18

19

20

64

65

66

69

70

71

72

73

74

75

76

77

78

79

83

84

/**

C++ Solution

1 #include <vector>

2 #include <string>

using std::vector;

6 using std::string;

using std::pair;

10 class LogSystem {

#include <unordered_map>

7 using std::unordered_map;

LogSystem() {

range without complex parsing.

self.logs = []

self.granularity_to_index = {

def put(self, log_id: int, timestamp: str) -> None:

sliced according to the granularity level

28 # Example of how to use the LogSystem class

// A list to store all logs

// based on the granularity

public LogSystem() -

31 # results = obj.retrieve(start, end, granularity)

The retrieve method will:

2. Iterate through the logs:

3. Create a list of the IDs where the condition holds true. In this case, logs 1 and 2 match, while log 3 does not.

As a result, results will have the value [1, 2], meaning log entries with IDs 1 and 2 are the ones that fall within the desired time

This example demonstrates how the implementation allows us to efficiently filter logs based on a specified granularity and time

2 - For log (2, "2017:01:02:23:59:59"), the truncated timestamp is also "2017:01", which is within the range.

3 - For log (3, "2017:02:01:23:59:59"), the truncated timestamp is "2017:02", which is outside the range.

1 - For log (1, "2017:01:01:23:59:59"), the truncated timestamp is "2017:01" which is within the range "2017:01" to "2017:01".

Python Solution class LogSystem: def __init__(self):

Initialize logs list to hold all the log entries as tuples (id, timestamp)

Dictionary to map the granularity level to the index of timestamp string

"Year": 4, # Index where the Year value ends in the timestamp

"Day": 10, # Index where the Day value ends in the timestamp

"Hour": 13, # Index where the Hour value ends in the timestamp

"Minute": 16, # Index where the Minute value ends in the timestamp

"Second": 19, # Index where the Second value ends in the timestamp

Retrieve log_ids of logs where timestamp is within the start and end range,

// Class representing a system to store and retrieve logs based on time granularity

private List<LogEntry> logEntries = new ArrayList<>();

private Map<String, Integer> granularityMap = new HashMap<>();

// A map to store the length of the timestamp string

// Constructor for initializing the granularity map

granularityMap.put("Year", 4);

granularityMap.put("Day", 10);

this.timestamp = timestamp;

// Getter for ID

public int getId() {

// Getter for Timestamp

return timestamp;

public String getTimestamp() {

* The LogSystem class can be used as follows:

granularityMap_["Year"] = 4;

granularityMap_["Day"] = 10;

// considering the specified granularity.

auto startSubstr = start.substr(0, index);

auto endSubstr = end.substr(0, index);

for (const auto& logEntry : logs_) {

* logSystem.put(id, timestamp); // to store a log entry

* List<Integer> results = logSystem.retrieve(start, end, granularity); // to retrieve log entries

// Year granularity ends at the 4th character index (yyyy)

// Day granularity ends at the 10th character index (yyyy-MM-dd)

granularityMap_["Month"] = 7; // Month granularity ends at the 7th character index (yyyy-MM)

// Retrieves IDs of log entries that have timestamps between the range [start, end],

// Substring the timestamp of the current log to match granularity.

// If the log's timestamp is within the range, add its ID to the result.

if (startSubstr <= timestampSubstr && timestampSubstr <= endSubstr) {</pre>

int index = granularityMap_[granularity]; // Index determines the granularity level.

// Substrings of start and end based on the granularity to compare the timestamps.

vector<int> retrieve(string start, string end, string granularity) {

vector<int> result; // To store the matching log entry IDs.

// Iterate through all stored logs to find matching entries.

result.emplace_back(logEntry.first);

auto timestampSubstr = logEntry.second.substr(0, index);

granularityMap_["Hour"] = 13; // Hour granularity ends at the 13th character index (yyyy-MM-dd:HH)

granularityMap_["Minute"] = 16; // Minute granularity ends at the 16th character index (yyyy-MM-dd:HH:mm)

granularityMap_["Second"] = 19; // Second granularity ends at the 19th character index (yyyy-MM-dd:HH:mm:ss)

* LogSystem logSystem = new LogSystem();

return id;

granularityMap.put("Month", 7);

return [log_id for log_id, ts in self.logs if start[:index] <= ts[:index] <= end[:index]]</pre>

"Month": 7, # Index where the Month value ends in the timestamp

Store the log with its id and timestamp 17 self.logs.append((log_id, timestamp)) 18 19 20 def retrieve(self, start: str, end: str, granularity: str) -> list: # Find the index up to which we will compare timestamps based on the granularity 21 index = self.granularity_to_index[granularity]

Java Solution 1 import java.util.ArrayList; 2 import java.util.HashMap;

29 + obj = LogSystem()

30 # obj.put(log_id, timestamp)

import java.util.List;

import java.util.Map;

class LogSystem {

```
21
            granularityMap.put("Hour", 13);
22
            granularityMap.put("Minute", 16);
23
            granularityMap.put("Second", 19);
24
25
26
        // Method to store an individual log entry
27
        public void put(int id, String timestamp) {
28
            logEntries.add(new LogEntry(id, timestamp));
29
30
31
       // Method to retrieve log entries within the start and end
32
       // timestamp based on the given granularity
33
        public List<Integer> retrieve(String start, String end, String granularity) {
34
            List<Integer> result = new ArrayList<>();
35
            int granularityLength = granularityMap.get(granularity);
36
37
            // Extract the relevant portions of the start and end timestamps
38
           // based on the granularity
            String startPrefix = start.substring(0, granularityLength);
39
40
            String endPrefix = end.substring(0, granularityLength);
41
42
            // Loop over all log entries to find the logs within the time range
            for (LogEntry logEntry : logEntries) {
43
                // Extract the relevant portion of the log's timestamp
44
                String logTimestampPrefix = logEntry.getTimestamp().substring(0, granularityLength);
45
46
47
                // Check if the log's timestamp is within the range
                if (startPrefix.compareTo(logTimestampPrefix) <= 0 && logTimestampPrefix.compareTo(endPrefix) <= 0) {</pre>
48
49
                    result.add(logEntry.getId());
50
51
52
            return result;
53
54
55
   // Class representing an individual log entry
   class LogEntry {
58
        private int id;
                                // The ID of the log entry
59
        private String timestamp; // The timestamp of the log entry
60
61
       // Constructor for log entry
62
        public LogEntry(int id, String timestamp) {
63
            this.id = id;
```

19 20 21 // Stores a log entry with a unique ID and timestamp. 22 void put(int id, string timestamp) { 23 logs_.emplace_back(id, timestamp); 24

11 public:

12

13

14

15

16

17

18

25

26

27

28

29

30

31

32

33

34

35

36

37

38

39

40

41

42

43

44

26

27

28

29

30

31

32

33

34

35

36

37

38

40

41

43

42 }

```
45
 46
             return result;
 47
 48
 49
    private:
         vector<pair<int, string>> logs_; // Vector storing log entries as <ID, timestamp> pairs.
 50
 51
         unordered_map<string, int> granularityMap_; // Maps granularity strings to substring index limits.
 52 };
 53
Typescript Solution
  1 // TypeScript code for a log system using global variables and functions
    // Define available granularities and corresponding character index limits in a timestamp
    const granularityMap: { [key: string]: number } = {
         "Year": 4, // Year granularity ends at the 4th character index (yyyy)
         "Month": 7, // Month granularity ends at the 7th character index (yyyy-MM)
  6
         "Day": 10, // Day granularity ends at the 10th character index (yyyy-MM-dd)
         "Hour": 13, // Hour granularity ends at the 13th character index (yyyy-MM-dd:HH)
  8
         "Minute": 16, // Minute granularity ends at the 16th character index (yyyy-MM-dd:HH:mm)
  9
 10
         "Second": 19 // Second granularity ends at the 19th character index (yyyy-MM-dd:HH:mm:ss)
 11 };
 12
 13 // Initialize an array to store log entries as {id, timestamp} objects
    let logs: { id: number; timestamp: string }[] = [];
 15
 16 // Function to store a log entry with a unique ID and timestamp
 17 function put(id: number, timestamp: string): void {
 18
         logs.push({ id, timestamp });
 19 }
 20
 21 // Function to retrieve IDs of log entries that have timestamps between the range [start, end],
 22 // considering the specified granularity
    function retrieve(start: string, end: string, granularity: string): number[] {
         const result: number[] = []; // Array to store the matching log entry IDs
 24
 25
         const index: number = granularityMap[granularity]; // Index determines the granularity level
```

// Substrings of start and end based on the granularity to compare timestamps

// Substring the timestamp of the current log to match granularity

const timestampSubstr: string = logEntry.timestamp.substring(0, index);

// If the log's timestamp is within the range, add its ID to the result

if (startSubstr <= timestampSubstr && timestampSubstr <= endSubstr) {</pre>

const startSubstr: string = start.substring(0, index);

// Iterate through all stored logs to find matching entries

const endSubstr: string = end.substring(0, index);

put method: The put operation has a time complexity of 0(1) since it only involves appending a tuple (ID and timestamp) to the end of the self.logs list which is a constant time operation.

Time and Space Complexity

for (const logEntry of logs) {

return result;

result.push(logEntry.id);

being compared. The complexity of the comparison operation is assumed to be 0(1) since the length of the substring depends on

retrieve can also be O(n).

proportional to the number of entries.

Time Complexity

the granularity and doesn't change with the size of self.logs. **Space Complexity** Overall: The space complexity of the LogSystem class is O(n), where n is the number of log entries stored in self.logs. Each log entry requires a fixed amount of space for the integer ID and a fixed-length string for the timestamp, resulting in space directly

retrieve method: The retrieve operation involves iterating over each log entry in self.logs and comparing the substrings of their

timestamps. The time complexity is 0(n * s) where n is the number of log entries and s is the length of the timestamp substring

retrieve method: The retrieve method creates a new list of IDs that match the criteria, the space complexity of which is O(k) where k is the number of matching log entries returned. Since k can be at most n in the worst case, the worst-case space complexity for

put method: For the put method, it only appends the log data without needing any extra space besides the self.logs, so the space complexity incurred by this method is 0(1) in addition to the overall 0(n) space needed for storing the logs themselves.