636. Exclusive Time of Functions

Problem Description

Medium Stack Array

respective timestamps. The logs are formatted with each entry looking like this: "function_id:start_or_end:timestamp". An important detail is that functions may start and end multiple times, as recursion is possible in the execution.

In this problem, we're dealing with a single-threaded CPU that executes a series of function calls. Each function has a unique ID

ranging from 0 to n-1, where n is the total number of functions. We have logs of when each function starts and ends, with their

The task is to calculate the exclusive time for each function, which is the total time that function spent executing in the CPU. By "exclusive," we mean that we're only interested in the time periods when that function was the one currently being executed by

the CPU (top of the call stack). The outcome should be an array of integers where each index i corresponds to the function with i and the value is that function's exclusive time.

To solve this problem, we can simulate the call stack of the CPU with a stack data structure. We iterate through the log entries,

keeping track of the current timestamp, we can calculate how much time has passed since we last encountered a log entry. The solution involves a few key steps: Initialize an answer array to hold the exclusive times for each function, all set initially to zero. • Use a stack to keep track of active function calls (function IDs) where the most recent call is at the top.

processing start and end signals. When a function starts, we add its ID to the stack. When it ends, we pop it from the stack. By

processing any function yet. Iteratively process each log entry:

- Split the log into parts to extract the function ID, the action (start or end), and the timestamp.

• Define a variable to keep track of the current time (curr). Initially, we set it to an invalid timestamp (e.g., -1) because we haven't started

- new timestamp.
- If the action is start, update the exclusive time of the function at the top of the stack by adding the difference between the current timestamp and the previous curr value (if the stack is not empty). Then push the new function ID onto the stack and update curr to the

If the action is end, pop the function ID from the stack, add the time it took to execute to the corresponding index in the answer array

- (inclusive of start and end timestamp), and update curr to one more than the current timestamp to account for the passage of time. By the end of processing, the answer array will contain the exclusive times for each function, fulfilling the requirements of the problem.
 - The solution employs a stack data structure and a list to maintain the exclusive times for each function. The stack is used to simulate the behavior of a call stack in a single-threaded CPU, where only the top function is currently executing, and each
- nested function call pushes a new frame onto it. Here's the step-by-step implementation walkthrough:

Initialize an array called ans with n zeros, where n is the total number of functions. This array will store the exclusive time for

Set a variable curr to -1. This will be used to store the most recent timestamp we've processed.

each function.

Solution Approach

Loop through each log entry in logs. For each log: a. Split the string using log.split(':') to get the function ID (fid), the type of log ('start' or 'end'), and the timestamp (ts). b. Convert fid and ts to integers.

If the type of log is 'start': a. If the stack is not empty, update the top function's exclusive time by adding the difference

If the type of log is 'end': a. Pop the last function ID off the stack with fid = stk.pop() since this function has finished

executing. b. Add the difference between the current timestamp and the previous curr value, inclusive of the end timestamp,

to ans[fid]. This adds the time for the function execution to ans[fid]. c. Update curr to ts + 1 as the next function (if

between the current timestamp and the previous curr value to ans[stk[-1]]. The top function of the stack would have been

any) would start executing at the next timestamp.

logs = ["0:start:0", "1:start:2", "1:end:5", "0:end:6"]

To calculate the exclusive times of these functions, follow these steps:

Since the stack is empty, we don't update any function's exclusive time.

Push fid onto the stack, so stk becomes [0].

Process the second log entry "1:start:2".

Split the log into fid = 1, type = 'start', and ts = 2.

Push fid onto the stack, so stk becomes [0, 1].

Initialize ans = [0, 0] to store the exclusive times for the 2 functions.

Create an empty list named stk to represent the call stack.

- executing so far, and we now have to pause its timer. b. Push the new function ID onto the stack. stk.append(fid) c. Update curr to the current timestamp ts.
- After processing all the logs, return the ans array, which now contains the exclusive times for each function. With this approach, we're able to simulate the execution of the functions on the CPU and calculate the exclusive times as needed. The use of the stack is crucial here because it correctly models the nested nature of function calls. Additionally, carefully
- **Example Walkthrough** Let's consider a small example to clearly illustrate the solution approach. Suppose we have n = 2 functions, and the logs for their execution on the single-threaded CPU are given as follows:
- Create an empty stack stk.

Since the stack is not empty, update function 0's exclusive time. ans [0] += 2 - 0, so ans becomes [2, 0].

managing the curr timestamp allows us to measure execution times correctly as we process the logs.

Set $\frac{1}{1}$ as the initial timestamp before any functions have started. Process the first log entry "0:start:0". Split the log into fid = 0, type = 'start', and ts = 0.

```
    Split the log into fid = 1, type = 'end', and ts = 5.
```

Process the third log entry "1:end:5".

Update curr to 0.

Update curr to 2.

class Solution:

- Pop fid from the stack (1 is popped, leaving stk as [0]). ○ Update function 1's exclusive time. ans [1] += 5 - 2 + 1, so ans becomes [2, 4]. ○ Update curr to 5 + 1.
- Process the fourth log entry "0:end:6".
- Split the log into fid = 0, type = 'end', and ts = 6. Pop fid from the stack (0 is popped, leaving stk as []).

exclusive_times = [0] * n

Process each log entry.

parts = log.split(':')

if call stack:

function_id = int(parts[0])

Split the string log into parts.

call_stack = []

 $current_time = -1$

for log in logs:

 Update function 0's exclusive time. ans[0] += 6 - 5 + 1, so ans becomes [4, 4]. ○ Update curr to 6 + 1.

of 4 units of time as well. Therefore, the final answer is [4, 4].

def exclusiveTime(self, n: int, logs: List[str]) -> List[int]:

Use a stack to keep track of the function call hierarchy.

Push the current function onto the stack.

Pop the last function from the stack.

Update the current time to the new start time.

Update the exclusive time for this function.

exclusive_times[function_id] += timestamp - current_time + 1

// Split the log into parts: functionId, event type, and timestamp.

// If the stack is not empty, update the exclusive time of the function on the top.

exclusiveTimes[functionStack.peek()] += timestamp - prevTimestamp;

call_stack.append(function_id)

else: # Otherwise, it's an end event.

function_id = call_stack.pop()

Return the list of calculated exclusive times.

int[] exclusiveTimes = new int[numFunctions];

// Stack to keep track of the function calls.

String[] parts = log.split(":");

if ("start".equals(parts[1])) {

// Check if the event type is start.

if (!functionStack.isEmpty()) {

Deque<Integer> functionStack = new ArrayDeque<>();

int functionId = Integer.parseInt(parts[0]);

} else { // If the event is a function ending

times[functionId] += timestamp - previousTime + 1;

functionId = functionStack.top();

function exclusiveTime(n: number, logs: string[]): number[] {

// Iterate through each log in the logs array.

if (callStack.length !== 0) {

if (callStack.length !== 0) {

callStack.push([functionId, timestamp]);

const startedFunction = callStack.pop();

const parts = log.split(':');

if (state === 'start') {

for (const log of logs) {

} else {

functionStack.pop();

// Pop the top function, add the time since the last event to it,

// and add one because an "end" logs the time at the end of the unit

return times; // Return the vector with calculated exclusive times for each function

const callStack: [number, number][] = []; // Stack to keep track of function calls.

const previousFunction = callStack[callStack.length - 1];

// If there's another function that was paused, resume its timer.

callStack[callStack.length - 1][1] = timestamp + 1;

const functionExecTimes = new Array(n).fill(0); // Array to store exclusive times for each function.

const [functionId, state, timestamp] = [Number(parts[0]), parts[1], Number(parts[2])];

functionExecTimes[previousFunction[0]] += timestamp - previousFunction[1];

// Start the new function by pushing it onto the stack with its start time.

functionExecTimes[startedFunction[0]] += timestamp - startedFunction[1] + 1;

Update the current time to the end time + 1 since the next time unit

will indicate the start of the next event.

current_time = timestamp + 1

return exclusive_times

Time and Space Complexity

Return the list of calculated exclusive times.

// Splitting the log string into useful parts: function id, state ('start' or 'end'), and timestamp.

// If there is a function currently being executed, pause it and update its execution time.

// Function is ending, pop from stack to get its start time and calculate the time spent.

int timestamp = Integer.parseInt(parts[2]);

// Previous timestamp to calculate the duration.

public int[] exclusiveTime(int numFunctions, List<String> logs) {

// This array holds the exclusive time for each function.

current time = timestamp

Solution Implementation **Python**

Initialize a list to hold the exclusive time for each function.

Store a pointer to the current time (initially set to an invalid time).

The function id is the first part of the log, converted to an integer.

exclusive_times[call_stack[-1]] += timestamp - current_time

The timestamp is the third part of the log, converted to an integer. timestamp = int(parts[2]) # If the log entry indicates a start event: if parts[1] == 'start': # If there's an ongoing function, update its exclusive time.

After processing all log entries, the ans array [4, 4] contains the exclusive execution times for functions 0 and 1. Function 0

was active from time 0 to 2 and from 6 to 6 for a total of 4 units of time, and function 1 was active from time 2 to 5 for a total

Update the current time to the end time + 1 since the next time unit # will indicate the start of the next event. current_time = timestamp + 1

Java

class Solution {

return exclusive_times

int prevTimestamp = -1;

for (String log : logs) {

// Iterate over each log entry.

```
// Push the current function to the stack.
                functionStack.push(functionId);
                // Update the 'prevTimestamp' with the current timestamp.
                prevTimestamp = timestamp;
            } else { // The event type is end.
                // Pop the function from the stack as it ends.
                functionId = functionStack.pop();
                // Update the exclusive time for the popped function.
                exclusiveTimes[functionId] += timestamp - prevTimestamp + 1;
                // Move to the next timestamp as this event marks the end of the function.
                prevTimestamp = timestamp + 1;
       // Return the computed exclusive times for all functions.
       return exclusiveTimes;
C++
#include <vector>
#include <string>
#include <stack>
using namespace std;
class Solution {
public:
   vector<int> exclusiveTime(int n, vector<string>& logs) {
        vector<int> times(n, 0); // Vector to store exclusive times for each function
        stack<int> functionStack; // Stack to keep track of function calls
        int previousTime = 0; // The time of the last event
        for (string& log : logs) {
            char type[10]; // To store whether the event is a 'start' or 'end'
            int functionId, timestamp;
            // Parse the log string: "<function id>:<start/end>:<timestamp>"
            sscanf(log.c_str(), "%d:%[^:]:%d", &functionId, type, &timestamp);
            if (type[0] == 's') { // If the event is a function starting
                if (!functionStack.empty()) {
                    // Add the time since the last event to the currently running function
                    times[functionStack.top()] += timestamp - previousTime;
                previousTime = timestamp: // Update the time of the last event
                functionStack.push(functionId); // Push functionId onto the stack
```

previous Time = timestamp + 1; // Update the time of the last event, moving past the end timestamp

};

TypeScript

```
// Return the array containing the exclusive times for each function.
   return functionExecTimes;
class Solution:
   def exclusiveTime(self, n: int, logs: List[str]) -> List[int]:
       # Initialize a list to hold the exclusive time for each function.
       exclusive_times = [0] * n
       # Use a stack to keep track of the function call hierarchy.
       call_stack = []
       # Store a pointer to the current time (initially set to an invalid time).
       current_time = -1
       # Process each log entry.
       for log in logs:
           # Split the string log into parts.
           parts = log.split(':')
           # The function id is the first part of the log, converted to an integer.
            function_id = int(parts[0])
           # The timestamp is the third part of the log, converted to an integer.
            timestamp = int(parts[2])
           # If the log entry indicates a start event:
           if parts[1] == 'start':
               # If there's an ongoing function, update its exclusive time.
               if call stack:
                   exclusive_times[call_stack[-1]] += timestamp - current_time
               # Push the current function onto the stack.
                call_stack.append(function_id)
               # Update the current time to the new start time.
               current time = timestamp
           else: # Otherwise, it's an end event.
               # Pop the last function from the stack.
               function_id = call_stack.pop()
               # Update the exclusive time for this function.
               exclusive_times[function_id] += timestamp - current_time + 1
```

The time complexity of the code is O(N) where N is the number of logs. This is because the code iterates through the list of logs only once, and for each log, it does a constant amount of work (splitting the log into parts, accessing the last element from the stack, adding time to the answer list).

Time Complexity

Space Complexity

The space complexity of the code is O(N) as well. In the worst case, where all function calls are started before any is finished, the stack (named stk in the code) would grow to contain all N function ids. Additionally, the answer list (named ans) requires space for n integers, where n is the number of functions (which is different from N, the number of logs).