644. Maximum Average Subarray II Binary Search Prefix Sum Hard Array

Problem Description You're given an integer array nums which contains n elements, and an integer k. The task is to find a contiguous subarray whose

allowed a small margin for error. Specifically, any answer within 10^-5 (0.00001) of the actual answer will be considered correct. A subarray, in this context, is a sequence of elements from the array nums that lies next to each other, without any gaps. For example, if nums is [1, 3, 5, 7, 9] and k is 2, we have to find a contiguous sequence of length 2 or more that gives the highest average. The

length is at least k which has the maximum possible average value. It's important to note that the answer needs to be accurate but is

Leetcode Link

returned value should be the average of that subarray. The essence of the problem is to find the optimal section of the list - it can be from k to n elements long - and calculate the average of the numbers within this range. The difficulty lies in doing this efficiently, as brute force methods that check all possible subarrays

will be too slow when the array nums is large. Intuition

The intuition behind the solution involves binary search and prefix sums. Instead of directly searching for the maximum average, we

flip the problem on its head: for a given average v, determine if there's a subarray with an average greater than or equal to v. Then

tricks:

we use binary search on v, narrowing down the range until we find the maximum average to satisfy the condition. Here's how we arrive at the solution approach: Initially, we know that the maximum average lies between the minimum and maximum values in our array nums. These are our

We define a check function to verify if a subarray exists that has an average at least as large as v. This involves some smart

 Calculate the prefix sum of the array elements minus v times k. This transforms the problem into finding a subarray sum that is non-negative.

initial lower (1) and upper (r) bounds for binary search.

- Use a running sum (s) to keep track of the current sum of the last k elements, and a minimum sum (mi) to keep the smallest sum of any k elements we've seen. We update these as we move the running sum forward through nums. If, at any point, our running sum minus the minimum sum we've seen so far (m1) is non-negative, it means there is a subarray
- with an average of at least v. The binary search operates by repeatedly checking the midpoint of our 1 and r range against the check function. If the function
- returns true, we know an average of at least mid is possible, so we move the lower bound 1 to mid. Otherwise, we move the upper bound r to mid.
- We continue the binary search until our range is less than 10^-5, at which point we can return either 1 or r as our result, since they will be sufficiently close to the maximum average. This solution approach is much more efficient than checking every possible subarray, as it takes advantage of the properties of

averages and the structure of contiguous subarrays to reduce the problem to a binary search over possible average values.

The solution uses a combination of binary search and prefix sums to efficiently calculate the maximum average subarray. Here is a step-by-step explanation of how the implementation works:

1. Binary Search Initialization: We start by establishing our search range for the possible maximum average, with 1 being the

minimum value in nums and r being the maximum value. These represent the absolute possible bounds for the average.

2. Helper Function - check(): This function is crucial; it checks if there exists a contiguous subarray with an average value of at least v. Here's how it operates: Sum up the first k elements of nums, then subtract k*v from this sum to get s. This operation shifts the problem from finding

a maximum average subarray to finding a non-negative sum subarray.

o and mi is the smallest possible value.

this loop:

Calculate the midpoint mid as (1 + r) / 2.

Solution Approach

 Iterate over the array starting from the k-th element. For each element at index i, add nums[i] - v to s and nums[i - k] v to t, effectively moving the k-length window one step forward.

Then, update mi to be the minimum of itself and t, which is the sum of the first i-k elements of the array minus k*v. This

Initialize two variables, t and mi, that represent the total sum and minimum sum encountered so far, respectively. Initially, t is

∘ If s - mi >= 0, the function returns True, indicating that it's possible to have an average of at least v. This is because we found a contiguous subarray where the sum of the numbers minus k*v is at least zero, which means the average of that subarray is at least v.

3. Binary Search Loop: The binary search loop keeps iterating while r - 1 is greater than a very small epsilon value (1e-5). Within

Use the check() function with this midpoint value. If check(mid) returns True, we know that it's possible to have a subarray

with an average greater than or equal to mid, so we move the lower bound up to mid. If the result is False, we make the

tracks the smallest sum we have seen up to that point, before considering the current k-length window.

upper bound r equal to mid. This continuously narrows the range of possible averages.

subarray individually, moving the solution from possibly quadratic time complexity to linearithmic.

 \circ Add nums[i] - v to s (running sum), so s += 2 - 6.5, and we get s = -4.5.

Add nums[i - k] - v to t. For i = 2, it's 1 - 6.5 (since k=2), so t = -5.5.

mi was updated along the way and it's min(mi, t); we'll consider it -5.5.

or r would serve as the result for the maximum average, which is approximately 6.5.

Helper function to check if average value v can be the result

If the current average is already >= 0, return True

Keep track of the minimum sum encountered so far

Initialize the sum of the first k elements adjusted by subtracting v

If the current window sum is greater than any seen before, return True

Setting the lower and upper bounds of binary search to the min and max of nums

If the current mid value can be an average, update the lower bound

// After the loop, lowerBound is the maximum average within specified precision

// Calculate the initial sum of the first k elements reduced by the averageValue

// Helper method to check if we can find an average larger than 'averageValue'

private boolean canFindLargerAverage(int[] nums, int k, double averageValue) {

double prevSum = 0; // Sum of the values from the start to i - k

sum += nums[i] - averageValue; // Increment current sum

double minPrevSum = 0; // Minimum sum encountered so far for prevSum

prevSum += nums[i - k] - averageValue; // Increment previous sum

// then there exists a subarray with an average ≥ averageValue

// If the current sum — the smallest sum of prevSum is non-negative,

// Function to find the maximum average subarray of size 'k' in the given vector 'nums'

// Lambda function to check if there's a subarray with average greater than 'value'

// Initial sum of the first 'k' elements with 'value' subtracted from each one

double sum_excluding_first_i_elements = 0, min_sum_excluding_first_i_elements = 0;

// If the sum is greater than or equal to the minimum sum, we can return true

// Check the rest of the array to find if any subarray can have an average larger than 'value'

sum_excluding_first_i_elements += nums[i - k] - value; // Update the sum excluding the first i elements

min_sum_excluding_first_i_elements = min(min_sum_excluding_first_i_elements, sum_excluding_first_i_elements); // Up

minPrevSum = Math.min(minPrevSum, prevSum); // Update the minimum of previous sums

The result is the left bound after the binary search loop ends

def findMaxAverage(self, nums: List[int], k: int) -> float:

def can_be_average(v: float) -> bool:

if current_sum >= 0:

return True

prev_sum = min_sum = 0

 $current_sum = sum(nums[:k]) - k * v$

Iterate over the rest of the elements

min_sum = min(min_sum, prev_sum)

Defining precision for the binary search result

Binary search routine to find maximum average

Otherwise, update the upper bound

if current_sum >= min_sum:

return True

left, right = min(nums), max(nums)

while right - left >= precision:

mid = (left + right) / 2

if can_be_average(mid):

left = mid

right = mid

return lowerBound;

double sum = 0;

if (sum >= 0) {

return true;

for (int i = 0; i < k; ++i) {

sum += nums[i] - averageValue;

// Iterate through the rest of the elements

for (int i = k; i < nums.length; ++i) {</pre>

double findMaxAverage(vector<int>& nums, int k) {

double epsilon = 1e-5;

double sum = 0;

if (sum >= 0) {

return false;

left = mid;

} else {

return true;

// Epsilon value to control the precision of our answer

// Initialize left and right boundaries for binary search

// If the sum is already non-negative, we can return true

if (sum >= min_sum_excluding_first_i_elements) {

sum += nums[i] - value; // Add the next element to 'sum'

// If no subarray has an average greater than 'value', return false

double left = *min_element(nums.begin(), nums.end());

auto canFindLargerAverage = [&](double value) {

for (int i = k; i < nums.size(); ++i) {</pre>

for (int i = 0; i < k; ++i) {

return true;

sum += nums[i] - value;

double right = *max_element(nums.begin(), nums.end());

if (sum >= minPrevSum) {

return true;

return false;

return False

precision = 1e-5

else:

return left

Java Solution

class Solution {

t has been updated as we moved along, let's say it is now -1.5.

accurate representation of the maximum average, within an error margin of 10^-5. The solution elegantly combines binary search, which is a classic divide and conquer algorithm, with a manipulation of sums to

min(nums))), as we have to run our check once for each iteration of binary search, and the range of the binary search is dictated by

the values within nums. The linear pass in the check function, where we iterate over the array updates running sums, takes O(n) time.

reframe the problem into one that is far more computationally efficient. The running time complexity is 0(n * log(max(nums) -

The constant factor has been reduced significantly as compared to a brute force method that might consider every possible

4. Getting the Result: Once the binary search loop exits, the value of 1 (or r, as they are very close to each other) is a sufficiently

Let's illustrate the solution approach with a small example. Assume we have an array nums = [1, 12, 2, 6, 7] and k = 2. Step 1: Binary Search Initialization

• The minimum value in nums is 1, and the maximum value is 12. So we initialize our binary search range with l = 1 and r = 12.

Step 2: Helper Function - check() For illustration purposes, let's assume the binary search is at a point where it's testing v = 6.5. We would initialize s as the sum of the first k elements minus k*v, which is 1 + 12 - 2*6.5 = 0. Initialize t and mi as well. Here, t will start at 0 and mi will be 0, as we have not seen any sum yet.

• Start iterating from the k-th element of nums (index 1): For each element nums [i], where i >= k, do the following (example for i

Update mi to be the minimum of mi and t. As this is the first update and t is less than our initial mi, we now set mi = t =

\circ s += 7 - 6.5, now s = 0.5 - 4.5 = -4.

Continuing this process, when i = 4 (nums [i] = 7):

set l = 6.5. If it returned False, we'd set r = 6.5.

average subarray to within the accepted error margin.

Continue until r - 1 < 10^-5.

-5.5.

Step 3: Iterate and Check

= 2):

Example Walkthrough

>= 0. So, a subarray with an average at least as high as 6.5 exists. Step 4: Binary Search Loop

When we check s - mi for each i, we are looking for it to be >= 0. In this case, at i = 4, s - mi is -4 - (-5.5) = 1.5, which is

We'd repeat the above process for different values of v, adjusting 1 and r with each iteration. If check(6.5) returned True, we'd

Step 5: Result • Once 1 and r are within 10^-5 of each other, the binary search terminates, and either value gives us the maximum possible

In this example, let's say the binary search concluded with l = 6.49995 and r = 6.50000, with an error margin below $l0^-5$, either l

for i in range(k, len(nums)): 14 # Update the sum for the new window by including the new element and excluding the old 15 current_sum += nums[i] - v 16 # Update the sum for the previous window 17 $prev_sum += nums[i - k] - v$

Python Solution

class Solution:

10

12

13

19

20

21

23

24

25

26

27

28

29

30

31

32

33

34

35

36

37

38

39

40

41

42

21

22

23

24

25

26

27

28

29

30

31

32

33

34

35

36

37

38

39

40

41

42

43

44

45

46

47

48

49

50

51

52

54

55

57

8

9

10

11

12

13

14

15

16

17

18

19

20

21

22

23

24

25

26

27

28

29

30

31

32

33

34

35

36

37

38

45

46

56 }

C++ Solution

public:

1 class Solution {

from typing import List

```
// Calculates the maximum average of any subarray of length k
        public double findMaxAverage(int[] nums, int k) {
            double precision = 1e-5;
            double lowerBound = 1e10;
            double upperBound = -1e10;
           // Finding the lowest and highest numbers in the input array nums
            for (int num : nums) {
 9
                lowerBound = Math.min(lowerBound, num);
                upperBound = Math.max(upperBound, num);
12
13
14
            // Binary search to find the maximum average subarray within the precision range
           while (upperBound - lowerBound >= precision) {
15
                double mid = (lowerBound + upperBound) / 2;
16
17
                if (canFindLargerAverage(nums, k, mid)) {
18
                    lowerBound = mid;
19
                } else {
20
                    upperBound = mid;
```

39 // Perform a binary search to find the maximum average within an epsilon range while (right - left >= epsilon) { 40 41 // Calculate mid-point of the range double mid = (left + right) / 2; 42 // Use the lambda function to check if we can find a larger average 43 if (canFindLargerAverage(mid)) { 44

};

```
47
                     right = mid;
 48
 49
             // Return the maximum average which is on the 'left' due to the way we updated it in binary search
 50
 51
             return left;
 52
 53 };
 54
Typescript Solution
     function findMaxAverage(nums: number[], k: number): number {
         const epsilon = 1e-5; // This is the precision for the floating-point comparison.
         let left = Math.min(...nums); // Initialize 'left' to the smallest element in the array.
         let right = Math.max(...nums); // Initialize 'right' to the largest element in the array.
         // 'check' function checks if there exists a subarray of length 'k' with average value greater than or equal to 'targetAverage'
         const check = (targetAverage: number): boolean => {
             // Initial sum of the first 'k' elements adjusted by 'targetAverage'.
             let sum = nums.slice(0, k).reduce((a, b) => a + b) - targetAverage * k;
 10
             if (sum >= 0) {
                 return true; // The average of the first 'k' elements is already greater than 'targetAverage'.
 11
 12
 13
 14
             // 'totalSum' stores the sum of previous elements adjusted by 'targetAverage' in the sliding window.
 15
             // 'minPrevSum' stores the minimum sum encountered in the sliding window.
 16
             let totalSum = 0;
             let minPrevSum = 0;
 17
 18
 19
             for (let i = k; i < nums.length; ++i) {</pre>
 20
                 sum += nums[i] - targetAverage; // Include the next element into 'sum' while shifting the window.
 21
                 totalSum += nums[i - k] - targetAverage; // Include the element that is exiting the window into 'totalSum'.
 22
 23
                 // Update 'minPrevSum' to the minimum sum we have seen so far.
 24
                 minPrevSum = Math.min(minPrevSum, totalSum);
 25
 26
                 // If the current sum adjusted by the minimum previous sum we've seen is non-negative,
 27
                 // it means there exists a subarray with an average at least 'targetAverage'.
 28
                 if (sum >= minPrevSum) {
 29
                     return true;
 30
 31
 32
             return false;
         };
 33
 34
 35
         // Use binary search to find the maximum average to a precision of 'epsilon'.
 36
         while (right - left >= epsilon) {
 37
             const mid = (left + right) / 2; // Calculate the mid point between left and right.
```

left = mid; // If there exists a subarray with average greater than 'mid', move 'left' to 'mid'.

The time complexity of the findMaxAverage function consists of two main parts: the binary search over the range of values and the sliding window check within the check function.

Time and Space Complexity

• Within each iteration of the binary search, the check function is called once. The check function uses a sliding window approach that takes O(n) time, where n is the length of the nums list.

The binary search runs in O(log((max(nums) - min(nums)) / eps)) time because it narrows down the range [min(nums),

Space Complexity

Combining these two parts, the overall time complexity of the function is 0(n * log((max(nums) - min(nums)) / eps)).

input size is required. Together, the function achieves linearithmic time complexity and constant space complexity.

if (check(mid)) {

right = mid; // Otherwise, move 'right' to 'mid'.

max(nums)] by half in each iteration until the range is smaller than eps (= 1e-5).

// 'left' will contain the maximum average subarray value within the desired precision.

} else {

return left;

Time Complexity

38

39

40

41

42

43

44

45

47

46 }

The space complexity of the function is 0(1).

• There are only a constant number of extra variables used (s, t, mi, eps, l, r, mid), and no additional space that scales with the