658. Find K Closest Elements

Two Pointers Binary Search Sorting

Problem Description

Medium Array

Given a sorted array of integers arr, and two integers k and x, the task is to find the k closest integers to x in the array. The result should be returned in ascending order. To determine which integers are the closest, we follow these rules:

Sliding Window

Heap (Priority Queue)

- 1. An integer a is considered closer to x than an integer b if the absolute difference |a x| is less than |b x|. 2. If the absolute differences are the same, then a is closer to x than b if a is smaller than b.
- In essence, the problem asks us to find a subsequence of the array that contains integers that are nearest to x, with a special

emphasis on the absolute difference and the value of the integers themselves when the differences are the same.

To solve the problem, we need an efficient way to find the subsequence of length k out of the sorted array that is closest to the

Intuition

target integer x.

Method 1: Sort A brute force approach would be to sort the elements of the array based on their distance from x. After sorting, we can simply take the first k elements. However, this method is not optimal in terms of time complexity because sorting would

take O(n log n) time. Method 2: Binary search A more efficient approach utilizes the fact that the array is already sorted. We know that the k closest integers we are looking for must form a consecutive subsequence in the array. Instead of looking for each integer separately, we

should look for the starting index of this subsequence. We can use a binary search strategy to find this starting index. The highlevel steps are: 1. Set two pointers: left at the start of the array and right at the position len(arr) - k because any subsequence starting beyond this point

does not have enough elements to be k long. 2. Perform a binary search between left and right: Calculate the middle index mid. Check if x is closer to arr[mid] or to arr[mid + k].

- o If x is closer to arr[mid] or at the same distance to both, we discard the elements to the right of mid + k because the best starting point for our subsequence must be to the left or inclusive of mid.
- If x is closer to arr[mid + k], we discard the elements to the left of mid since the best starting point must be to the right. 3. Keep narrowing down until left is equal to right, indicating that we have found the starting index of the k closest elements.

step-by-step breakdown of the approach and the code implementation:

between x and the potential upper bound for our sequence at arr[mid + k].

start of the subarray containing the k closest numbers to x.

Solution Approach The solution implements a binary search algorithm to efficiently find the starting index for the k closest elements to x. Here's a

consecutive entrants in the array, and starting any further to the right would leave insufficient elements to reach a count of k.

+ 1.

yielding the correct answer.

We enter a loop that continues until left is no longer less than right, indicating that we have narrowed down to the exact starting index for the k closest elements. Inside the loop, we find the middle index between left and right using the formula mid = (left + right) >> 1. The >> 1 is a

We begin by initializing two pointers, left at 0 and right at len(arr) - k. This is because the k elements we seek must form

- bitwise right shift that effectively divides the sum by 2. With mid established, we check the distance between x and the current middle value at arr[mid] compared to the distance
- ∘ If the element at mid is closer to x (or the elements are equally distant), we know that the k closest elements can't start any index higher than mid. Thus we set right to mid, moving our search space leftward.

Conversely, if the element at mid + k is closer to x, it means the sequence starts further to the right, so we move our left pointer up to mid

By repeatedly halving our search space, we eventually arrive at a point where left equals right. This index represents the

- Once we have the starting index (left), we slice the original array from left to left + k to select the k closest elements. This slice is then returned, and since the original array was sorted, this resulting slice is guaranteed to be sorted as well,
- optimization over a complete sort-based method that would require O(n log n) time. This optimization is significant for large datasets where the difference in time complexity could be substantial.

Using this binary search approach, we achieve a time complexity of O(log(n - k)) for finding the subsequence, which is an

Let us walk through a simple example to illustrate the solution approach.

Assume we have an array arr = [1, 3, 4, 7, 8, 9], and we want to find k = 3 closest integers to x = 5. **Initial Setup**

• We begin by initializing two pointers, left is set to 0 and right is set to len(arr) - k, which in this case is 6 - 3 = 3. The subarray [7, 8, 9] is

• The binary search begins. We calculate the middle index between left and right. The initial mid is (0 + 3) / 2 = 1.5, rounded down to 1. • We then compare the distance of x = 5 from arr[mid] = 3 and from arr[mid + k] = arr[4] = 8. The distance from 3 is |5 - 3| = 2 and from 8

Binary Search

Narrowing Down

Example Walkthrough

Example

is |5 - 8| = 3. • Since 3 is closer to 5 than 8 is, we set right to mid, which is now 1. Our search space for starting indices is now [0, 1].

the last possible sequence of k numbers, and we will not consider starting indices beyond this.

• Since 7 is closer to 5 than 1 is, we move the left pointer to mid + 1, which is 1.

• We slice the original array from left to left + k, which yields [3, 4, 7].

• We perform another iteration of the binary search. Our new mid is (0 + 1) / 2 = 0.5, rounded down to 0. • Again, we compare distances from x = 5: arr[mid] = 1 has a distance of 4, and arr[mid + k] = arr[3] = 7 has a distance of 2.

Convergence • Because left now equals right, we have converged to the starting index for the k closest elements. Our starting index is 1.

From our example using the array arr = $\begin{bmatrix} 1 & 3 & 4 & 7 & 8 & 9 \end{bmatrix}$ and the values k = 3 and x = 5, we have walked through the binary

search approach and determined that the k closest integers to x in this array are [3, 4, 7], as illustrated by the binary search

logic to find the starting index and the subsequent slice of the array. This method efficiently finds the correct subsequence without fully sorting the array based on proximity to x.

Solution Implementation

from typing import List

class Solution:

Conclusion

Obtaining the Result

Python

left, right = 0, len(arr) - k

mid = (left + right) // 2

right = mid

left = mid + 1

if x - arr[mid] <= arr[mid + k] - x:</pre>

if $(x - arr[mid] \le arr[mid + k] - x)$ {

// Create a list to store the k closest elements.

List<Integer> result = new ArrayList<>();

for (int i = left; i < left + k; ++i) {</pre>

// Return the list of k closest elements.

// Initialize the binary search bounds

int mid = left + (right - left) / 2;

if $(x - arr[mid] \le arr[mid + k] - x) {$

// move the right bound to mid

// move the left bound to mid + 1

* Finds the k closest elements to a given target x in a sorted array.

* @param $\{number\} \times - The target number to find the closest elements to.$

function findClosestElements(arr: number[], k: number, x: number): number[] {

// Binary search to find the start index of the k closest elements.

// then the closer elements are to the left side of the array.

if (x - arr[midPointer] <= arr[midPointer + k] - x) {</pre>

* @returns {number[]} - An array of k closest elements to the target.

* @param {number[]} arr - The sorted array of numbers.

// Initialize two pointers for binary search.

let rightPointer = arr.length - k;

while (leftPointer < rightPointer) {</pre>

rightPointer = midPointer;

leftPointer = midPointer + 1;

* @param {number} k - The number of closest elements to find.

int right = arr.size() - k;

right = mid;

left = mid + 1;

} else {

let leftPointer = 0;

} else {

while (left < right) {</pre>

// Otherwise, move right (higher indices).

// Add the closest elements to the list, starting from the left pointer.

std::vector<int> findClosestElements(std::vector<int>& arr, int k, int x) {

// Perform binary search to find the start index of the k closest elements

// Calculate mid index (avoid potential overflow by using left + (right-left)/2)

// Compare the differences between x and elements at mid index and mid+k index

// If the element at mid index is closer to x, or equally close

// Create and return a vector of k closest elements starting from the 'left' index

// as the element at mid+k index (prefer the smaller element),

// Otherwise, if the element at mid+k index is closer to x,

// The goal is to find the smallest window such that the elements are closest to x

right = mid;

result.add(arr[i]);

left = mid + 1;

} else {

return result;

int left = 0;

while left < right:</pre>

else:

def findClosestElements(self, arr: List[int], k: int, x: int) -> List[int]:

The right pointer is set to the highest starting index for the sliding window.

Check the distance from the x to the middle element and the element at mid + k position.

we move the right pointer to mid. Otherwise, we adjust the left pointer to mid + 1.

Extract the subarray from the left index of size k, which will be the k closest elements.

If the element at mid is closer to x or equal in distance compared to the element at mid + k,

Perform binary search to find the left bound of the k closest elements.

Initialize the left and right pointers for binary search.

Calculate the middle index between left and right.

```
return arr[left:left + k]
# Example usage:
# sol = Solution()
# result = sol.findClosestElements([1, 2, 3, 4, 5], k=4, x=3)
# print(result) # Output should be [1, 2, 3, 4]
Java
import java.util.List;
import java.util.ArrayList;
class Solution {
    public List<Integer> findClosestElements(int[] arr, int k, int x) {
       // Initialize the left and right pointers for binary search.
        int left = 0;
        int right = arr.length - k;
       // Continue searching until the search space is reduced to a single element.
       while (left < right) {</pre>
            // Calculate the middle index.
            int mid = left + (right - left) / 2; // Using left + (right - left) / 2 to avoid potential overflow.
            // If the distance to the left is less than or equal to the distance to the right,
           // we need to move towards the left (lower indices).
```

C++

public:

#include <vector>

class Solution {

```
return std::vector<int>(arr.begin() + left, arr.begin() + left + k);
};
TypeScript
```

/**

*/

```
// Slice the array from the left pointer to get the k closest elements.
      return arr.slice(leftPointer, leftPointer + k);
from typing import List
class Solution:
   def findClosestElements(self, arr: List[int], k: int, x: int) -> List[int]:
       # Initialize the left and right pointers for binary search.
       # The right pointer is set to the highest starting index for the sliding window.
        left, right = 0, len(arr) - k
       # Perform binary search to find the left bound of the k closest elements.
       while left < right:</pre>
           # Calculate the middle index between left and right.
            mid = (left + right) // 2
           # Check the distance from the x to the middle element and the element at mid + k position.
            # If the element at mid is closer to x or equal in distance compared to the element at mid + k,
            # we move the right pointer to mid. Otherwise, we adjust the left pointer to mid + 1.
            if x - arr[mid] <= arr[mid + k] - x:</pre>
                right = mid
            else:
                left = mid + 1
```

Extract the subarray from the left index of size k, which will be the k closest elements.

const midPointer = (leftPointer + rightPointer) >> 1; // Equivalent to Math.floor((left + right) / 2)

// If the element at the middle is closer to or as close as the element k away from it,

// Otherwise, they are to the right, and we move the left pointer one step past the middle.

```
The provided algorithm is a binary search approach for finding the k closest elements to x in the sorted array arr. Here's the
analysis:
```

result = sol.findClosestElements([1, 2, 3, 4, 5], k=4, x=3)

return arr[left:left + k]

Time and Space Complexity

print(result) # Output should be [1, 2, 3, 4]

Example usage:

sol = Solution()

is because the binary search is performed within a range that's reduced by k elements (since we're always considering k contiguous elements as a potential answer).

The algorithm performs a binary search by repeatedly halving the search space, which is initially N - k. It does not iterate through all N elements, but narrows down the search to the correct starting point for the sequence of k closest elements.

Time Complexity: The time complexity of this algorithm is O(log(N - k)) where N is the number of elements in the array. This

After finding the starting point, it returns the subarray in constant time, as array slicing in Python is done in 0(1) time. Hence, the iterative process of the binary search dominates the running time, which leads to the O(log(N - k)) time complexity.

Space Complexity: The space complexity of this algorithm is 0(1). The code uses a fixed amount of extra space for variables

left, right, and mid. The returned result does not count towards space complexity as it is part of the output. Please note that array slicing in the return statement does not create a deep copy of the elements but instead returns a reference

to the same elements in the original array, thus no extra space proportional to k or N is used in this operation.