787. Cheapest Flights Within K Stops

Breadth-First Search Graph

Problem Description

Depth-First Search

Medium

This LeetCode problem requires finding the cheapest flight route from one city to another with certain constraints. We have n cities and a list of flights where each flight is represented by [from_i, to_i, price_i]. This indicates a flight from city from_i to city to_i at the cost price_i. Given additional parameters src, dst, and k, the goal is to find the minimum cost to travel from the source city src to the destination city dst with at most k stops in between. If no such route exists, the function should return -1.

Dynamic Programming

Shortest Path

Heap (Priority Queue)

ntuition

the same city.

To solve this problem, we use a variation of the Bellman-Ford algorithm which is suitable for finding shortest paths in graphs with edge weights that could be negative (not the case here) and can handle queries for paths with a limited number of edges.

We initialize a list dist of size n with a high value representing infinity since we are looking for the minimum cost. This list will keep track of the minimum cost to reach each city. The entry for the source city src is set to 0 since it costs nothing to stay at

The algorithm will go through the flights at most k+1 times (the number of stops plus the initial city). On each iteration, we make a copy of the dist list called backup. This is important to ensure that updates in the current round do not affect the other updates in

the same round since we are only allowed to use k stops.

For each flight [f, t, p] in flights, we check if the current known cost to city t is greater than the cost to city f plus the price p of the flight from f to t. If it is, we update the cost for city t in the dist list with the new lower cost. This represents a relaxation step where we relax the cost to get to t via f if it is possible to get there more cheaply.

After k+1 iterations, the dist list contains the minimum costs of reaching each city from the source city with up to k stops. The value in dist[dst] is the cheapest price to get to the destination dst. However, if the value at dist[dst] is still infinity (represented by INF), it means there was no possible route within the given number of stops, hence we return -1.

The implementation of the solution uses <u>dynamic programming</u> to iteratively update the cheapest prices for reaching each city

within the constraint of at most k stops. Here's a step-by-step walkthrough of the algorithm: Initialization: An array dist is initialized to store the cheapest prices to each city. It is filled with INF (a large number indicating infinity) to represent that initially, all destinations are unreachable. The price to reach the src city is set to 0

iteration is because a journey from src to dst with k stops consists of k+1 flights. Backup Array: In each iteration, a backup of the dist array is created. This is crucial because updates to the dist array within a single iteration must be based on the state of the array from the previous iteration, not the current one where some values

Algorithm: The core part of the algorithm runs for k+1 iterations, where k is the maximum number of stops allowed. The extra

- have already been updated. **Relaxation:** For each flight [f, t, p], where f is the starting city, t is the target city, and p is the price of the flight, the algorithm checks if the current known price to city t can be improved by taking the flight from city f. The relaxation condition
- dist[t] = backup[f] + p If the price to reach city t (dist[t]) is higher than the price to reach city f (backup[f]) plus the price of the flight p, the cheaper price is updated in dist[t].
- **Returning the result**: After k+1 iterations, the algorithm looks at the price recorded in dist[dst], which is the minimum cost to reach destination dst from source src within k stops. If this price is still INF, it means no such route was found within the limit

if dist[t] > backup[f] + p:

is:

because it costs nothing to start from there.

This approach efficiently utilizes the Bellman-Ford algorithm's concept of relaxing edges but modifies it to account for the number of stops constraint by limiting the number of relaxation iterations.

src = 0 (source city), dst = 3 (destination city), k = 1 (maximum 1 stop allowed), and the flights are represented by a list of

[from_i, to_i, price_i] as follows: flights = [[0, 1, 100], [1, 3, 100], [0, 2, 500], [2, 3, 100]].

Algorithm: We will iterate k+1 times, which in this case is 2 times (1 stop + 1 initial trip).

Let's use a small example to illustrate the solution approach described above. Consider the following scenario where n = 4 cities,

In this example, there are direct flights from city 0 to city 1 and city 2, and from city 1 and city 2 to city 3 with the respective prices of 100 and 500. Initialization: We start by initializing the dist array with INF values for all cities except the src city: dist = [0, INF, INF,

Backup Array: In each iteration, a backup of dist is made. Relaxation:

■ Flight [1, 3, 100] (from city 1 to city 3): Checks if dist[3] > backup[1] + 100. Since INF > INF + 100, no update is made. Flight [0, 2, 500] (from city 0 to city 2): Checks if dist[2] > backup[0] + 500. Since INF > 0 + 500, update dist[2] to 500. ■ Flight [2, 3, 100] (from city 2 to city 3): Checks if dist[3] > backup[2] + 100. Since INF > INF + 100, no update is made.

Flight [0, 1, 100] (from city 0 to city 1): Checks if dist[1] > backup[0] + 100. Since INF > 0 + 100, update dist[1] to 100.

 Second iteration (with updated dist from first iteration; new backup is made): ■ Flight [0, 1, 100]: Since dist[1] is already at its lowest from the first iteration, no changes occur.

of stops, and -1 is returned. Otherwise, the cheapest price found is returned.

■ Flight [0, 2, 500]: Since dist[2] is already at its lowest from the first iteration, no changes occur.

INF].

First iteration:

```
■ Flight [2, 3, 100]: Checks if dist[3] > backup[2] + 100. Since 200 > 500 + 100, no update is made. (Showing that we found a
```

cheaper way to get to city 3 during this iteration) Returning the result: After 2 iterations, we look at dist[dst], which is dist[3]. The value is 200, which is the cheapest price

Flight [1, 3, 100]: Checks if dist[3] > backup[1] + 100. Since INF > 100 + 100, update dist[3] to 200.

def findCheapestPrice(self, n: int, flights: List[List[int]], src: int, dst: int, k: int) -> int:

If the destination is still at 'infinity', it means it's unreachable within K stops

public int findCheapestPrice(int n, int[][] flights, int source, int destination, int k) {

// Initialize distance array with infinity to represent no flights booked yet

return -1, indicating no possible route within the given stop constraint. Solution Implementation

to reach city dst from src within 1 stop. This value is returned as the result. If dist[dst] had still been INF, then we would

Use a high value to represent 'infinity', since there is no direct representation in Python INF = float('inf') # Initialize distance array, with 'infinity' for all nodes except the source node distance = [INF] * ndistance[src] = 0

Perform (K+1) iterations of Bellman—Ford algorithm to find the cheapest path with at most K stops

Make a copy of distance array to prevent using updated values within the same iteration

Update the distance for each edge (flight) if a cheaper price is found for from_node, to_node, price in flights: # The new price is considered only if the previous node was reached within K stops distance[to_node] = min(distance[to_node], distance_backup[from_node] + price)

Java

class Solution {

for _ in range(k + 1):

distance_backup = distance.copy()

// Represent an unreachable cost with a high value

for (const auto& flight : flights) {

int from = flight[0], to = flight[1], price = flight[2];

if (previousIterationDistances[from] < INF) {</pre>

// otherwise, return the shortest distance to the destination.

distances[src] = 0; // The distance from the source to itself is always 0.

// Make a copy of the current state of distances before this iteration.

return distances[dst] == INF ? -1 : distances[dst];

// which translates to K+1 edges in the shortest path.

for (let i = 0; i <= K; ++i) {

private static final int INFINITY = 0x3f3f3f3f;

int[] currentDist = new int[n];

return -1 if distance[dst] == INF else distance[dst]

// Finds the cheapest price for a flight with up to k stops

from typing import List

class Solution:

```
// Backup array used during relaxation to avoid overwriting information prematurely
       int[] prevDist = new int[n];
       // Fill the distance array with infinity
       Arrays.fill(currentDist, INFINITY);
       // The cost to get to the source from the source is 0
       currentDist[source] = 0;
       // Perform relaxation k+1 times (since 0 stops means just 1 flight)
        for (int i = 0; i < k + 1; ++i) {
           // Copy current distances to the backup array
           System.arraycopy(currentDist, 0, prevDist, 0, n);
           // Iterate through each flight
           for (int[] flight : flights) {
               int from = flight[0], to = flight[1], cost = flight[2];
               // Relax the distance if a shorter path is found
               // Only update the distance using values from previous iteration (prevDist)
               currentDist[to] = Math.min(currentDist[to], prevDist[from] + cost);
       // If destination is reachable, return the cost, otherwise return -1
       return currentDist[destination] == INFINITY ? -1 : currentDist[destination];
C++
class Solution {
public:
   int findCheapestPrice(int n, vector<vector<int>>& flights, int src, int dst, int K) {
       const int INF = 0x3f3f3f3f; // Define an "infinity" value for initial distances.
       vector<int> distances(n, INF); // Initialize all distances to "infinity" except the source.
       vector<int> previousIterationDistances; // Used to store distances from the previous iteration.
       distances[src] = 0; // The distance from the source to itself is always 0.
       // Run the Bellman—Ford algorithm for K+1 iterations because you can have at most K stops in between,
       // which translates to K+1 edges in the shortest path.
        for (int i = 0; i < K + 1; ++i) {
           // Make a copy of the current state of distances before this iteration.
           previousIterationDistances = distances;
```

// For each edge in the graph, try to relax the edge and update the distance to the destination node

// Relaxation step: If the current known distance to 'from' plus the edge weight

distances[to] = min(distances[to], previousIterationDistances[from] + price);

// After K+1 iterations, if the distance to the destination is still "infinity", no such path exists;

let distances: number[] = new Array(n).fill(INF); // Initialize all distances to "infinity" except the source.

// to 'to' is less than the currently known distance to 'to', update it.

function findCheapestPrice(n: number, flights: number[][], src: number, dst: number, K: number): number {

let previousIterationDistances: number[]; // Used to store distances from the previous iteration.

// Run the Bellman-Ford algorithm for K+1 iterations because you can have at most K stops in between,

const INF = Number.POSITIVE_INFINITY; // Define an "infinity" value for initial distances.

```
previousIterationDistances = distances.slice();
// For each edge in the graph, try to relax the edge and update the distance to the destination node
for (const flight of flights) {
    const [from, to, price] = flight;
```

};

TypeScript

```
// Relaxation step: if the current known distance to 'from' plus the edge weight
              // to 'to' is less than the currently known distance to 'to', update it.
              if (previousIterationDistances[from] < INF) {</pre>
                  distances[to] = Math.min(distances[to], previousIterationDistances[from] + price);
      // After K+1 iterations, if the distance to the destination is still "infinity", no such path exists;
      // otherwise, return the shortest distance to the destination.
      return distances[dst] === INF ? -1 : distances[dst];
from typing import List
class Solution:
   def findCheapestPrice(self, n: int, flights: List[List[int]], src: int, dst: int, k: int) -> int:
       # Use a high value to represent 'infinity', since there is no direct representation in Python
        INF = float('inf')
       # Initialize distance array, with 'infinity' for all nodes except the source node
       distance = [INF] * n
       distance[src] = 0
       # Perform (K+1) iterations of Bellman—Ford algorithm to find the cheapest path with at most K stops
        for _ in range(k + 1):
           # Make a copy of distance array to prevent using updated values within the same iteration
            distance_backup = distance.copy()
           # Update the distance for each edge (flight) if a cheaper price is found
            for from_node, to_node, price in flights:
               # The new price is considered only if the previous node was reached within K stops
               distance[to_node] = min(distance[to_node], distance_backup[from_node] + price)
       # If the destination is still at 'infinity', it means it's unreachable within K stops
        return -1 if distance[dst] == INF else distance[dst]
```

The given code defines a function that finds the cheapest price for a flight from a source to a destination with at most k stops. The main operations in this code are as follows:

Time Complexity

Time and Space Complexity

complexity of O(n). We then run a loop k+1 times to account for the fact that we can make at most k stops, which means we are considering paths with at most k+1 edges. The loop operation involves iterating through the list of flights.

We initialize the dist array of size n with INF to represent the minimum cost to reach each node. This operation has a time

Then, in the nested loop, for each flight (f, t, p), we update the dist of the destination t if a cheaper price is found by considering the flight from f to t with price p. With m being the number of flights, this nested loop has a time complexity of

Inside the loop, a copy() operation is performed on the dist array, which again takes O(n) time.

O(m) per iteration of the outer loop. Putting these together, the total time complexity of the algorithm becomes the combined complexity of the loop executed k+1

times and the nested iteration over all the flights, resulting in O((k+1) * (n + m)) = O(k * (n + m)).

In the worst case, we might have to consider every flight for every iteration, so the worst-case time complexity is O(k * m) when m is significantly larger than n.

Space Complexity

The space complexity of the algorithm is determined by: The dist array, which consumes O(n) space.

- The backup array, which also consumes O(n) space and is created anew in each iteration of the loop but does not add to the space complexity asymptotically as only one extra array exists at any given time.
- Thus, the total space complexity is O(n), due to the space required for the dist and backup arrays to store the cost of reaching

each node.