847. Shortest Path Visiting All Nodes

Breadth-First Search Graph Dynamic Programming

## **Problem Description**

Bit Manipulation

Hard

In this LeetCode problem, we are given an undirected, connected graph with n nodes, where nodes are labeled from 0 to n - 1. The graph is represented as an adjacency list: graph[i] contains all the nodes connected to the node i by an edge. Our goal is to determine the shortest path length that visits every node in the graph. We are allowed to start and end at any node, revisit

**Bitmask** 

nodes, and traverse any edge multiple times if needed. The key challenge is to visit all nodes in the most efficient way possible, keeping the path as short as we can manage. It is

important to note that the problem isn't asking for the shortest path between two particular nodes, but rather a path that covers all nodes.

## To solve this problem, we'll use Breadth-First Search (BFS) algorithm paired with state compression to keep track of the nodes

undirected graphs.

Intuition

State compression is used here to efficiently represent the set of nodes visited at any point in time. This is crucial because revisiting nodes is allowed, hence the state of our search isn't just the current node, but also the collection of nodes that have been visited so far. By using a bitmask to represent visited nodes (where the i-th bit corresponds to whether the i-th node has been visited), we can easily update and check which nodes have been seen as we process the BFS queue.

visited. The intuition here is that since all edges are equally weighted, BFS is a natural choice to find the shortest path in

For each node, we start a BFS traversal, where each element in the BFS queue is a pair (current\_node, visited\_state). We define the initial state as 1 << i, which means only the node i is visited. If the visited state ever becomes (1 << n) - 1, it indicates all nodes have been visited, and we have our solution: the length of the path taken to reach this state. The BFS ensures we find the shortest such path, and the visited set (vis) prevents us from processing the same state (node +

visited combination) more than once, which is essential to avoid infinite loops and reduce computation. **Solution Approach** 

The solution leverages the BFS algorithm and bitmasking for state tracking to efficiently determine the shortest path that visits

starts traversing from a node and explores all its neighbors before moving onto the neighbors of those neighbors, thus

every node in the graph. The approach can be broken down into key components that work synchronously to find the solution: Breadth-First Search (BFS): BFS algorithm is a perfect fit to find the shortest path in a graph with edges of equal weight. It

### Queue: The algorithm uses a queue data structure to process nodes in FIFO (first-in-first-out) order, which is the essence of BFS. In Python, this is implemented using the deque data structure from the collections module for efficient pops from the

1) straightforward.

<< i)).

visited\_state.

**Example Walkthrough** 

[1]

The implementation steps are as follows:

ensuring that the shortest paths to all nodes are identified.

• Set ans to zero, which will keep track of the number of steps from the start.

• Begin the while loop, which will run until we break out of it when the solution is found.

front of the list. State Compression with Bitmasking: To track the nodes that have been visited during the traversal without actually storing

all the nodes themselves, state compression through bitmasking is used. For instance, if five nodes have been visited in a

graph with n = 5 nodes, the visited state can be represented as 11111 in binary, which is 31 in decimal. This approach

drastically reduces the memory footprint and makes the check for the condition of all nodes being visited (st == (1 << n) -

Visited Set: A set named vis is used to keep track of all the (node, state) pairs that have been visited. This is to avoid repeating searches that have already been done and to ensure we do not go into cycles, which can be a common issue in graph problems.

• Initialize the queue and visited set, and populate them with the first step of the BFS, which includes every node visited on its own first ((i, 1

• Iterate over the queue. Remove (popleft) the front of the queue to process the current state. Each state contains a current\_node and the

have been visited). If so, this is the shortest path, and we return ans. • If all nodes have not been visited, for each neighbor j of the current\_node, calculate the new state nst by setting the corresponding bit (1 << j) in the visited\_state. • If this new state has not been visited (if (j, nst) not in vis), add it to the visited set and queue. • If the inner loop completes without returning, increment ans to indicate that another level of BFS is completed.

• Check if the visited state signifies that all nodes have been visited by comparing it to (1 << n) - 1 (which represents a state where all nodes

[0, 2, 3], // Node 1 is connected to Node 0, 2, and 3

// Node 3 is connected to Node 1

[0, 1], // Node 2 is connected to Node 0 and 1

2. ans is set to 0, marking the starting point of the BFS levels.

Processing current\_node = 0, visited\_state = 1:

• For neighbor 1, nst becomes 1 | (1 << 1) which is 3 in binary it's 011.

• We check each neighbor of node 0, which is 1 and 2.

Update the visited\_state to include these neighbors.

- Let us consider a small graph represented as an adjacency list for the clarity of our example: graph = [[1, 2], // Node 0 is connected to Node 1 and 2
- This is an undirected graph with n = 4 nodes. The goal is to find the shortest path that visits all nodes.
- Queue and vis will have the following entries: [(0, 1), (1, 2), (2, 4), (3, 8)] where each pair has the format (current\_node, visited\_state).

Let's walk through the steps of the BFS algorithm and state compression starting from Node 0:

1. Initialize the BFS queue with our starting nodes and initial state and the visited set vis.

1, 2, 4, 8 are the bitmask states for visiting nodes 0, 1, 2, and 3 respectively.

# 3. We start our while loop, processing the states in our queue one by one.

Queue now looks like:

Solution Implementation

queue = deque()

while True:

for node in range(num\_nodes):

for in range(len(queue)):

public int shortestPathLength(int[][] graph) {

state = 1 << node

class Solution:

• For neighbor 2, nst becomes 1 | (1 << 2) which is 5 in binary it's 101. • Since these new states have not been visited, we add them to vis and queue: [(1, 3), (2, 5)].

```
After processing all initial states and their neighbors, we increment ans to 1, which reflects moving to a new level in our BFS
search. The queue and vis have been updated with the new states.
```

This process continues, applying BFS. Each node visits its neighbors, creates new states, adds them to the queue if not already

def shortestPathLength(self, graph: List[List[int]]) -> int:

num nodes = len(graph) # Number of nodes in the graph

# Perform a BFS until a path visiting all nodes is found

# Check if the current state has all nodes visited

return steps # Return the number of steps taken so far

boolean[][] visited = new boolean[numNodes][1 << numNodes]; // Visited states</pre>

int bitmask = 1 << i; // Binary representation where only the ith bit is set</pre>

int levelSize = nodesQueue.size(); // Number of elements at the current level of BFS

auto [currentNode, stateBitmask] = nodesQueue.front(); // Fetch the front element

// If the state bitmask represents all nodes visited, return the current path length

int nextStateBitmask = stateBitmask | (1 << nextNode); // Update the visited mask</pre>

if (!visited[nextNode][nextStateBitmask]) { // If this state has not been visited

new Array(1 << nodeCount).fill(false));</pre>

const startState = 1 << i; // Initial state for node i (only node i has been visited).</pre>

visited[nextNode][nextStateBitmask] = true; // Mark it as visited

nodesQueue.emplace(nextNode, nextStateBitmask); // Add to the queue

visited[i][bitmask] = true; // Set the initial state as visited

nodesQueue.pop(); // Remove the element from the queue

while (levelSize--) { // Loop over all nodes in the current BFS level

// Loop continually, increasing the path length in each iteration

if (stateBitmask == (1 << nodeCount) - 1) {</pre>

for (int nextNode : graph[currentNode]) {

nodesQueue.emplace(i, bitmask);

for (int pathLength = 0; ; ++pathLength) {

return pathLength;

// Explore adjacent nodes

function shortestPathLength(graph: number[][]): number {

const nodeCount = graph.length;

const queue: number[][] = [];

queue.push([i, startState]);

// The 'n' represents the total number of nodes in the graph.

// The queue to perform BFS. Elements are pairs: [node, state].

# Explore neighbors of the current node

new state = state | (1 << neighbor)</pre>

if (neighbor, new state) not in visited:

visited.add((neighbor, new state))

queue.append((neighbor, new\_state))

# After exploring all nodes at the current depth, increment steps

for neighbor in graph[current node]:

// Visited states represented by node index and state as a bitmask.

const visited: boolean[][] = Array.from({ length: nodeCount }, () =>

# If the new state has not been visited, add it to the queue and set

# Iterate over nodes in the current layer

if state == (1 << num nodes) - 1:</pre>

current node, state = queue.popleft()

# Explore neighbors of the current node

new state = state | (1 << neighbor)</pre>

int numNodes = graph.length; // Number of nodes in the graph

// Initialize by adding each node as a starting point in the queue

Deque<int[]> queue = new ArrayDeque<>(); // Queue for BFS

// and marking it visited with its own unique state

if (neighbor, new state) not in visited:

visited.add((neighbor, new state))

queue.append((neighbor, new\_state))

for neighbor in graph[current node]:

visited = set() # Set to store visited (node, state) tuples

# Initialize the queue and visited set with all individual nodes and their bitmasks

(1, 3), (2, 5) // New states after processing Node 0

(1, 2), (2, 4), (3, 8), // Remaining initial states

// More states would be added as we continue the BFS

**Python** from collections import deque

visited, and checks if the visited state equates to (1 << n) - 1. If at any point the visited state becomes 1111 (15 in decimal),

this means all nodes have been visited, and the value of ans at this point will give us our shortest path length.

Upon finding this condition, the BFS stops, and we return ans as the length of the shortest path to visit all nodes.

queue.append((node, state)) visited.add((node, state)) # Initialize the number of steps taken to reach all nodes steps = 0

#### # After exploring all nodes at the current depth, increment steps steps += 1Java

class Solution {

```
for (int i = 0; i < numNodes; ++i) {
            queue.offer(new int[] {i, 1 << i}); // Node index, state as a bit mask</pre>
            visited[i][1 << i] = true;</pre>
        // BFS traversal
        for (int steps = 0; ; ++steps) { // No terminal condition because the answer is quaranteed to exist
            for (int size = queue.size(); size > 0; --size) { // Process nodes level by level
                int[] pair = queue.poll(); // Get pair (node index, state)
                int node = pair[0], state = pair[1];
                // If state has all nodes visited (all bits set), return the number of steps
                if (state == (1 << numNodes) - 1) {</pre>
                    return steps;
                // Check the neighbors
                for (int neighbor : graph[node]) {
                    int newState = state | (1 << neighbor); // Set bit for the neighbor</pre>
                    // If this state has not been visited before, mark it visited and add to queue
                    if (!visited[neighbor][newState]) {
                        visited[neighbor][newState] = true;
                        queue.offer(new int[] {neighbor, newState});
C++
#include <vector>
#include <queue>
#include <cstring>
class Solution {
public:
    // Function to find the shortest path that visits all nodes in an undirected graph.
    int shortestPathLength(vector<vector<int>>& graph) {
        int nodeCount = graph.size(); // Number of nodes in the graph
        queue<pair<int, int>> nodesQueue; // Queue to maintain the state
        // Visited array to keep track of visited states: node index and visited nodes bitmask
        bool visited[nodeCount][1 << nodeCount];</pre>
        memset(visited, false, sizeof(visited)); // Initializing all elements as false
        for (int i = 0; i < nodeCount; ++i) {
            // Initialize the queue with all nodes as starting points
```

## // Initialize queue and visited with the starting nodes and their respective bitmasks. for (let i = 0; i < nodeCount; i++) {</pre>

**TypeScript** 

```
visited[i][startState] = true;
   // BFS to find the shortest path length.
    for (let steps = 0; ; steps++) {
        // Iterate through the current level of the queue.
        for (let size = queue.length; size > 0; size--) {
            const [currentNode, state] = queue.shift()!; // Get next element from the queue.
            // If all nodes have been visited, return the number of steps taken.
            if (state === (1 << nodeCount) - 1) {</pre>
                return steps;
            // Check all neighbors of the current node.
            for (const nextNode of graph[currentNode]) {
                const nextState = state | (1 << nextNode); // Mark nextNode as visited in the state.</pre>
                // If the state is not yet visited for the nextNode, mark it and enqueue.
                if (!visited[nextNode][nextState]) {
                    visited[nextNode][nextState] = true;
                    queue.push([nextNode, nextState]);
from collections import deque
class Solution:
   def shortestPathLength(self, graph: List[List[int]]) -> int:
        num nodes = len(graph) # Number of nodes in the graph
       queue = deque()
        visited = set() # Set to store visited (node, state) tuples
       # Initialize the queue and visited set with all individual nodes and their bitmasks
        for node in range(num_nodes):
            state = 1 << node
            queue.append((node, state))
            visited.add((node, state))
       # Initialize the number of steps taken to reach all nodes
        steps = 0
       # Perform a BFS until a path visiting all nodes is found
       while True:
            # Iterate over nodes in the current layer
            for in range(len(queue)):
                current node, state = queue.popleft()
                # Check if the current state has all nodes visited
                if state == (1 << num nodes) - 1:</pre>
                    return steps # Return the number of steps taken so far
```

### **Time Complexity** The time complexity of the provided code can be considered as $0(N * 2^N)$ where N is the number of nodes in the graph. Here's the breakdown:

BFS.

steps += 1

**Time and Space Complexity** 

• The algorithm uses Breadth-First Search (BFS) to traverse through every possible state of visited nodes. • There are 2^N possible states since each node can either be visited or not, marked by bitmasks. • For each state, we go through all the N nodes and perform constant-time operations (a bitwise OR to update the state and some checks). • The outer loop iterates until the queue is empty, and since we enqueue each state at most once, it iterates N \* 2^N times across all layers of

**Space Complexity** 

# If the new state has not been visited, add it to the queue and set

- The space complexity is  $O(N * 2^N)$ , which arises from the following: • A vis set is maintained to keep track of visited states to ensure that each state is processed exactly once. Since there are N nodes and 2^N
- possible states, the set size can go up to  $N * 2^N$ . • A queue q is used to perform BFS, this queue will also store at most N \* 2^N elements, which are pairs of current node and the state of visited nodes up to that point.