898. Bitwise ORs of Subarrays

Bit Manipulation Array

Problem Description

Medium

The problem presents us with a task: given an array arr of integers, we are required to find the total number of distinct bitwise OR results that can be obtained from all possible non-empty contiguous subarrays of arr.

To understand the problem better, let's clarify some concepts:

Dynamic Programming

Bitwise OR (1): A binary operation that takes two bit patterns of equal length and performs the logical inclusive OR operation on each pair of corresponding bits. The result in each position is 1 if at least one of the bits is 1.

Subarray: A sequence of elements within an array that is contiguous (elements are consecutive without gaps) and non-

- empty. Distinct: Each unique value should be counted only once in the final tally.
- The problem, therefore, requires us to explore and assess every possible subarray that can be generated from the given array, perform the bitwise OR operation on its elements, and count the unique outcomes.

ntuition The intuitive approach to solving this problem might involve a brute force strategy - calculating the bitwise OR for every subarray

and then utilizing a set to count the distinct values. However, this approach would result in a high time complexity due to the

multiple nested loops required (each subarray is recursively built starting from each element).

The solution code, however, uses a clever strategy that reduces the amount of redundant computation:

Iterative Building: It takes advantage of the fact that the bitwise OR of a subarray ending at position i depends upon the bitwise OR result of the subarray ending at i-1. The prev variable is used to accumulate the bitwise OR result till the current element, so that for the next element, we don't have to start from scratch. Early Breaking: It recognizes that if a new subarray's bitwise OR equals the accumulated OR (prev), we can stop early

because all subsequent subarrays will only give the same or greater OR results, which we would already have encountered

- Set for Uniqueness: The set s is used to store unique bitwise OR results. This way, every time we add a new OR result, duplicate values are inherently avoided. The main algorithm goes as follows:
- Initialize a set s to keep track of distinct bitwise OR results. Loop over each element v in the array: • Start a prev accumulator that holds the bitwise OR up to the current element.

 Calculate the bitwise OR for the current subarray and add it to the set s. Once the current subarray OR equals the accumulated prev OR, break early.

- Finally, return the length of set s, which represents the number of distinct bitwise ORs.
- By using the set and early breaking, the solution immensely reduces the number of calculations compared to the brute force
- approach, and therefore, is much more efficient.

Loop backwards from the current element to the beginning of the array:

due to the nature of the OR operation building on previous subarrays.

of all subarrays. Here's how the approach works in detail: We initialize an empty set s that will hold the distinct bitwise OR results. Apart from this, we have a prev variable that is used

The implementation of the reference solution makes use of sets and two pointers to efficiently compute the distinct bitwise ORs

to keep track of the bitwise OR up to the current element in the outer loop, which goes through each element i in the array. As we iterate over each element v in the array using an index i, we first update the prev variable to hold the bitwise OR

from all the subarrays of arr. This is the result that the function returns.

new variable curr to hold the result of bitwise ORs for the subarrays ending in the current element i.

between itself and the current element (prev |=v). This new prev will be the bitwise OR of all elements from the beginning of

duplicate elements.

Solution Approach

the array up to the current element i. In the inner loop, we start from the current element and go backwards through the array using another index j. We create a

- During each iteration of the inner loop, we update curr to be the bitwise OR of itself and the current jth element (curr |= arr[j]). After updating curr, we add it to the set s. This operation ensures that only distinct OR results are kept, as sets do not store
- computation. Lastly, once both loops are done, the length of set s will represent the number of distinct bitwise ORs that can be formed

The above steps form an efficient algorithm as it reduces the number of subarrays we need to check, leveraging the properties of

the bitwise OR operation and making use of sets to maintain distinct entries. Such an optimization is essential to pass all test

cases on platforms like LeetCode, where the input size can be large and brute force methods would result in a timeout error.

value and have been considered before. When this condition is met, the inner loop breaks, avoiding unnecessary

The early breaking condition is checked (if curr == prev), which is based on the understanding that once the current

subarray's OR matches the overall OR up to the current element (prev), all subsequent larger subarrays will yield the same OR

Example Walkthrough To illustrate the solution approach, let's take a small example:

Initialize a set and variables: Set s = {}

\circ For element 3 (i=2), prev becomes prev | 3 = 3 | 3 = 3. Inner loop - Backward iteration from current element:

At i=0: Only one subarray [1]: curr = 1

Outer loop - Iterate over each element in the array:

 \circ For element 2 (i=1), prev becomes prev | 2 = 1 | 2 = 3.

At i=1: Subarrays [2], [1, 2]: • curr = 2, add to set $s \rightarrow \{1, 2\}$

Early Breaking:

Python

Java

class Solution:

Set s becomes {1}

Variable prev is initialized.

For element 1 (i=0), set prev = 1.

Let's say we have an array arr = [1, 2, 3].

We will apply the solution approach step by step:

 Since curr now equals prev, inner loop breaks early. At i=2: Subarrays [3], [2, 3], but we stop at [2, 3] because:

curr = 3, already in set s, no need to add.

■ curr = 2 | 1 = 3, add to set $s \rightarrow \{1, 2, 3\}$

- Utilized at each step in the inner loop when curr matches prev. Final Result: The length of set s is 3, representing the distinct bitwise OR results: {1, 2, 3}.
- Solution Implementation

def subarrayBitwiseORs(self, arr: List[int]) -> int:

unique_or_results.add(current)

Return the number of unique bitwise OR results found

for index, value in enumerate(arr):

if current == prev:

for (int i = 0; i < arr.length; ++i) {</pre>

int subarrayBitwiseORs(vector<int>& arr) {

int currentOr = 0;

// Iterate over each element in the array.

for (int i = 0; i < arr.size(); ++i) {</pre>

int subarrayOr = 0;

for (int $j = i; j >= 0; ---j) {$

return uniqueOrValues.size();

// Iterate over each element in the array.

let subarrayOr: number = 0;

def subarrayBitwiseORs(self, arr: List[int]) -> int:

for index, value in enumerate(arr):

for j in range(index, -1, -1):

current |= arr[j]

Initialize a set to store unique bitwise OR results

Iterate over the input array with both value and index

for (let i = 0; i < arr.length; i++) {</pre>

currentOr |= arr[i];

unique_or_results = set()

prev |= value

Time and Space Complexity

current progression in curr.

Time Complexity

current = 0

subarrayOr |= arr[j];

if (subarrayOr == currentOr) break;

// The TypeScript standard library already includes Set, so we don't need

break

return len(unique_or_results)

int aggregate = 0;

unique_or_results = set()

prev |= value

Initialize a set to store unique bitwise OR results

Iterate over the input array with both value and index

'prev' will hold the cumulative OR result of the current iteration

Update 'prev' by taking the OR with the current value

■ Loop attempts to compute curr = 3 | 2, but since prev is already 3, the inner loop breaks.

current = 0 # Iterate backwards from the current index to the start of the array for j in range(index, -1, -1): # Update 'current' by taking the OR with the value at j current |= arr[j] # Add 'current' to the set of unique OR results

'current' will hold the cumulative OR result starting from 'index' going back to the start

If 'current' equals 'prev', no new unique values can be found by continuing; break

// 'aggregate' will hold the cumulative bitwise OR value up to the current element

// We iterate from the current element down to the start of the array

unordered_set<int> uniqueOrValues; // To store unique OR values of subarrays.

// Iterate from the current element to the beginning of the array.

uniqueOrValues.insert(subarrayOr); // Store the calculated OR in the set.

// The size of the set represents the number of distinct OR values of all subarrays.

// Break the loop if the subarray OR equals the currently calculated OR (all bits already set).

Each bitwise OR operation only computes new results, while duplicates are ignored due to the set data structure. By breaking

elements present in the solution approach, demonstrating its efficiency and how it leads to the final answer.

early from the inner loop, we avoid unnecessary computation, optimizing the process. This compact example covers all the main

```
public int subarrayBitwiseORs(int[] arr) {
   // We use a set to store unique values of bitwise ORs for all subarrays
   Set<Integer> uniqueBitwiseORs = new HashSet<>();
   // We iterate through each element in the array
```

class Solution {

```
for (int j = i; j >= 0; --j) {
                // We calculate the bitwise OR from the current element to the 'jth' element
                aggregate |= arr[j];
                // Add the current subarray's bitwise OR to the set
                uniqueBitwiseORs.add(aggregate);
                /* If the current aggregate value is the same as the previous
                   aggregate value, all future aggregates will also be the same
                   due to the properties of bitwise OR, so we break out early. */
                if (aggregate == (aggregate | arr[i])) {
                   break;
       // Return the number of unique bitwise ORs found
       return uniqueBitwiseORs.size();
C++
#include <vector>
#include <unordered_set>
class Solution {
public:
   // Function to count the number of distinct bitwise OR values of all subarrays.
```

// To store the running OR value of the current subarray.

// Used to calculate OR for each possible subarray ending at 'i'.

// Update the OR for the subarray ending at 'i' starting at 'j'.

// To store the running OR value of the current subarray.

// Used to calculate OR for each possible subarray ending at 'i'.

// Update the running OR with the current element.

```
// a separate import for unordered_set as we would in C++.
// Function to count the number of distinct bitwise OR values of all subsequences.
function subarrayBitwiseORs(arr: number[]): number {
    let uniqueOrValues: Set<number> = new Set(); // To store unique OR values of subarrays.
    let current0r: number = 0;
```

class Solution:

prev = 0

};

TypeScript

```
// Iterate from the current element to the beginning of the array.
    for (let j = i; j >= 0; j--) {
        subarrayOr |= arr[j];
                                          // Update the OR for the subarray ending at 'i' starting at 'j'.
       uniqueOrValues.add(subarrayOr); // Store the calculated OR in the set.
       // Break the loop if the subarray OR equals the currently calculated OR (all bits already set).
       if (subarrayOr === currentOr) break;
// The size of the set represents the number of distinct OR values of all subarrays.
return uniqueOrValues.size;
```

'current' will hold the cumulative OR result starting from 'index' going back to the start

The given code aims to find the number of distinct subarray bitwise ORs. To do this, it iterates over the given array and computes

the OR of elements from the current element to all previous elements by keeping a record of the previous OR in prev and the

from the inner loop. Specifically, the sequence of ORs will eventually stablize into a set of values that does not grow with each

export { subarrayBitwiseORs }; // Export the function to be available for import in other modules.

```
# Add 'current' to the set of unique OR results
        unique_or_results.add(current)
        # If 'current' equals 'prev', no new unique values can be found by continuing; break
        if current == prev:
           break
# Return the number of unique bitwise OR results found
return len(unique_or_results)
```

Update 'current' by taking the OR with the value at j

Iterate backwards from the current index to the start of the array

'prev' will hold the cumulative OR result of the current iteration

Update 'prev' by taking the OR with the current value

• The outer loop runs exactly n times where n is the number of elements in arr. • The inner loop runs up to i+1 times in the worst case (when curr never equals prev early). However, due to the properties of the bitwise OR operation, repetitions are likely to occur much earlier, resulting in earlier breaks

additional OR operation. The actual number of unique elements in these OR sequences across all iterations is bounded by a factor much smaller than n^2.

The time complexity of this algorithm mainly depends on the number of iterations within the double-loop structure.

- While it's difficult to put a precise bound on this without specifics about the input distribution, let's denote the average unique
- sequence length as k (which is considerably smaller than n due to the saturation of OR operations). Therefore, the total number of operations is approximately 0(n*k).

The space complexity is due to the set s that is used to store the unique subarray OR results.

time complexity argument, this won't actually occur due to the saturation of bitwise ORs.

number of unique OR values across all subarrays.

implying that in the worst case the time complexity could tend towards 0(n^2), but in practical scenarios, it is expected to perform significantly better. **Space Complexity**

However, it is important to note that k is not guaranteed to be a constant and its relation with n can depend heavily on the input,

Let m represent the maximum possible unique OR values which can be much less than the total subarray count of roughly n* (n+1)/2. Therefore, the space complexity can be approximated as 0(m). In conclusion, the time complexity of the code is approximately 0(n*k) (with k being influenced by the input nature and much

smaller than n) and space complexity is around O(m) for storing the unique OR results set, where m represents the maximum

• In the worst case, each subarray OR could be unique, which means the set could grow to the size of the sum of all subarray counts. As with the