937. Reorder Data in Log Files

be sorted by their identifiers in lexicographic order.

Sorting

String]

Problem Description

Array

Medium

either a string of lowercase English letters (a letter-log) or a string of digits (a digit-log). The identifier and the log content are separated by a space. Our task is to reorder these logs based on certain rules: 1. Letter-logs must be sorted and come before all digit-logs. 2. Letter-logs should be ordered lexicographically (i.e., alphabetically) by their content. If two letter-logs have identical content, they should then

In this problem, we are dealing with an array of strings called logs, where each log entry is composed of an identifier followed by

- 3. Digit-logs should retain their original order as they appeared in the input. The expected output is the reordered array of the logs following these rules.
- Intuition

To solve this, we need a way to differentiate between letter-logs and digit-logs and then sort them according to the given criteria.

We can achieve this through a custom sorting function. When we look at the identifier and the log content, we observe that letter-logs contain alphabetical characters while digit-logs

content is a letter or a digit. Given the sorting requirements, we know that all letter-logs should be sorted lexicographically by their content, and if these are

contain numbers. We can exploit this characteristic to split logs into two categories by checking if the first character of the log

equal, by their identifiers. Digit-logs do not need sorting among themselves, but they need to be placed after all letter-logs. To keep them in their original order, we can just ignore their content and identifiers during the sorting process. With these observations, we craft a sorting function that:

1. Splits each log into its identifier and content. 2. Checks whether it's a letter-log or digit-log by inspecting the first character of the content. 3. If it's a letter-log, the function returns a tuple (0, content, identifier). The zero indicates that it's a letter-log for sort precedence, and sorting by content and then identifier follows naturally. 4. If it's a digit-log, the function returns a tuple (1,). The '1' ensures that all digit-logs are ordered after letter-logs, and since no other sorting keys

Finally, by calling the sorted function and passing our custom comparison function, we obtain a sorted list that meets all the

Here's the step-by-step approach within the cmp function:

leading of ensures these come before any digit-logs.

which are the content and, if necessary, the identifier.

algorithms, making it an efficient and elegant solution to the problem.

"let1 art can" splits into identifier "let1" and content "art can"

("let1 art can".split(" ", 1)) => ("let1", "art can") => (0, "art can", "let1")

("let2 own kit dig".split(" ", 1)) => ("let2", "own kit dig") => (0, "own kit dig", "let2")

For digit-logs, cmp will return a tuple just containing '1', which places all digit-logs after the letter-logs.

"let2 own kit dig" splits into identifier "let2" and content "own kit dig"

"dig1 8 1 5 1" splits but we only need to recognize it as a digit-log

After applying the cmp function, the sorted function sorts the logs based on these tuples:

"dig2 3 6" also splits but we only need to mark it as a digit-log

function along with a custom comparison key to effectively reorder the logs.

Split each log into two parts: identifier and content

If it's a digit-log (i.e., content starts with a digit),

we use (1,) so that it appears after all the letter-logs,

and within the digit-logs subgroup, it maintains the original order

The logs are sorted according to the keys provided by the get_log_key function.

(since we don't provide any further sorting keys for digit-logs).

return (0, content, identifier) if content[0].isalpha() else (1,)

Sort logs using sort() function. The key parameter takes a function

def reorderLogFiles(self, logs: List[str]) -> List[str]:

identifier, content = log.split(' ', 1)

that extracts sorting keys from the log entries.

// Sort the given logs with specified ordering rules

// Split the logs into two parts: identifier and content

// Check if log1 and log2 contents start with a digit

// One log is a digit-log and the other is a letter-log

std::vector<std::string> reorderLogFiles(std::vector<std::string>& logs) {

std::sort(logs.begin(), logs.end(), [&](const std::string& log1, const std::string& log2) {

// If log1 is a digit log, then it should come after log2 if log2 is a letter log

// Find the first space in both logs to separate identifier from content

// Check if the first character of the content of log1 and log2 is a digit

boolean isDigitLog1 = Character.isDigit(splitLog1[1].charAt(0));

boolean isDigitLog2 = Character.isDigit(splitLog2[1].charAt(0));

return sorted(logs, key=get_log_key)

public String[] reorderLogFiles(String[] logs) {

private int compareLogs(String log1, String log2) {

String[] splitLog1 = log1.split(" ", 2);

String[] splitLog2 = log2.split(" ", 2);

if (isDigitLog1 && isDigitLog2) {

// Function to reorder log files based on specified rules

size t firstSpaceLog1 = log1.find(' ');

size_t firstSpaceLog2 = log2.find(' ');

if (isDigitLog1 && isDigitLog2) {

function reorderLogFiles(logs: string[]): string[] {

// Custom sort function for log files

return logs.sort((log1, log2) => {

// Function to check if a character is a digit

const lastCharLog1 = log1[log1.length - 1];

const lastCharLog2 = log2[log2.length - 1];

const isDigit = (char: string): boolean => char >= '0' && char <= '9';</pre>

// If both logs are digit logs, maintain their relative order

if (isDigit(lastCharLog1) && isDigit(lastCharLog2)) {

const contentLog1 = log1.split(' ').slice(1).join(' ');

const contentLog2 = log2.split(' ').slice(1).join(' ');

if (contentLog1 === contentLog2) {

// Check last character of log1 and log2 to determine if they are digit logs

// If contents are equal, sort based on lexicographical order of the entire log

If it's a letter-log (i.e., content starts with an alphabetic character),

we want it sorted by (0, content, identifier) to ensure that it appears

and within the digit-logs subgroup, it maintains the original order

The logs are sorted according to the keys provided by the get_log_key function.

(since we don't provide any further sorting keys for digit-logs).

return (0, content, identifier) if content[0].isalpha() else (1,)

Sort logs using sort() function. The key parameter takes a function

and secondly by identifier if there's a tie in content.

If it's a digit-log (i.e., content starts with a digit),

we use (1,) so that it appears after all the letter-logs,

before any digit-logs and that letter-logs are sorted firstly by content

return false;

if (isDigitLog1) {

TypeScript

return false;

// Lambda function to check if a character is a digit

auto isDigit = [](char c) { return c >= '0' && c <= '9'; };</pre>

bool isDigitLog1 = isDigit(log1[firstSpaceLog1 + 1]);

bool isDigitLog2 = isDigit(log2[firstSpaceLog2 + 1]);

// If both logs are digit logs, maintain their relative order

return isDigitLog1 ? 1 : -1;

// Custom sort function for log files

return 0;

Arrays.sort(logs, this::compareLogs);

def get log key(log: str) -> tuple:

("dig1 8 1 5 1".split(" ", 1)) => ("dig1", "8 1 5 1") => (1,)

("dig2 3 6".split(" ", 1)) => ("dig2", "3 6") => (1,)

criteria outlined in the problem description. The sorted function rearranges the elements according to the tuples returned by our comparison function cmp, effectively pushing letter-logs to the front and sorting them while keeping digit-logs in their original

follow, they maintain their relative input order.

- order at the end of the resulting list.
- Solution Approach

The implementation of the solution primarily revolves around the Python sorted function and a custom sorting key function,

which is the cmp function in the reference solution. The sorted function is an efficient sorting algorithm that uses the TimSort algorithm, which is a hybrid sorting algorithm derived from merge sort and insertion sort. The cmp function is where the distinction between letter-logs and digit-logs is made and the sorting logic is customized. It takes a log entry as an argument and uses the split method to divide it into two parts, the identifier and the content, with a space

being the delimiter. Since we need both parts separately for the sorting criteria, splitting them does the job efficiently.

Split the log into identifier a and content b: a, b = x.split(' ', 1)

Example Walkthrough

Step 1: Differentiating logs

Letter-logs

code:

Digit-logs

contain numbers after the first space.

Check whether the log is a letter-log or a digit-log: if b[0].isalpha(): else:

For letter-logs, return a tuple (0, b, a), where b is the content of the log without the identifier and a is the identifier. The

For digit-logs, return a tuple (1,), which only has the digit 1. It ensures digit-logs are placed after letter-logs and since no other values are in the tuple, their original order (relative to other digit-logs) is maintained. The sorted function uses this cmp key function to sort the entire list of logs. The logs are compared based on the tuples

returned by cmp; if the first elements of these tuples are different (which differentiate between letter and digit logs), it sorts

based on this. If they are the same (as with all letter-logs), it proceeds to sort based on the subsequent elements in the tuple,

To sum up, the algorithm utilizes a sorting key function to create sortable tuples for each log entry, which Python's sorted

function then utilizes to sort the entire list accordingly. This approach does not require any additional data structures or complex

Let's walk through a small example to illustrate the solution approach. Consider the following logs array: logs = ["dig1 8 1 5 1", "let1 art can", "dig2 3 6", "let2 own kit dig", "let3 art zero"]

We are tasked with reordering these logs based on the rules provided in the problem description.

Step 2: Applying the custom sorting function We implement a custom sorting key function, cmp, to sort the logs. This function will handle logs differently based on their type.

For letter-logs, cmp will return a tuple of the form (0, content, identifier). Here's how it will look for our letter-logs in python

Firstly, we want to differentiate the logs into letter-logs and digit-logs. We can see that "let1 art can", "let2 own kit dig", and "let3"

art zero" are letter-logs because they contain letters after the first space. "dig1 8 1 5 1" and "dig2 3 6" are digit-logs because they

"let3 art zero" splits into identifier "let3" and content "art zero" ("let3 art zero".split(" ", 1)) => ("let3", "art zero") => (0, "art zero", "let3")

Step 3: Sorting logs

Letter-logs sorted by content and then by identifier if needed sorted_letter_logs = sorted(["let1 art can", "let2 own kit dig", "let3 art zero"], key=cmp) # Output for letter-logs is: ["let3 art zero", "let1 art can", "let2 own kit dig"]

Solution Implementation

from typing import List

Digit-logs maintain their original order

Therefore, after letter-logs they are just appended: ["dig1 8 1 5 1", "dig2 3 6"] # Combined sorted loas sorted_logs = sorted_letter_logs + ["dig1 8 1 5 1", "dig2 3 6"]

This final output respects all the sorting rules mentioned in the problem description and uses the power of Python's sorted

Final output: ["let3 art zero", "let1 art can", "let2 own kit dig", "dig1 8 1 5 1", "dig2 3 6"]

We return a tuple that will help us sort the logs: # If it's a letter-log (i.e., content starts with an alphabetic character), # we want it sorted by (0, content, identifier) to ensure that it appears # before any digit-logs and that letter-logs are sorted firstly by content # and secondly by identifier if there's a tie in content.

Java

C++

#include <vector>

#include <string>

#include <algorithm>

class Solution {

return logs;

class Solution:

Python

```
// Both logs are letter-logs
if (!isDigitLog1 && !isDigitLog2) {
    // Compare the contents of the logs
    int contentComparison = splitLog1[1].compareTo(splitLog2[1]);
    if (contentComparison != 0) {
        return contentComparison;
    // If contents are identical, compare the identifiers
    return splitLog1[0].compareTo(splitLog2[0]);
// Both logs are digit-logs. in which case we return 0 to retain their original order
```

// Letter—logs should come before digit—logs, thus returning 1 if the first log is a digit—log, else —1

```
// If log2 is a digit log, then it should come after log1 if log1 is a letter log
    if (isDigitLog2) {
        return true;
    // If both logs are letter logs, compare their contents excluding identifier
    std::string contentLog1 = log1.substr(firstSpaceLog1 + 1);
    std::string contentLog2 = log2.substr(firstSpaceLog2 + 1);
    // If the contents are equal, sort based on the lexicographical order of the entire log
    if (contentLog1 == contentLog2) {
        return log1 < log2;</pre>
    // Otherwise, sort based on the lexicographical order of just the contents
    return contentLog1 < contentLog2;</pre>
});
// Return the reordered logs
return logs;
```

return 0; // If log1 is a digit log, it should come after log2 if log2 is a letter log if (isDigit(lastCharLog1)) { return 1; // If log2 is a digit log, it should come after log1 if log1 is a letter log if (isDigit(lastCharLog2)) { return -1; // If both logs are letter logs, compare based on content excluding identifier

return log1.localeCompare(log2); // Changed to use standard localeCompare for string comparison // Otherwise, sort based on the lexicographical order of just the content return contentLog1.localeCompare(contentLog2); // Changed to use standard localeCompare for string comparison }); from typing import List class Solution: def reorderLogFiles(self, logs: List[str]) -> List[str]: def get log key(log: str) -> tuple: # Split each log into two parts: identifier and content identifier, content = log.split(' ', 1) # We return a tuple that will help us sort the logs:

Time and Space Complexity The time complexity of the given code is O(NlogN), where N is the number of logs. This arises from the use of the sorted

return sorted(logs, key=get_log_key)

that extracts sorting keys from the log entries.

function, which typically has O(NlogN) time complexity for sorting an array. Specifically, every comparison between two logs requires string comparison, and in the worst case, each comparison can take O(M) time, where M is the maximum length of a log file. However, assuming M is relatively small and therefore the comparison time is constant, the overall time complexity still remains O(NlogN). The space complexity of the code is O(N). This is because the sorting algorithm used in Python (Timsort) is a hybrid stable sorting algorithm derived from merge sort and insertion sort characterizes O(N) space complexity. This space is needed to hold

the intermediate results for the sorted array. Furthermore, the extra space needed for the lambda function's return values is not significant compared to the space needed for these intermediate results as they do not depend on the number of logs N.