947. Most Stones Removed with Same Row or Column Medium Depth-First Search Union Find Graph Hash Table

Problem Description

stones can occupy the same coordinate point. A stone can be removed if there is at least one other stone sharing the same row or column on the grid, and that other stone is still on the plane (i.e., it hasn't been removed). Our task is to find out the maximum number of stones that can be removed while following the stated rule. The stones are

Leetcode Link

represented in an array named stones, with each element stones [i] being a list of 2 integers: [xi, yi], which corresponds to the location of the i-th stone on the 2D plane.

Intuition

The intuition behind the solution begins by understanding that if two stones share the same row or the same column, they are part of

the same connected component. In the context of the problem, connected components are groups of stones that are all connected

by rows or columns in such a way that if you pick any stone in a connected component, you'd be able to trace to any other stone in the same connected component through shared rows or columns. The last stone in such a connected component can never be

one.

removed as there would be no other stone in the same row or column. Now, the key observation is that the maximum number of stones that can be removed is equal to the total number of stones minus the number of these connected components. To illustrate, imagine you have three stones forming a straight line either horizontally or vertically. You can remove the two endpoint stones but have to leave the middle one, making the connected components count as

The provided Python solution employs the union-find (or disjoint-set) data structure to efficiently find connected components. This structure helps to keep track of elements that are connected in some way, and in this case, is used to identify stones that are in the same row or column. Here's the breakdown of how the union-find algorithm is applied:

1. Each stone is represented by a node, whose initial "parent" is itself, indicating that it is initially in its own connected component. 2. Iterate over all stones. For each stone (x, y), unify the component containing x with the component containing y + n (we offset y by n to avoid collisions between row and column indices as they could be the same). 3. After all stones have been iterated over, we count unique representatives of the connected components which are the parent

4. The answer is the total number of stones minus the number of unique representatives (connected components).

By applying union-find, the solution achieves a merging process of the stones that lie in the same connected component quickly, leading to an efficient way to calculate the maximum number of stones that can be removed.

each node, where a node is either a row or a column index.

Initially, every element is set to be its own parent.

Here's a step-by-step explanation of the solution:

nodes for the rows and columns.

- Solution Approach The solution implements a Union-Find data structure to keep track of connected components. This data structure is particularly useful in dealing with problems that involve grouping elements into disjoint sets where the connectivity between the elements is an
- essential attribute.

Initialization An array p is created with size 2*n. This is because we have n possible x-coordinates (0 to n-1) and n possible y-coordinates (0

to n-1), but we need to separate the representation to avoid collision, hence n is doubled. The array p represents the parent of

The find function is a standard function in Union-Find, which returns the representative (root parent) of the disjoint set that x

If x does not point to itself, we recursively find the parent of x until we reach the root while applying path compression. Path

compression means we set p[x] to its root directly to flatten the structure, which speeds up future find operations on elements

The find function

belongs to.

in the same set.

Unification Process

coordinate.

The main loop iterates over every stone.

Counting Connected Components

rows and columns. The unification is done by setting the parent of the set containing x to be the parent of the set containing y + n using the find function.

A set s is used to store unique parents of the stones, which comes from applying the find function on every stone's x-

By implementing Union-Find, the solution approach efficiently identifies and unifies stones into connected components and

len(stones) - len(s) gives us the count of removable stones, satisfying the problem's requirement.

calculates the maximum number of stones that can be removed, leading to an effective strategy for this problem.

The number of unique parents indicates the number of connected components.

Let's consider a small example with 5 stones to illustrate the solution approach.

like this: stones = [[0,0], [0,2], [1,1], [2,0], [2,2]].

For stones [2] = [1,1], we unify 1 with 1 + 5.

Next, stones[3] = [2,0], we unify 2 with 0 + 5.

Finally, stones [4] = [2,2], we unify 2 with 2 + 5.

def removeStones(self, stones: List[List[int]]) -> int:

unique_roots = {find_root(x) for x, _ in stones}

return len(stones) - len(unique_roots)

public int removeStones(int[][] stones) {

for (int[] stone : stones) {

private int find(int x) {

return parent[x];

if (parent[x] != x) {

// of the stone's adjusted column index

uniqueRoots.add(find(stone[0]));

return stones.length - uniqueRoots.size();

parent[x] = find(parent[x]);

parent[find(stone[0])] = find(stone[1] + n);

// since each root represents a connected component

// Function to find the root of the element with path compression

parents[x] = find(parents[x]); // Path compression step

// Function to initialize 'parents' whereby each element is its own parent.

// Function to find the root of element 'x' with path compression.

22 // Function to remove as many stones as possible while ensuring

// Define a large enough value for the maximum coordinate.

// Initialize the 'parents' collection to twice the value of 'maxCoord'.

// The number of stones that can be removed is the total number of stones

// minus the number of unique roots in the disjoint-set forest.

// at least one stone is left in the same row or column.

function removeStones(stones: Point[]): number {

// Perform union operations for the stones.

const colParent = find(col + maxCoord);

parents[x] = find(parents[x]); // Path compression step

// Global variable to hold parent pointers for the disjoint-set (Union-Find) data structure.

(which have to remain) from the total number of stones.

29 # Note: This code assumes that the List class has been imported from the typing module:

int n = 10010; // Representative value for scaling row and column indices

// Add the representative of each stone's row to the set of unique roots

// the structure of the tree, effectively speeding up future `find` operations.

if parent[root_id] != root_id:

return parent[root_id]

Recursive function that finds the root of a disjoint set.

parent[root_id] = find_root(parent[root_id])

Path compression is applied to flatten the structure for efficiency.

Given the constraints of the problem, there can be at most 20000 nodes

Calculate how many stones can be removed by subtracting the number of unique sets

private int[] parent; // Array to represent the parent of each element in the Disjoint Set Union (DSU)

Set<Integer> uniqueRoots = new HashSet<>(); // Set to store unique roots after unions have been performed

// The number of stones that can be removed is the total number of stones minus the number of unique roots,

// Recursively find the root representative of the element `x`. Path compression is applied here to flatten

because stones could be placed in rows and columns 0 through 9999.

We keep a set s and insert the root parent for each x-coordinate after unification.

the roots for our example, meaning we have 3 connected components.

considering both x and y coordinates.

applying path compression.

3. Unification Process:

• For a given stone at coordinates (x, y), we unify the set containing x with the set containing y + n to ensure we don't mix up

Calculating the Answer Finally, the solution is the total number of stones minus the number of connected components. To see why this is the case, for each connected component, one stone will inevitably remain, making all others removable.

Example Walkthrough

Steps According to the Solution Approach: 1. Initialization: We create an array p of size 2*n = 10. This array will help us keep track of the parent of each coordinate

2. The find function: Let's define our find function which takes an index x and recursively finds the representative of x's set, while

Looking at stones[0] = [0,0], we unify the x-coordinate (0) with the y-coordinate (0 + 5) to avoid collision with the x-

After the unification, we have root parents for the x-coordinates: [0, 1, 2]. Let's assume union-find identifies 0, 1, and 2 as

Assume we have n = 5 stones at the following coordinates on a 2D plane: [[0,0], [0,2], [1,1], [2,0], [2,2]]. So our stones array looks

- coordinates. • Then, stones[1] = [0,2], we unify 0 with 2 + 5.
- We have a total of 5 stones, and we've identified 3 connected components. The maximum number of stones that we can remove is len(stones) - len(s) which is 5 - 3 = 2. Thus, we can remove a maximum of 2 stones while ensuring that no

def find_root(root_id):

 $num_nodes = 10000$

from typing import List

5. Calculating the Answer:

stone is isolated.

removable stones.

Python Solution

class Solution:

8

9

10

11

12

13

14

15

23

24

25

26

27

31

12

13

14

15

16

17

18

21

23

24

25

26

27

28

29

30

31

32

33

34

35

36

37

38

39

41

43

42 };

12 }

13

16

18

19

21

25

26

28

29

30

31

33

34

35

36

43

44

45

46

48

47 }

O(m).

20 }

int find(int x) {

Typescript Solution

if (parents[x] != x) {

return parents[x];

// Type definition for a 2D point.

function initParents(size: number): void {

for (let i = 0; i < size; i++) {

type Point = [number, number];

const parents: number[] = [];

parents[i] = i;

function find(x: number): number {

if (parents[x] !== x) {

const maxCoord = 10010;

initParents(maxCoord << 1);</pre>

stones.forEach(([row, col]) => {

const rowParent = find(row);

parents[rowParent] = colParent;

return stones.length - uniqueRoots.size;

Time and Space Complexity

complexity analysis of this code is as follows:

return parents[x];

4. Counting Connected Components:

This walkthrough demonstrates the solution approach using the union-find algorithm to efficiently compute the maximum number of stones that can be removed according to the given rules. By identifying the connected components where stones share the same

row or column, union-find allows us to easily count these connected components and therefore calculate the maximum number of

- # Create an array representing the disjoint set forest with an initial parent of itself. 16 parent = list(range(num_nodes * 2)) # Iterate through each stone and unify their row and column into the same set. 19 for x, y in stones: parent[find_root(x)] = find_root(y + num_nodes) 20 21 22 # Use a set comprehension to store the unique roots of all stones' rows and columns.
- parent = new int[n << 1]; // Initialize the parent array for the DSU for (int i = 0; i < parent.length; ++i) { parent[i] = i; // Initially, each element is its own parent for (int[] stone : stones) { 10 // Perform union of stone's row and column by setting the parent of the stone's row to the representative 11

Java Solution

class Solution {

```
33 }
34
C++ Solution
 1 class Solution {
2 public:
       // Class member to hold the parent pointers for the Union-Find data structure
       vector<int> parents;
       // Function to remove as many stones as possible while ensuring at least one
 6
       // stone is left in the same row or column
       int removeStones(vector<vector<int>>& stones) {
           // Define a large enough value for the maximum coordinate
           int maxCoord = 10010;
10
           // Resize the parents vector to double the value since we are mapping
11
12
           // 2-D coordinates to a 1-D array
           parents.resize(maxCoord << 1);</pre>
13
           // Initialize the parent of each element to be itself
14
           for (int i = 0; i < parents.size(); ++i) {</pre>
15
               parents[i] = i;
16
17
18
           // Perform union operations for the stones
19
           for (auto& stone : stones) {
21
               parents[find(stone[0])] = find(stone[1] + maxCoord);
23
24
           // Using a set to store unique roots after path compression
25
           unordered_set<int> uniqueRoots;
            for (auto& stone : stones) {
26
                uniqueRoots.insert(find(stone[0]));
27
28
29
30
           // The number of stones that can be removed is the total number of stones
           // minus the number of unique roots in the disjoint-set forest
31
32
           return stones.size() - uniqueRoots.size();
```

37 // Using a set to store unique roots after path compression, // which indicates the connected components. 38 const uniqueRoots = new Set<number>(); 39 stones.forEach(([row, _]) => { 40 uniqueRoots.add(find(row)); 41 42 });

Time Complexity:

});

1. The find function is at most traversing the depth of the union-find tree, which, with path compression, results in nearly constant time per operation. However, worst-case without any optimizations can be O(n), but realistically, with path compression, it is closer to $O(\alpha(n))$, where $\alpha(n)$ is the Inverse Ackermann function, which grows very slowly and is practically considered constant time.

2. There is a loop that runs once for each stone, and inside this loop, two find operations are performed, one for the x-coordinate

Considering there are m stones, each union-find operation is $O(\alpha(n))$, the overall time complexity across all stones will be $O(m * \alpha(n))$.

Since the Inverse Ackermann function α(n) is practically constant for all reasonable values of n, the time complexity simplifies to

1. An array p of size n << 1 (or 2 * n) is created to represent the union-find structure, where n is a constant representing the

maximum number of different row/column values to expect, set to 10010 in the code. Hence, the space used by p is O(n).

and one for the y-coordinate shifted by n to keep the x and y coordinates distinct in the union-find structure.

The given code snippet uses the Union-Find algorithm to remove as many stones as possible on a 2D grid while ensuring at least one

stone is left in a connected group (a group of stones connected by either the same row or the same column). The time and space

Space Complexity:

2. A set s is created to store unique roots after path compression. In the worst case, this set will contain all the stones if none share the same row or column, hence an O(m) space complexity at worst. Combining these, the total space complexity is O(n + m) which, given the constant size of n, simplifies to O(m), where m is the

number of stones.

- In summary: Time Complexity: O(m), where m is the number of stones.
- Space Complexity: O(m), where m is the number of stones.

In this problem, we're given n stones positioned at integer coordinate points in a 2-dimensional plane. The key rule is that no two