1188. Design Bounded Blocking Queue

should block the calling thread until there is space available to add the new element.

Medium Concurrency

Problem Description

queue is a queue with a fixed maximum size, and it has the capability to block or wait when operations like enqueue (adding to the queue) and dequeue (removing from the queue) cannot be performed because the queue is full or empty, respectively. The queue supports three main operations: • enqueue(int element): This method adds an element to the end of the queue. If the queue has reached its capacity, the method

The problem is about creating a data structure which is a bounded blocking queue with thread-safety in mind. A bounded blocking

Leetcode Link

- dequeue(): This removes and returns the element at the front of the queue. If the queue is empty, this operation should block until there is an element available to dequeue. • size(): This returns the current number of elements in the queue.
- It is especially noted that the implementation will be tested in a multithreaded environment, where multiple threads could be calling

these methods simultaneously. Therefore, it is crucial that the implementation ensures that all operations on the bounded blocking

queue are thread-safe (i.e., function correctly when accessed from multiple threads).

another thread. This behavior fits perfectly for our problem's requirements.

Lastly, the use of any built-in bounded blocking queue implementations is prohibited as the goal is to understand how to create such a data structure from scratch, potentially in a job interview. Intuition

The solution to the problem involves coordinating access to the queue to ensure that only one thread can perform an enqueue or

that could lead to incorrect behavior or data corruption. To achieve this, we employ two synchronization primitives called Semaphores. A semaphore maintains a set of permits, and a thread that wants to perform an operation must first acquire a permit. If no permit is available, the thread blocks until a permit is released by

dequeue operation at a time to maintain thread safety. This means protecting the internal state of the queue from race conditions

1. s1, which starts with a number of permits equal to the capacity of the queue. This semaphore controls access to enqueue operation. A thread can enqueue if it can acquire a permit from \$1, which signifies that there is room in the queue. After the enqueue operation, the thread releases a permit to \$2, signaling that there is an item available to be dequeued. 2. s2, which starts with zero permits as the queue is initially empty. This semaphore controls access to the dequeue operation. For

a thread to dequeue, it must acquire a permit from \$2, which signifies that there is at least one item in the queue to be

dequeued. After the dequeue operation, the thread releases a permit to \$1, indicating that there is now additional space available in the queue.

We use two semaphores in the solution:

- The queue itself (q) is represented by a deque (double-ended queue) from Python's collections module, which allows for fast appending and popping from both ends. Note that accessing len(q) to get the size of the queue does not need to be serialized by semaphores, as it is not modifying the queue. This approach enables us to limit the number of elements in the queue to the defined capacity, and to ensure that enqueue and
- dequeue operations wait for the queue to be not full or not empty, respectively, before proceeding, fulfilling the conditions for a bounded blocking queue in a thread-safe manner. **Solution Approach**

In the provided solution, the implementation of the BoundedBlockingQueue class is done with two semaphores and a deque. Here's a

walkthrough of the pattern and algorithms used: • Semaphores: Two instances of the Semaphore class are used, s1 and s2, each serving a distinct purpose. s1 governs the ability to insert an item into the queue, and begins with permits equal to the capacity. 52 reflects the number of items in the queue available to be dequeued and starts with no permits. This is a classic application of the "producer-consumer" problem's solution,

• Deque: A deque (double-ended queue) from the collections module is used to represent the queue's data structure. This

2. self.q.append(element): Once a permit has been acquired (meaning there is space in the queue), the element is safely

1. self.s2.acquire(): This acquires a permit from s2, which signifies that there is at least one element in the queue to be

dequeued. If the queue is empty, this call will block until a thread calls enqueue and releases a permit on s2.

where one semaphore is used to signal "empty slots" and another semaphore is used to signal "available items".

provides efficient FIFO (first-in-first-out) operations needed for enqueueing (append) and dequeueing (popleft).

call will block until another thread calls dequeue and releases a permit on \$1.

the queue, so it does not require interaction with semaphores.

1. It acquires a permit from s1, which now has 1 permit left.

1. It acquires the remaining permit from s1, and s1 now has 0 permits.

2. It appends 2 to the deque, so the queue state becomes [1,2].

1. It acquires the permit from s1, and again s1 has 0 permits.

dequeuing operations using semaphores to manage its capacity and state.

self.semaphore_empty_slots = Semaphore(capacity)

self.semaphore_filled_slots = Semaphore(0)

Add an element to the end of the queue.

def enqueue(self, element: int) -> None:

element = self.queue.popleft()

When enqueue(element: int) is called, the following steps are performed:

not empty.

Example Walkthrough

• Imagine thread A calls enqueue(1):

signifying the queue is at full capacity.

Meanwhile, if thread D calls dequeue():

enqueued into the queue.

For dequeue(), the steps are the mirror image:

2. ans = self.q.popleft(): Removes the oldest (front) element from the queue safely because it's been ensured that the queue is

3. self.s2.release(): Releasing a permit on s2 to signal that an element has been enqueued and is now available for dequeuing.

1. self.s1.acquire(): Attempt to acquire a permit from s1, which represents a free slot in the queue. If there are no free slots, this

3. self.sl.release(): Releasing a permit on s1 to signal that an element has been dequeued and there is now a free slot in the queue.

size() is a straightforward operation as it simply returns the current number of items in the queue, len(self.q). It does not modify

The solution effectively serializes access to the mutable shared state (self.q), preventing race conditions by using semaphores to

coordinate enqueue and dequeue actions. This guarantees that the queue never exceeds its capacity and avoids dequeue

operations being called on an empty queue, thus satisfying thread safety and other requirements of the problem.

- Let's consider a small example to illustrate the solution approach for a BoundedBlockingQueue with a capacity of 2. • Initially, we create the queue with a capacity of 2, initializing semaphore s1 with 2 permits and semaphore s2 with 0 permits.
- 2. It then appends 1 to the deque, and the queue state becomes [1]. 3. Finally, it releases a permit to \$2, indicating that there is one item available for dequeuing. Now, thread B calls enqueue(2):

• At this state, the queue is full. If another thread, say thread C, tries to enqueue(3), it will be blocked as \$1 has no permits left,

3. It releases a permit to s1, increasing the number of permits back to 1, signaling that there is now space for one more item in

3. It releases another permit to \$2, now \$2 has 2 permits indicating there are two items available to be dequeued.

the queue.

needing to acquire a permit.

1 from threading import Semaphore

self.queue = deque()

def dequeue(self) -> int:

return element

def size(self) -> int:

1 import java.util.ArrayDeque;

import java.util.concurrent.Semaphore;

public BoundedBlockingQueue(int capacity) {

dequeueSemaphore = new Semaphore(0);

queue = new ArrayDeque<>();

enqueueSemaphore.acquire();

queue.offer(element);

synchronized (this) {

enqueueSemaphore = new Semaphore(capacity);

public class BoundedBlockingQueue {

2 import java.util.Deque;

Python Solution

8

9

10

11

12

13

14

15

16

17

18

19

20

21

22

23

24

13

14

15

16

17

18

19

20

21

22

23

24

26

27

28

13

14

15

16

17

18

19

20

21

22

23

24

25

26

27

28

29

30

31

32

33

34

35

36

37

38

39

40

41

43

44

45

46

47

48

50

49 };

private:

void enqueue(int element) {

queue_.push(element);

++count_;

int dequeue() {

queue_.pop();

return value;

return count_;

--count_;

int size() {

Typescript Solution

std::unique_lock<std::mutex> lock(mutex_);

// Wait until there is space in the queue

std::unique_lock<std::mutex> lock(mutex_);

// Wait until there is an item to dequeue

std::lock_guard<std::mutex> lock(mutex_);

std::queue<int> queue_; // The queue that holds the elements

std::mutex mutex_; // Mutex to protect access to the queue

int capacity_; // Maximum number of items in the queue

1 const queue: number[] = []; // The queue that holds the elements

2 let capacity: number; // Maximum number of items in the queue

int count_; // Current number of items in the queue

not_empty_condition_.notify_one();

int value = queue_.front();

not_full_condition_.notify_one();

// Get the current size of the queue.

not_full_condition_.wait(lock, [this] { return count_ < capacity_; });</pre>

// Notify one waiting thread (if any) that an item was dequeued

not_empty_condition_.wait(lock, [this] { return count_ > 0; });

// Notify one waiting thread (if any) that space is now available

1. It acquires a permit from \$2 (which has 2 permits at this point), leaving 1 permit left in \$2. 2. It dequeues an element from the deque which is 1 (FIFO order), leaving the queue state as [2].

• If thread C is still waiting to enqueue(3), it can now proceed as a permit became available in s1.

2. It appends 3 to the deque, so the queue state becomes [2,3]. 3. It releases a permit to \$2, which now has 2 permits, reflecting the two items in the queue.

This example demonstrates how the BoundedBlockingQueue enforces its bounds and provides thread-safe enqueueing and

At any time, calling size() returns the number of items currently in the queue, which can be accessed by any thread without

Semaphore to track empty slots

Semaphore to track filled slots

Use deque for queue operations

from collections import deque # Ensure deque is imported class BoundedBlockingQueue: def __init__(self, capacity: int): # Initialize the queue with given capacity.

self.queue.append(element) # Add the element to the queue

// This class represents a thread-safe bounded blocking queue with a fixed capacity.

// Constructor initializes the semaphores and queue with specified capacity.

// Acquire a permit from enqueueSemaphore discarding it, if available capacity is 0 waits.

// Enqueues an element into the queue if there's available capacity.

public void enqueue(int element) throws InterruptedException {

// Adds the element to the end of the queue.

Remove and return an element from the front of the queue.

Get the current number of elements in the queue.

26 return len(self.queue) # Return the size of the queue 27

self.semaphore_empty_slots.acquire() # Decrease the counter of empty slots, wait if no empty slots

self.semaphore_filled_slots.acquire() # Decrease the counter of filled slots, wait if no filled slots

Return the dequeued element

self.semaphore_filled_slots.release() # Increase the counter of filled slots, signaling dequeue if slots are filled

Remove the element from the queue

self.semaphore_empty_slots.release() # Increase the counter of empty slots, signaling enqueue if slots are available

// Semaphore to control the number of elements that can be added (based on capacity). private final Semaphore enqueueSemaphore; // Semaphore to control the number of elements that can be removed (starts at 0). private final Semaphore dequeueSemaphore; 10 // The queue to store elements. 11 private final Deque<Integer> queue; 12

Java Solution

```
29
           // Release a permit to dequeueSemaphore, increasing the number of available elements to dequeue.
30
           dequeueSemaphore.release();
31
33
       // Dequeues an element from the front of the queue.
34
       public int dequeue() throws InterruptedException {
35
           // Acquire a permit from dequeueSemaphore, waiting if necessary until an element is available.
36
           dequeueSemaphore.acquire();
           int element;
37
           synchronized (this) {
               // Remove and return the front element of the queue.
40
               element = queue.poll();
41
           // Release a permit to enqueueSemaphore, increasing the available capacity.
43
           enqueueSemaphore.release();
           return element;
45
46
47
       // Returns the current number of elements in the queue.
       public int size() {
48
           synchronized (this) {
49
               // The size of the queue is returned.
               return queue.size();
52
53
54 }
55
C++ Solution
    #include <queue>
    #include <mutex>
     #include <condition_variable>
    class BoundedBlockingQueue {
     public:
         // Constructor initializes the queue with a capacity limit.
         BoundedBlockingQueue(int capacity): capacity_(capacity), count_(0) {
             // No need to initialize semaphores since we will use condition_variable and mutex
 11
 12
         // Enqueue adds an element to the queue. If the queue is full, blocks until space is available.
```

// Dequeue removes and returns an element from the queue. If the queue is empty, blocks until an element is available.

std::condition_variable not_full_condition_; // Condition variable to block enqueue when queue is full

std::condition_variable not_empty_condition_; // Condition variable to block dequeue when queue is empty

let count: number = 0; // Current number of items in the queue

```
// Utilizing these for synchronization could be complex without encapsulation
  6 const mutex = new Mutex(); // A pretend Mutex since JavaScript/TypeScript doesn't have one
    const notFullCondition = new ConditionVariable(); // A pretend condition variable for when the queue is not full
     const notEmptyCondition = new ConditionVariable(); // A pretend condition variable for when the queue is not empty
    // Initialization of the queue with a capacity limit.
    function initializeQueue(initialCapacity: number): void {
 12
         capacity = initialCapacity;
 13
 14
    // Adds an element to the queue. If the queue is full, it is supposed to block until space is available.
    async function enqueue(element: number): Promise<void> {
         await mutex.lock();
 17
 18
         try {
 19
             // Wait until there is space in the queue
 20
             while (count >= capacity) {
                 await notFullCondition.wait(mutex);
 22
 23
             queue.push(element);
 24
             count++;
 25
             // Notify one waiting thread (if any) that an item was enqueued
             notEmptyCondition.notifyOne();
 26
 27
         } finally {
             mutex.unlock();
 28
 29
 30 }
 31
    // Removes and returns an element from the queue. If the queue is empty, it is supposed to block until an element is available.
    async function dequeue(): Promise<number> {
         await mutex.lock();
 35
 36
             // Wait until there is an item to dequeue
             while (count === 0) {
 37
                 await notEmptyCondition.wait(mutex);
 38
 39
             const value = queue.shift();
 40
 41
             count--;
             // Notify one waiting thread (if any) that space is now available
 42
             notFullCondition.notifyOne();
 43
             return value;
 44
         } finally {
 45
 46
             mutex.unlock();
 47
 48
 49
     // Returns the current size of the queue.
     function size(): number {
         return count; // No need for synchronization in a single-threaded environment
 52
 53
 54
    // Note that Mutex and ConditionVariable are not native JS/TS classes.
     // These would need to be implemented or a library would have to be used to simulate them.
 57
Time and Space Complexity
For this BoundedBlockingQueue implementation using semaphores, we will analyze the time and space complexities of its operations.
```

• enqueue: The enqueue operation involves two semaphore operations (acquire and release) and an append operation on a deque. The semaphore operations are generally 0(1) assuming they don't block; if they do block, the time complexity is dependent on external factors such as contention from other threads. Appending to the deque is an 0(1) operation. Therefore, the combined

again 0(1) per call in the absence of blocking.

Time Complexity

- time complexity is 0(1) per call in the absence of blocking. dequeue: Like enqueue, dequeue also has two semaphore operations and a popleft operation on the deque. Since deque's popleft is designed to be 0(1) and semaphore operations are 0(1) without blocking, the overall time complexity for dequeue is
- size: This simply returns the number of items in the deque, which is maintained internally and is thus an 0(1) operation. **Space Complexity** The space complexity revolves around the deque that stores elements. Since the capacity of the queue is fixed, the maximum space it will use is O(capacity), corresponding to the maximum number of elements that can be enqueued at any given time.

_init__: Initializing the queue involves setting up two semaphores and the underlying deque. This operation is constant time,

0(1), as it involves only a fixed number of operations, regardless of the capacity of the queue.

In summary, except for the potential blocking on semaphores (which can't be quantified in standard complexity analysis), all fundamental operations (__init__, enqueue, dequeue, size) on the BoundedBlockingQueue class have a time complexity of 0(1). The space complexity is O(capacity) based on the fixed maximum size of the underlying deque.