1371. Find the Longest Substring Containing Vowels in Even Counts

Medium Bit Manipulation String] **Prefix Sum** Hash Table

In this problem, you're given a string s. Your task is to find the length of the longest substring of s where each of the vowels - 'a',

Problem Description

'e', 'i', 'o', 'u' - appears an even number of times. Substrings are continuous parts of the original string, and in this case, any substring that satisfies the vowel condition could potentially be the longest. The goal is to figure out the maximum length possible while adhering to the even occurrence condition for vowels.

Intuition

The intuition behind the solution is to use a bitwise representation to track the count of vowels in the substring efficiently.

Because we only care about even or odd occurrences, we don't need to keep a count of each vowel; we only need to know if a vowel has appeared an even or odd number of times. Thus, each vowel can be represented with 1 bit, resulting in a 5-bit integer for all five vowels (where 0 means even and 1 means odd number of occurrences). The key insight is that if at two different points in the string, the 5-bit integer (or state) is the same, it means that we have a substring (between these two points) where all vowels have appeared an even number of times.

than the previous maximum. The pos array keeps track of the first occurrence of each state to calculate lengths of valid substrings. We start by initializing the pos array with inf to denote that those states have not been seen yet, except for the

state 0, which is initialized to -1 to handle the edge case where a valid substring starts at the beginning of the input string.

The solution iterates through the string and toggles the respective bit in the state when a vowel is found. If the state has been

seen before, we calculate the length of the new substring ending at the current character and update the answer if it's longer

By using this bitwise state and the pos array, we can efficiently find the longest substring where all vowels appear an even number of times. Solution Approach

The solution utilizes a bitwise approach combined with a hash map (implemented as a list called post in the code) to keep track of

the indices at which each state occurs for the first time. **Key Components of the Solution:**

only 'a' and 'i' have been seen an odd number of times, the state would be 10100 in binary, which is 20 in decimal.

state is o since we start with an even count (zero occurrences) for all vowels.

Vowels Bitmasking: Each vowel has a corresponding bit in a 5-bit integer, where the order is 'a', 'e', 'i', 'o', 'u'. For example, if

Index Memory (pos): An array called pos of size 32 (since there are 2^5 possible states due to 5 vowels) is used to

Tracking States: The state variable is a 5-bit integer representing the current state of vowels' counts (even/odd). The initial

Iterating Through the String: As we iterate through the string, we check each character. If it's a vowel, we flip the corresponding bit in the state using the exclusive OR (XOR) operation: state ^= 1 << j, where j is the index of the vowel

remember the earliest occurrence of every state. This array is initialized with inf, except for pos[0] which is set to -1.

- in the string 'aeiou'. **Updating the Answer**: With each new character processed, we check if the current state has been encountered before: If it's been seen, we calculate the length of the substring from the first occurrence of this state to the current position.
- First Occurrence: We also check if the current state's first occurrence needs to be updated in the post array. If the current index is less than the stored value in pos[state], we update pos[state] with the current index.

By the end of the string traversal, ans will hold the length of the longest substring where all vowels appear an even number of

• If the calculated length is greater than the ans (which keeps track of the maximum length seen so far), we update ans.

• We start at state = 0 and pos[0] = -1 because we have an even count (zero) for all vowels at the start.

Let's walk through a quick example: s = "eleetminicoworoep"

• If state is 3 again at a later point, we know that between these two indices, 'a' and 'e' must have appeared an even number of times. Therefore, we calculate the length of this substring and check if it's the maximum. • We continue this process until the end of the string, constantly updating ans with longer valid substrings as we find them.

The implementation is efficient with a time complexity of O(n), where n is the length of the string, because we only need a single

Initialize the pos array with length 32 to represent all possible states of vowel occurrences (2^5 for 5 vowels). Set all values

We start with state = 0 (since no vowels have been seen yet) and ans = 0 (since no substrings have been found yet).

• When at index 0, s[0] = 'a': We see our first vowel 'a'. The bit for 'a' in the state is toggled from 0 to 1. Now, state = 00001.

pass through the string to compute the result, and each operation within that pass is of constant time complexity.

• Iterating over s, whenever we encounter a vowel, we update state. Suppose state becomes 3 after some operations; this means 'a' and 'e'

Example Walkthrough

s = "aeiobcdf"

Example Walkthrough:

• At index 1, s[1] = 'e': We toggle the bit for 'e'. Now, state = 00011. • At index 2, s[2] = 'i': Toggling the 'i' bit: state = 00111.

At index 3, s[3] = 'o': Toggling 'o': state = 01111.

At index 4, s[4] = 'b': 'b' is not a vowel; state remains 01111.

At index 5, s[5] = 'c': 'c' is not a vowel; state remains 01111.

At index 6, s[6] = 'd': 'd' is not a vowel; state remains 01111.

At index 7, s[7] = 'f': 'f' is not a vowel; state remains 01111.

As we iteratively check each character in s, we use the following steps:

to inf, except pos[0] to -1.

Let's illustrate the solution approach using the string:

times, and that's what we return as the solution.

have been seen an odd number of times so far.

```
final state is 01111.
```

Solution Implementation

positions = [inf] * 32

for index, char in enumerate(s):

if char == vowel:

state **^= 1 <<** j

public int findTheLongestSubstring(String s)

// String of vowels to check against

// Loop through characters of input string

for (int i = 0; i < s.length(); i++) {</pre>

char currentChar = s.charAt(i);

Check if the current character is a vowel

for j, vowel in enumerate(vowels):

def findTheLongestSubstring(self, s: str) -> int:

Initialize a list to store the first position we encounter a given state

The starting state (all vowels have even counts) is at index 0

Toggle the corresponding bit if we encounter a vowel

max length = max(max length, index - positions[state])

// pos array to keep track of the earliest index of each state

positions[state] = min(positions[state], index)

State is represented as a bitmask integer of size 32 (for 5 vowels, each can be on/off)

Calculate max length using current index and first index where current state was seen

Update the position for this state if it's the first time we're seeing this state

// Function to find the length of the longest substring containing vowels in even counts

// 'state' will represent the binary value of the vowels seen odd number of times

// 'maxLength' stores the length of the longest substring found so far

Python

from math import inf

class Solution:

vowel, we checked pos to see if the current state had been recorded before. In this instance, since no state repeated other than the initial state 0, the length of the longest valid substring is zero because we have not encountered a state twice where the state would be 00000 again (meaning that all vowels have an even count).

The key takeaway is the tracking of vowel occurrences through bit manipulation, using XOR to toggle between even and odd

counts, and using the pos array to store the first occurrence of a state. This approach allows for quick lookups and updates while

maintaining an efficient assessment of the longest valid substring. However, in our example, there were no repeated states to

determine a valid substring. In a longer string with repeated vowel occurrences, the ans would likely be greater than zero.

So, the ans remains 0 in this example; there are no substrings where each vowel occurs an even number of times.

We've reached the end of the string, and throughout, each vowel has appeared exactly once, so their counts are odd. Our

Throughout our iteration, the state has changed from 00000 to 01111. After toggling the bits whenever we encountered a

positions[0] = -1# List of vowels for reference vowels = 'aeiou' # `state` to keep track of the count of vowels encountered (even/odd as a bitmask) # `max length` to keep track of the length of the longest valid substring so far state = max_length = 0 # Iterate over the string characters with their index

int[] earliestPos = new int[32]; // 32 possible states for 5 vowels (2^5) // Initialize all positions to max value except for state 0 Arrays.fill(earliestPos, Integer.MAX VALUE); // State 0 (no vowels seen or all seen even times) starts at index -1 earliestPos[0] = -1;

int state = 0;

int maxLength = 0;

String vowels = "aeiou";

return max_length

Java

class Solution {

```
// Check and flip the corresponding bit for the vowel
            for (int i = 0; i < 5; i++) {
                if (currentChar == vowels.charAt(j)) {
                    state ^= (1 << j);
            // If state has been seen before, update the maxLength
            maxLength = Math.max(maxLength, i - earliestPos[state]);
            // If state has not been seen before, set it to the current index
            if (earliestPos[state] == Integer.MAX_VALUE) {
                earliestPos[state] = i;
        // Return the length of the longest substring
        return maxLength;
C++
class Solution {
public:
    int findTheLongestSubstring(string s) {
        // Initialize a vector to keep track of the first occurrence of all states.
        vector<int> firstOccurrence(32, INT MAX);
        // Setting the base condition that for state '0', the first occurrence is at -1
        first0ccurrence[0] = -1;
        // Define a string of vowels for easy indexing.
        string vowels = "aeiou";
        // This will keep track of the current state of vowels encountered.
        int currentState = 0;
        // The result variable to store the length of the longest substring.
        int longestSubstringLength = 0;
        // Iterate over the characters of the string.
        for (int i = 0; i < s.size(); ++i) {
            // Check if the current character is a vowel and update the state accordingly.
            for (int j = 0; j < 5; ++j) { // There are 5 vowels.
                if (s[i] == vowels[i]) {
                    // Toggle the i-th bit of the state to represent the occurrence of vowel.
                    currentState ^= (1 << j);</pre>
```

```
// The variable to store the length of the longest substring.
let longestSubstringLength: number = 0;
// Iterate over the characters of the string.
for (let i = 0; i < s.length; i++) {
   // Check if the current character is a vowel and update the state accordingly.
    for (let j = 0; j < vowels.length; j++) {</pre>
        if (s[i] === vowels[i]) {
            // Toggle the i-th bit of the state to represent the occurrence of the vowel.
            currentState ^= (1 << j);</pre>
    // If the current state hasn't occurred before, update the occurrence of the current state.
```

// Calculate the length if the current state has occurred before

// The length of the valid string would be difference of indices.

// Update the occurrence of current state if it's the first time.

// Return the length of the longest valid substring found.

// This function calculates and returns the length of the longest substring

// This will keep track of the current state of vowels encountered.

if (firstOccurrence[currentState] === Number.MAX_SAFE_INTEGER) {

// Calculate the length if the current state has occurred before

// The length of the valid substring would be the difference of indices.

firstOccurrence[currentState] = i;

// Initialize an array to keep track of the first occurrence of all states.

let firstOccurrence: number[] = new Array(32).fill(Number.MAX SAFE INTEGER);

// Setting the base condition that for state '0', the first occurrence is at -1

return longestSubstringLength;

// Define a string of vowels for easy indexing.

function findTheLongestSubstring(s: string): number {

// where the count of each vowel is even.

const vowels: string = "aeiou";

firstOccurrence[\emptyset] = -1;

let currentState: number = 0;

};

TypeScript

firstOccurrence[currentState] = min(firstOccurrence[currentState], i);

longestSubstringLength = max(longestSubstringLength, i - firstOccurrence[currentState]);

```
longestSubstringLength = Math.max(longestSubstringLength, i - firstOccurrence[currentState]);
    // Return the length of the longest valid substring found.
    return longestSubstringLength;
from math import inf
class Solution:
   def findTheLongestSubstring(self, s: str) -> int:
        # Initialize a list to store the first position we encounter a given state
        # State is represented as a bitmask integer of size 32 (for 5 vowels, each can be on/off)
        positions = [inf] * 32
       # The starting state (all vowels have even counts) is at index 0
        positions[0] = -1
       # List of vowels for reference
        vowels = 'aeiou'
       # `state` to keep track of the count of vowels encountered (even/odd as a bitmask)
       # `max length` to keep track of the length of the longest valid substring so far
        state = max_length = 0
       # Iterate over the string characters with their index
        for index, char in enumerate(s):
            # Check if the current character is a vowel
            for i, vowel in enumerate(vowels):
                # Toggle the corresponding bit if we encounter a vowel
                if char == vowel:
                    state ^= 1 << j
           # Calculate max length using current index and first index where current state was seen
           max length = max(max length, index - positions[state])
           # Update the position for this state if it's the first time we're seeing this state
            positions[state] = min(positions[state], index)
        return max_length
```

Time Complexity The given code iterates over each character of the string s, which has a length n. It performs a constant amount of work for

Time and Space Complexity

pos[state]. Because these operations all have a constant time complexity, the overall time complexity of the function is O(n). **Space Complexity** The space complexity is determined by the fixed-size array pos, which has 32 elements (one for each possible state of the 5

each character: checking if the character is a vowel, possibly flipping a bit in the state variable, and updating ans and

vowel bits), and a few additional integer variables (state, ans, and the loop variables). Therefore, the space complexity of the algorithm is 0(1) since the space required does not grow with the size of the input string s.