1438. Longest Continuous Subarray With Absolute Diff Less Than or Equal to Limit Medium Queue Ordered Set **Monotonic Queue** Heap (Priority Queue) **Sliding Window Leetcode Link** Array

The problem provides an array of integers named nums and an integer called limit. The task is to find the length of the longest non-

Problem Description

empty contiguous subarray (a sequence of adjacent elements from the array) where the absolute difference between any two

elements in the subarray does not exceed the limit value.

For example, if the input array is [10, 1, 2, 4, 7, 2] and the limit is 5, the longest subarray where the absolute difference between any two elements is less than or equal to 5 is [1, 2, 4, 7, 2], which has a length of 5. The two key aspects of the problem are:

Working with contiguous elements (subarray), not just any subsets of the array.

Ensuring that every pair of elements in the subarray has an absolute difference of at most limit.

- The solution approach involves using a data structure that maintains a sorted order of elements. This allows the efficient retrieval of
- the smallest and largest elements in the current window (subarray) to check if their absolute difference is within the limit.

1. Maintain a sliding window that expands and contracts as we iterate through the nums array.

2. In each iteration, add the current element to the SortedList, which is our window.

arrive at the solution:

Intuition

3. Check if the absolute difference between the smallest and largest elements in the SortedList exceeds the limit. 4. If it does, we remove the leftmost element from our window (which was first added when the window was last valid) to try and bring the difference back within limit.

A suitable data structure for this problem is a SortedList, provided by the sortedcontainers library in Python. Here's how we can

5. We keep track of the maximum size of the window that satisfied the condition of staying within the limit. This approach ensures that at any given point, we have the longest valid subarray ending at the current position, and we keep

- updating the answer with the maximum size found so far.
- The reason we use a SortedList instead of sorting the window array in each iteration is the time complexity—SortedList maintains

left one. Also, initialize a variable ans to store the longest length of the subarray found so far.

largest (sl[-1]) and smallest (sl[0]) elements in the SortedList does not exceed the limit.

smallest and largest elements up to the current position.

the order of elements with a far lesser time complexity for insertion and removal compared to sorting an array at each step.

The implementation is based on the sliding window pattern, which is an optimization technique to reduce repeated work and maintain a range of elements that fulfill certain criteria. Here is a detailed explanation: 1. Initialize a SortedList named sl and two pointers for the window indices i and j with i being the right pointer and j being the

2. Loop through the elements of nums with index i and value v. o Add the new element v to the SortedList. Because the list is always sorted, doing this helps us quickly reference the

3. Inside the loop, check if the current window (from j to i inclusively) is valid, meaning that the absolute difference between the

If the limit is exceeded, we need to contract the window by removing the leftmost element. This is done by removing

for i, v in enumerate(nums):

sl.remove(nums[j])

evaluating the entire subarray every time.

Following the proposed solution:

that will hold the answer.

2. We start iterating through the elements of nums.

sl.add(v)

9 return ans

end of the window and j is the beginning.

1 sl = SortedList() # Instantiate a SortedList data structure.

2 ans = j = 0 # Initialize the answer and the left pointer of the window to 0.

sorting a list which would take O(N log N) time for each change in the window.

The updated window [2, 2, 5] is now valid. Update ans to 3.

maximum within the current window to decide if the subarray satisfies the condition.

from typing import List # Import List from typing module for type annotation

Initialize a SortedList which allows us to maintain a sorted collection of numbers

Remove the leftmost value from the sorted list as we're shrinking the window

Initialize variables for the answer and the start index of the window

Shrink the window from the left if the condition is violated

max_length = max(max_length, window_end - window_start + 1)

Return the length of the longest subarray after examining all windows

Calculate the length of the current window and compare with the max

def longest_subarray(self, nums: List[int], limit: int) -> int:

while sorted_list[-1] - sorted_list[0] > limit:

sorted_list.remove(nums[window_start])

Move the start of the window to the right

// Iterate over the array using the right pointer 'right'

if (frequencyMap.get(nums[left]) == 0) {

frequencyMap.put(nums[right], frequencyMap.getOrDefault(nums[right], 0) + 1);

// Shrink the sliding window until the absolute difference between the max

// and min within the window is less than or equal to 'limit'

while (frequencyMap.lastKey() - frequencyMap.firstKey() > limit) {

// Decrease the frequency of the number at the left pointer

frequencyMap.put(nums[left], frequencyMap.get(nums[left]) - 1);

// If the frequency drops to zero, remove it from the frequency map

for (int right = 0; right < nums.length; ++right) {</pre>

// Update the frequency of the current number

1 # We import SortedList from the sortedcontainers module

2 from sortedcontainers import SortedList

max_length = 0

window_start = 0

return max_length

sorted_list = SortedList()

window_start += 1

36 # result = sol.longest_subarray([10,1,2,4,7,2], 5)

Update the max_length as needed

while sl[-1] - sl[0] > limit: # If current window is invalid, contract it.

ans = max(ans, i - j + 1) # Store the largest size of the valid window.

Solution Approach

nums[j] from SortedList since j corresponds to the leftmost index of the window. Increment i to move the start of the window to the right.

5. After the loop ends, return ans, which holds the size of the longest subarray found that satisfies the condition. Here is a code snippet to illustrate the solution's core logic:

4. Update the answer ans with the maximum length found so far. We compute the current window size with i - j + 1, as i is the

The SortedList is efficient because it keeps elements sorted at all times. Hence, we can always get the smallest and largest element

in O(1) time and remove elements in O(log N) time, where N is the number of elements in the list. This is much more optimal than

Overall, the use of a sliding window algorithm with SortedList allows us to efficiently solve this problem with a time complexity that

depends on inserting and removing each element into the sorted list (generally O(log N) for each operation), rather than re-

Let's take a small example to illustrate the solution approach. Consider the input array nums = [4, 2, 2, 5, 4] with limit = 2.

1. We initialize an empty SortedList named sl, set two pointers for the window indices j = 0 and i = 0, and ans = 0, the variable

 \circ For i = 0 (v = 4): Add 4 to s1, so s1 = [4]. The window [4] is valid because there is only one element. Update ans to 1.

Example Walkthrough

3. Move to i = 1 (v = 2): ○ Add 2 to sl, leading to sl = [2, 4]. The window [4, 2] is valid as 4 - 2 = 2, which does not exceed the limit. Update ans to 2. 4. Proceed to i = 2 (v = 2):

Add another 2 to sl, so sl = [2, 2, 4]. The window [4, 2, 2] is still valid for the same reasons. Update ans to 3.

○ Add 5 to s1, which yields s1 = [2, 2, 4, 5]. The window [4, 2, 2, 5] is invalid because 5 - 2 = 3, which is larger than

the limit. We remove the leftmost element (nums[j] which is 4) from sl, resulting in sl = [2, 2, 5], and increment j to 1.

 \circ We add 4 to s1 to get s1 = [2, 2, 4, 5]. The window [2, 2, 5, 4] is again invalid because 5 - 2 = 3 exceeds the limit.

We remove nums[j] (which is now the leftmost 2) from sl, making it sl = [2, 4, 5], and increment j to 2. The new window

[2, 5, 4] is valid and its size (3) is used to update ans if it's larger than the current ans. 7. Having iterated through all elements, we find that the longest subarray where the absolute difference between any two elements

Python Solution

class Solution:

10

11

12

19

20

22

23

24

25

26

27

28

29

30

31

32

33

38

8

9

10

11

12

13

14

15

16

17

18

19

21

22

23

24

25

26

27

28

29

30

31

32

33

34

36

35 };

8

9

10

11

16

18

20

19 }

}

5. Move on to i = 3 (v = 5):

6. Finally, for i = 4 (v = 4):

- does not exceed the limit is 3. Thus, we return ans = 3. This example shows how the sliding window moves through the array and adjusts by adding new elements and potentially removing the leftmost element to maintain a valid subarray within the limit. The SortedList makes it efficient to find the minimum and
- 13 14 # Iterate through the array with index and value 15 for window_end, value in enumerate(nums): # Add the current value to the sorted list 16 17 sorted_list.add(value) 18

The condition being if the absolute difference between the max and min values in the window exceeds the limit

public int longestSubarray(int[] nums, int limit) { // Create a TreeMap to keep track of the frequency of each number TreeMap<Integer, Integer> frequencyMap = new TreeMap<>(); int maxLength = 0; // Stores the maximum length of the subarray int left = 0; // The left pointer for our sliding window

Java Solution

class Solution {

34 # Example usage

35 # sol = Solution()

37 # print(result) # Output: 4

```
frequencyMap.remove(nums[left]);
20
21
22
                   // Move the left pointer to the right, shrinking the window
23
                   ++left;
24
25
26
               // Update the maximum length found so far
27
               maxLength = Math.max(maxLength, right - left + 1);
28
29
           // Return the maximum length of the subarray that satisfies the condition
           return maxLength;
30
31
32 }
33
C++ Solution
1 #include <vector>
2 #include <set>
   #include <algorithm>
  class Solution {
6 public:
       // Function to calculate the length of the longest subarray with the absolute difference
       // between any two elements not exceeding `limit`.
       int longestSubarray(vector<int>& nums, int limit) {
9
           // Initialize a multiset to maintain the elements in the current sliding window.
10
11
           multiset<int> window_elements;
12
           int longest_subarray_length = 0; // Variable to keep track of the max subarray length.
13
           int window_start = 0; // Starting index of the sliding window.
14
           // Iterate over the array using `i` as the end of the sliding window.
15
16
           for (int window_end = 0; window_end < nums.size(); ++window_end) {</pre>
17
               // Insert the current element into the multiset.
18
               window_elements.insert(nums[window_end]);
19
20
               // If the difference between the largest and smallest elements in the multiset
```

// exceeds the `limit`, shrink the window from the left until the condition is satisfied.

while (*window_elements.rbegin() - *window_elements.begin() > limit) {

int current_subarray_length = window_end - window_start + 1;

// Return the length of the longest subarray found.

return longest_subarray_length;

type CompareFunction<T> = (a: T, b: T) => number;

function getSize(node: ITreapNode<any> | null): number {

function getFac(node: ITreapNode<any> | null): number {

Typescript Solution

value: T;

interface ITreapNode<T> {

priority: number;

left: ITreapNode<T> | null;

12 let compareFn: CompareFunction<any>;

return node?.size ?? 0;

right: ITreapNode<T> | null;

count: number;

size: number;

let leftBound: any;

let rightBound: any;

let root: ITreapNode<any>;

// Erase the leftmost element from the multiset and shrink the window.

// Calculate the length of the current subarray and update the maximum length.

longest_subarray_length = max(longest_subarray_length, current_subarray_length);

window_elements.erase(window_elements.find(nums[window_start++]));

```
26
27
28
29
```

```
22
        return node?.priority ?? 0;
23
   }
24
   function createTreapNode<T>(value: T): ITreapNode<T> {
        const node: ITreapNode<T> = {
            value: value,
            count: 1,
            size: 1,
            priority: Math.random(),
30
31
            left: null,
32
            right: null,
33
       };
34
35
        return node;
36 }
38 function pushUp(node: ITreapNode<any>): void {
39
        let tmp = node.count;
40
        tmp += getSize(node.left);
        tmp += getSize(node.right);
41
42
       node.size = tmp;
43
44
   function rotateRight<T>(node: ITreapNode<T>): ITreapNode<T> {
46
       const left = node.left;
47
       node.left = left?.right ?? null;
       if (left) {
48
            left.right = node;
49
50
51
        if (node.right) pushUp(node.right);
52
        pushUp(node);
53
        return left ?? node;
54 }
55
   function rotateLeft<T>(node: ITreapNode<T>): ITreapNode<T> {
57
        const right = node.right;
58
       node.right = right?.left ?? null;
59
       if (right) {
60
            right.left = node;
61
62
       if (node.left) {
63
            pushUp(node.left);
64
65
        pushUp(node);
66
       return right ?? node;
67 }
68
   // ... (Other methods would be similarly defined as global functions, but not included here for brevity)
70
71 // Initialize the global Treap
72 function initTreap<T>(
        compFn: CompareFunction<T>,
73
        leftBnd: T = -Infinity as unknown as T,
74
        rightBnd: T = Infinity as unknown as T,
75
76
   ): void {
77
        compareFn = compFn as unknown as CompareFunction<any>;
78
        leftBound = leftBnd;
        rightBound = rightBnd;
79
80
        const treapRoot: ITreapNode<any> = createTreapNode<any>(rightBound);
        treapRoot.priority = Infinity;
81
```

Time Complexity

treapRoot.left = createTreapNode<any>(leftBound);

// Example usage of initializing and using the treap

addNode(root, 10); // This assumes the addNode function is implemented globally as mentioned above.

treapRoot.left.priority = -Infinity;

pushUp(treapRoot.left);

initTreap<number>((a, b) => a - b);

Time and Space Complexity

similar data structure used for keeping the list sorted.

pushUp(treapRoot);

root = treapRoot;

remove(nums[j]) is called, which is also an O(log n) operation because the list must be searched for the value to remove it, and it might need restructuring to keep it sorted.

82

83

84

85

86

88

92

87 }

The variable ans is updated using a max() function which is an O(1) operation. Since the for loop iterates over each element in nums once, and the inner while loop only processes each element once due to the

The provided code utilizes a sliding window approach within a loop and a SortedList to keep track of the order of elements. For each

element in nums, it is added to the SortedList, which is typically an O(log n) operation due to the underlying binary search tree or

The while loop inside the for loop is executed only when the current subarray does not meet the limit condition. Within the loop,

Space Complexity The additional space used by the algorithm consists of the SortedList and the variables used for iteration and storing the current longest subarray's length. The space complexity of the SortedList depends on the number of unique elements inserted. In the

sliding window mechanism, the overall time complexity is $O(n \log n)$, where n is the number of elements in the nums list.

0(n).

worst-case scenario, all elements of nums are different, and the SortedList will contain n elements, leading to a space complexity of The space for the other variables is comparatively negligible (0(1)), so the overall space complexity is 0(n).