Depth-First Search

Tree

Medium

The problem involves finding the diameter of an N-ary tree. The diameter is defined as the length (number of edges) of the longest

path between any two nodes within the tree. A key point to note is that this path is not required to go through the tree's root. Understanding how to handle N-ary trees, which are trees where a node can have any number of children, is crucial in solving this problem. Level order traversal serialization (where each group of children is separated by the null value) is used for input representation.

Intuition When finding the diameter of an N-ary tree, a depth-first search (DFS) approach can be applied effectively. The diameter of the tree could potentially be the distance between two nodes that are below the root of the tree, which means it might not include the root itself. Therefore, for each node, we need to consider the two longest paths that extend from this node down its children, because

these paths could potentially contribute to the maximum diameter if they are connected through the current node.

found. The key steps in our approach are: 1. Traverse each node with DFS, starting from the root. 2. For each node visited, calculate the lengths of the longest (max) and second-longest (smax) paths among its child nodes.

To implement this, we use a recursive DFS function that computes and returns the height of the current subtree while simultaneously

calculating the potential diameter that passes through the current node (the sum of the two longest child paths connected at the

node). We track the overall longest path seen so far with a nonlocal variable 'ans', which gets updated when a larger diameter is

4. At each node, return the height of this subtree to its parent, which is 1 plus the longest path length among the child nodes.

By doing this for each node, once the traversal is complete, 'ans' will hold the length of the longest path, which is the diameter of the

contribute to height or diameter.

height among the children.

path through the node).

N-ary tree.

The solution approach for calculating the diameter of an N-ary tree involves a depth-first search (DFS) algorithm. The DFS is

Solution Approach

3. Update the overall potential diameter 'ans' if the sum of max and smax is greater than the current 'ans'.

customized to perform two operations at once - calculate the height of the current subtree and use this information to evaluate the longest path that passes through each node (potential diameters). Here's a step-by-step breakdown of the implementation using the provided solution code:

1. We define a DFS helper function dfs that takes a node of the tree as an argument. This function returns the height of the tree rooted at the given node.

2. Inside the dfs function: • We start by checking if the current root node is None (base case). If it is, we return 0 since a non-existent node does not

∘ We then introduce two variables m1 and m2 initialized to 0, which will store the lengths of the longest (m1) and secondlongest (m2) paths found within the children of the current node.

edge connecting to its parent and m1 as the height of its longest subtree).

search pattern efficiently to solve the problem of finding the tree's diameter.

For child node 3, the recursive call to dfs returns 0 (as it has no child).

5. The DFS continues recursively for all nodes but finds no longer path than the current ans.

7. The diameter method then returns ans, which is 3, as the diameter of the tree.

self.children = children if children is not None else []

base case: if the current node is None, return 0

iterate over all the children of the current node

elif path_length > second_longest_path:

Initialize max_diameter to 0 before starting DFS

Call the dfs function starting from the root node

Once DFS is complete, return the max_diameter found

second_longest_path = path_length

longest_path = second_longest_path = 0

initialize the two longest paths from the current node to 0

recursively find the longest path for each child

max_diameter = max(max_diameter, longest_path + second_longest_path)

5. Finally, the diameter method returns the ans variable, which now contains the diameter of the tree.

• We loop through each child of the current node and recursively call dfs(child). The returned value t is the height of the subtree rooted at child. We use this value to potentially update m1 and m2.

■ If t is greater than m1, we set m2 to m1 and then m1 to t, thus keeping m1 as the max height and m2 as the second max

with the maximum of its current value and the sum of m1 + m2, representing the potential diameter at this node (the longest

- If t is not greater than m1 but is greater than m2, we update m2 with t. After considering all children, we update the ans variable (which is kept in the outer scope of dfs and declared as nonlocal)
- 4. In the diameter method of the Solution class, we initialize the ans variable to 0 (to globally keep track of the longest path seen) and call dfs(root) to kick off the recursive DFS from the root of the tree.

3. The dfs function concludes by returning 1 + m1, which represents the height of the subtree rooted at the current node (1 for the

Example Walkthrough Let's illustrate the solution approach with a small example. Suppose we have an N-ary tree represented using level order serialization

The use of recursion to traverse the tree, combined with updating the two longest paths at each node, harnesses the depth-first

2 / | \ 3 2 3 4 4 / | \ 5 5 6 7

We want to find the diameter of this tree, which is the longest path between any two nodes. We will walk through the DFS approach using the provided solution template:

2. When we process the root (node 1), we initialize m1 and m2 both to 0. These variables will track the longest and second-longest

For child node 2, the recursive call to dfs returns 1 (since node 5 is its only child and adds one edge to the height).

∘ For child node 4, the DFS must consider both children. Nodes 6 and 7 contribute a height of 1 each, and since these are the

longest and second-longest heights for this subtree, after considering both, node 4 returns 2 as the height (1 for the longest

1. We start the DFS with the root, which is node 1. We initialize ans to 0. The dfs function will traverse the tree and track the

3. We recursively call dfs for each child of the root.

longest path at each node.

heights among the children of each node.

child plus 1 for the edge to node 4).

without necessarily passing through the root.

def dfs(node: 'Node') -> int:

nonlocal max_diameter

for child in node.children:

return 1 + longest_path

max_diameter = 0

dfs(root)

if node is None:

return 0

self.val = val

as follows:

becomes 1 (from node 2). Therefore, the potential diameter passing through root at this stage is m1 + m2 = 2 + 1 = 3. We update ans to 3.

4. Now, with the heights from the children of the root, we update m1 and m2 for the root. m1 becomes 2 (from node 4), and m2

6. After the DFS is complete and all nodes have been visited, ans holds the diameter of the tree. In this case, it remains 3, representing the length of the longest path, which happens to be from node 5 to node 7 via the root $(5 \rightarrow 2 \rightarrow 1 \rightarrow 4 \rightarrow 6 \text{ or } 7)$.

This walkthrough illustrates the algorithm's efficiency in computing the diameter by determining the longest path through each node

6 class Solution: def diameter(self, root: 'Node') -> int: # Helper function to perform depth-first search

21 path_length = dfs(child) 22 # check if the current path is longer than the longest recorded path 23 if path_length > longest_path: 24 # update the second longest and longest paths accordingly 25 second_longest_path, longest_path = longest_path, path_length 26 # else if it's only longer than the second longest, update the second longest

accessing the non-local variable 'max_diameter' to update its value within this helper function

update the maximum diameter encountered so far based on the two longest paths from this node

return the longer path increased by 1 for the edge between this node and its parent

```
Python Solution
   # Definition for a Node.
   class Node:
      def __init__(self, val=None, children=None):
```

9

10

11

12

13

14

15

16

17

18

19

20

27

28

29

30

31

32

33

34

35

36

37

38

43

44

45

46

47

48

49

50

51

52

53

54

55

56

57

58

59

60

61

62

63

64

65

66

67

/**

* @param root the current node being traversed

int maxDepth1 = 0; // Tracks the longest path

int depth = depthFirstSearch(child);

private int depthFirstSearch(Node root) {

for (Node child : root.children) {

maxDepth2 = maxDepth1;

} else if (depth > maxDepth2) {

maxDepth1 = depth;

maxDepth2 = depth;

if (depth > maxDepth1)

// Leaf nodes have a depth of 0

if (root == null) {

return 0;

* @return the maximum depth emanating from the current node

int maxDepth2 = 0; // Tracks the second longest path

// Recursively obtain the depth for children nodes to find the longest paths

// Check and set the longest and second—longest distances found

```
39
            return max_diameter
40
Java Solution
   /**
    * Definition for a N-ary tree node.
    */
   class Node {
       public int val;
       public List<Node> children;
       // Constructor initializes value and empty children list
8
       public Node() {
           children = new ArrayList<Node>();
10
12
13
       // Constructor initializes node with a value and empty children list
       public Node(int val) {
14
           this.val = val;
15
           children = new ArrayList<Node>();
16
18
       // Constructor initializes node with a value and given children list
19
       public Node(int val, ArrayList<Node> children) {
20
           this.val = val;
21
           this.children = children;
22
23
24
25
   public class Solution {
27
       // This class variable keeps track of the diameter of the tree
       private int diameter;
28
29
30
       /**
31
        * Computes the diameter of an N-ary tree.
32
        * The diameter of an N-ary tree is the length of the longest path between any two nodes in a tree.
33
34
        * @param root the root node of the tree
        * @return the diameter of the tree
35
36
37
       public int diameter(Node root) {
38
           diameter = 0;
           depthFirstSearch(root);
39
           return diameter;
40
41
42
```

* Uses Depth First Search (DFS) to find the length of the longest path through the N-ary tree from the current node.

68 69

```
70
71
           // Update the diameter if the sum of two longest paths through the current node is greater than the current diameter
72
           diameter = Math.max(diameter, maxDepth1 + maxDepth2);
73
           // Return the maximum depth plus one for the current path
74
75
           return 1 + maxDepth1;
76
77 }
78
C++ Solution
  1 // Definition for a Node is provided as per the question context
  2 class Node {
    public:
                                      // The value contained within the node.
         int val;
         vector<Node*> children;
                                      // A vector containing pointers to its children.
        Node() {}
  8
        Node(int _val) {
  9
 10
             val = _val;
 11
 12
 13
        Node(int _val, vector<Node*> _children) {
 14
             val = _val;
 15
             children = _children;
 16
 17 };
 18
 19 class Solution {
 20
    public:
 21
         int maxDiameter; // Class variable to store the maximum diameter found.
 22
 23
         // Public method which is the starting point for finding the diameter of the tree.
         int diameter(Node* root) {
 24
             maxDiameter = 0; // Initialize the max diameter to 0.
 25
 26
                           // Call the depth-first search helper method.
             dfs(root);
             return maxDiameter; // Return the maximum diameter calculated.
 27
 28
 29
 30
         // Private helper method for DFS traversal which calculates the depths and updates the diameter.
 31
        // It returns the maximum depth from the current node to its furthest leaf node.
 32
         int dfs(Node* root) {
 33
             if (!root) return 0; // Base case: If the node is null, return 0 (no depth).
 34
             int maxDepth1 = 0; // To store the maximum length of the paths in the children.
 35
             int maxDepth2 = 0; // To store the second maximum length of the paths in the children.
 37
 38
             // Iterate through each child node of the current root.
             for (Node* child : root->children) {
 39
                 int depth = dfs(child); // Recursive call to get the depth for each child.
 40
                 // Update the two maximum depths found among children nodes
 41
 42
                 if (depth > maxDepth1) {
 43
                     maxDepth2 = maxDepth1; // Update the second max if the new max is found
 44
                                          // Update the new max depth
                     maxDepth1 = depth;
 45
                 } else if (depth > maxDepth2) {
 46
                     maxDepth2 = depth; // Update the second max if greater than it but less than the max depth
 47
 48
 49
 50
             // Update the maximum diameter if the sum of the two largest depths is greater than the current diameter.
 51
             maxDiameter = max(maxDiameter, maxDepth1 + maxDepth2);
 52
 53
             // Return the maximum depth of this subtree to its parent caller, which is 1 plus the max depth of its children.
 54
             return 1 + maxDepth1;
 55
 56 };
 57
Typescript Solution
1 // Define the type for a Node that includes a value and an array of Node references as children
2 type Node = {
     val: number;
     children: Node[];
5 };
```

23 24 // Iterate through each child node of the current root. for (let child of root.children) { 25 let depth = dfs(child); // Recursive call to get the depth for each child. 26 27 // Update the two maximum depths found among children nodes if (depth > maxDepthOne) { 28

return 1 + maxDepthOne;

Time and Space Complexity

11

15

20

21

29

30

33

34

35

36

37

38

39

40

42

41 }

14 }

let maxDiameter: number; // Global variable to store the maximum diameter found.

// This function is the entry point for finding the diameter of the n-ary tree.

16 // Helper function for DFS traversal that calculates depths and updates the diameter.

// It returns the maximum distance from the current node to its furthest leaf node.

maxDepthTwo = maxDepthOne; // Second max updated if a new max is found

if (root === null) return 0; // Base case: If the node is null, return 0 (no depth).

let maxDepthOne = 0; // To store the maximum depth among the paths in the children.

let maxDepthTwo = 0; // To store the second maximum depth among the paths in the children.

maxDepthTwo = depth; // Update the second max if it's greater than the current second max

// Update the maximum diameter if the sum of the two largest depths is greater than the current maximum diameter.

// Return the maximum depth of this subtree to its parent, which is 1 plus the maximum depth of its children.

function diameter(root: Node | null): number {

function dfs(root: Node | null): number {

} else if (depth > maxDepthTwo) {

maxDiameter = 0; // Initialize the max diameter to 0.

dfs(root); // Call the depth-first search helper function.

maxDepthOne = depth; // Set as the new max depth

maxDiameter = Math.max(maxDiameter, maxDepthOne + maxDepthTwo);

return maxDiameter; // Return the maximum diameter calculated.

The code provided defines a function dfs(root) that is used to compute the diameter of an N-ary tree. The diameter of an N-ary tree is defined as the length of the longest path between any two nodes in the tree. **Time Complexity:**

function is called recursively for each node exactly once. Within each call to dfs(), it performs a constant amount of work for each

child. Since the sum of the sizes of all children's lists across all nodes is N - 1 (there are N - 1 edges in a tree with N nodes), the

The time complexity of the provided code is O(N), where N is the total number of nodes in the tree. This is because the dfs()

overall time complexity is linear relative to the number of nodes in the tree. **Space Complexity:**

The space complexity of the code can be analyzed in two parts: the space required for the recursion call stack and the space required for the dfs() function execution.

- The worst-case space complexity for the recursion call stack is O(H), where H is the height of the tree. In the worst case, the tree can be skewed, resembling a linked list, thus the height can become N, leading to a worst-case space complexity of O(N). The additional space used by the dfs() function is constant, as it only uses a few integer variables (m1, m2, and t).
- Considering both parts, the total space complexity is O(H), which is O(N) in the worst case (when the tree is a linear chain), but can be better (such as $O(\log N)$) if the tree is balanced.