1570. Dot Product of Two Sparse Vectors

Medium Design Hash Table Two Pointers Array

The task is to create a class, SparseVector, which encapsulates the concept of a sparse vector and implements a method to

out the zero values.

Problem Description

should not be stored in a typical dense format as that would waste space. The goal is to store such vectors efficiently and perform operations on them.

compute the dot product of two sparse vectors. A sparse vector is defined as a vector that contains mostly zeros and therefore

To efficiently store the sparse vector, we can use a dictionary to hold only the non-zero elements, where the keys are the indices of these elements, and the values are the elements themselves. This way we aren't storing all the zero values, which can dramatically reduce memory usage for very sparse vectors. The SparseVector class has two methods:

1. SparseVector(nums): The constructor takes a list of integers nums and initializes the sparse vector using a dictionary comprehension that filters

another sparse vector, vec.

The dot product of two vectors is computed by multiplying corresponding entries and summing those products. In the case of sparse vectors, most of these products will be zero, because they involve multiplications by zero, so we only need to consider the

2. dotProduct(vec): This method computes the dot product between the vector represented by the current instance of SparseVector and

non-zero entries. The follow-up question asks about efficiently computing the dot product if only one of the two vectors is sparse.

Intuition For the solution, the idea is to leverage the sparsity of the vectors to optimize the dot product computation. Given that most of the elements in the vectors are zeros, we want to perform multiplications only for the non-zero elements. By converting the

vectors into a dictionary structure with non-zero elements, we can quickly identify which elements actually need to be multiplied.

dictionary comprehension for building a structure to represent the sparse vectors.

vector and only iterate over the potential non-zero counterparts.

we ignore all zero-product cases which would contribute nothing to the final sum.

In the dotProduct method, we iterate over the items in the smaller dictionary (to optimize the number of iterations) and multiply values by the corresponding values in the other dictionary, if they exist. If an index does not exist in the other dictionary, it means

get method to handle such cases, which allows us to specify a default value of 0 when an index is not found. When dealing with one sparse and one non-sparse vector, the current approach still works efficiently because the dot product will focus on iterating over the non-zero elements of the sparse vector and lookup the corresponding values in the non-sparse vector. **Solution Approach**

that the value for that index in the other vector is zero and thus does not contribute to the dot product. Therefore, we use the

Algorithm and Data Structure Dictionary for Sparse Representation: A Python dictionary is an ideal data structure for representing a sparse vector. It

non-zero element itself. This structure is memory efficient since we only store entries for non-zero elements.

• The resulting dictionary self.d holds only the elements of the input list that are non-zero, along with their indices.

allows storing key-value pairs where the key is the index of a non-zero element in the original vector, and the value is the

The implementation of the SparseVector solution makes use of two main aspects: Python dictionaries and the concept of

Constructor __init__:

dotProduct Method:

not in the dictionary would be zero.

 We use a dictionary comprehension in the constructor to iterate over nums using enumerate to get both the index i and the value v together. • The condition if v ensures that we're only storing the non-zero values (as zero is considered False in Python).

- non-zero entries. By comparison of the lengths of these two dictionaries, we choose to iterate over the smaller dictionary (here represented as a). This is an optimization step; since the dot product will be zero for all indices not present in both vectors, we can skip the zero values of the larger
 - We then use a generator expression to iterate over the items of a: for i, v in a.items(). • For each element, we calculate the product v * b.get(i, 0). The get method of the dictionary is very handy in this case, as it will return 1 if the index i doesn't exist in b—a common occurrence with sparse vectors, and also safe considering the default value for any index

• We receive another SparseVector object, vec, and we want to compute the dot product with the current sparse vector instance.

• We access the internal dictionaries of the current instance (a) and the vec (b) — these dictionaries store the indices and values of the

- Finally, the sum function accumulates all the products to give us the result of the dot product. **Pattern Used**
 - Iterating Over a Smaller Set: Choosing to iterate over the smaller set of elements to reduce the number of operations is a common optimization strategy in algorithms involving collection processing.

Combining these algorithms and patterns, the SparseVector class efficiently implements the computation of a dot product

between two sparse vectors, considering only the meaningful, non-zero elements, and thus avoiding unnecessary computations.

By using a compressed representation for the vectors with dictionaries and strategically leveraging the sparsity, the

sparseB. When initializing these objects, our dictionary comprehension will filter out the zeroes and store only the non-zero

• sparseA's internal dictionary will have the elements {0: 1, 3: 2}, corresponding to indices and values (index 0 has value 1, and index 3 has value

find the dot product of sparseA and sparseB, we invoke the dotProduct method on one of them, let's say

two elements, we can choose either one to iterate over, but for the sake of this example, we will iterate over sparseA because

• Efficient Computation with Sparse Representation: By representing the vectors in a sparse form, computations are made more efficient since

Example Walkthrough

implementation maximizes efficiency in terms of both time and space complexity.

Let's walk through a small example to illustrate the solution approach.

Suppose we have two sparse vectors represented as follows:

Using the given solution approach, we first convert these lists into SparseVector objects. Let's call these objects sparseA and

2).

Vector A: [1, 0, 0, 2, 0]

Vector B: [0, 3, 0, 4, 0]

elements and their indices:

sparseA.dotProduct(sparseB).

Here's a step-by-step explanation:

we called the method on it.

computations.

from typing import List

class SparseVector:

Example usage:

v1 = SparseVector(nums1)

v2 = SparseVector(nums2)

ans = v1.dotProduct(v2)

class SparseVector {

- Inside the dotProduct method, we compare the sizes of the internal dictionaries of sparseA and sparseB. Since both have
- We now loop over the items in sparseA's dictionary. For each item, we look up whether the corresponding index is in

sparseA) * 4 (from sparseB) which equals 8.

init (self, nums: List[int]):

Return the dot product of two sparse vectors

def dotProduct(self, vec: "SparseVector") -> int:

Store the non-zero elements with their indices as keys

For efficiency, iterate through the smaller vector

overlapping elements of the two sparse vectors

else (vec.non_zero_elements, self.non_zero_elements)

Calculate the dot product by summing the product of the

self.non_zero_elements = {i: v for i, v in enumerate(nums) if v != 0}

if len(self.non zero elements) < len(vec.non zero elements) \</pre>

// Using a HashMap to efficiently store non-zero elements and their positions

// Constructor to populate the map with non-zero elements from the input array

// Reference to the smaller of the two maps to iterate over for efficiency

private Map<Integer, Integer> nonZeroElements = new HashMap<>();

for (int i = 0; i < nums.length; ++i) {</pre>

// Return the dotProduct of two sparse vectors

public int dotProduct(SparseVector vec) {

smallerMap = largerMap;

for (var entry : smallerMap.entrySet()) {

int index = entry.getKey();

// SparseVector v1 = new SparseVector(nums1);

// int product = v1.dotProduct(v2);

const sparseVectorOperations = {

// Object to encapsulate sparse vector data and operations

sparseVectorData: new Map<number, Map<number, number>>(),

for (let index = 0; index < nums.length; ++index) {</pre>

sparseData.set(index, nums[index]);

this.sparseVectorData.set(vectorId, sparseData);

const sparseData = new Map<number, number>();

if (nums[index] !== 0) {

let smallerMap = vec1Data;

smallerMap = vec2Data;

largerMap = vec1Data;

if (vec1Data.size > vec2Data.size) {

if (largerMap.has(index)) {

for (const [index, value] of smallerMap) {

sparseVectorOperations.createSparseVector([1, 0, 0, 2, 3], 1);

sparseVectorOperations.createSparseVector([0, 3, 0, 4, 0], 2);

let largerMap = vec2Data;

let result = 0;

// Function to initialize a sparse vector from a given array of numbers

// Function to calculate the dot product of two sparse vectors identified by their IDs

calculateDotProduct: function(vectorId1: number, vectorId2: number): number {

// Swap if vec1Data has more elements than vec2Data to minimize iterations

// Calculate dot product by iterating through the map with fewer elements

const dotProductResult = sparseVectorOperations.calculateDotProduct(1, 2); // Should calculate the dot product

result += value * (largerMap.get(index) || 0);

console.log(dotProductResult); // Outputs the result of the dot product calculation

const vec1Data = this.sparseVectorData.get(vectorId1) || new Map();

const vec2Data = this.sparseVectorData.get(vectorId2) || new Map();

createSparseVector: function(nums: number[], vectorId: number): void {

TypeScript

},

int value = entry.getValue();

largerMap = temp;

nonZeroElements.put(i, nums[i]);

Map<Integer, Integer> smallerMap = nonZeroElements;

int productSum = 0; // The result of the dot product operation

productSum += value * largerMap.getOrDefault(index, 0);

return productSum; // Return the computed dot product

// Iterating through the smaller map and multiplying the matching values

smaller vector, larger vector = (self.non zero elements, vec.non zero_elements) \

return sum(value * larger_vector.get(index, 0) for index, value in smaller_vector.items())

sparseB's dictionary. We have two iterations in our loop:

• sparseB's internal dictionary will have the elements {1: 3, 3: 4}.

sum and thus the dot product of A and B is 8. In conclusion, the SparseVector class successfully computes the dot product of sparseA and sparseB as 8 using an efficient approach, taking advantage of the sparse representation by only considering non-zero elements and their indices in the

• First iteration: For index 0 in sparseA, there is no corresponding index in sparseB. When we attempt to multiply 1 (from sparseA) with

• Second iteration: For index 3 in sparseA, there is a matching index in sparseB which has the value 4. We perform the multiplication 2 (from

We sum the results of the multiplications. In this example, the only non-zero result came from the second iteration (8), so the

sparseB.get(0, 0) (using .get to specify a default value of 0), the result is 0 since index 0 is not present in sparseB.

- Solution Implementation **Python**
- Java import java.util.HashMap; import java.util.Map;

Map<Integer, Integer> largerMap = vec.nonZeroElements; // Swap if 'vec's map has fewer elements to iterate over the smaller map if (largerMap.size() < smallerMap.size()) {</pre> Map<Integer. Integer> temp = smallerMap;

// Example of usage:

SparseVector(int[] nums) {

if (nums[i] != 0) {

```
// SparseVector v2 = new SparseVector(nums2);
// int ans = v1.dotProduct(v2);
C++
#include <vector>
#include <unordered map>
using namespace std;
// Class to represent a Sparse Vector
class SparseVector {
public:
    // This map will store the non-zero values of the sparse vector associated with their indices
    unordered_map<int, int> indexToValueMap;
    // Constructor which takes a vector of integers and populates the indexToValueMap
    SparseVector(vector<int>& nums) {
        for (int i = 0; i < nums.size(); ++i) {</pre>
            if (nums[i] != 0) {
                indexToValueMap[i] = nums[i];
    // Return the dot product of this sparse vector with another sparse vector vec
    int dotProduct(SparseVector& vec) {
        // Using references to the internal maps for easier access
        auto& thisVectorMap = indexToValueMap;
        auto& otherVectorMap = vec.indexToValueMap;
        // Optimize by iterating over the smaller map
        if (thisVectorMap.size() > otherVectorMap.size()) {
            swap(thisVectorMap, otherVectorMap);
        int total = 0;
        // Compute dot product by only considering non-zero elements
        for (auto& [index, value] : thisVectorMap) {
            if (otherVectorMap.count(index)) {
                total += value * otherVectorMap[index];
        return total;
};
// Usage example:
// vector<int> nums1 = { ... };
// vector<int> nums2 = { ... };
// SparseVector v1(nums1):
// SparseVector v2(nums2);
```

return result; // Example usage:

from typing import List

class SparseVector: def init (self, nums: List[int]): # Store the non-zero elements with their indices as keys self.non_zero_elements = {i: v for i, v in enumerate(nums) if v != 0} # Return the dot product of two sparse vectors def dotProduct(self, vec: "SparseVector") -> int: # For efficiency, iterate through the smaller vector smaller vector, larger vector = (self.non zero elements, vec.non zero elements) \ if len(self.non zero elements) < len(vec.non zero elements) \</pre> else (vec.non_zero_elements, self.non_zero_elements) # Calculate the dot product by summing the product of the # overlapping elements of the two sparse vectors return sum(value * larger_vector.get(index, 0) for index, value in smaller_vector.items()) # Example usage: # v1 = SparseVector(nums1) # v2 = SparseVector(nums2) # ans = v1.dotProduct(v2)Time and Space Complexity **Time Complexity:** The constructor $__{init}$ has a time complexity of O(n) where n is the number of elements in nums, as it needs to iterate through all elements to create the dictionary with non-zero values.

The dotProduct function has a time complexity of O(min(k, l)) where k and l are the number of non-zero elements in the two SparseVectors, respectively. This is because the function iterates over the smaller of the two dictionaries (after ensuring a has the smaller length, swapping if necessary) and attempts to find matching elements in the larger one. The get operation on a

dictionary has an average case time complexity of 0(1).

Space Complexity: The space complexity of the $_init_$ function is 0(k), where k is the number of non-zero elements in nums, since the space required depends on the stored non-zero elements.

The dotProduct function operates in 0(1) space complexity because it calculates the sum on the fly and does not store intermediate results or allocate additional space based on input size, other than a few variables for iteration and summing.