1714. Sum Of Special Evenly-Spaced Elements In Array

Problem Description

Array

Hard

Dynamic Programming

The problem requires us to answer a set of queries on an array of non-negative integers, nums. Each query is represented by a pair $[x_i, y_i]$, where x_i is the starting index to consider in the array and y_i is the stride for selecting elements. Specifically, for each query, we want to find the sum of every nums [j] such that j is greater than or equal to x_i and the difference $(j - x_i)$ is divisible by y_i . Importantly, the answers to the queries should be returned modulo $10^9 + 7$.

Intuition The intuition behind the solution is based on the observation that direct computation of every query could be inefficient,

especially when there are many queries or the array is very large. Thus, we need to optimize the process. For each y_i that is small (specifically, not larger than the square root of the size of nums), we can precompute a suffix sum array suf that helps answer the queries quickly. This exploits the fact that with smaller strides, there's more repetition and structure to use to our advantage.

This approach balances between precomputation for frequent, structured queries, and on-the-fly computation for less structured and less frequently occurring query patterns.

sparsity of the elements we'd be summing. Therefore, for such strides, it's better to compute the sum directly per query.

For larger strides (y_i greater than the square root of the size of nums), it might be less efficient to use precomputation due to the

Solution Approach

First, we identify the threshold for small strides, which is the square root of n, the length of nums. We prepare a 2D array suf

smaller strides (small y_i) and larger strides (large y_i). Precomputation for smaller strides

where each row corresponds to a different stride length upto the threshold. The idea is to calculate suffix sums starting from

The solution approach uses a combination of precomputation and direct calculation to handle the two scenarios efficiently:

each index j. For stride i, suf[i][j] represents the sum of elements nums[j], nums[j+i], nums[j+2i], and so on till the end of the array.

The precomputation occurs like this: 1. Iterate over each stride length i from 1 to m, where m is the square root of n. 2. For each stride length i, iterate backwards through the nums array starting from the last index.

3. Calculate the sum incrementally, adding nums[j] to the next value in the prefix already computed which is suf[i][min(n, j + i)]. This is

because the next element in the sequence we're summing would be j+i elements ahead considering the stride. This process essentially fills up the suffix sum array with all the necessary sums for smaller strides.

1. Starting at index x, generate a range of indices by slicing the nums array from x to the end of the array with step y. This selects every yth

Direct calculation for larger strides

element starting at x.

2. Sum the selected elements.

By combining these two approaches, we obtain an efficient method to answer all types of queries. The algorithm only uses precomputation for scenarios where it significantly reduces complexity, and falls back to direct summation when precomputation

3. Apply modulo 10^9 + 7 on the sum to avoid integer overflow and to conform with the problem requirements.

For each query with stride y > m, the direct calculation takes place as follows:

• Each answer is then appended to the ans list, after applying the modulo operation.

would not offer a speedup. The main algorithm consists of: Precomputing the suf array for smaller strides. Iterating over each query.

• Checking if the stride y of the query is smaller or equal to m. If it is, use the precomputed suf value to answer the query. If it's larger, directly

This hybrid approach is particularly efficient for handling a mix of query types on potentially large datasets, as it minimizes unnecessary computation while making use of precomputation wherever beneficial.

Example Walkthrough

• Finally, return the ans list as the result.

calculate the sum using slicing on the nums array.

Let's consider m to be the threshold for small strides, which is the square root of the length of nums. Here, n = 11, so m = sqrt(11) \approx 3. Thus, any stride y_i <= 3 will involve precomputation.

For stride 1: suf[1][...] = [35, 32, 31, 30, 29, 24, 22, 20, 14, 8, 5]

Precompute the suffix sum array suf for the strides less than or equal to m.

Let's walk through an example to illustrate the solution approach.

Suppose we have the following nums array and queries:

nums = [3, 1, 4, 1, 5, 9, 2, 6, 5, 3, 5]

queries = [[1, 2], [0, 3], [2, 1]]

For stride 2: suf[2][...] = [36, 1, 25, 4, 21, 9, 11, 6, 8, 3, 5]

For stride 3: suf[3][...] = [22, 1, 6, 1, 14, 9, 2, 6, 5, 3, 5]

```
We now have precomputed sums for all smaller strides.
```

Answering Queries

• Final Answer List: [1 mod (10^9 + 7), 22 mod (10^9 + 7), 31 mod (10^9 + 7)] or [1, 22, 31] since all values are already less than 10^9 + 7.

This completes the example, which demonstrates how the algorithm efficiently answers queries by using a combination of

Query [1, 2]: Since the stride is 2 (which is less than or equal to m), we use the precomputed suf array.

Query [0, 3]: Since the stride is 3 (which is less than or equal to m), we also use the precomputed suf array.

Answer: suf[2][1] = 1 (which is sum of nums[1], nums[1+2], nums[1+4], etc., up to the end of array).

Answer: suf[3][0] = 22 (which is sum of nums[0], nums[0+3], nums[0+6], etc., up to the end of array).

Query [2, 1]: As the stride is 1, we refer again to the precomputed suf array.

Finally, for each answer, we apply the modulo 10^9 + 7 operation.

precomputation for smaller strides and direct computation for larger ones.

def solve(self, nums: List[int], queries: List[List[int]]) -> List[int]:

Fill the prefix sums matrix for all blocks with size up to sqrt(n)

result = prefix_sums[interval][start] % MOD

result = sum(nums[start::interval]) % MOD

return results # Return the results list for all queries

int squareRootOfNumLength = (int) Math.sqrt(numLength);

n = len(nums) # The total number of elements in nums

for block_size in range(1, sqrt_n + 1):

for start, interval in queries:

if interval <= sqrt_n:</pre>

results.append(result)

int numLength = nums.length;

final int mod = (int) 1e9 + 7;

public int[] solve(int[] nums, int[][] queries) {

// Define the modulo value to avoid overflow

// Create a 2D array for storing suffix sums

else:

○ Answer: suf[1][2] = 31 (which is sum of nums[2], nums[3], nums[4], ..., until the end of the array).

Python

class Solution:

Initial Setup

- Solution Implementation
 - from typing import List from math import sqrt

MOD = 10**9 + 7 # Define the modulus for result as per problem statement to avoid large integers

sqrt_n = int(sqrt(n)) # The square root of the length of nums, which determines the threshold

prefix_sums = [[0] * (n + 1) for _ in range(sqrt_n + 1)] # Initialize the prefix sums matrix

If the interval is under the square root threshold, use precomputed sums

// Calculate the length of the nums array and the square root of that length

int[][] suffixSums = new int[squareRootOfNumLength + 1][numLength + 1];

const int mod = 1e9 + 7; // The modulo value to prevent integer overflow.

suffix[i][j] = (suffix[i][std::min(numsSize, j + i)] + nums[j]) % mod;

int start = query[0], step = query[1]; // Start index and step for the current query.

// If the step is less than or equal to the block size, use the precomputed suffix sums.

int suffix[blockSize + 1][numsSize + 1]; // Suffix sums matrix.

std::vector<int> ans; // Vector to store the answers to the queries.

// Otherwise, perform a brute-force sum calculation.

for (let startIndex = arrayLength - 1; startIndex >= 0; --startIndex) {

let nextIndex = Math.min(arrayLength, startIndex + blockSize);

for (int i = start; i < numsSize; i += step) {</pre>

// Iterate over each query and calculate the sum accordingly.

// Initialize the suffix sums matrix with zeros.

for (int $j = numsSize - 1; j >= 0; --j) {$

ans.push_back(suffix[step][start]);

// Pre-compute the suffix sums for all possible blocks.

memset(suffix, 0, sizeof(suffix));

for (auto& query : queries) {

if (step <= blockSize) {</pre>

int sum = 0;

// Array to hold answer for each query

if (stepSize <= sqrtLength) {</pre>

for (const [startIndex, stepSize] of queries) {

answers.push(suffixSums[stepSize][startIndex]);

sum = (sum + nums[i]) % modulus;

next_index = min(n, i + block_size)

for start, interval in queries:

if interval <= sqrt_n:</pre>

results.append(result)

Time and Space Complexity

else:

Space Complexity

results = [] # This will store the results of each query

result = prefix_sums[interval][start] % MOD

result = sum(nums[start::interval]) % MOD

return results # Return the results list for all queries

For each query in queries, there are two cases to consider:

// For larger step sizes, calculate the sum manually

return answers; // Return the array containing sums for each query

def solve(self, nums: List[int], queries: List[List[int]]) -> List[int]:

for (let i = startIndex; i < arrayLength; i += stepSize) {</pre>

answers.push(sum); // Append the sum to the answers array

const answers: number[] = [];

let sum = 0;

// Process each query

} else {

from typing import List

from math import sqrt

class Solution:

} else {

for (int i = 1; i <= blockSize; ++i) {</pre>

For intervals larger than square root of n, calculate the sum on the fly

for i in range(n - 1, -1, -1): # Start from the end to compute prefix sums next_index = min(n, i + block_size) prefix_sums[block_size][i] = prefix_sums[block_size][next_index] + nums[i] results = [] # This will store the results of each query

```
Java
```

class Solution {

```
// Calculate suffix sums for blocks with size up to the square root of numLength
        for (int i = 1; i <= squareRootOfNumLength; ++i) {</pre>
            for (int j = numLength - 1; j >= 0; --j) {
                suffixSums[i][j] = (suffixSums[i][Math.min(numLength, j + i)] + nums[j]) % mod;
        // Get the number of queries and initialize an array to store the answers
        int queryCount = queries.length;
        int[] answers = new int[queryCount];
       // Process each query
        for (int i = 0; i < queryCount; ++i) {</pre>
            int startIndex = queries[i][0];
            int stepSize = queries[i][1];
            // If the step size is within the computed suffix sums, use the precomputed value
            if (stepSize <= squareRootOfNumLength) {</pre>
                answers[i] = suffixSums[stepSize][startIndex];
            } else {
                // If the step size is larger, calculate the sum on the fly
                int sum = 0;
                for (int j = startIndex; j < numLength; j += stepSize) {</pre>
                    sum = (sum + nums[j]) % mod;
                answers[i] = sum;
       // Return the array of answers
        return answers;
C++
#include <vector>
#include <cstring>
#include <cmath>
class Solution {
public:
    std::vector<int> solve(std::vector<int>& nums, std::vector<std::vector<int>>& queries) {
        int numsSize = nums.size();
        int blockSize = static_cast<int>(sqrt(numsSize)); // The size of each block for the sqrt decomposition.
```

```
sum = (sum + nums[i]) % mod;
                ans.push_back(sum); // Add the computed sum to the answers vector.
        return ans; // Return the vector of answers.
};
TypeScript
// Defining a function to solve the query based on the array of numbers and list of queries
function solve(nums: number[], queries: number[][]): number[] {
    const arrayLength: number = nums.length;
                                                               // Length of the nums array
    const sqrtLength: number = Math.floor(Math.sqrt(arrayLength)); // Sqrt decomposition length
    const modulus: number = 1e9 + 7;
                                                              // Define modulus for large numbers
    // Suffix arrays to store pre-computed sums. These are used to answer queries efficiently for small y values
    const suffixSums: number[][] = Array(sqrtLength + 1)
        .fill(0)
        .map(() => Array(arrayLength + 1).fill(0));
    // Pre-compute the suffix sums for each possible block size up to the square root of the length of the array
    for (let blockSize = 1; blockSize <= sqrtLength; ++blockSize) {</pre>
```

suffixSums[blockSize][startIndex] = (suffixSums[blockSize][nextIndex] + nums[startIndex]) % modulus;

// If stepSize is within the pre-computed range, use pre-computed sum for efficiency

MOD = 10**9 + 7 # Define the modulus for result as per problem statement to avoid large integers

prefix sums[block_size][i] = prefix_sums[block_size][next_index] + nums[i]

For intervals larger than square root of n, calculate the sum on the fly

If the interval is under the square root threshold, use precomputed sums

```
n = len(nums) # The total number of elements in nums
sqrt_n = int(sqrt(n)) # The square root of the length of nums, which determines the threshold
prefix_sums = [[0] * (n + 1) for _ in range(sqrt_n + 1)] # Initialize the prefix sums matrix
# Fill the prefix sums matrix for all blocks with size up to sqrt(n)
for block_size in range(1, sqrt_n + 1):
    for i in range(n - 1, -1, -1): # Start from the end to compute prefix sums
```

Time Complexity The time complexity of the precomputation step (building the suf array) is 0(m * n), where m is int(sqrt(n)) and n is the length of the input array nums. This is because the outer loop runs m times and the inner loop runs n times.

Given that there are q queries, if we denote the number of queries where y > m as q1 and where y <= m as q2, then the total time for all queries is 0(q1 * (n / y) + q2). However, in the worst case, all queries could be such that y > m, which makes it 0(q * (n + (n + y)))/ y)).

2. When y > m: The complexity is O(n / y) because the sum is computed using slicing with a step y, so it touches every yth element.

1. When $y \ll m$: The complexity for this case is O(1) because the result is directly accessed from the precomputed suf array.

The total time complexity is therefore 0(m * n + q * (n / y)). In the worst case scenario where y is just above m, this could be approximated as 0(m * n + q * n / m), which simplifies to 0(m * n + q * sqrt(n)).

The space complexity is primarily due to the additional 2D list suf. Since suf has a size of (m+1) * (n+1), its space complexity is 0(m * n). Additionally, the space used by ans to store the results grows linearly with the number of queries q, hence 0(q). Therefore, the total space complexity is 0(m * n + q).