1723. Find Minimum Time to Finish All Jobs Bitmask Bit Manipulation Hard Array **Dynamic Programming** Backtracking

Leetcode Link

Problem Description

In this problem, we are given an array called jobs, where each entry jobs[i] represents the time it takes to complete the i-th job. Additionally, there are k workers available to assign these jobs to. Each job must be assigned to one and only one worker. The

"working time" for a worker is the total time they must spend to complete all jobs assigned to them. The objective is to distribute the jobs among the workers in such a way that the maximum working time for any single worker is minimized. This means we want to find the least amount of time any worker has to work if the jobs are assigned optimally. The function should return this minimum maximum working time.

Intuition To solve this problem, we use a backtracking (or depth-first search) approach. Since we want to minimize the maximum working time

of any worker, a brute force approach would involve trying all possible assignments of jobs to workers and finding the assignment with the minimum maximum working time, which would be very inefficient. This is where backtracking becomes helpful, as it allows

chance to deal with difficult cases early and potentially prune the search space more effectively. Using the dfs function (depth-first search), we recursively explore different job assignments, where we assign the i-th job to each worker one by one and recursively assign the next job. If adding the current job to a worker's existing workload would exceed the current best answer (which is stored in the global variable ans), we skip that assignment to prune the search space.

us to explore possibilities but backtrack as soon as we know a certain assignment cannot be better than the best one found so far.

We start by sorting the jobs in decreasing order. This ensures that we are dealing with the largest jobs first, which gives us the

When we assign a job to a worker whose current workload is 0 (meaning they have not been assigned a job yet), and after this assignment, their workload becomes non-zero, we don't need to try assigning the job to other workers with 0 workloads. This is

because the order in which workers without any assigned jobs start their first job doesn't matter. The ans variable keeps track of the best (minimum) maximum working time found so far, and the cnt array tracks the current working time for each worker. The algorithm terminates when all jobs have been assigned and updates the ans with the minimum of the

current maximum working times if it's better than the previously recorded maximum.

Solution Approach The solution provided follows the backtracking approach to solve the problem efficiently. Here's a step-by-step walkthrough of the implementation of the solution based on the reference approach provided above:

1. Sorting the jobs array: Sorting the jobs in descending order helps in handling the largest jobs first during the backtracking

process. This early handling of more significant jobs can lead to earlier pruning of the branches in the search tree, hence

1 jobs.sort(reverse=True)

1 def dfs(i):

1 if cnt[j] == 0:

1 cnt = [0] * k

optimizing the performance.

2. Depth-First Search (DFS): The dfs function is the core of the solution, which performs depth-first search (or backtracking). When a job is to be assigned, the function iterates over all k workers to try and assign the job to each one of them.

It uses a counter array cnt, where cnt[j] is the total amount of time worker j has been assigned jobs.

4. Exploring and backtracking: After assigning a job to a worker, the recursion continues with the next job (1 + 1).

total working time that exceeds the best solution found so far (ans). If so, this branch of the recursion is abandoned; otherwise, the job is added to the worker's total time. 1 if cnt[j] + jobs[i] >= ans:

3. Pruning the search space: Within the dfs function, there is a check to see if assigning the current job to a worker will result in a

1 cnt[j] += jobs[i] 3 cnt[j] -= jobs[i] # Backtrack 5. Avoiding identical workers' permutations: If a worker has not yet been assigned any job (cnt[j] == 0), there's no need to

6. Updating the answer: Once all jobs have been considered (i = len(jobs)), the function updates the minimum maximum

1 ans = min(ans, max(cnt)) 7. Initialization: Before diving into the backtracking, initialize the search with initial values.

continue to the next worker after assigning a job to them because all workers are identical.

working time found so far (ans) if the current assignment is better.

8. **Kick off the DFS**: Start the DFS with the first job (i=0).

1 dfs(0) 9. Return the result: After the backtracking is complete, the ans variable holds the desired result, which is the minimum maximum working time across all workers.

The algorithm makes use of recursive backtracking and pruning techniques to navigate the space of possible allocations, aiming to

minimize the maximum working time of any worker. The key here is to explore different combinations of job assignments in a depth-

first search manner while trimming branches that exceed the current best solution, which efficiently leads to the minimum possible

Example Walkthrough

maximum working time.

• jobs = [5,3,2,1]

• k = 2 workers available

1 jobs.sort(reverse=True) # jobs = [5, 3, 2, 1]

1 cnt = [0] * k # cnt = [0, 0]

with other workers with 0 workload.

def dfs(curr_job_index):

nonlocal min_time

for worker in range(k):

continue

break

Return the minimum time found

Assign job to worker

dfs(curr_job_index + 1)

2 ans = float('inf')

1 return ans

Now let's walk through the backtracking solution: 1. Sorting the jobs array: We sort the jobs array in descending order to deal with the larger jobs first.

• We start our DFS from the first job; dfs(0) means we are considering how to assign the job with a time requirement of 5. 4. Exploring job assignments: We try assigning the job to each worker.

3. First DFS call

efficiently.

15

16

17

18

20

21

22

23

24

25

26

27

28

29

30

31

32

33

40

43

44

4

5

6

8

9

10

11

12

13

14

15

16

dfs(0)

Java Solution

1 class Solution {

return min_time

this.workers = k;

int temp = jobs[i];

jobs[i] = jobs[j];

jobs[j] = temp;

Python Solution

class Solution:

from typing import List

 Assign job 2 to worker 1, cnt = [8, 2], then dfs(3). We continue because it's still less than ans. Assign job 1 to worker 0, cnt = [9, 2], then we check if all jobs are assigned & update ans to 9.

Backtrack and assign job 1 to worker 1, cnt = [8, 3], update ans to 8.

Backtrack to when job 3 was assigned to worker 0.

def minimumTimeRequired(self, jobs: List[int], k: int) -> int:

Helper function for the depth-first search algorithm.

Accessing 'min_time' as a nonlocal variable to modify it

Try assigning current job to each worker one by one

workers_load[worker] += jobs[curr_job_index]

workers_load[worker] -= jobs[curr_job_index]

Backtrack: Remove job from worker (undo assignment)

Recursively assign the next job

if workers_load[worker] == 0:

Start the DFS with the first job (0-index)

Assign job 5 to worker 0. cnt = [5, 0], then make a recursive DFS call dfs(1) and go to the next job.

Assign job 2 to worker 0, cnt = [10, 0], then dfs(3). Since the sum exceeds the best so far, we backtrack.

6. Skip identical workers' states: When a job is assigned to a worker with a 0 workload, we don't have to consider permutations

After running through all possible combinations and backtracking when necessary, we find that the best distribution that minimizes

the maximum working time for a worker is [5, 3] and [2, 1] with a maximum working time of 8. So the function will return 8.

This example demonstrates the depth-first search combined with pruning to avoid unnecessary work, leading us to the solution

Assign job 3 to worker 0, cnt = [8, 0], then dfs(2). Since the sum is less than ans, we continue.

Let's use a small example to illustrate the solution approach. Suppose we have the following input:

2. Depth-First Search (DFS): Initialize the dfs function and the global variables cnt and ans.

- Try assigning job 3 to worker 1 instead, cnt = [5, 3], and proceed with the next jobs similarly. 5. Backtracking and pruning: The DFS algorithm will backtrack whenever it encounters a sum that cannot be better than the current ans, which is constantly being updated.
- # Base case: all jobs have been assigned if curr_job_index == len(jobs): # Update the minimum time if a better solution is found min_time = min(min_time, max(workers_load)) 14 return

If worker hadn't been assigned any job before, then this is the least loaded they can be

No need to try assigning jobs to other workers with the same load as it'll be equivalent

34 # Initialize each worker's load as 0 35 workers_load = [0] * k 36 # Sort the jobs in descending order to optimize the search (using greedy aproach) 37 jobs.sort(reverse=True) # Initialize the minimum time to infinity 38 39 min_time = float('inf')

private int[] workerLoads; // array to keep track of current load on each worker

// Entry method to find the minimum time required to complete all jobs with k workers

private int minimumTime; // minimum time for completing all jobs

Arrays.sort(jobs); // sort the jobs in ascending order

for (int i = 0, j = jobs.length - 1; i < j; ++i, --j) {

// Reverse the jobs array to have jobs in descending order

int minimumJobTime; // Renamed from 'ans' to 'minimumJobTime' for clarity

int minimumTimeRequired(vector<int>& jobs, int k) {

dfs(0, k, jobs, workerTimes);

for (int j = 0; j < k; ++j) {

return minimumJobTime;

// Helper method to perform DFS

return;

sort(jobs.begin(), jobs.end(), greater<int>());

// Start the Depth-First Search (DFS) to assign jobs

// Update the minimum job time if necessary

// Iterate through each worker to assign the job

// Assign the job to the worker

if (workerTimes[j] == 0) break;

// Recur for the next job

workerTimes[j] += jobs[currentIndex];

workerTimes[j] -= jobs[currentIndex];

1 let minimumJobTime: number; // Holds the minimum job time found

// Start the recursive search to assign jobs

// Iterate through each worker to assign the job

function minimumTimeRequired(jobs: number[], k: number): number {

dfs(currentIndex + 1, k, jobs, workerTimes);

// Method to find the minimum time required to complete all jobs given 'k' workers

vector<int> workerTimes(k); // Holds the current total time for each worker

// Sort jobs in descending order to try the largest jobs first for optimization

minimumJobTime = INT_MAX; // Initialize with the maximum possible value

void dfs(int currentIndex, int k, vector<int>& jobs, vector<int>& workerTimes) {

// Skip if adding this job exceeds the current minimum time

// Function to find the minimum time required to complete all jobs given 'k' workers

minimumJobTime = Math.min(minimumJobTime, Math.max(...workerTimes));

if (workerTimes[j] + jobs[currentIndex] >= minimumJobTime) continue;

let workerTimes: number[] = new Array(k).fill(0); // Holds the current total time for each worker

// Skip if adding this job would make the worker's time exceed the current minimum job time

// Sort jobs in descending order to try the largest jobs first for better optimization

if (currentIndex == jobs.size()) { // Base case: all jobs have been assigned

if (workerTimes[j] + jobs[currentIndex] >= minimumJobTime) continue;

// Backtrack: remove the job from the worker for the next iteration

minimumJobTime = min(minimumJobTime, *max_element(workerTimes.begin(), workerTimes.end()));

// Optimization: If the worker time is zero, no need to try assigning this job to other workers

minimumJobTime = Number.MAX_SAFE_INTEGER; // Initialize with the largest safe integer value due to the absence of INT_MAX

private int workers; // number of workers available

public int minimumTimeRequired(int[] jobs, int k) {

private int[] jobs; // array to store jobs (sorted in descending order)

Pruning: skip if current assignment exceeds the known minimum time

if workers_load[worker] + jobs[curr_job_index] >= min_time:

```
45
46
47
48
```

C++ Solution

1 class Solution {

2 public:

9

10

12

13

14

15

16

17

20

21

22

24

25

26

28

30

31

32

33

34

35

36

38

10

11

19

20

21

22

23

24

25

26

27

28

29

30

31

32

33

35

37 };

58

```
17
18
            this.jobs = jobs; // store the sorted jobs array
19
            workerLoads = new int[k]; // initialize the worker loads array
20
            minimumTime = Integer.MAX_VALUE; // initialize the minimum time to a large number
21
22
            // Start the depth-first search to assign jobs to workers
23
            dfs(0);
24
            return minimumTime;
25
26
27
       // Depth-first search method to assign jobs to workers and find the minimum time
28
        private void dfs(int jobIndex) {
29
            if (jobIndex == jobs.length) { // If all jobs have been assigned
                int maxLoad = 0; // find out the maximum load on any worker
30
31
                for (int load : workerLoads) {
32
                    maxLoad = Math.max(maxLoad, load);
33
34
                minimumTime = Math.min(minimumTime, maxLoad); // update minimum time if necessary
35
                return;
36
37
            // Iterate over workers and try to assign the current job to each worker
38
39
            for (int j = 0; j < workers; ++j) {</pre>
40
               // Skip assignment if job addition exceeds current minimum time
41
                if (workerLoads[j] + jobs[jobIndex] >= minimumTime) {
42
                    continue;
43
44
                workerLoads[j] += jobs[jobIndex]; // assign the job to the worker
                dfs(jobIndex + 1); // recurse to assign the next job
                // Backtrack: remove the job from the worker
                workerLoads[j] -= jobs[jobIndex];
49
50
               // If the current worker had 0 load, then there's no need to try further workers
51
                // since all are identical at this point
52
                if (workerLoads[j] == 0) {
53
                    break;
54
55
56
57 }
```

12 } // Recursive function to perform Depth-First Search (DFS) for job assignments 15 const dfs = (currentIndex: number, k: number, jobs: number[], workerTimes: number[]): void => { if (currentIndex === jobs.length) { // Base case: all jobs have been assigned // Update the minimum job time if a new minimum is found

return;

explored states.

jobs.sort((a, b) => b - a);

return minimumJobTime;

dfs(0, k, jobs, workerTimes);

for (let j = 0; j < k; ++j) {

Typescript Solution

```
// Assign the job to the worker
           workerTimes[j] += jobs[currentIndex];
           // Recurse for the next job
           dfs(currentIndex + 1, k, jobs, workerTimes);
           // Backtrack: remove the job from the worker for the next iteration
           workerTimes[j] -= jobs[currentIndex];
           // Optimization: if the current worker has no jobs, don't assign this job to other workers
           if (workerTimes[j] === 0) break;
34 };
Time and Space Complexity
Time Complexity
The time complexity of the minimumTimeRequired function is determined by the depth-first search (DFS) as it tries to assign jobs to
workers in every possible combination to find the minimum possible time.

    Since jobs are being sorted initially, this contributes 0(n log n) to the time complexity, where n is the number of jobs.

    The DFS process has a worst-case time complexity of O(k^n) where k is the number of workers, and n is the number of jobs,

    since each job can be assigned to any of the k workers.

    However, due to pruning:
```

If cnt[j] + jobs[i] >= ans, the DFS does not continue down that path, which can significantly reduce the number of

If cnt[j] == 0, it breaks the inner loop, avoiding redundant assignments to empty slots which are equivalent.

The exact time complexity is hard to characterize due to these pruning strategies, but the worst-case without pruning is O(k^n),

Space Complexity For space complexity, the following considerations are taken into account:

factoring in the sorting of jobs, we have $0(n \log n + k^n)$.

number of jobs). An array cnt of size k is used to keep track of the current sum of job times for each worker, contributing O(k) space complexity.

Therefore, the overall space complexity is O(n + k) due to the recursive stack and the cnt array. However, if we consider that the depth of the recursive stack dominates as n can potentially be larger than k, the space complexity simplifies to 0(n).

• The recursive call stack of DFS contributes O(n) space complexity since the maximum depth of the recursion stack is n (the