1763. Longest Nice Substring

Hash Table

String

Problem Description

Bit Manipulation

Easy

which contains every letter in both uppercase and lowercase forms. For example, a string "aA" is nice because it contains both 'a' and 'A'. If there are several such nice substrings, we are to return the one that occurs first. If no nice substrings exist, we should return an empty string.

The challenge in this problem is to find the longest substring of the given string, s, where a "nice" substring is defined as one

Sliding Window

Divide and Conquer

Intuition

To solve this problem, we iterate through the string character by character, starting from each character in turn. We use two

bitmask integers, lower and upper, to represent the presence of lowercase and uppercase letters, respectively. For each character in 'a' to 'z', a bit is set to 1 in lower if the lowercase letter is present, and correspondingly in upper if the uppercase letter is present. For instance, if the lowercase letter 'b' is encountered, the second bit (assuming 0-indexing) of lower is set to 1. Similarly, if the uppercase letter 'B' is spotted, the second bit of upper is turned on.

We move through the string with a nested loop, checking consecutive characters starting from each index pointed by the outer loop and ending at the length of the string. While doing this, we keep updating our lower and upper bitmasks for each character we see. A substring is nice if the lower and upper bitmasks are equal at some point, meaning every lowercase character

observed so far has a corresponding uppercase version present, and vice versa. When we find such a substring, we check if it's longer than any previous "nice" substring found (initially none, as denoted by an empty string ans). If it's longer, we update our answer with the current substring.

implement, its time complexity is not optimal for very long strings. However, it is perfectly viable for strings of moderate length and guarantees that the longest "nice" substring will be found.

This brute force approach checks all possible substrings starting at each point in the string, and while it's easy to understand and

Solution Approach To implement the solution, we use a simple brute force approach which may not be the most efficient in terms of time complexity but is straightforward to understand and guarantees to find the correct answer. Here's a step-by-step breakdown:

checks every subsequent character up to the end of the string, indexed by j.

For each character, two bitmasks lower and upper are maintained. If the character is lowercase (checked using s[j].islower()), we set the corresponding bit in lower. This is done by shifting 1 to the left by ord(s[j]) - ord('a') places (since 'a' is the ASCII starting point for lowercase letters, the difference gives us the correct bit position).

Iterate through the string s using a nested loop. The outer loop starts from each character indexed by i and the inner loop

Similarly, if the character is uppercase, the corresponding bit in upper is set by shifting 1 to the left by ord(s[j]) - ord('A') places (as 'A' is the ASCII starting point for uppercase letters).

right we naturally prioritize earlier substrings over later ones of the same length.

Initialize an empty string ans to keep track of the longest nice substring found.

- We compare lower and upper to check if the current substring is nice, which would be true if lower equals upper. This comparison realizes if for every lowercase letter there's a matching uppercase letter and vice versa.
- If a nice substring is found and its length is greater than the length of the current ans, the ans string is updated to this substring. We achieve substring extraction using Python's slice notation s[i : j + 1]. We repeat this process, expanding our current substring check from all possible starting points (1) to include all following

characters (j). Since the answer should be of the earliest occurrence of the longest nice substring, by iterating from left to

Finally, after all iterations, and holds the longest nice substring, and it is returned. Throughout this implementation, we rely on basic bitwise operations, simple string manipulation, and nested loops. The algorithm's complexity is O(n^2) where n is the number of characters in the string. Each nested loop iteration checks one

Although not efficient for very large strings, for smaller strings, this simple and direct method effectively locates the desired nice

Consider the string s = "aAbBcC". We are looking for the longest "nice" substring, that is, a substring where each letter exists in

We begin by iterating over the string s with index i from 0 to the length of s. The first character is 'a', and we start the inner

Continuing this process of updating lower and upper for each character in the inner loop, we reach the end of the string. By

now, lower and upper should both be 111111 in binary, which corresponds to 63 in decimal, standing for the presence of 'a',

We check if lower equals upper after each inner loop iteration and update ans only if the current substring (s[i : j + 1]) is

longer than ans. In this case, after the first full iteration (i = 0 to j = the end of s), the ans will become "aAbBcC", which is the

Completing the loops without finding a longer nice substring, we would still have ans = "aAbBcC", which is indeed the longest

Example Walkthrough Let's walk through an example to illustrate how the described solution approach is applied to the problem.

We start with an empty string ans to hold the longest nice substring we find.

substring, and there are O(n^2) substrings in total.

loop from i to check subsequent characters.

lower = 1 as 'a' is the first lowercase letter.

'b', and 'c' in lowercase and uppercase.

entire string, since lower == upper is true.

nice substring that meets the criteria.

Iterate over the string with two pointers

lower case flags = 0 # Bit flags for lowercase letters

upper_case_flags = 0 # Bit flags for uppercase letters

Set the bit corresponding to the lowercase letter

Set the bit corresponding to the uppercase letter

lower_case_flags |= 1 << (ord(s[j]) - ord('a'))</pre>

upper_case_flags |= 1 << (ord(s[j]) - ord('A'))

// 'start' will keep the index at which the longest nice substring begins

// 'lowerCaseBitmask' and 'upperCaseBitmask' are bitmasks to keep track of

// 'maxLength' is the length of the longest nice substring found so far

// Iterate over each character in the string as the starting point

Explore the substring starting from index i

Solution Implementation

for i in range(n):

else:

for j in range(i, n):

return longest_nice_substring

// Length of the input string

int stringLength = inputString.length();

for (int i = 0; i < stringLength; ++i) {</pre>

if s[i].islower():

Python

both uppercase and lowercase form.

substring without additional data structures or complex algorithms.

For each character s[j] we encounter as we move through the inner loop: ∘ If s[j].islower() is true, for example s[j] = 'a', we modify lower bitmask. We do lower |= (1 << (ord('a') - ord('a'))), resulting in

 \circ If s[j].isupper() is true instead, for example s[j] = 'A', we modify the upper bitmask in a similar fashion, so upper = 1.

entire string, no longer substring is possible.

At the end of the algorithm, the ans string, which equals "aAbBcC", is returned as the correct answer.

- To continue our process, we would next start with i = 1, but given that we already found a nice substring that includes the
- This example demonstrates the scenario where the entire string meets the conditions to be a nice string. Thus, the first and longest nice substring is found on the very first iteration of the outer loop.
- class Solution: def longestNiceSubstring(self, s: str) -> str: n = len(s) # length of the input string longest_nice_substring = '' # Initialize the longest nice substring

Update the longest nice substring if the current one is longer longest_nice_substring = s[i : j + 1] # Return the longest nice substring found

Check if the current substring is nice (lowercase and uppercase bits match)

if lower case flags == upper case flags and len(longest nice substring) < j - i + 1:</pre>

```
Java
class Solution {
    public String longestNiceSubstring(String inputString) {
```

int start = -1:

int maxLength = 0;

// and its length.

int startIndex = -1, maxLength = 0;

for (int i = 0; i < strSize; ++i) {</pre>

char c = s[j];

if (islower(c))

else

};

/**

TypeScript

int lowerBitmask = 0, upperBitmask = 0;

for (int i = i; i < strSize; ++i) {</pre>

// Get the current character.

// Explore substrings starting at index i.

lowerBitmask |= 1 << (c - 'a');

upperBitmask |= 1 << (c - 'A');

// Also check if it's the longest so far.

// If no nice substring is found, return an empty string.

* @param {string} s - The input string to search for the nice substring.

// Iterate through the string to find all possible substrings

for (let start = 0; start < length0fString; start++) {</pre>

function longestNiceSubstring(s: string): string {

// Return the longest nice substring found

return longestSubstring;

const lengthOfString = s.length;

let longestSubstring = '';

* @return {string} - The longest nice substring found in the input string.

longestSubstring = s.substring(start, end + 1);

return startIndex == -1 ? "" : s.substr(startIndex, maxLength);

// Otherwise, return the longest nice substring found.

// Iterate through each character as starting point of the nice substring.

// Bitmask to represent lowercase and uppercase letters encountered.

// Set the bit for this character in the appropriate bitmask.

if (lowerBitmask == upperBitmask && maxLength < i - i + 1) {</pre>

* Finds the longest substring where each character appears in both lower and upper case.

// Check if the current substring is nice: it has both cases for each letter.

maxLength = i - i + 1; // Update max length to the new longest nice substring.

startIndex = i; // Update start index to the starting index of the new longest nice substring.

```
// lowercase and uppercase characters encountered
            int lowerCaseBitmask = 0, upperCaseBitmask = 0;
            // Try extending the substring from the starting point 'i' to 'j'
            for (int j = i; j < stringLength; ++j) {</pre>
                // Get the current character
                char currentChar = inputString.charAt(j);
                // If it's lowercase, set the corresponding bit in the bitmask
                if (Character.isLowerCase(currentChar)) {
                    lowerCaseBitmask |= 1 << (currentChar - 'a');</pre>
                // If it's uppercase, set the corresponding bit in the bitmask
                else {
                    upperCaseBitmask |= 1 << (currentChar - 'A');
                // Check if the substring from 'i' to 'i' is a nice string
                // A nice string has the same set of lowercase and uppercase characters
                if (lowerCaseBitmask == upperCaseBitmask && maxLength < i - i + 1) {</pre>
                    // Update the maxLength and the starting index 'start'
                    maxLength = j - i + 1;
                    start = i;
        // If 'start' was updated (meaning a nice substring was found), return it
        // Otherwise, if no nice substring exists, return an empty string
        return (start == -1) ? "" : inputString.substring(start, start + maxLength);
C++
class Solution {
public:
    string longestNiceSubstring(string s) {
        // Initialize the size of the string.
        int strSize = s.size();
        // These will keep track of the start index of the longest nice substring
```

```
let lowerCaseMask = 0; // Bitmap for tracking lowercase letters
let upperCaseMask = 0; // Bitmap for tracking uppercase letters
// Explore the substrings starting from 'start' index
for (let end = start; end < lengthOfString; end++) {</pre>
    const charCode = s.charCodeAt(end);
    // If the character is lowercase, update the lowerCaseMask
    if (charCode > 96) {
        lowerCaseMask |= 1 << (charCode - 97);</pre>
    // If the character is uppercase, update the upperCaseMask
    else {
        upperCaseMask |= 1 << (charCode - 65);
```

// Check if the current substring is "nice" and if it is longer than the current longest

if (lowerCaseMask === upperCaseMask && end - start + 1 > longestSubstring.length) {

```
class Solution:
   def longestNiceSubstring(self, s: str) -> str:
       n = len(s) # length of the input string
        longest_nice_substring = '' # Initialize the longest nice substring
       # Iterate over the string with two pointers
        for i in range(n):
            lower case flags = 0 # Bit flags for lowercase letters
           upper_case_flags = 0 # Bit flags for uppercase letters
           # Explore the substring starting from index i
            for j in range(i, n):
               if s[i].islower():
                   # Set the bit corresponding to the lowercase letter
                    lower_case_flags |= 1 << (ord(s[j]) - ord('a'))
               else:
                   # Set the bit corresponding to the uppercase letter
                   upper_case_flags |= 1 << (ord(s[j]) - ord('A'))
               # Check if the current substring is nice (lowercase and uppercase bits match)
               if lower case flags == upper case flags and len(longest nice substring) < j - i + 1:
                   # Update the longest nice substring if the current one is longer
                    longest nice substring = s[i : j + 1]
       # Return the longest nice substring found
       return longest_nice_substring
Time and Space Complexity
  The given Python code snippet is designed to find the longest substring of a given string s such that the substring is "nice". A
```

Time Complexity The time complexity of the code is determined by the nested for-loops. The outer loop runs n times where n is the length of

string s. The inner loop runs at most n times for each iteration of the outer loop as it starts from the current position of i to the end of the string s. This gives us a quadratic number of iterations in the worst case.

substring is considered "nice" if it contains both the uppercase and lowercase forms of the same letter.

output, is O(n).

The operations inside the inner loop are constant time operations, such as checking if a character is lower or uppercase and setting bits in an integer. Hence, they don't affect the time complexity's order. Therefore, the overall time complexity of the code is $O(n^2)$. **Space Complexity**

The space complexity includes the variables to store the bitmasks for lowercase (lower) and uppercase (upper) letters, a

variable for the answer (ans), and two loop variables (i and j). The bitmasks lower and upper use fixed space since there are

exactly 26 lowercase and 26 uppercase English letters, and they can be represented within a fixed-size integer type. Since the algorithm's space usage does not scale with the size of the input string s, besides the output string ans, the space complexity is 0(1) for the working variables. However, if we consider the space used by the output ans, it could be 0(n) in the case when the whole string s is a "nice" substring itself. Therefore, the overall space complexity, including the space for the