1817. Finding the Users Active Minutes

`Medium` `Array` `Hash Table`

Problem Description

In this problem, we are dealing with logs of user actions on LeetCode. The logs are given as a 2-dimensional array where each element represents a user's action with two pieces of information: the user's ID and the time the action occurred, recorded by the minute. We are instructed to compute a user's active minutes (UAM), which are the unique minutes during which they performed any actions. Importantly, even if a user performs multiple actions within a single minute, that minute is only counted once for their UAM.

Our goal is to create an array of size `k`, where each entry `i` indicates the number of users whose UAM is exactly `j`. To clarify, the array we are constructing is 1-indexed, meaning that the index of each item starts at 1, not 0. Therefore, for every number of active minutes that a user could have (from 1 to `k`), we want to count how many users have that amount of UAM.

Intuition

The natural approach to solve this problem relies on representing each user's actions without duplications and counting their active minutes efficiently. Since a user can perform many actions, possibly at the same times, we need a way to store these actions such that we can easily determine the unique minutes.

Using a hash table, or in Python, a dictionary that maps each user ID to a set of unique times, provides an effective solution. Sets are useful here because they naturally eliminate duplicate entries, which aligns perfectly with counting unique active minutes.

After populating the hash table with each user's unique active minutes, we then need to determine the number of users for each possible UAM value up to `k`. We accomplish this by initiating an answer array `ans` filled with zeroes to counts for each UAM from 1 to `k`. We loop through the hash table and increment the count in `ans` at the index that corresponds to the user's UAM (adjusted by subtracting 1 for 1-indexing).

The reason behind this approach is two-fold:

1. By mapping users to sets of timestamps, we efficiently record user activity without duplication, directly giving us the UAM for each user.
2. By using an array indexed by active minutes, we can simply tally up the number of users for each UAM, which the problem requires us to report.

Solution Approach

The solution's strategy is based on the hash table data structure which is very common for storing relationships between key-value pairs. In Python, this is typically implemented using a `defaultdict` which is a subclass of the dictionary (`dict`). A `defaultdict` allows us to specify a default value for the dictionary – in this case, a `set`. The `set` is chosen as the default value because we are interested in keeping a collection of unique active minutes for each user.

Here's the detailed implementation flow:

- **Initialize the `defaultdict`:** We create `d` as a defaultdict of sets. Each user ID will be a key, and the corresponding value will be a set that contains all unique minutes during which the user performed actions.

- **Populate the `defaultdict`:**
  - Iterate over each entry in the `logs` array.
  - For each `[user_id, timestamp]` pair, add `timestamp` to the set mapped to by `user_id`. The set ensures that if the same `timestamp` is inserted multiple times, it will only be stored once, preserving the unique nature of the active minutes.

- **Initiate the `ans` array:** We create an array (list in Python) called `ans` with length `k`, filled with zeros. Each index `i` in this array represents the number of users whose UAM is `i + 1`, because the problem asks for a 1-indexed array.

- **Count the UAM for each user:**
  - Loop through the sets in the `defaultdict` values, which each represent the set of unique active minutes for a particular user ID.
  - The size of each set is the UAM count for that user. Subtracting 1 from the set size gives the index that corresponds to that UAM in the `ans` array.
  - Increment the value at the calculated index in `ans` by 1.

By following this solution approach, we effectively tabulate the distribution of UAMs across all users. Finally, we return the populated `ans` array as our result, which contains the count of users for each UAM value from 1 to `k`.

Here's an illustration of the reference solution approach in a concise form:

```
1  class Solution:
2      def findingUsersActiveMinutes(self, logs: List[List[int]], k: int) -> List[int]:
3          d = defaultdict(set) # Create defaultdict of sets
4          for i, t in logs:
5              d[i].add(t) # Populate the defaultdict
6          ans = [0] * k # Initialize ans list
7          for ts in d.values():
8              ans[len(ts) - 1] += 1 # Count UAM and populate ans
9          return ans
```

Example Walkthrough

Let's consider an example to illustrate the solution approach with the logs array and `k` value as follows:

```
1  logs = [[1, 1], [2, 2], [1, 2], [1, 3], [3, 4], [3, 4], [3, 5]], k = 3
```

With this data, we'll walk through the implementation steps:

Initialize the defaultdict

We create a `defaultdict` named `d`. Each key will be unique user IDs, and each value will be a set of unique timestamps:

```
1  d = defaultdict(set)
```

Initially, `d` will be empty.

Populate the defaultdict

We iterate through each log entry and update `d` with unique active minutes for each user:

- First log entry is `[1, 1]`—we add minute `1` to user `1`'s set.
- Next, we process three log entries for user `2`. They are all in minutes `2` and `3`, resulting in user `2`'s set containing `{2, 3}`.
- Lastly, user `3` has entries at minutes `4` and `5`, but duplicates, so user `3`'s set will only contain `{4, 5}`.

After processing, `d` looks like this:

```
1  {1: {1}, 2: {2, 3}, 3: {4, 5}}
```

Initiate the ans array

K is `3`, so we create an answer list `ans` of length `k`, with all elements initialized to `0`:

```
1  ans = [0] * 3 # Which is [0, 0, 0]
```

Count the UAM for each user

We'll iterate through the sets in `defaultdict` values and update the `ans` list based on the number of unique active minutes for each user:

- User 1 has a UAM count of `1`, so we increment `ans[0]` from `0` to `1` (because index `0` corresponds to UAM `1`).
- User 2 has a UAM count of `2`, so we increment `ans[1]` to `1`.
- User 3 also has a UAM of `2`, so we increase `ans[1]` again, from `1` to `2`.

After processing, the `ans` array is:

```
1  [1, 2, 0]
```

This array means there is one user with a UAM of `1`, two users with a UAM of `2`, and zero users with a UAM of `3`.

The final output of the reference solution for our given `logs` and `k` values would be `[1, 2, 0]`.

Python Solution

```python
1  from collections import defaultdict
2  from typing import List
3
4  class Solution:
5      def findingUsersActiveMinutes(self, logs: List[List[int]], k: int) -> List[int]:
6          # Dictionary mapping user IDs to a set of their active minutes
7          user_active_minutes = defaultdict(set)
8
9          # Iterate through the logs, adding unique minutes to each user's set
10         for user_id, minute in logs:
11             user_active_minutes[user_id].add(minute)
12
13         # Initialize an answer array to store counts of users with k active minutes
14         # where k is the index + 1
15         answer = [0] * k
16
17         # Iterate through the user_active_minutes to update the answer array
18         for minutes in user_active_minutes.values():
19             # Increase the count for the number of users (i) with a specific
20             # We do -1 because the index is 0-based and minutes is 1-based
21             answer[len(minutes) - 1] += 1
22
23         return answer
```

Java Solution

```java
1  class Solution {
2      public int[] findingUsersActiveMinutes(int[][] logs, int k) {
3          // A hashmap to store unique time stamps for each user
4          Map<Integer, Set<Integer>> userActiveTimes = new HashMap<>();
5
6          // Iterate over each log entry
7          for (int[] log : logs) {
8              int userId = log[0]; // Extract the user ID
9              int timestamp = log[1]; // Extract the timestamp
10
11             // If the user ID doesn't exist in the map, create a new HashSet for that user.
12             // Then, add the timestamp to the user's set of active times.
13             userActiveTimes.computeIfAbsent(userId, key -> new HashSet<>()).add(timestamp);
14         }
15
16         // Initialize an array to store the UAM count for each possible count 1 through k
17         int[] answer = new int[k];
18
19         // Iterate over each user's set of timestamps
20         for (Set<Integer> timeStamps : userActiveTimes.values()) {
21             // Count the number of unique timestamps for the user and
22             // increment the respective count in the answer array.
23             // Subtract 1 from the size to convert the count into a zero-based index.
24             answer[timeStamps.size() - 1]++;
25         }
26
27         // Return the filled answer array
28         return answer;
29     }
30 }
```

C++ Solution

```cpp
1  #include <vector>
2  #include <unordered_map>
3  #include <unordered_set>
4  using namespace std;
5
6  class Solution {
7  public:
8      // This function takes a vector of vector of integers which contain user ID and
9      // login time, along with an integer k representing the total number of minutes.
10     // It returns a vector with the count of users who have a specific number of
11     // active minutes ranging from 1 to k.
12     vector<int> findingUsersActiveMinutes(vector<vector<int>>& logs, int k) {
13         // Create a map that associates each user (identified by their ID) with a set of unique active minutes.
14         unordered_map<int, unordered_set<int>> userToActiveMinutesMap;
15         for (const auto& logEntry : logs) {
16             int userID = logEntry[0];
17             int time = logEntry[1];
18             // Insert the time into the set for this user, duplicates are automatically handled by the set.
19             userToActiveMinutesMap[userID].insert(time);
20         }
21
22         // Initialize a vector to store the counts of users having k number of active minutes,
23         // where k ranges from 1 to k.
24         vector<int> activeMinutesCounts(k, 0);
25         for (const auto& userToTimeEntry : userToActiveMinutesMap) {
26             // Increment the count for the appropriate number of active minutes (decrement by 1 to match 0-indexing).
27             int uniqueActiveMinutes = userToTimeEntry.second.size();
28             // Increment the count for the appropriate number of active minutes (decrement by 1 to match 0-indexing).
29             if (uniqueActiveMinutes <= k) {
30                 activeMinutesCounts[uniqueActiveMinutes - 1]++;
31             }
32         }
33         // Return the resulting vector of counts.
34         return activeMinutesCounts;
35     }
36 };
```

Typescript Solution

```typescript
1  // This function calculates the number of unique users who have exactly 1 to k active minutes.
2  // Logs are represented as an array of [userId, minute] and k is the upper limit of active minutes.
3  function findingUsersActiveMinutes(logs: number[][], k: number): number[] {
4      // Initialize a map to store unique minutes for each user.
5      const userActivityMap: Map<number, Set<number>> = new Map();
6
7      // Iterate through all the log entries.
8      for (const [userId, minute] of logs) {
9          // If the user does not have a set yet, initialize it.
10         if (!userActivityMap.has(userId)) {
11             userActivityMap.set(userId, new Set<number>());
12         }
13         // Add the minute to the set corresponding to the user.
14         userActivityMap.get(userId)!.add(minute);
15     }
16
17     // Initialize an array to count users with exactly 1 to k active minutes.
18     const userCount: number[] = Array(k).fill(0);
19
20     // Iterate through the map to count the number of active minutes for each user.
21     for (const userMinutesSet of userActivityMap.values()) {
22         // Determine the count for the number of users with a specific count of active minutes.
23         // (size - 1) is used because indices are 0-based and our minute counts are 1-based.
24         ++userCount[userMinutesSet.size - 1];
25     }
26
27     // Return the array with counts of users having 1 to k active minutes.
28     return userCount;
29 }
```

Time and Space Complexity

The time complexity of the code is $O(n)$, where n is the length of the `logs` array. This is because we iterate over each element of the `logs` array exactly once, adding the timestamp to a set in the dictionary. Since set operations such as `.add()` have a time complexity of $O(1)$, iterating over the `logs` array contributes $O(n)$ to the time complexity.

Further, we also iterate over the values of the dictionary, with the number of values being at most n. For each of the values, we increment the appropriate count in the `ans` array, which is another $O(1)$ operation. Hence, the total time complexity remains $O(n)$.

The space complexity of the code is also $O(n)$ because in the worst-case scenario, each log entry could be for a unique user and a unique timestamp, leading to n unique entries in the dictionary. As such, both the dictionary and the set of timestamps could potentially grow linearly with the number of log entries. Therefore, the space used by the dictionary is proportional to the size of the `logs` array, n.