1817. Finding the Users Active Minutes

**Array** 

Hash Table

Medium

**Problem Description** 

any actions. Importantly, even if a user performs multiple actions within a single minute, that minute is only counted once for their UAM. Our goal is to create an array of size k, where each entry j indicates the number of users whose UAM is exactly j. To clarify, the array we are constructing is 1-indexed, meaning that the index of each item starts at 1, not 0. Therefore, for every number of

In this problem, we are dealing with logs of user actions on LeetCode. The logs are given as a 2-dimensional array where each

element represents a user's action with two pieces of information: the user's ID and the time the action occurred, recorded by the

minute. We are instructed to compute a user's active minutes (UAM), which are the unique minutes during which they performed

active minutes that a user could have (from 1 to k), we want to count how many users have that amount of UAM. Intuition

The natural approach to solve this problem relies on representing each user's actions without duplications and counting their

active minutes efficiently. Since a user can perform many actions, possibly at the same times, we need a way to store these

actions such that we can easily determine the unique minutes. Using a hash table, or in Python, a dictionary that maps each user ID to a set of unique times, provides an effective solution. Sets are useful here because they naturally eliminate duplicate entries, which aligns perfectly with counting unique active minutes.

possible UAM value up to k. We accomplish this by initiating an answer array ans filled with zeroes to correspond to counts for each UAM from 1 to k. We loop through the hash table and increment the count in ans at the index that corresponds to the

After populating the hash table with each user's unique active minutes, we then need to determine the number of users for each

The reason behind this approach is two-fold: 1. By mapping users to sets of timestamps, we efficiently record user activity without duplication, directly giving us the UAM for each user. 2. By using an array indexed by active minutes, we can simply tally up the number of users for each UAM, which the problem requires us to report.

value pairs. In Python, this is typically implemented using a defaultdict which is a subclass of the dictionary (dict). A

defaultdict allows us to specify a default value for the dictionary – in this case, a set. The set is chosen as the default value

Initialize the defaultdict: We create d as a defaultdict of sets. Each user ID will be a key, and the corresponding value will

because we are interested in keeping a collection of unique active minutes for each user.

The solution's strategy is based on the hash table data structure which is very common for storing relationships between key-

Here's the detailed implementation flow:

Iterate over each entry in the logs array.

**Solution Approach** 

user's UAM (adjusted by subtracting 1 for 1-indexing).

be a set that contains all unique minutes during which the user performed actions. Populate the defaultdict: •

• For each [user\_id, timestamp] pair, add timestamp to the set mapped to by user\_id. The set ensures that if the same timestamp is

inserted multiple times, it will only be stored once, preserving the unique nature of the active minutes. Initiate the ans array: We create an array (list in Python) called ans with length k, filled with zeros. Each index i in this array

- represents the number of users whose UAM is i + 1, because the problem asks for a 1-indexed array. Count the UAM for each user: •
- The size of each set is the UAM count for that user. Subtracting 1 from the set size gives the index that corresponds to that UAM in the ans array. Increment the value at the calculated index in ans by 1.

By following this solution approach, we effectively tabulate the distribution of UAMs across all users. Finally, we return the

Loop through the sets in the defaultdict values, which each represent the set of unique active minutes for a particular user ID.

populated ans array as our result, which contains the count of users for each UAM value from 1 to k.

ans[len(ts) - 1] += 1 # Count UAM and populate ans

logs = [[1, 1], [2, 2], [2, 2], [2, 3], [3, 4], [3, 4], [3, 4], [3, 5]], k = 3

We iterate through each log entry and update d with unique active minutes for each user:

• Lastly, user 3 has entries at minutes 4 and 5, all duplicates, so user 3's set will only contain {4, 5}.

Here's an illustration of the reference solution approach in a concise form:

return ans

Initialize the defaultdict

d = defaultdict(set)

Populate the defaultdict

Initially, d will be empty.

After processing, d looks like this:

*{*1: *{*1*}*, 2: *{*2, 3*}*, 3: *{*4, 5*}}* 

class Solution:

def findingUsersActiveMinutes(self, logs: List[List[int]], k: int) -> List[int]: d = defaultdict(set) # Create defaultdict of sets for i, t in logs: d[i].add(t) # Populate the defaultdict ans = [0] \* k # Initialize ans list for ts in d.values():

**Example Walkthrough** Let's consider an example to illustrate the solution approach with the logs array and k value as follows:

We create a defaultdict named d. Each key will be unique user IDs, and each value will be a set of unique timestamps:

First log entry is [1, 1] —we add minute 1 to user 1's set.

• User 2 has a UAM count of 2, so we increment ans [1] to 1.

• User 3 also has a UAM of 2, so we increase ans [1] again, from 1 to 2.

With this data, we'll walk through the implementation steps:

**Initiate the ans array** 

[1, 2, 0]

**Python** 

Java

C++

class Solution {

**Count the UAM for each user** 

After processing, the ans array is:

from collections import defaultdict

user\_active\_minutes = defaultdict(set)

user\_active\_minutes[user\_id].add(minute)

public int[] findingUsersActiveMinutes(int[][] logs, int k) {

// A hashmap to store unique time stamps for each user

int userId = log[0]; // Extract the user ID

int timestamp = log[1]; // Extract the timestamp

Map<Integer, Set<Integer>> userActiveTimes = new HashMap<>();

// Then, add the timestamp to the user's set of active times.

// Count the number of unique timestamps for the user and

// increment the respective count in the answer array.

answer[timeStamps.size() - 1]++;

// Return the filled answer array

for user id, minute in logs:

# where X is the index + 1

// Iterate over each log entry

for (int[] log : logs) {

answer = [0] \* k

return answer

from typing import List

class Solution:

K is 3, so we create an answer list ans of length k, with all elements initialized to 0: ans = [0] \* 3 # Which is <math>[0, 0, 0]

• Next, we process three log entries for user 2. They are all in minutes 2 and 3, resulting in user 2's set containing {2, 3}.

We'll iterate through the sets in defaultdict values and update the ans list based on the number of unique active minutes for each user:

This array means there is one user with a UAM of 1, two users with a UAM of 2, and zero users with a UAM of 3.

• User 1 has a UAM count of 1, so we increment ans [0] from 0 to 1 (because index 0 corresponds to UAM 1).

Solution Implementation

The final output of the reference solution for our given logs and k values would be [1, 2, 0].

for minutes in user active minutes.values(): # Increase the count for the number of active minutes # We do -1 because the index is 0-based and minutes is 1-based answer[len(minutes) - 1] += 1

// If the user ID doesn't exist in the map, create a new HashSet for that user.

userActiveTimes.computeIfAbsent(userId, key -> new HashSet<>()).add(timestamp);

// Subtract 1 from the size to convert the count into a zero-based index.

// This function calculates the number of unique users who have exactly 1 to k active minutes.

function findingUsersActiveMinutes(logs: number[][], k: number): number[] {

// Initialize a map to store unique minutes for each user.

// Iterate through all the log entries.

if (!userActivitvMap.has(userId)) {

userActivityMap.get(userId)!.add(minute);

for (const userMinutesSet of userActivityMap.values()) {

++userCount[userMinutesSet.size - 1];

const userCount: number[] = Array(k).fill(0);

user\_active\_minutes = defaultdict(set)

answer[len(minutes) - 1] += 1

user\_active\_minutes[user\_id].add(minute)

for user id, minute in logs:

# where X is the index + 1

for (const [userId, minute] of logs) {

const userActivityMap: Map<number, Set<number>> = new Map();

// If the user does not have a set yet, initialize it.

userActivityMap.set(userId, new Set<number>());

// Add the minute to the set corresponding to the user.

// Initialize an array to count users with exactly 1 to k active minutes.

if (userMinutesSet.size > 0 && userMinutesSet.size <= k) {</pre>

// Return the array with counts of users having 1 to k active minutes.

# Dictionary mapping user IDs to a set of their active minutes

// Iterate through the map to count the number of active minutes for each user.

def findingUsersActiveMinutes(self, logs: List[List[int]], k: int) -> List[int]:

# Initialize an answer array to store counts of users with X active minutes

# We do -1 because the index is 0-based and minutes is 1-based

complexity of O(1), iterating over the logs array contributes O(n) to the time complexity.

# Iterate through the logs, adding unique minutes to each user's set

// Increment the count for the number of users with a specific count of active minutes.

// (size - 1) is used because indices are 0-based and our minute counts are 1-based.

// Logs are represented as an array of [userId, minute] and k is the upper limit of active minutes.

# Initialize an answer array to store counts of users with X active minutes

def findingUsersActiveMinutes(self, logs: List[List[int]], k: int) -> List[int]:

# Iterate through the logs, adding unique minutes to each user's set

# Iterate through the user active minutes to update the answer array

# Dictionary mapping user IDs to a set of their active minutes

## // Initialize an array to store the UAM count for each possible count 1 through k int[] answer = new int[k]; // Iterate over each user's set of timestamps for (Set<Integer> timeStamps : userActiveTimes.values()) {

return answer;

```
#include <vector>
#include <unordered map>
#include <unordered set>
using namespace std;
class Solution {
public:
    // This function takes a vector of vector of integers which contain user ID and
    // login time, along with an integer k representing the total number of minutes.
    // It returns a vector with the count of users who have a specific number of
    // active minutes ranging from 1 to k.
    vector<int> findingUsersActiveMinutes(vector<vector<int>>& logs, int k) {
        // Create a map that associates each user (identified by their ID) with a set of unique active minutes.
        unordered map<int. unordered set<int>> userToActiveMinutesMap;
        for (const auto& logEntry : logs) {
            int userID = logEntry[0];
            int time = logEntry[1];
            // Insert the time into the set for this user, duplicates are automatically handled by the set.
            userToActiveMinutesMap[userID].insert(time);
        // Initialize a vector to store the counts of users having x number of active minutes,
        // where x ranges from 1 to k.
        vector<int> activeMinutesCounts(k, 0);
        for (const auto& userToTimeEntry : userToActiveMinutesMap) {
            // The size of the set for each user gives the total number of unique active minutes.
            int uniqueActiveMinutes = userToTimeEntry.second.size();
            // Increment the count for the appropriate number of active minutes (decrement by 1 to match 0-indexing).
            if (uniqueActiveMinutes <= k) {</pre>
                activeMinutesCounts[uniqueActiveMinutes - 1]++;
        // Return the resulting vector of counts.
        return activeMinutesCounts;
};
TypeScript
```

## answer = [0] \* k# Iterate through the user active minutes to update the answer array for minutes in user active minutes.values(): # Increase the count for the number of active minutes

return answer

the logs array, n.

from collections import defaultdict

return userCount;

from typing import List

class Solution:

Time and Space Complexity The time complexity of the code is O(n), where n is the length of the logs array. This is because we iterate over each element of the logs array exactly once, adding the timestamp to a set in the dictionary. Since set operations such as ladd() have a time

Further, we also iterate over the values of the dictionary, with the number of values being at most n. For each of the values, we increment the appropriate count in the ans array, which is another 0(1) operation. Hence, the total time complexity remains 0(n). The space complexity of the code is also 0(n) because in the worst-case scenario, each log entry could be for a unique user and

a unique timestamp, leading to n unique entries in the dictionary. As such, both the dictionary and the set of timestamps could

potentially grow linearly with the number of log entries. Therefore, the space used by the dictionary is proportional to the size of