

1855. Maximum Distance Between a Pair of Values

Medium

Greedy

Array

Two Pointers

Binary Search

Leetcode Link

Problem Description

In this problem, we are given two non-increasing integer arrays, `nums1` and `nums2`. We define a valid pair as a pair of indices (i, j) such that:

- i is an index from `nums1` and j is an index from `nums2`.
- i must be less than or equal to j ($i \leq j$).
- The value at `nums1[i]` must be less than the value at `nums2[j]` (`nums1[i] < nums2[j]`).

The distance for a valid pair (i, j) is calculated by the difference $j - i$. The goal of the problem is to return the maximum possible distance for any valid pair of indices. If there are no valid pairs, we should return `0`.

Note that an array is non-increasing if each element is less than or equal to the previous element. This implies that the numbers in both `nums1` and `nums2` are sorted in non-increasing order (from largest to smallest).

Intuition

The intuition behind the solution involves recognizing that if a pair (i, j) is valid, then any pair (i, k) where $k < j$ is also valid due to the non-increasing order of the arrays. This property allows us to use a two-pointer or a binary search approach to efficiently find the maximum distance. We choose the two-pointer approach here as it simplifies the implementation.

We start by initializing two pointers, i and j , at the beginning of `nums1` and `nums2` respectively. We then iterate through `nums1` with i , trying to find the furthest j in `nums2` that will form a valid pair. As we go, we keep track of the maximum distance `ans`.

The crucial observation is that since both `nums1` and `nums2` are non-increasing, once we find a j for a particular i such that `nums1[i] <= nums2[j]`, we don't need to reset j back to i for the next i . Instead, we can continue from the current j , because if `nums1[i] <= nums2[j]` and `nums1` is non-increasing, then `nums1[i + 1]` will be less than or equal to `nums1[i]`, and hence also less than or equal to `nums2[j]`. We keep incrementing j until the condition `nums1[i] <= nums2[j]` is no longer met.

However, we need to calculate the distance $j - i$ and store the maximum. Since j is incremented in the inner loop until `nums1[i] <= nums2[j]` is false, at this point j is actually one index past the valid pair, so we compute $j - i - 1$ to get the correct distance for considering the current index i .

As a result, the algorithm efficiently calculates the maximum distance of any valid pair by scanning through both arrays only once.

Solution Approach

The given reference solution approach indicates using a Binary Search algorithm. However, the provided solution code actually uses a two-pointer technique, which intuitively exploits the sorted nature of the arrays to find the maximum distance. Below, we'll walk through the implementation of the two-pointer solution given in the code.

The code defines a class `Solution` with a method `maxDistance` which accepts two arrays, `nums1` and `nums2`.

- Initialization:** We begin by initializing a few variables:
 - `m`, the length of `nums1`
 - `n`, the length of `nums2`
 - `ans`, which will keep track of the maximum distance found among valid pairs. It is set to `0` initially.
 - Two pointers i and j , both set to `0`, to iterate through `nums1` and `nums2` respectively.
- Iterating with Two Pointers:** We use a while loop to iterate through `nums1` with i as long as $i < m$ (ensuring we don't go out of bounds).
- Finding the Furthest j :** For every index i , we have an inner while loop that seeks the furthest j such that `nums1[i] <= nums2[j]`. This loop runs as long as $j < n$. If the condition is met, j is incremented by `1`, moving the second pointer ahead in `nums2`.
- Tracking the Maximum Distance:** Once we are either out of bounds in `nums2` or the condition is no longer satisfied, the inner loop breaks and we use the maximum function `max(ans, j - i - 1)` to update our answer with the largest distance found so far. The `-1` is there to correct for the j index that is now at an invalid pair, having just moved past the last valid `nums2[j]`.
- Incrementing i :** After checking and possibly updating the `ans` with the maximum distance for the current i , we increment i by `1` and repeat the above steps until all elements of `nums1` have been considered.
- Returning the Answer:** When the loop terminates, `ans` will hold the maximum distance of any valid pair (i, j) , or it will remain `0` if no valid pairs existed. This value is then returned as the result of the method.

The two-pointer approach is efficient because it leverages the sorted properties of `nums1` and `nums2` to avoid redundant checks. It doesn't require additional data structures and runs in linear time, as each pointer only passes through its respective array once.

Thus, the implementation uses a methodical two-pointer strategy to measure the maximum distance between valid pairs across two sorted arrays.

Example Walkthrough

Let's walk through an example to illustrate the solution approach in practice.

Suppose we have the following inputs: `nums1 = [4, 2, 2]` `nums2 = [3, 2, 0]`

We want to find the maximum possible distance for any valid pair of indices.

- We start with $i = 0$ in `nums1` and $j = 0$ in `nums2`, with `ans = 0` as we haven't found any valid pairs yet.
- In the outer loop, i is now at index `0` and `nums1[0]` is `4`. The inner while loop starts iterating j through `nums2`. `nums2[j]` is `3`, which is less than `nums1[i]`, so we don't have a valid pair. We continue to the next j without changing `ans`.
- Now, $j = 1$ and `nums2[j]` is `2`. It's still not greater than `nums1[i] = 4`, so we move to the next j .
- At $j = 2$, we have `nums2[j]` equal to `0`, which is also not greater than `4`. Having reached the end of `nums2` without finding a valid pair for $i = 0$, we increment i to `1` and start over with $j = 0$.
- Now $i = 1$ and `nums1[i]` is `2`. Starting again from the beginning of `nums2`, we find that `nums2[0]` is `3`, which is greater than `nums1[i]`. This is a valid pair ($i=1, j=0$). Hence, we update `ans = max(ans, 0 - 1 - 1)`, but this gives us `-2`, which is not an increase over `0`, so `ans` stays at `0`.
- We increment j to see if there is a pair with a larger distance. Now j is `1` and since `nums1[i] = 2` is already less than `nums2[j] = 2`, we stop the inner while loop and update `ans`. Now `ans = max(0, 1 - 1 - 1)`, which remains `0` as the indices are the same.
- Finally, we increment i to `2`. `nums1[i]` is still `2`, and we try with `nums2[j]` starting from the beginning. At $j = 0$, we once again have a valid pair ($i=2, j=0$) and update `ans` to `max(0, 0 - 2 - 1)`, giving us `-3`, which doesn't increase `ans`. Since all values are non-increasing and there's no greater element in `nums2` than `3`, we can conclude there's no need to iterate j further for $i = 2$.

After considering all elements of `nums1` and finding no valid pair with a positive distance, the function returns the `ans`, which in this case remains `0` as there were no valid pairs that satisfied the condition for a positive maximum distance.

This example demonstrates how the two-pointer approach navigates through both arrays, cleverly using the condition of non-increasing order to find the valid pairs and calculating the maximum possible distance.

Python Solution

```
1 # Import the List type from the typing module for type annotations
2 from typing import List
3
4 class Solution:
5     def maxDistance(self, nums1: List[int], nums2: List[int]) -> int:
6         # Initialize the lengths of the two input lists
7         len_nums1, len_nums2 = len(nums1), len(nums2)
8
9         # Initialize the maximum distance and the indices for nums1 and nums2
10        max_distance = idx_nums1 = idx_nums2 = 0
11
12        # Loop through the elements of nums1
13        while idx_nums1 < len_nums1:
14            # For the current idx_nums1, increase idx_nums2 as long as
15            # the conditions are satisfied (nums1[idx_nums1] <= nums2[idx_nums2])
16            while (idx_nums2 < len_nums2 and nums1[idx_nums1] <= nums2[idx_nums2]):
17                idx_nums2 += 1
18            # Update the maximum distance if a larger one is found
19            # Decrement by 1 because idx_nums2 is increased one more time before the condition fails
20            max_distance = max(max_distance, idx_nums2 - idx_nums1 - 1)
21            # Move to the next index in nums1
22            idx_nums1 += 1
23
24        # Return the maximum distance found
25        return max_distance
26
```

Java Solution

```
1 class Solution {
2     public int maxDistance(int[] nums1, int[] nums2) {
3         int lengthNums1 = nums1.length; // length of the first array
4         int lengthNums2 = nums2.length; // length of the second array
5         int maxDist = 0; // variable to keep track of the maximum distance
6
7         // Initialize two pointers for both arrays
8         for (int indexNums1 = 0, indexNums2 = 0; indexNums1 < lengthNums1; ++indexNums1) {
9             // Move the indexNums2 pointer forward as long as the condition holds
10            while (indexNums2 < lengthNums2 && nums1[indexNums1] <= nums2[indexNums2]) {
11                ++indexNums2;
12            }
13            // Update maxDist with the maximum distance found so far
14            // We subtract 1 because indexNums2 has moved one step further than the true distance
15            maxDist = Math.max(maxDist, indexNums2 - indexNums1 - 1);
16        }
17        return maxDist; // return the maximum distance found
18    }
19 }
20
```

C++ Solution

```
1 class Solution {
2 public:
3     int maxDistance(vector<int>& nums1, vector<int>& nums2) {
4         int nums1Size = nums1.size(); // Represents the size of the first vector nums1
5         int nums2Size = nums2.size(); // Represents the size of the second vector nums2
6         int maxDist = 0; // Initialize maximum distance to 0
7
8         // Using two pointers, iterate through both arrays
9         // i is the pointer for nums1 and j is the pointer for nums2
10        for (int i = 0, j = 0; i < nums1Size; ++i) {
11            // While j is within the bounds of nums2 and the value at nums1[i]
12            // is less than or equal to the value at nums2[j], increment j
13            while (j < nums2Size && nums1[i] <= nums2[j]) {
14                ++j;
15            }
16            // Calculate the current distance, and update maxDist if it's greater
17            // than the current maxDist. Subtract one because j has been incremented
18            // in the last iteration of the inner loop where the condition was still valid.
19            maxDist = max(maxDist, j - i - 1);
20        }
21
22        return maxDist; // Return the computed maximum distance
23    }
24 };
25
```

Typescript Solution

```
1 function maxDistance(nums1: number[], nums2: number[]): number {
2     let maximumDistance = 0; // This will store the maximum distance found.
3     const nums1Length = nums1.length; // Length of the first array
4     const nums2Length = nums2.length; // Length of the second array
5
6     // Start with two pointers, 'i' for array nums1 and 'j' for array nums2
7     for (let i = 0, j = 0; i < nums1Length; ++i) {
8         // Increment 'j' as long as it's within the bounds of nums2
9         // and the element in nums1 is less than or equal to the element in nums2.
10        // This ensures that the 'j' finds the furthest distance it can go for each 'i'.
11        while (j < nums2Length && nums1[i] <= nums2[j]) {
12            j++;
13        }
14
15        // Update the maximum distance. Since 'j' moved one step extra, we subtract 1.
16        // i.e., j - i is the total distance from current 'i' to the last valid 'j'.
17        // but we need to subtract 1 since 'j' has moved one step ahead of valid position.
18        maximumDistance = Math.max(maximumDistance, j - i - 1);
19    }
20
21    // Return the maximum distance found
22    return maximumDistance;
23 }
24
```

Time and Space Complexity

Time Complexity

The time complexity of the given code is $O(m + n)$. Here, m is the length of `nums1` and n is the length of `nums2`. This is because the two pointers i and j only move forward through their respective arrays, with i incrementing in every outer loop, and j potentially moving forward during each inner loop until `nums1[i] <= nums2[j]` no longer holds. Since each pointer goes through its respective array at most once, the time complexity is linear with respect to the sizes of the input arrays.

Space Complexity

The space complexity of the code is $O(1)$. No additional space is used that grows with the input size. The variables `ans`, i , j , `m`, and `n` each use a constant amount of space.