

1879. Minimum XOR Sum of Two Arrays

Hard Bit Manipulation Array Dynamic Programming Bitmask

Leetcode Link

Problem Description

In this problem, you have two integer arrays `nums1` and `nums2`, each of the same length `n`. You need to calculate the XOR sum of these arrays which is obtained by taking the XOR of each corresponding pair of elements from the two arrays and then summing those results up. The XOR operation is a bitwise operation where the result is 1 if the two bits are different and 0 if they are the same.

The XOR sum is computed as follows:

```
(nums1[0] XOR nums2[0]) + (nums1[1] XOR nums2[1]) + ... + (nums1[n - 1] XOR nums2[n - 1])
```

The challenge is to rearrange the elements in `nums2` in such a way that the resulting XOR sum is as small as possible. In other words, you want to find an optimal permutation of `nums2` that when paired with the elements of `nums1`, yields the minimum possible XOR sum.

For example, if `nums1` is `[1,2,3]` and `nums2` is `[3,2,1]`, the XOR sum of these arrays before any rearrangement is `(1 XOR 3) + (2 XOR 2) + (3 XOR 1) = 2 + 0 + 2 = 4`. By rearranging `nums2` properly, you might get a smaller XOR sum.

Your task is to determine this minimum XOR sum after rearranging `nums2`, and to return it.

Intuition

The solution to this problem is using a dynamic programming approach that involves trying out different combinations of pairings between elements in `nums1` and `nums2`. Specifically, we use a bitmask to represent the elements from `nums2` that have been paired up with elements in `nums1`. The bitmask is a binary number where each bit corresponds to an element in `nums2`. If a bit is set (i.e., it is 1), it means the corresponding element in `nums2` has been used in a pairing.

We start with an array `f` to keep track of the minimum XOR sum that can be achieved with each possible bitmask. Initially, `f[0]` (representing no elements paired) is set to 0 and all other entries are set to infinity because we haven't computed them yet.

Then, we iterate over all possible bitmasks. For each bitmask, we figure out how many bits are set; this tells us the position (`k`) in `nums1` that we are considering pairing up. Now, for each bit that is set in the current bitmask, we try unsetting it (which means we're considering that the `j`-th element in `nums2` could be paired with the `k`-th element in `nums1`) and calculate the XOR of `nums1[k]` and `nums2[j]`, and add it to the minimum XOR sum stored in `f` for the bitmask without the `j`-th bit set.

This is done for each set bit in the bitmask to find the minimum XOR sum possible for that bitmask and store it back in `f`. Finally, `f[-1]` (which corresponds to all elements in `nums2` being paired up) will contain the minimum XOR sum we can achieve, and that's the value we return as the solution.

Solution Approach

The solution approach for minimizing the XOR sum of two integer arrays is a dynamic programming approach that utilizes bit masking to explore the state space of possible pairings. The algorithm uses the following concepts:

- Dynamic Programming (DP):** The idea is to break the problem into overlap subproblems and use DP to remember results of already solved subproblems, such that each subproblem is solved only once. This significantly reduces the computation time as compared to a naive approach that might solve the same subproblems repeatedly.

- Bitmasking:** This technique is used to represent the pairing state of elements from the second array `nums2`. A bitmask is an array that uses each individual bit to represent a binary state (on/off, used/not used). In this case, we are using it to keep track of which elements in `nums2` have already been paired up with elements in `nums1`.

Now, let's go through the code:

For every masked state (ranging from `1` to `2n - 1`, where `n` is the size of the arrays), the bit count is retrieved. Here `1.bit_count() - 1` calculates the number of set bits in the bitmask and decrements it by one to get the correct index in `nums1` since we are zero-indexed.

Next, the algorithm iterates over all possible indices `j` in the range of `0` to `n-1` to find the index where the bit is set in the current mask (checked by `i >> j & 1`) which means that the `j`-th element of `nums2` is considered for pairing up with the `k`-th element of `nums1`.

Then, the XOR of `nums1[k]` and `nums2[j]` is computed and added to the previously computed value of `f[i ^ (1 << j)]` (which represents the minimum XOR sum for the state where `j`-th bit is not included in the mask). The algorithm selects the minimum of these sums and updates the DP table at `f[i]`.

The DP table `f` is a one-dimensional array with `2n` elements (since this is the number of possible states for `n` bits), initialized with infinity (`inf`) to represent that those XOR sums have not been computed yet, except for the base case `f[0]` which is set to 0 because no elements are paired and thus the XOR sum is 0.

Finally, after the algorithm iterates through all subproblems, `f[-1]` will hold the minimum XOR sum which is returned as the solution. Note that `f[-1]` is Python's way of accessing the last element of the list, which corresponds to the mask where all bits are set and all elements in `nums2` have been paired with elements in `nums1`.

Here is a simplified outline of the algorithm applied in the code:

- Initialize the DP table `f` with `inf` and set `f[0]` to 0.
- Iterate over all possible bit masks from `1` to `2n - 1`.
- Calculate the current position `k` in `nums1` based on the bitmask using bit count.
- For each set bit `j` in the current bitmask, calculate the new XOR sum and update `f[i]` with the minimum value obtained.
- Return `f[-1]` as the minimum XOR sum after considering all pairings.

This dynamic programming solution ensures that the minimum XOR sum for any possible pairing is calculated efficiently and accurately by considering all possible combinations without redundant computations.

Example Walkthrough

Let's take `nums1 = [1, 3]` and `nums2 = [2, 4]` as an example to illustrate the solution approach.

- Initialization**

First, we initialize our DP table `f` with values set to infinity and `f[0]` to 0 because at the start, no elements are paired, so the XOR sum is 0.

 - So `f = [0, inf, inf, inf]`, corresponding to the 2-bit masks from `00` to `11`.
- Iterating Over Bit Masks**

We iterate over all possible bit masks from `1` to `2n - 1` to calculate the minimum XOR sum for all pairings, where `n` is the size of the input arrays. Since our arrays have 2 elements, we iterate over masks `'01'` and `'10'`, representing the different pairings.

 - `f[1]` will track the results for mask `01` (where the second element of `nums2` is used, but not the first).
 - `f[2]` will track results for mask `10` (where the first element of `nums2` is used, but not the second).
- Updating DP Table `f` with Subproblem solutions**
 - Consider the bit mask `01` (binary for 1):

The bit count minus one is `0` (`1.bit_count() - 1`) which means we are looking to pair the first element of `nums1` with elements of `nums2`.

 - Pair `nums1[0]` with `nums2[1]` (bit set at position 1 in mask `01`)

XOR sum is `1 XOR 4 = 5`. With the base case `f[0] = 0`, the result is `0 + 5 = 5`.

So we update `f[1]` to `min(inf, 5)` which is 5.
 - Consider the bit mask `10` (binary for 2):

The bit count minus one is `0` (`1.bit_count() - 1`) which means we're again pairing the first element of `nums1`.

 - Pair `nums1[0]` with `nums2[0]` (bit set at position 0 in mask `10`)

XOR sum is `1 XOR 2 = 3`. With the base case `f[0] = 0`, the result is `0 + 3 = 3`.

So we update `f[2]` to `min(inf, 3)` which is 3.
- Finding the Minimum XOR Sum**

Now, we consider all elements as paired (mask `11` represents all bits set), which in this simplified example means we've already made our optimal pairings. We'd use `f[1]` and `f[2]` to calculate the XOR sum for mask `11`, but as it's beyond the size of input arrays, we stick to the subproblem results.

Since we compared all possible combinations, we already have our answer and do not need this step for input arrays of size 2.
- Returning the Result**

The minimum XOR sum after rearranging `nums2` is the minimum of `f[1]` and `f[2]`. In this case, it is `min(f[1], f[2])` which equates to `min(5, 3) = 3`.

So, the answer for `nums1 = [1, 3]` and `nums2 = [2, 4]` after rearranging `nums2` for the minimum XOR sum is 3.

This walk-through covers the dynamic programming solution approach involving bit masks to optimize the XOR sum between two arrays by finding the best possible rearrangement of elements in `nums2`.

Python Solution

```
1 class Solution:
2     def minimum_xor_sum(self, nums1: List[int], nums2: List[int]) -> int:
3         # Determine the length of the second list
4         length = len(nums2)
5
6         # Initialize a memoization table with infinity values,
7         # representing the minimum XOR sum for each subset
8         memo = [float('inf')] * (1 << length)
9
10        # Base case: the minimum XOR sum for an empty subset is 0
11        memo[0] = 0
12
13        # Iterate over all possible subsets of nums2
14        for bitmask in range(1, 1 << length):
15            # Count the number of bits set in bitmask to determine
16            # the index k in nums1 that is being considered
17            k = bin(bitmask).count('1') - 1
18
19            # Check all elements of nums2 by iterating over bits of bitmask
20            for j in range(length):
21                # If the j-th bit of bitmask is set, calculate the potential XOR sum
22                # and update the memo table accordingly
23                if bitmask & (1 << j):
24                    # Clear the j-th bit to find the XOR sum of the previous subset
25                    previous_bitmask = bitmask ^ (1 << j)
26
27                    # Update the memo table entry for the current bitmask with the minimum
28                    # XOR sum obtained by either taking the current element from nums2 or not
29                    memo[bitmask] = min(memo[bitmask],
30                                       memo[previous_bitmask] + (nums1[k] ^ nums2[j]))
31
32        # The last element of memo contains the minimum XOR sum for the full set
33        return memo[-1]
34
```

Java Solution

```
1 class Solution {
2     public int minimumXORSum(int[] nums1, int[] nums2) {
3         // Get the length of the array.
4         int n = nums1.length;
5         // Initialize the 'dp' array with the max possible values (using shift to get 2^30).
6         int[] dp = new int[1 << n];
7         Arrays.fill(dp, 1 << 30);
8         // The starting state has a minimum XOR sum of 0.
9         dp[0] = 0;
10
11        // Iterate over all possible combinations of pairs.
12        for (int i = 0; i < (1 << n); ++i) {
13            // Find the current number of bits set to 1 in the bitmask 'i'.
14            int bitCount = Integer.bitCount(i) - 1;
15            for (int j = 0; j < n; ++j) {
16                // Check if the j-th bit in the mask 'i' is set to 1.
17                if ((i & (1 << j)) != 0) {
18                    // This is done by removing the j-th element from the current state (using XOR)
19                    int prevState = i ^ (1 << j);
20                    // Calculate the minimum XOR sum by comparing the previous state with the new masked value.
21                    dp[i] = Math.min(dp[i], dp[prevState] + (nums1[bitCount] ^ nums2[j]));
22                }
23            }
24        }
25        // Return the minimum XOR sum for all pairs by examining the last element in 'dp' array.
26        return dp[(1 << n) - 1];
27    }
28 }
29
```

C++ Solution

```
1 class Solution {
2 public:
3     int minimumXORSum(vector<int>& nums1, vector<int>& nums2) {
4         int n = nums1.length(); // Size of the input vectors
5         vector<int> dp(1 << n, INT_MAX); // Initialize the dp array with maximum integer values
6
7         dp[0] = 0; // Initial state: no numbers are paired, so the XOR sum is 0
8
9         // Iterate over all possible states
10        for (int i = 0; i < (1 << n); ++i) {
11            // k represents the number of elements already included from nums1
12            int k = __builtin_popcount(i) - 1;
13            // Iterate over all elements in nums2
14            for (int j = 0; j < n; ++j) {
15                // Check if the j-th element in nums2 has already been paired
16                if (i & (1 << j)) {
17                    // If paired, calculate the new value for the dp state
18                    // This is done by removing the j-th element from the current state (using XOR)
19                    // Then, add the XOR of nums1[k] and nums2[j] to the dp value of the previous state
20                    dp[i] = min(dp[i], dp[i ^ (1 << j)] + (nums1[k] ^ nums2[j]));
21                }
22            }
23        }
24        // Return the result for the state where all elements are included
25        return dp[(1 << n) - 1];
26    }
27 };
28
29
```

Typescript Solution

```
1 function minimumXORSum(nums1: number[], nums2: number[]): number {
2     const n = nums1.length; // length of the arrays
3     const dp: number[] = Array(1 << n).fill(1 << 30); // dynamic programming array initialized with high values
4     dp[0] = 0; // base case: XOR sum is 0 when there are no numbers to pair
5
6     // Iterate over all possible subsets of pairs created from nums2
7     for (let i = 0; i < (1 << n); ++i) {
8         const bitsSet = bitCount(i) - 1; // calculate how many bits are set in i
9
10        // Try matching each element in nums2 with nums1 based on bits set
11        for (let j = 0; j < n; ++j) {
12            if (((i >> j) & 1) === 1) { // if the j-th bit is set
13                // Calculate new minimum XOR for the new subset by toggling j-th bit
14                dp[i] = Math.min(dp[i], dp[i ^ (1 << j)] + (nums1[bitsSet] ^ nums2[j]));
15            }
16        }
17    }
18
19    // dp[(1 << n) - 1] contains the answer for the full set
20    return dp[(1 << n) - 1];
21 }
22
23 // Helper function that returns the count of set bits in the binary representation of i
24 function bitCount(i: number): number {
25     // Binary magic to count number of 1s
26     i = i - ((i >> 1) & 0x55555555);
27     i = (i & 0x33333333) + ((i >> 2) & 0x33333333);
28     i = (i + (i >> 4)) & 0x0f0f0f0f;
29     i = i + (i >> 8);
30     i = i + (i >> 16);
31     return i & 0x3f;
32 }
33
```

Time and Space Complexity

The code provided is a solution to the Minimum XOR Sum problem using dynamic programming with bit masking to represent different combinations of pairings between elements in `nums1` and `nums2`.

Time Complexity

The time complexity can be analyzed by looking at the two nested loops in which the outer loop iterates over all subsets of `nums2` and the inner loop iterates over every individual element in `nums2`.

- The outer loop runs for `2n` iterations because it loops over all possible subsets of a set with `n` elements, which are represented as bit masks.
- The inner loop runs for `n` iterations because it checks each position in the bit mask to see if it is set (which corresponds to `nums2[j]` being selected).

Since the inner loop operates within the outer loop, the total time complexity is `O(n * 2n)`.

Space Complexity

The space complexity is primarily determined by the storage of the array `f`, which has a length of `2n` to represent all possible combinations of matching `nums2` elements with elements in `nums1`.

So, the space complexity is `O(2n)` as that is the size of the array `f`.