Medium **Bit Manipulation** Game Theory Brainteaser Array Math Dynamic Programming Leetcode Link

Problem Description In this game-themed problem, we have two players, Alice and Bob, who are playing a turn-based game involving piles of stones.

1908. Game of Nim

player can't make a move (which happens when all piles are empty), they lose. The twist in the problem comes from the assumption that both players will make the best possible move at any point in the game, meaning they play optimally. The task is to determine if Alice, who always goes first, will win the game given the initial number of stones in each pile represented by an integer array piles, where piles [i] specifies the number of stones in the ith pile. Intuition

There are n piles of stones, and on each player's turn, that player must remove any positive number of stones from any one pile. If a

sizes.

understanding the optimal strategies that players can employ and the winning positions in such a game. A fundamental concept in playing Nim optimally is the Nim-sum, or the bitwise XOR sum of the number of stones in each pile. If the Nim-sum is 0 at the beginning of a player's turn, that player is in a losing position (assuming both players play perfectly). This is

The problem is a variant of the classical game theory problem known as Nim. The key to solving this sort of problem lies in

because no matter what move they make, the opponent can always return the game to a state where the Nim-sum is 0, putting the initial player back into a losing position on their next turn. Therefore, if the Nim-sum of all the pile sizes is not 0 when it's Alice's turn (which is the case at the start of the game since she goes first), then Alice can make a move that forces a 0 Nim-sum on Bob's turn, putting her in a winning position from the start.

As the problem assumes optimal play from both Alice and Bob, we can deduce the winner by calculating the Nim-sum of all the pile

Although the provided solution code uses a depth-first search (DFS) approach, which can handle variations of the problem where the rules of play diverge from classical Nim, for this specific problem, a simplified solution would compute the Nim-sum, like so:

class Solution: def nimGame(self, piles: List[int]) -> bool: $xor_sum = 0$

```
for pile in piles:
               xor_sum ^= pile
           return xor_sum != 0
This code simply iterates over the piles and computes the cumulative XOR. If the result is non-zero, Alice wins, since she can always
```

make a move to set the Nim-sum to 0 for Bob; if it is zero, Bob wins by maintaining the 0 Nim-sum position through each of his turns.

```
Solution Approach
```

Here's a walk-through of the provided DFS solution:

rules where the winning strategy might not be straightforward.

from the initial state; otherwise, if it's False, Bob will win.

best sequence of actions when players are acting optimally.

5]. We will use this example to illustrate the solution approach.

The binary representation of the Nim-sum 7: 111

game to a state with a Nim-sum of 0 at the end of her turn.

def nimGame(self, piles: List[int]) -> bool:

@lru_cache(maxsize=None)

return False

return can_win(tuple(piles))

piles_list = list(st)

Iterate over each pile

def can_win(st):

Apply memoization to avoid recomputation of the same game states

Try removing 1 to pile stones from the current pile

Remove the stones for the current move

31 # print(result) # Outputs True or False depending on whether you can win the Nim game

Convert tuple to list to modify the piles

for stones_removed in range(1, pile + 1):

Start the game with the initial piles configuration

30 # result = sol.nimGame([1, 2, 3]) # Pass the initial piles as a list

// Initialize a memoization map to store computed game states

private Map<Integer, Boolean> memoization = new HashMap<>();

powersOfEight[i] = powersOfEight[i - 1] * 8;

private int[] powersOfEight = new int[8];

for (int i = 1; i < 8; ++i) {

public boolean nimGame(int[] piles) {

return depthFirstSearch(piles);

// Precompute the powers of 8 up to 8^7

// 'powersOfEight' array will store the powers of 8 (1, 8, 64, ...) for encoding the game state

// Method to determine if the current player can win the nim game given the piles

// Function to encode the current state of piles as a single integer.

// Depth-First Search (DFS) algorithm to determine if current player can win.

piles[i] -= j; // Remove stones from the current pile.

piles[i] += j; // Backtrack to restore pile size.

piles[i] += j; // Backtrack to restore pile size.

return memo[currentState]; // Return precomputed result if available.

for (int j = 1; j <= piles[i]; ++j) { // Try removing 1 to all stones from pile.

if (!dfs(piles)) { // If opponent loses from the resulting state...

return memo[currentState] = true; // Current player can win.

return memo[currentState] = false; // If no winning move found, current player loses.

auto encodeState = [&](vector<int>& piles) {

for (int i = 0; i < piles.size(); ++i) {</pre>

int currentState = encodeState(piles);

if (memo.count(currentState)) {

state += piles[i] * powerOfEight[i];

function<bool(vector<int>&)> dfs = [&](vector<int>& piles) {

return dfs(piles); // Start the game with the initial piles.

for (int i = 0; i < piles.size(); ++i) { // Try each pile.</pre>

int state = 0;

return state;

};

};

int main() {

for i, pile in enumerate(piles_list):

The binary representation of the piles: 1 -> 001, 3 -> 011, 5 -> 101

The reference solution provided uses a recursive DFS approach to simulate every possible state of the game. Although for the specific game rules provided, there is a more efficient solution as explained earlier, this approach can handle a wider variety of game

1. The dfs function is defined within the context of the nimGame method. It is decorated with @cache, which is a Python decorator

2. The dfs function is intended to be called with a tuple st representing the current state of the game, i.e., the number of stones in

each pile. It converts st into a list 1st for easier manipulation since tuples are immutable.

that automatically memoizes the results of the function calls. Memoization ensures that repeated states of the game are not reevaluated, which optimizes the solution by reducing the number of recursive calls.

it returns True.

3. The dfs function then iterates over every pile lst[i] and for each pile tries removing every possible positive number of stones j (from 1 to the pile's size inclusive). This simulates all potential moves a player could make. 4. After simulating the move (removing j stones from pile i), dfs recursively checks whether this state leads to a losing state for

the opponent by calling not dfs(tuple(lst)). If a losing state for the opponent is found, then the current state is winning, hence

5. If no move leads to an immediate winning state, the dfs function returns False, indicating that any opponent's response to the current state can at most prolong the game, but eventually the current player will lose.

6. nimGame calls dfs using the initial state tuple(piles) and returns the result. If the result is True, it means Alice can win starting

The beauty of this approach is that it considers the entire space of possible moves and outcomes, and it ensures that we do not overlook any strategy that could lead to a win. However, as mentioned, the specific problem of Nim has a well-known mathematical solution, which is more efficient than the general DFS approach.

Nonetheless, the DFS approach is a powerful technique when dealing with complex decision-making problems in games or situations

where optimal strategies are not known in advance. It provides a methodical way to explore all possible outcomes and determine the

Example Walkthrough Let's consider a small example where there are 3 piles of stones with the following number of stones in each pile: piles = [1, 3,

1. At the beginning of the game, the Nim-sum is calculated by performing a bitwise XOR operation of the numbers of stones in

2. Performing the calculation gives us 1 ^ 3 = 2 and then 2 ^ 5 = 7. Since 7 is not equal to 0, Alice can make a move to leave the Nim-sum at 0 for Bob. 3. To find Alice's optimal move, we can look at the binary representation of the pile sizes and compare it with the Nim-sum. Let's do

4. Alice will try to make a move that results in a new state with a Nim-sum of 0. She can do this by choosing a pile and removing enough stones to turn the current Nim-sum (111) to 000. For example, Alice can remove 4 stones from the third pile (5 stones), which gives us a new state of [1, 3, 1].

Python Solution

12

13

14

15

23

24

25

26

27

32

8

9

10

11

12

13 14

15

16

17

18

19

20

15

16

17

18

19

20

21

22

23

24

25

26

27

28

29

30

31

32

33

34

35

36

37

38

39

40

41

42

43

44

45

46

47

};

this:

each pile: 1 ^ 3 ^ 5.

5. The new Nim-sum after Alice's move is 1 ^ 3 ^ 1, which equals 0, since 1 ^ 1 = 0 and 0 ^ 3 = 3 and finally 3 ^ 0 = 3, and in binary, 3 is 011, which will result back to 0 once XOR'ed with itself during Bob's turn.

6. Now it is Bob's turn, and since the Nim-sum is 0, any move he makes will leave a non-zero Nim-sum for Alice to exploit.

leaving a single stone in two piles: [1, 1, 1], which has a Nim-sum of 0 again. 8. Ultimately, no matter how Bob plays, if Alice starts with a non-zero Nim-sum and both play optimally henceforth, Alice will always be able to return the game to a state where the Nim-sum is 0 at Bob's turn.

Using this example, we can see that the decision on the first move is crucial: given a non-zero Nim-sum, there is always a winning

strategy for the player making the first move. Alice wins this game if she follows the strategy of always making a move to return the

7. If Bob removes 1 stone from any pile, the piles may look like [1, 2, 1]. Alice can then remove 2 stones from the second pile,

- from typing import List from functools import lru_cache class Solution:
- piles_list[i] -= stones_removed 16 17 # Check if opponent would lose from this state if not can_win(tuple(piles_list)): return True # If opponent loses, current player wins 19 20 # Undo the move piles_list[i] += stones_removed 21 22 # If no winning move, return False

```
Java Solution
   import java.util.Map;
    import java.util.HashMap;
```

class Solution {

public Solution() {

powersOfEight[0] = 1;

28 # Example usage:

29 # sol = Solution()

```
21
 22
 23
         // Private method for a recursive depth-first search to simulate all possible moves
 24
         private boolean depthFirstSearch(int[] piles) {
 25
             // Calculate a unique state hash for the current state of the piles
 26
             int stateHash = computeStateHash(piles);
 27
             // If this state has been encountered before, return the stored result
 28
             if (memoization.containsKey(stateHash)) {
 29
                 return memoization.get(stateHash);
 30
 31
 32
             // Try reducing each pile by a number of stones and perform DFS on the new state
 33
             for (int i = 0; i < piles.length; ++i) {</pre>
 34
                 for (int stonesToRemove = 1; stonesToRemove <= piles[i]; ++stonesToRemove) {</pre>
 35
                     piles[i] -= stonesToRemove;
 36
                     // If the opponent cannot win after our move, we found a winning move
 37
                     if (!depthFirstSearch(piles)) {
 38
                         piles[i] += stonesToRemove; // Undo the move
 39
                         memoization.put(stateHash, true); // Memoize the winning result
                         return true;
 40
 41
 42
                     piles[i] += stonesToRemove; // Undo the move before trying next move
 43
 44
 45
 46
             // If no winning move is found, this is a losing state
 47
             memoization.put(stateHash, false);
 48
             return false;
 49
 50
         // Computes a unique hash for the current state of the piles using powers of 8
 51
 52
         private int computeStateHash(int[] piles) {
 53
             int stateHash = 0;
 54
             // Encode the piles' state into a single integer using base 8 representation
             for (int i = 0; i < piles.length; ++i) {</pre>
 55
 56
                 stateHash += piles[i] * powersOfEight[i];
 57
 58
             return stateHash;
 59
 60 }
 61
C++ Solution
    #include <vector>
  2 #include <unordered_map>
     #include <functional>
     using namespace std;
     class Solution {
     public:
         bool nimGame(vector<int>& piles) {
             unordered_map<int, int> memo; // Store calculated game results to avoid recomputation.
 10
             int powerOfEight[8] = {1}; // Precomputed powers of 8 for state encoding.
 11
 12
             for (int i = 1; i < 8; ++i) {
 13
                 powerOfEight[i] = powerOfEight[i - 1] * 8;
 14
```

```
vector<int> piles = {3, 4, 5};
 49
         Solution solution;
 50
 51
         bool canWin = solution.nimGame(piles);
 52
         // canWin will be true or false depending on whether the player can win or not.
 53
         return 0;
 54 }
 55
Typescript Solution
  1 // Represents the Nim game initial state
    function nimGame(piles: number[]): boolean {
         // Base for representing the state
         const base: number[] = Array(8).fill(1);
         for (let i = 1; i < base.length; ++i) {</pre>
             base[i] = base[i - 1] * 8;
  6
  8
  9
         // Converts a pile array into a unique state representation
         const toState = (piles: number[]): number => {
 10
 11
             let state = 0;
 12
             for (let i = 0; i < piles.length; ++i) {</pre>
 13
                 state += piles[i] * base[i];
 14
 15
             return state;
 16
         };
 17
 18
         // Memoization cache to store known game states
 19
         const memo: Map<number, boolean> = new Map();
 20
         // Recursive function to determine if the current player can win from the given pile configuration
 21
 22
         const canWin = (piles: number[]): boolean => {
 23
             const currentState = toState(piles);
 24
             if (memo.has(currentState)) {
 25
                 return memo.get(currentState)!;
 26
 27
             for (let i = 0; i < piles.length; ++i) {</pre>
 28
                 for (let j = 1; j <= piles[i]; ++j) {
 29
                     piles[i] -= j; // Try reducing the current pile by 'j'
                     if (!canWin(piles)) { // If the opponent can't win from here, current player wins
 30
 31
                         piles[i] += j; // Restore the pile before memoizing and returning
 32
                         memo.set(currentState, true);
 33
                         return true;
 34
 35
                     piles[i] += j; // Restore the pile for the next iteration
 36
 37
             memo.set(currentState, false); // Current player can't win from this state
 38
 39
             return false;
         };
 40
 41
 42
         // Start the game by calling the recursive function with the initial piles configuration
         return canWin(piles);
 43
 44 }
 45
Time and Space Complexity
The provided Python code implements a recursive solution to the Nim Game problem with memoization using the cache decorator,
```

Time Complexity The time complexity of the recursive solution is determined by the number of possible states and the transitions between these

which caches results of subproblems to avoid repetitive calculations.

states. Since the state is represented by the st tuple, which is essentially the current configuration of piles, the number of distinct states can be vast. For each state, the dfs function iterates over each pile and attempts to reduce the number of stones in each pile by every possible number from 1 to \times (the size of the pile).

m is the maximum size of a single pile. In the worst case, each pile can be reduced to zero in m different ways, which gives us m^n possible states (assuming piles can have different maximum sizes, in practice it's less due to combining identical pile sizes). The dfs function is called once for each transition

n is the number of piles.

Let's suppose:

between these states, and because of memoization, each state is computed only once. Hence, the time complexity is O(m^n) but in practice, it can be significantly less due to memoization and the fact that not all transitions occur due to pruning (when a winning state is found, further exploration of that branch is stopped).

Space Complexity The space complexity is determined by the maximum size of the recursion stack and the space required to store the memoization cache. The recursion depth can go as deep as the number of possible transitions m^n, in the worst case. The cache will need to store

a result for each unique state of the piles, thus requiring 0(m^n) space as well. Therefore, the total space complexity is 0(m^n). In summary:

- Time Complexity: O(m^n) Space Complexity: 0(m^n)
- Please note that in practice, the actual complexities may be lower due to memoization and pruning as explained above.