2350. Shortest Impossible Sequence of Rolls

`Hard` `Greedy` `Array` `Hash Table`

Problem Description

In this problem, we are given two inputs: an integer array `rolls` of size `n`, which represents the outcomes of rolling a `k`-sided dice `n` times, with `k` being the second input — a positive integer. The dice have sides numbered from `1` to `k`.

We need to find the length of the **shortest sequence of rolls** that is not present in the given `rolls` array. A sequence of rolls is defined by the numbers that are obtained when rolling the `k`-sided dice some number of times. The key point is that we're not looking for a sequence that must be in consecutive order in `rolls`, but the numbers must appear in the same order as they would in the dice rolls.

To put it in simple words: imagine you are writing down each result of rolling a `k`-sided dice on a piece of paper, one after another. You now have to find the smallest length of a list of numbers that you could never write down only by using the numbers in the exact same order they appear in the `rolls` array.

Intuition

The solution to this problem relies on the observation that the length of the shortest absent rolls sequence depends directly on the variety of numbers in each segment of the array. Since we want to find the shortest sequence that can't be found in `rolls`, we can look at the input array as a series of segments where, once we encounter all `k` unique dice numbers, we can start looking for a new sequence.

The approach for the solution is to track the unique numbers we will roll with a set `s`. As we iterate over the `rolls`:

1. We add each number to the set `s`.
2. We check if the size of the set `s` has reached `k`. If it has, it means we've seen all possible outcomes from the dice in this segment.
3. Once we see that all numbers have appeared, it implies that any sequence of length equal to the current answer `ans` can be constructed from the segment, so we begin searching for the next longer sequence that cannot be constructed, by incrementing `ans` by 1.
4. We then clear the set `s` to start checking for the next segment.

With this strategy, each time we complete a full set of `k` unique numbers, the length of the shortest missing sequence increases, since we're able to construct any shorter sequence up to that point. The increment in `ans` represents the sequential nature of sequences that can't be found in `rolls`. When we have gone through all the elements in `rolls`, the value stored in `ans` will be the length of the shortest sequence that didn't appear in the `rolls`.

Solution Approach

The solution utilizes a greedy approach which aims to construct the shortest missing sequence incrementally.

Algorithm:

1. We initialize an empty set `s`. The purpose of this set is to store unique dice roll outcomes from the `rolls` array as we iterate through it.

2. We also define an integer `ans`, which is initialized at `1`. This variable represents the length of the current shortest sequence that cannot be found in `rolls`.

3. We iterate through every integer `v` in the `rolls` array and add the roll outcome `v` to the set `s` every time. This is our way of keeping track of the outcomes we have seen so far. By using a set, we automatically ensure that each outcome is only counted once. If a particular number repeats in `rolls`, it does not affect our counting in the set `s`.

4. After each insertion, we check if the size of the set `s` has reached the size `k`. This step is key to the implementation, since reaching a set size of `k` implies that we have seen all possible outcomes that could be the result of a dice roll.

   - If and when the set size is `k`, it means we have found a sequence of rolls of length `ans` that can be created using `rolls`. Consequently, we can't be looking for sequences of this length anymore; we need to look for a longer sequence that cannot be generated. Thus, we increment `ans` by `1`.

   - To start looking for the next shortest missing sequence, we need to clear the set `s`. By doing so, we reset our tracking of unique dice outcomes.

5. The loop repeats until all dice rolls in `rolls` have been examined. By now, `ans` will be one more than the length of the longest sequence of rolls that can be found in `rolls`. Therefore, `ans` will be the shortest sequence length that cannot be constructed from `rolls`.

Data Structures:

- The use of a **set** is essential in this approach. The set allows us to maintain a collection of unique items, which is perfect for keeping track of which dice numbers we have encountered. As a set doesn't store duplicates, it's also very efficient for this use case.

- The use of an **integer**, `ans`, to represent the length of the sequence we're currently looking for.

Patterns:

- The pattern in the solution approach follows a **greedy algorithm**. Greedy algorithms make the locally optimum choice at each step with the intent of finding the global optimum.

In summary, while iterating over the array `rolls`, we are greedily increasing the length of the sequence `ans` whenever we confirm that a sequence of that length can indeed be created using the rolls seen so far. The final value of `ans` when we have completed our iteration is the length of the shortest sequence that cannot be constructed from `rolls`.

Example Walkthrough

Let's illustrate the solution approach with a small example:

Assume we are given an array `rolls = [1, 2, 3, 3, 2]`, and the dice has `k = 3` sides. Our goal is to find the length of the shortest sequence of dice rolls that is not represented in `rolls`.

1. First, we initialize an empty set `s` and an integer `ans` with a value of `1`. The set `s` will track unique dice roll outcomes, and `ans` represents the length of the current shortest sequence not found in `rolls`.

2. We begin iterating over each element in `rolls`:

   - We add `1` to the set `s`, which becomes `{1}`.

   - No increment to `ans` is needed since the set does not yet contain all `k` rolls.

   - We move to the next roll, adding `2` to the set. Now `s = {1, 2}`.

   - We still don't need to increment `ans` as `s` does not contain all `k` unique values.

   - We add the next roll, `3`, to the set. Now `s = {1, 2, 3}`.

   - Since `s` now contains `k` unique numbers (all possible outcomes of the dice), we have seen at least one of each possible roll. At this point, we can construct any sequence of length `1` with the numbers in `s`, so we increment `ans` to `2` and clear the set `s` to start tracking the next sequence.

   - The next roll is a `3`. We add it to the now-empty set `s`, resulting in `s = {3}`.

   - The set still doesn't contain all `k` unique numbers, so we don't increment `ans`.

   - Finally, we add the last roll, `2`, to the set. Now `s = {2, 3}`.

   - Since we have reached the end of the `rolls` array and the set `s` does not contain `k` unique numbers, we stop here.

3. After completing the iteration, the value of `ans` stands at `2`.

Therefore, the length of the shortest sequence of rolls that cannot be found in `rolls` is `2`. This means there is no sequence of two rolls in the order they were rolled that we did not see at least once in the given `rolls` array.

Python Solution

```python
from typing import List

class Solution:
    def shortestSequence(self, rolls: List[int], k: int) -> int:
        # `shortestSeqLength` will hold the length of the shortest sequence that is not a subsequence of `rolls`
        shortestSeqLength = 1

        # `uniqueNumbers` will keep track of the unique numbers we have seen in the current subsequence
        uniqueNumbers = set()

        # Iterate over each number in the `rolls` list
        for number in rolls:
            # Add the number to the set of unique numbers
            uniqueNumbers.add(number)

            # Check if we have seen all `k` different numbers
            if len(uniqueNumbers) == k:
                # If true, we can form a new subsequence which will not be a subsequence of the current `rolls`
                shortestSeqLength += 1

                # Clear the set to start tracking a new sequence
                uniqueNumbers.clear()

        # Return the length of the shortest sequence that is not a subsequence of `rolls`
        return shortestSeqLength
```

Java Solution

```java
class Solution {
    public int shortestSequence(int[] rolls, int k) {
        // Initialize a set to keep track of unique elements
        Set<Integer> set = new HashSet<>();
        // Initialize the answer variable which represents the shortest sequence
        int answer = 1;

        // Iterate through each number in the rolls array
        for (int number : rolls) {
            // Add the current number to the set
            set.add(number);
            // Check if the set size equals k, meaning all numbers are present
            if (set.size() == k) {
                // Reset the set for the next sequence
                set.clear();
                // Increment the answer value, as we've completed a sequence
                answer++;
            }
        }
        // Return the count of the shortest sequence
        return answer;
    }
}
```

C++ Solution

```cpp
#include <vector>
#include <unordered_set>

class Solution {
public:
    // Function to find the shortest sequence that contains every number from 1 to k
    int shortestSequence(vector<int>& rolls, int k) {
        unordered_set<int> numbers; // Set to store unique numbers
        int sequenceLength = 1; // Initialize the sequence length as 1

        // Iterate over the roll values
        for (int roll : rolls) {
            numbers.insert(roll); // Insert the current roll value into the set
            // Check if the set of roll equals k, we have found a full sequence
            if (numbers.size() == k) {
                numbers.clear(); // Clear the set for the next sequence
                ++sequenceLength; // Increment the sequence length counter
            }
        }

        return sequenceLength; // Return the shortest sequence length
    }
};
```

Typescript Solution

```typescript
// Importing Set from the ES6 standard library
import { Set } from "typescript-collections";

// Function to find the shortest sequence that contains every number from 1 to k
function shortestSequence(rolls: number[], k: number): number {
    const numbers: Set<number> = new Set(); // Set to store unique numbers
    let sequenceLength: number = 1; // Initialize the sequence length as 1

    // Iterate over the roll values
    rolls.forEach(roll => {
        numbers.add(roll); // Insert the current roll value into the set
        // If the size of the set equals k, we have found a full sequence
        if (numbers.size === k) {
            numbers.clear(); // Clear the set for the next sequence
            sequenceLength++; // Increment the sequence length counter
        }
    });

    return sequenceLength; // Return the shortest sequence length
}
```

Time and Space Complexity

Time Complexity

The time complexity of the given code is $O(n)$ where `n` is the length of the input list `rolls`. This is because the code iterates through each roll exactly once. Inside the loop, adding an element to the set and checking its length are both $O(1)$ operations. When `s` reaches the size `k`, it is cleared, which is also an $O(1)$ operation because it happens at most `n/k` times and does not depend on the size of the set when cleared.

Space Complexity

The space complexity is $O(k)$ because the set `s` is used to store at most `k` unique values from the `rolls` list at any given time. The other variables, `ans` and `v`, use a constant amount of space, hence they contribute $O(1)$, which is negligible in comparison to $O(k)$.