2387. Median of a Row Wise Sorted Matrix

```
Binary Search
Medium Array
                             Matrix
```

## **Problem Description**

number of columns), which contains an odd number of integers. It's stated that the integers in each row of the grid are sorted in non-decreasing order. Our task is to find the median of all the integers in the matrix. But there's a catch: we must find the median in less than 0(m \* n) time complexity. This condition tells us that we cannot simply sort all the elements of the matrix and pick the middle one, as this would exceed the allowed time complexity. Intuition

The problem presents us with a matrix grid of size m x n (where 'm' represents the number of rows and 'n' represents the

The intuition behind the solution lies in binary search and the understanding of how the median is positioned within a sorted list of numbers. In a sorted list of odd numbers, the median is the middle element. In a matrix form, however, we cannot directly access the middle element since the numbers are spread across rows, but we know that exactly half of the numbers are less than or equal to it, and half are greater than it.

Since each row is sorted, we can make use of this property to apply binary search, but not on the elements directly. Instead, we use binary search on a range of possible values (since we don't have a flat sorted list). We use bisect library in Python which offers a way to determine the position where a number would be inserted to maintain order, which is bisect\_right, and **bisect\_left** to return the smallest number such that a given condition is met.

Here's how it works: **Define a search space:** The problem hints at the range of elements by suggesting the use of range(10\*\*6 + 1). This means each element can be in the range from 0 to 10\*\*6. **Define the target:** Because there's an odd number of elements in the matrix, the median is the (m\*n + 1) / 2th smallest element (since Python uses 0-based indexing, we use target = (m \* n + 1) >> 1 to get the index).

Counting function: We create a function count(x) that tells us the number of elements in the matrix which are less than or

- equal to x. This uses bisect\_right which counts how many elements in each row are less than or equal to x, and we sum these counts for all rows. Binary search: We perform a binary search across the possible values of elements and use our count(x) function to
- determine if a value is too low or too high. We're looking for the smallest number x where count(x) is at least the target number. This  $\mathbf{x}$  is the median. In the code:

• bisect\_left(range(10\*\*6 + 1), target, key=count) is the binary search command.

To get a better understanding of how the bisect\_left method operates here:

the matrix into a list, which would result in a time complexity greater than the allowed 0(m \* n).

Let's illustrate the solution approach with an example. We are given a 3x3 matrix with non-decreasing rows:

largest elements of the matrix as the bounds for binary search. Here, the search space is from 1 to 7.

 $\circ$  In the third row, there are 2 elements (3, 4) less than or equal to 4. The total count for x = 4 is 2 + 2 + 2 = 6.

• It looks through the sorted search space (range(10\*\*6 + 1)) for the value such that at least target numbers in the matrix are less than or equal to it. count function is used as a key to guide the binary search using the count of the elements. The intuition for why this works is that bisect\_left will hone in on the value with enough smaller or equal elements in the rows,

thanks to the sorted property of the rows, and the odd number of total elements ensures there's a definite middle element which

is the median.

- **Solution Approach** The implementation relies heavily on the bisect module from Python's standard library, which is designed for binary searching
- within a sorted list. By conceptualizing the problem in terms of finding the median via binary search, we avoid having to sort or merge all of the matrix's elements, which enables us to achieve a solution that meets the required time complexity bound.

established as the binary search space, where bisect\_left will search the number that is the median.

Binary Search Space: The range of possible values each element in the matrix can take is from 0 to 10\*\*6. This is

Count Function: A helper function count(x) is designed to compute the count of elements that are less than or equal to x.

Target: Since there is an odd total number of elements in the matrix, the median is at the ((m \* n) + 1) / 2th smallest

element. We right-shift by one bit to find the target index target = (m \* n + 1) >> 1 more efficiently (equivalent to integer

This uses bisect\_right, which, when used on each row of the matrix, returns the index at which x should be inserted to keep the row sorted. Summing these indices across all rows gives the total count of elements less than or equal to x.

division by 2).

**Example Walkthrough** 

grid = [

[1, 3, 5],

Here are the key components of the implementation:

bisect\_left: The binary search is performed by bisect\_left(range(10\*\*6 + 1), target, key=count). The range(10\*\*6 + 1) gives us the space in which to look for our median value, target is the number of elements that must be smaller than or equal to our median, and count as the key-function verifies whether the current middle value in our binary search meets the condition of having target number of elements less than or equal to it.

It will go through the whole range (10\*\*6 + 1) one interval at a time, zeroing in on where the median would be if we had one giant sorted list of the elements (but without actually creating this list). At each step of the binary search, bisect\_left looks at the count of numbers less than or equal to the middle number it's currently considering and determines whether to search higher or lower.

The combination of these components leads to a solution that efficiently finds the median without directly sorting or flattening

number such that there would be target numbers less than or equal to it in the sorted array representation of the matrix.

bisect\_left searches for the insertion point so as to maintain the sorted order. In this case, we are trying to find the smallest

[2, 4, 6], [3, 5, 7] In this grid, m = 3 and n = 3. There are m \* n = 9 elements, and we are looking for the median element, which is the (9 + 1) / 2 = 5 th smallest element due to the odd number of elements.

**Define a search space:** The possible values for elements range from 0 to 10\*\*6. In practice, we can use the smallest and

Define the target: We calculate the target as the median's position. Since there are 9 elements, the median will be the 5th

Binary search: We use bisect\_left to perform a binary search within our search space (from 1 to 7), looking for the

smallest number x such that the count of numbers less than or equal to x is at least 5 (our target). We start the binary

Next we check the interval from 1 to 4, with a midpoint of 2. With x = 2, the total count of elements less than or equal to 2 is 1 + 1 + 1

 $\circ$  Now we check the interval from 3 to 4, with midpoint 3. With x = 3, the total count of elements less than or equal to 3 is 2 + 1 + 2 = 5.

Now we've found exactly 5 elements, which meets our target, and since we're looking for the smallest such number, 3 is our median.

Thus, using the solution approach, we've found that the median of the matrix is 3 without having to flatten and sort the entire

 In the first row, there are 2 elements (1, 3) less than or equal to 4. In the second row, there are 2 elements (2, 4) less than or equal to 4.

**Counting function:** We need to count the number of elements less than or equal to some value x. If x = 4, then:

## Midpoint of 1 and 7 is 4. Counting using our previous counting function, we find 6 elements less than or equal to 4. Therefore, 4 may be too high—we need at least 5 elements, and we found 6—so we will now look for lower numbers.

Solution Implementation

**Python** 

Java

class Solution {

import bisect

class Solution:

search:

smallest element, so target = 5.

= 3. This is too low, so we need to search higher.

def matrixMedian(self. grid: List[List[int]]) -> int:

num\_rows, num\_cols = len(grid), len(grid[0])

mid = left + (right - left) // 2

if count less or equal(mid) < target:</pre>

# Get the dimensions of the matrix

def find median(left, right, target):

left = mid + 1

right = mid

return find\_median(1, 10\*\*6, target)

while left < right:</pre>

else:

while (minVal < maxVal) {</pre>

} else {

maxVal = midVal;

private int countLessEqual(int x) {

for (int[] row : grid) {

minVal = midVal + 1;

// Iterate over each row in the grid.

return left

# Helper function to count numbers less or equal to x using binary search

# Function to determine the median by binary searching over the range of values

# Initialize the search range between the smallest and largest possible values

# based on the problem constraints, then use binary search to find the median

# which could range from 1 to 10\*\*6 considering constraints from the problem

# Calculate the target position for median in the sorted array

target = (num\_rows \* num\_cols + 1) >> 1 # Equivalent to // 2

int numRows = grid.length; // The number of rows in the grid.

int maxVal = 1000010; // Upper bound for the value of median.

int minVal = 0; // Lower bound for the value of median.

// Binary search to find the median value in the matrix.

// If not, narrow down the lower half.

int left = 0: // Left pointer of the binary search.

if (countLessEqual(midVal) >= target) {

int numCols = grid[0].length; // The number of columns in the grid.

int midVal = (minVal + maxVal) / 2; // Calculate mid value.

// If count of elements less than or equal to midVal

return minVal; // minVal is the median after the end of binary search.

// is greater than or equal to target, narrow down the upper half.

// Helper method to count the number of elements less than or equal to x in the matrix.

int count = 0; // Counter for the number of elements less than or equal to x.

- matrix. The time complexity of this approach is 0(m \* log(max-min)), which is significantly less than 0(m \* n) for large matrices.
- def count less or equal(x): # Using bisect right to find the insertion point which gives us # the count of elements less than or equal to 'x' return sum(bisect\_bisect\_right(row, x) for row in grid)
- private int[][] grid; // Initialize a private grid variable to store the input grid. // Method to find the median in a row-wise sorted matrix. public int matrixMedian(int[][] grid) { this grid = grid; // Assign the passed grid to the class's grid variable.

int target = (numRows \* numCols + 1) / 2; // The median's position in the sorted order of elements.

```
int right = row.length; // Right pointer of the binary search.
            // Binary search to find the count of elements less than or equal to x in each row.
            while (left < right) {</pre>
                int mid = (left + right) / 2; // Calculate mid index.
                if (row[mid] > x) {
                    // If the current element is greater than x, narrow down the right part.
                    right = mid;
                } else {
                    // Else, narrow down the left part.
                    left = mid + 1;
            count += left; // Add the numbers less than or equal to x found in the current row to the count.
        return count; // Return the total count.
C++
#include <vector>
#include <algorithm> // Include algorithm for using upper_bound
class Solution {
public:
    // Function to find the median of the matrix
    int matrixMedian(vector<vector<int>>& matrix) {
        int rows = matrix.size(); // Number of rows in the matrix
        int cols = matrix[0].size(); // Number of columns in the matrix
        // Initializing search space for median
        int minValue = 0:
        int maxValue = 1000001; // Assuming 1e6 + 1 as the upper bound of the matrix values
        // Calculate the 'target' as the index after which we have the median
        int target = (rows * cols + 1) >> 1; // Use bitwise shift for division by 2
        // Lambda function to count numbers less than or equal to 'x' using 'upper_bound'
        auto countLessOrEqual = [&](int x) {
            int count = 0;
            for (const auto& row : matrix) {
                // 'upper bound' returns an iterator to the first element greater than 'x'
                count += (upper_bound(row.begin(), row.end(), x) - row.begin());
            return count;
        // Binary search to find the median
       while (minValue < maxValue) {</pre>
            int midValue = (minValue + maxValue) >> 1; // Finding the mid value for binary search
            // If count of elements less than or equal to 'midValue' is at least 'target'
            // then median is at 'midValue' or before it.
            if (countLessOrEqual(midValue) >= target) {
                maxValue = midValue; // Continue to search in the left half
            } else {
                minValue = midValue + 1; // Continue to search in the right half
```

}, 0); **}**; // Perform a binary search to find the median while (minValue < maxValue) {</pre> const midValue: number = (minValue + maxValue) >> 1; // Calculate the mid-value // If the count of elements less than or equal to 'midValue' is at least 'target' // then the median must be 'midValue' or smaller. if (countLessOrEqual(midValue) >= target) { maxValue = midValue; // Search in the left half } else {

// 'minValue' will hold the median value of the matrix after finishing the binary search

# Helper function to count numbers less or equal to x using binary search

# Function to determine the median by binary searching over the range of values

element can take in the grid, m is the number of rows and n is the number of columns in the grid.

# Using bisect right to find the insertion point which gives us

let maxValue: number = 1000001; // Assuming 1e6 + 1 as the upper limit for matrix values

const target: number = ((rows \* cols) + 1) >> 1; // Use bitwise shift for division by 2

// Find the index of the first element that is greater than 'x'

// upperBound from lodash acts similarly to 'upper\_bound' in C++

// Define a function to count how many numbers in the matrix are less than or equal to 'x'

// 'minValue' will hold the median value of the matrix after binary search

// Importing array utilities from lodash for operations like 'upperBound'

const rows: number = matrix.length; // Number of rows in the matrix

// Calculate the 'target' index that represents the median position

minValue = midValue + 1; // Search in the right half

def matrixMedian(self, grid: List[List[int]]) -> int:

num\_rows, num\_cols = len(grid), len(grid[0])

# the count of elements less than or equal to 'x'

return sum(bisect\_bisect\_right(row, x) for row in grid)

# Calculate the target position for median in the sorted array

target = (num\_rows \* num\_cols + 1) >> 1 # Equivalent to // 2

def count less or equal(x):

# Get the dimensions of the matrix

const cols: number = matrix[0].length; // Number of columns in the matrix

return minValue;

import { upperBound } from 'lodash';

let minValue: number = 0;

return count;

return minValue;

import bisect

class Solution:

// Define a function to find the median of the matrix

// Initializing the search space for the median

const countLessOrEqual = (x: number): number => {

return matrix.reduce((count, row) => {

count += upperBound(row, x);

function matrixMedian(matrix: number[][]): number {

**}**;

**TypeScript** 

# which could range from 1 to 10\*\*6 considering constraints from the problem def find median(left, right, target): while left < right:</pre> mid = left + (right - left) // 2 if count less or equal(mid) < target:</pre> left = mid + 1else: right = mid return left # Initialize the search range between the smallest and largest possible values # based on the problem constraints, then use binary search to find the median return find\_median(1, 10\*\*6, target) Time and Space Complexity

During each step of the binary search, a count of elements less than the current value (x) is performed across all rows, which takes 0(m \* n) time  $(m \text{ calls to bisect_right}, \text{ each potentially iterating over up to } n \text{ elements}).$ The space complexity of the code is 0(1), since no additional space is used that's based on the input size. The only variables in

The time complexity of the provided code is  $O(\log(MaxElement) * m * n)$  where MaxElement is the maximum value that an

The reason for this complexity is that the binary search is performed over a range of [0, 10^6] which requires log(10^6) time.

use are for counters and indices, which require a constant amount of space.