

2580. Count Ways to Group Overlapping Ranges

Medium

Array

Sorting

Leetcode Link

Problem Description

The given problem asks us to find different ways to partition a set of ranges into two distinct groups. Each range is a pair of integers `[start, end]`, which includes all integers from `start` to `end` inclusively. The two key rules for partitioning are: each range must belong to exactly one group, and overlapping ranges—that is, ranges which share at least one common integer—must be in the same group. Our task is to find the total number of distinct ways to achieve such a partition and return this number modulo $10^9 + 7$.

The problem is a combinatorial one which means that the solution typically involves counting certain arrangements or groupings. Since partitions are affected by whether ranges overlap, our approach must consider overlapping intervals as a single groupable entity.

Intuition

The first step towards the solution of the problem revolves around identifying overlapping ranges. To facilitate this identification, we sort the intervals based on their start and end points. Once sorted, we can iterate through the list of ranges and merge any overlapping ranges. This merging process is akin to finding connected components in a graph, where overlapping ranges are vertices connected by an edge.

As we merge overlapping ranges, we keep count of the non-overlapping intervals, or "chunks," since each of these can be assigned independently to one of the two groups. Here lies the crux of the problem: for each non-overlapping interval, we have two choices—the first group or the second group. Thus, if we end up with `cnt` non-overlapping intervals, there are 2^{cnt} ways to arrange them between the two groups.

The solution utilizes modular exponentiation to compute 2^{cnt} modulo $10^9 + 7$, which is a standard way to handle potentially large numbers in combinatorial problems, ensuring the result fits within typical integer size bounds in programming languages.

Ultimately, the given Python code keeps track of non-overlapping intervals by comparing the start of the current range to the maximum end point (`mx`) seen thus far. If the current start is greater than `mx`, we have encountered a new non-overlapping interval. The `max(mx, end)` update ensures we consider the entire merged interval for future comparisons. After counting all non-overlapping intervals, the `pow` function in Python computes the result of 2^{cnt} modulo $10^9 + 7$, thus providing the total number of ways to split the ranges into two groups that satisfy the constraints.

Solution Approach

The solution to this problem involves understanding the characteristics of interval partitioning and efficiently calculating the number of ways to split these intervals into two non-overlapping groups.

Here's an overview of the steps involved, further clarified with algorithms and data structures used:

1. **Sort the intervals:** Before we can count the non-overlapping intervals, we need to be able to identify them easily. Sorting the given `ranges` list by their starting points (and also by the end points in case of identical starts) helps us to handle them in a linear fashion. In Python, this is done by simply calling the `sort()` method on the `ranges` list which sorts in-place.

Sorting brings a pattern to how we explore the intervals—ranges that could possibly overlap are now next to each other, which simplifies the process of merging.

2. **Merging Overlapping Intervals:** As we iterate through the sorted intervals, we use a variable `mx` to keep track of the farthest end point we've seen so far. For each interval, if the start point is greater than `mx`, it signifies the beginning of a new non-overlapping interval, and we increment our count `cnt`. If the start is not greater than `mx`, the interval overlaps with the previous and does not contribute to increasing the count.

We then update `mx` to be the maximum of itself and the current interval's end point to reflect the most extended range of current overlapping intervals.

3. **Counting Non-Overlapping Intervals:** The variable `cnt` represents the count of these non-overlapped, merged intervals. Each of them can either go to the first group or the second, giving us two choices per interval.

4. **Calculating the Number of Ways:** Given `cnt` non-overlapping intervals, and knowing each interval has two choices of groups, the total number of ways to partition the set of ranges is 2^{cnt} .

To perform this calculation, we could multiply by 2 for each non-overlapping interval we find, but that might become computationally expensive and prone to overflow errors for very large numbers. Instead, we use modular exponentiation.

5. **Modular Exponentiation:** To find 2^{cnt} modulo $10^9 + 7$, we use the `pow` function in Python which takes three arguments—base 2, exponent `cnt`, and modulus $10^9 + 7$. This power function uses a fast exponentiation algorithm to efficiently compute large powers under a modulus, a process referred to as modular exponentiation.

In summary, the solution exploits the linearity of a sorted list to merge overlapping intervals, a count to track independent choices, and modular arithmetic to handle large numbers. The sorting incurs an $O(n \log n)$ time complexity, where `n` is the number of intervals. After sorting, the merge process and the counting that follows is done in a single pass, which contributes an $O(n)$ time complexity. The overall space complexity is $O(\log n)$ due to the recursive stack calls that could happen in sorting algorithm and the `pow` function.

Using these strategies, the code provided efficiently solves the problem while avoiding common issues associated with large number calculations in combinatorial problems.

Example Walkthrough

Let's say we are given the following ranges: `[[1, 3], [2, 5], [6, 8], [9, 10]]`.

1. **Sort the intervals:** We start by sorting these ranges. Sorted ranges: `[[1, 3], [2, 5], [6, 8], [9, 10]]` (Note: Since the ranges are already sorted, we don't need to do anything in this case.)

2. **Merging Overlapping Intervals:** Now, we need to merge overlapping ranges.

- Start with `mx = 0` (maximum end point seen so far).
 - For the first range `[1, 3]`, since `1 > mx`, it is a new non-overlapping interval, so increment `cnt` (count of non-overlapping intervals).
 - Update `mx` to `max(mx, end)` which is 3.
 - Move to the next range `[2, 5]`, since `2 <= mx`, it overlaps with the previous range, so we keep `cnt` the same.
 - Update `mx` to `max(mx, end)` which is 5.
 - The next range `[6, 8]` has a start greater than `mx`, so it's a new non-overlapping interval, increment `cnt`.
 - Update `mx` to `max(mx, end)` which is 8.
 - Finally, the last range `[9, 10]` also starts after `mx`, so increment `cnt` once more.
 - Update `mx` to 10.
- After merging, we have `cnt = 3` non-overlapping intervals.

3. **Counting Non-Overlapping Intervals:** We determined that there are 3 non-overlapping intervals after merging.

4. **Calculating the Number of Ways:** Since each non-overlapping interval can be assigned to either of two groups, the total number of distinct ways is 2^{cnt} , which is $2^3 = 8$ ways.

5. **Modular Exponentiation:** Finally, we calculate 2^{cnt} modulo $10^9 + 7$.

- Using the `pow` function in Python: `pow(2, 3, 10**9 + 7)` which equals 8.

In conclusion, for the given example with ranges `[[1, 3], [2, 5], [6, 8], [9, 10]]`, there are 8 distinct ways to partition these ranges into two groups that satisfy the given constraints.

Python Solution

```
1 from typing import List
2
3 class Solution:
4     def countWays(self, intervals: List[List[int]]) -> int:
5         # Sort intervals by their starting points
6         intervals.sort()
7
8         # Initialize the overlap counter and max_end variable to keep track of the furthest end point
9         overlap_count, max_end = 0, -1
10
11        # Iterate through each interval
12        for start, end in intervals:
13            # If the current start is greater than the max_end found so far,
14            # it means there is no overlap with the previous intervals
15            if start > max_end:
16                overlap_count += 1 # Increment the count for non-overlapping interval
17
18            # Update the max_end to the maximum of current max_end and the current interval's end
19            max_end = max(max_end, end)
20
21        # Define the modulo for the result as per the problem statement
22        mod = 10**9 + 7
23
24        # There are 2 options for each non-overlapping interval (include or exclude)
25        # Calculate the number of ways to arrange the intervals
26        # using power of 2 (as each non-overlapping interval doubles the number of ways)
27        # and take the result modulo mod
28        return pow(2, overlap_count, mod)
29
```

Java Solution

```
1 class Solution {
2     public int countWays(int[][] ranges) {
3         // Sort the ranges by their starting points.
4         Arrays.sort(ranges, (a, b) -> a[0] - b[0]);
5
6         int count = 0; // Initialize the count of distinct segments.
7         int maxEnd = -1; // Variable to keep track of the maximum endpoint seen so far.
8
9         // Iterate through the ranges.
10        for (int[] range : ranges) {
11            // If the current range starts after the maximum endpoint so far,
12            // it is a distinct segment which increases the count.
13            if (range[0] > maxEnd) {
14                count++;
15            }
16            // Update the maximum endpoint.
17            maxEnd = Math.max(maxEnd, range[1]);
18        }
19        // Calculate 2^count modulo 10^9+7 to get the number of ways.
20        return quickPow(2, count, (int) 1e9 + 7);
21    }
22
23    /**
24     * Calculate the power of a number modulo 'mod' efficiently using binary exponentiation.
25     *
26     * @param base The base number.
27     * @param exponent The exponent.
28     * @param mod The modulus for the operation.
29     * @return The result of (base^exponent) % mod.
30     */
31    private int quickPow(long base, int exponent, int mod) {
32        long result = 1; // Initialize result to 1.
33        // Loop until the exponent becomes zero.
34        for (; exponent > 0; exponent >= 1) {
35            // If the least significant bit of exponent is 1, multiply the result with base.
36            if ((exponent & 1) == 1) {
37                result = result * base % mod;
38            }
39            // Square the base for the next iteration.
40            base = base * base % mod;
41        }
42        return (int) result; // Return the result as an integer.
43    }
44 }
45
```

C++ Solution

```
1 #include <vector>
2 #include <algorithm>
3
4 class Solution {
5 public:
6     countWays(vector<vector<int>>& ranges) {
7         // Sort the intervals based on their starting points
8         sort(ranges.begin(), ranges.end());
9
10        // Initialize the counter for disjoint ranges and max end of intervals seen so far
11        int disjointRangeCount = 0, maxEnd = -1;
12
13        // Traverse through each range in the sorted ranges
14        for (const auto& range : ranges) {
15            // Increment the counter if the current range starts after the max end of previous ranges
16            if (range[0] > maxEnd) {
17                disjointRangeCount++;
18            }
19
20            // Update max end, if current range's end is greater
21            maxEnd = max(maxEnd, range[1]);
22        }
23
24        // Define long type for large number calculations
25        using ll = long long;
26
27        // Define a quick power function to calculate (a ^ n) mod
28        auto quickPowerMod = [&](ll base, int exponent, int modulus) {
29            ll result = 1;
30            // Iterate over each bit of the exponent
31            for (; exponent > 0; exponent >= 1) {
32                // If the current bit of exponent is set, multiply the result with base
33                if (exponent & 1) {
34                    result = (result * base) % modulus;
35                }
36                // Square the base for the next bit of exponent
37                base = (base * base) % modulus;
38            }
39            return result;
40        };
41
42        // Since there are 2 ways to cover each disjoint range, return 2^count of such ranges mod 10^9+7
43        return quickPowerMod(2, disjointRangeCount, 1e9 + 7);
44    };
45 };
46
```

Typescript Solution

```
1 function countWays(ranges: number[][]): number {
2     // Sort the ranges in increasing order based on their start values
3     ranges.sort((a, b) => a[0] - b[0]);
4
5     let maxEnd = -1; // Initialize the maximum end value seen so far
6     let totalWays = 1; // Initialize total ways to count combinations
7     const MODULO = 10 ** 9 + 7; // Define the modulo value for avoiding large numbers
8
9     // Iterate over each range
10    for (const [start, end] of ranges) {
11        // If the start of the current range is greater than the maxEnd seen so far,
12        // this indicates a new independent interval, so we double the ways and take modulo.
13        if (start > maxEnd) {
14            totalWays = (totalWays * 2) % MODULO;
15        }
16
17        // Update the maxEnd with the maximum of current range's end and maxEnd
18        maxEnd = Math.max(maxEnd, end);
19    }
20
21    // Return the total number of ways to cover all intervals modulated by MODULO
22    return totalWays;
23 }
24
```

Time and Space Complexity

Time Complexity

The provided code has two primary operations that contribute to its time complexity:

- Sorting the `ranges` list.
- Traversing the sorted `ranges` list to calculate the number of ways.

The `sort()` method on the list has a time complexity of $O(n \log n)$, where `n` is the number of elements in the `ranges` list. This is because the sorting operation is based on the TimSort algorithm, which is a hybrid sorting algorithm derived from merge sort and insertion sort.

After sorting, the for-loop iterates over the `ranges` exactly once. The iteration has a constant time complexity of $O(1)$ per element for basic operations like comparison and assignment. Thus, for `n` elements, this part of the code has a time complexity of $O(n)$.

Combining these two parts, the overall time complexity of the code is $O(n \log n) + O(n)$, which simplifies to $O(n \log n)$ because the $n \log n$ term dominates for large `n`.

Space Complexity

The space complexity of the code is related to the space used by the input and the internal mechanism of the sort operation. Since the input list `ranges` is sorted in place, no additional space proportional to the size of the input is required except for the internal space used by the sorting algorithm which is $O(n)$ in the worst case for TimSort.

However, the `cnt`, `mx` variables use a constant amount of space, and the `mod` is also constant, contributing a total of $O(1)$ space.

Therefore, the total space complexity of the code is $O(n)$ (space used by the sorting algorithm) plus $O(1)$ (space used by the variables), which is $O(n)$ overall.