Sorting

Problem Description

Greedy Array

Medium

In this problem, we are managing a set of tasks and processors in a way that optimizes the total time required to complete all tasks. We have n processors, each with 4 cores, meaning each processor can work on up to 4 tasks simultaneously. We also have 4 tasks that need to be completed, with the stipulation that each core can only execute one task at a time.

We are provided with two arrays: processorTime and tasks. The processorTime array tells us when each processor will become available to start executing tasks, with the index of the array representing each processor. The tasks array indicates the time each individual task takes to execute. The objective is to determine the minimum amount of time needed to finish all n * 4 tasks when they are distributed across the n

processors and their cores. It's a scheduling problem where we must figure out an optimal assignment of tasks to processors to minimize the overall completion time. Intuition

To develop the solution approach, consider if we had all processors available at the start, we would likely start with the longest

However, since processors become available at different times, we want to pair the longest tasks with the soonest available processors. This way, we can ensure that a task doesn't get unnecessarily delayed waiting for a busy processor when it could be

tasks to get them out of the way while all processors are available. We would then fill in the shorter tasks where possible.

started earlier by a free processor. Here is how we arrive at our solution approach: 1. Sort the processorTime array in ascending order, so we know the order in which processors become available. 2. Sort the tasks array in descending order, so we have the longest tasks at the beginning of the list.

Once we have these sorted lists, we begin assigning tasks to processors. Starting with the first processor in processorTime (the

one that becomes available earliest), we'll assign it the current longest task available in tasks. Since each processor has 4 cores, we can assign up to 4 tasks at a time to each processor. So, for every processor, we will assign up to 4 of our current longest

optimization problems. Here's how the provided solution translates our intuition into an algorithm:

decrement by 4 (the number of cores per processor) to always select the next set of tasks.

We decrement the index i by 4 to move on to the next set of tasks for the following processor.

Let's walk through this process with a little more detail with respect to the provided solution code snippet:

tasks, then move to the next processor that becomes available.

longest task time to the processor's available time. The minimum completion time is the maximum of these end times since all tasks need to be finished. By implementing this greedy algorithm, we ensure that we're always using the earliest available processor time optimally by pairing it with the longest remaining tasks.

To compute the minimum time all tasks are completed, we track the end time for each processor assigned tasks by adding the

Solution Approach The implementation of the solution follows the <u>Greedy</u> approach combined with <u>sorting</u>, which is a common pattern used in

Sorting - We sort both the processorTime and tasks arrays. The former is sorted in ascending order, so we know which processor will be available the soonest. The latter is sorted in descending order, ensuring we pick the longest tasks first.

Greedy Assignment - Once sorted, we assign the longest tasks to the earliest available processors - that is, each processor

after its assigned tasks are done.

the current ans to update the maximum completion time.

ans = 0 # Initialize the answer variable

starting from the earliest available will take four of the remaining longest tasks. Calculating Completion Time - For every processor, we calculate the end time which is when the processor becomes

available (processorTime) plus the time of the longest task assigned to it. This uses the max function to ensure we are only

considering the time at which the process will have completed its longest task, which is indicative of its actual availability

Data Structures - We use basic Python lists to represent the arrays and sort them. No advanced data structures are needed as the problem deals with direct array manipulation.

Iteration and Indexing - We iterate over the processorTime array, and for each processor, we keep an index i that we

- We start by sorting both the processorTime and the tasks arrays with the built-in .sort() method. • We initialize ans, which will keep track of the current maximum completion time of all tasks. • We iterate through the processorTime, for each processor, we add its available time to the time of the largest pending task and compare it with
- Note that this approach works since i is started at the end of the tasks array, saying that we always add the largest task to the next processor. It's also crucial to note that this calculation gives us the earliest possible end time for each processor to finish
- processor will not finish all of them at the same time, but the completion time will be bounded by the finish time of the longest one.

i = len(tasks) - 1 # Start with the last index of the sorted tasks (the largest task)

processorTime.sort() # Sort the processorTime array in ascending order

for t in processorTime: # Iterate through each processor's available time

Let's walk through a small example to illustrate the solution approach described above:

processorTime = [3, 6] and the individual tasks take times tasks = [5, 2, 3, 7, 8, 1, 4, 6].

Start with the first processor (processorTime[0] = 3) and assign the four longest tasks [8, 7, 6, 5].

Move to the second processor (processorTime[1] = 6) and assign the remaining four tasks [4, 3, 2, 1].

Here's the reference solution approach, further elucidated: class Solution: def minProcessingTime(self, processorTime: List[int], tasks: List[int]) -> int:

tasks.sort() # Sort the tasks array in descending order (requires reversing since Python's sort is ascending

one of its four tasks, not all four. This is because once we assign the four longest available tasks to a processor (if available), the

ans = max(ans, t + tasks[i]) # Update the answer with the latest finish time for the current processor i -= 4 # Move the index to the next set of 4 tasks return ans # Return the final calculated answer The implementation confirms the intuition that earlier available processors should be assigned the longest tasks in a way that minimizes empty processor time and maximizes task processing overlaps.

 Sort processorTime in ascending order: It's already [3, 6]. o Sort tasks in descending order: It becomes [8, 7, 6, 5, 4, 3, 2, 1]. **Greedy Assignment:**

Assume we have n = 2 processors, and hence n * 4 = 8 tasks. The processors become available at times given by

• For the second processor, starting at time 6, the maximum task time is 4, so it will finish at time 6 + 4 = 10. **Determine Final Answer:**

Example Walkthrough

Sorting:

Following the solution steps:

Calculating Completion Time:

minimize the total completion time.

from typing import List

tasks.sort()

min_time = 0

return min_time

import java.util.Collections;

import java.util.List;

Java

C++

public:

#include <vector>

class Solution {

class Solution:

• The largest among the completion times for the processors is 11, which corresponds to the first processor.

def minProcessingTime(self, processorTime: List[int], tasks: List[int]) -> int:

tasks.sort(reverse=True) # tasks becomes [8, 7, 6, 5, 4, 3, 2, 1]

i = len(tasks) - 1 # i starts at 7, index of the smallest task

i -= 4 # Decrease i to select the next set of 4 tasks

for t in processorTime: # Iterate through processor times [3, 6]

def minProcessingTime(self, processor_times: List[int], tasks: List[int]) -> int:

We select the max of the current min_time or the new completion time

min_time = max(min_time, processor_time + tasks[task_index])

This represents the earliest time at which all processors have finished processing

Decrement the task_index by 4 because the current processor is considered to

process a task every 4 cycles/time units (as per the given decrement step)

Therefore, with this approach, all tasks will be finished in 11 units of time. The solution code for this example would execute as follows: class Solution:

After running through this example, we have a clear view of how the algorithm assigns tasks to processors in a way that aims to

ans = max(ans, t + tasks[i]) # ans becomes max(0, 3+5), then max(11, 6+1)

• For the first processor, which starts at time 3, the maximum task time is 8, so it will finish its longest task at time 3 + 8 = 11.

Solution Implementation **Python**

Importing List from the typing module for type annotations

Sort the processor times in ascending order

Initialize the answer for the minimum processing time

Return the minimum time required to process all tasks

#include <algorithm> // Include the algorithm header for using the sort function

int minProcessingTime(vector<int>& processorTimes, vector<int>& tasks) {

// Initialize the answer to 0. This will track the minimum processing time.

// Iterate over processors and assign them the heaviest remaining task

function minProcessingTime(processorTimes: number[], tasks: number[]): number {

// Return the minimum time after all processors have been accounted for

def minProcessingTime(self, processor_times: List[int], tasks: List[int]) -> int:

// Sort the processor times in ascending order

// Initialize the answer and the index for the tasks array.

processorTimes.sort((a, b) => a - b);

let taskIndex: number = tasks.length - 1;

for (const processorTime of processorTimes) {

// Sort the tasks in ascending order

// Loop through each processor time

tasks.sort((a, b) => a - b);

let answer: number = 0;

taskIndex -= 4;

return answer;

if (taskIndex >= 0) { // Check if there are still tasks to process

// Sort the processor times in ascending order

// Start from the last task and work backwards

for (int processorTime : processorTimes) {

// Sort the tasks in ascending order

sort(tasks.begin(), tasks.end());

int taskIndex = tasks.size() - 1;

int minimumProcessingTime = 0;

sort(processorTimes.begin(), processorTimes.end());

Start from the last task, which has the longest processing time

Sort the tasks in ascending order

processorTime.sort() # processorTime becomes [3, 6]

return ans # Return 11 as the final calculated answer

ans = 0 # Initialize the answer variable

Iterate over each processor to assign the longest task(s) they can process for processor_time in processor_times: # Ensure the task_index is not less than 0 to avoid IndexError if task_index >= 0: # The completion time for the processor is its processing time # added to the task time it will process

task_index -= 4

task_index = len(tasks) - 1

processor_times.sort()

```
class Solution {
    /**
    * Calculates the minimum amount of time required to process all tasks given an array of processor times.
    * @param processorTimes A list of integers representing the times each processor requires to be ready for a task.
    * @param tasks A list of integers representing the times required to process each task.
    * @return The minimum processing time to complete all tasks.
    */
    public int minProcessingTime(List<Integer> processorTimes, List<Integer> tasks) {
       // Sort the processor times in ascending order
       Collections.sort(processorTimes);
       // Sort the tasks in ascending order
       Collections.sort(tasks);
        int minTime = 0; // Variable to store the minimum processing time required
        int taskIndex = tasks.size() - 1; // Start from the last task (which is the largest due to sorting)
       // Iterate over each processor time
        for (int processorTime : processorTimes) {
           // If there are no more tasks to allocate, break the loop
            if (taskIndex < 0) {</pre>
                break;
           // Calculate the total time for current processor by adding its ready time to the task time
           // and update minTime if this is larger than the current minTime
           minTime = Math.max(minTime, processorTime + tasks.get(taskIndex));
            // Move to the task which is 4 positions earlier in the list since there are 4 processors (0-based index)
            taskIndex -= 4;
       // Return the minimum time needed to complete all tasks
       return minTime;
```

```
minimumProcessingTime = max(minimumProcessingTime, processorTime + tasks[taskIndex]);
                // Move to the next set of tasks, assuming each processor can process 4 tasks simultaneously
                taskIndex -= 4;
            } else {
               // No more tasks to assign, break out of the loop.
               break;
        return minimumProcessingTime; // Return the calculated minimum processing time
};
TypeScript
/**
* Calculates the minimum time required to process all tasks by assigning them to processors.
 * @param {number[]} processorTimes - Array of processor times, where each element represents the time a processor takes to compl
 * @param {number[]} tasks - Array of task times, where each element represents the time a task requires for processing.
 * @returns {number} The minimum time required to process all tasks.
```

// Update the minimum processing time if it's less than the current processor's time plus task time

```
# Importing List from the typing module for type annotations
from typing import List
class Solution:
```

```
# Sort the processor times in ascending order
       processor_times.sort()
       # Sort the tasks in ascending order
       tasks.sort()
       # Initialize the answer for the minimum processing time
       min_time = 0
       # Start from the last task, which has the longest processing time
       task_index = len(tasks) - 1
       # Iterate over each processor to assign the longest task(s) they can process
       for processor_time in processor_times:
           # Ensure the task_index is not less than 0 to avoid IndexError
           if task_index >= 0:
               # The completion time for the processor is its processing time
               # added to the task time it will process
               # We select the max of the current min_time or the new completion time
               # This represents the earliest time at which all processors have finished processing
               min_time = max(min_time, processor_time + tasks[task_index])
               # Decrement the task_index by 4 because the current processor is considered to
               # process a task every 4 cycles/time units (as per the given decrement step)
               task_index -= 4
       # Return the minimum time required to process all tasks
       return min_time
Time and Space Complexity
```

// Calculate the potential time to process the current task with the current processor

answer = Math.max(answer, processorTime + (taskIndex >= 0 ? tasks[taskIndex] : 0));

// and update the answer with the maximum value of the current answer and this potential time

// Decrement the taskIndex by 4 for the next iteration, to simulate assignment of tasks to every 4th processor

log n).

After the sorting, there is a for loop that iterates through processorTime. The loop itself runs in O(m) time, where m is the number of processors. However, this does not affect the overall time complexity since it is assumed that m is much less than n and because the m loop is not nested within an n loop. Thus, the for loop's complexity does not exceed 0(n log n) of the sorting step. The space complexity of the sort operation depends on the implementation of the sorting algorithm. Typically, the sort method in Python (Timsort) has a space complexity of O(n). However, in the reference answer, they've noted a space complexity of $O(\log n)$

The time complexity of the code is primarily determined by the sorting operations performed on processorTime and tasks which

are 0(n log n) where n is the number of tasks. Each list is sorted exactly once, and therefore the time complexity remains 0(n

n). This can be considered correct under the assumption that the sorting algorithm used is an in-place sort like heapsort or inplace mergesort which has an O(log n) space complexity due to the recursion stack during the sort, but Python's default sorting algorithm is not in-place and actually takes O(n) space. If the sizes of the processorTime and tasks lists are immutable, and cannot be changed in place, the space complexity could indeed be O(log n) due to the space used by the sorting algorithm's recursion stack.