2944. Minimum Number of Coins for Fruits

Problem Description

Array

prices representing the cost of each fruit, indexed starting from 1. The special offer at the market is that if you buy the ith fruit, you can get the next i fruits for free. However, one can still choose to purchase a fruit that could be gotten for free, in order to activate the offer from that fruit and get additional fruits for free. The goal is to find the minimum number of coins required to buy all the fruits. It's important to note that the strategy might involve

The given problem presents a scenario where you're at a fruit market considering the purchase of exotic fruits. You have an array

<u>Dynamic Programming</u> <u>Monotonic Queue</u> <u>Heap (Priority Queue)</u>

sometimes purchasing a fruit, even when it could be taken for free, to benefit from the subsequent offers. Intuition

Solution Approach

Medium

subproblems. The main idea is to approach the problem from the start and, for each fruit we consider purchasing, explore all the possibilities of the subsequent offers it might unlock. Then, we choose the path that leads to the minimum cost.

The function dfs(i) defines the cost of purchasing fruits starting from the ith fruit onward. If choosing to buy fruit i at prices[i - 1] coins (since the array is 1-indexed), you can then skip i fruits (because they are free) and consider the next i fruits, thus we recursively calculate dfs(j) for j ranging from i+1 to i*2+1. We are looking for the minimal cost among all these

To find the minimum coins needed, we can use a recursive strategy together with memoization to avoid recalculating

possible j paths. When i * 2 is greater or equal to the length of the prices array, it means buying the ith fruit would give us all remaining fruits for free, so in that case, we just return the price of the ith fruit. By recording our results with @cache, we save the results of subproblems, speeding up the recursive process since we won't recompute the cost for any starting index more than once.

This solution leverages dynamic programming to break the problem down into overlapping subproblems and ensures that we only

calculate each subproblem once, thus efficiently arriving at the overall minimum cost.

The solution implements a Depth-First Search (DFS) strategy with dynamic programming and memoization for optimization. Let's go through each piece of the solution to understand how it's implemented:

The core of the solution is a recursive DFS function that explores all possible options to minimize the cost. The function dfs(i)

represents the cost of acquiring all fruits starting from the ith index. The dfs function is defined to operate on the assumption

that, if we purchase fruit i, we instantly receive the next i fruits for free as per the given market offer. We then call dfs recursively to accumulate the minimum cost needed for the remaining fruits.

Depth-First Search (DFS)

Solution Approach

Memoization To ensure that we don't recompute the minimum cost for the same starting index more than once, the solution uses the @cache decorator from Python. This decorator automatically saves the results of the dfs function calls into a cache. Subsequent calls

with the same arguments will return the stored result instead of recomputing it. Thus, it converts our DFS into a dynamic

• It checks if i is such that buying the ith fruit would result in acquiring all remaining fruits for free (when i * 2 >= len(prices)). In this case, it

The solution starts by calling dfs(1) which indicates that we begin by considering the first fruit and recursively determine the

• Otherwise, it calculates the sum of the price of the ith fruit and the minimum cost among all possible next steps. The next steps are given by the range i + 1 to i * 2 + 1. It does this by calling dfs(j) for each j in this range and choosing the minimum result via the min function.

Implementation Details

Algorithm Complexity

minimum cost to acquire all fruits from there.

i next fruits for free when we purchase the ith fruit.

Calling dfs(1): We are considering buying the first fruit, which costs 3 coins.

• This will offer us the next 2 fruits for free (3rd and 4th fruits get skipped).

The algorithm begins by defining the dfs function with the following logic:

returns the price of the ith fruit since no further purchases are necessary.

programming approach that reuses previously computed values.

of unique recursive calls. The exact time complexity can be hard to determine due to the varying nature of recursive calls depending on the input; however, it is significantly optimized by avoiding redundant calculations.

dynamic programming concepts to efficiently solve the problem of acquiring all the fruits for a minimum cost.

• If we purchase it, we get to pick the next fruit for free, which means we can skip considering the second fruit.

Now, we calculate the cost of starting from the 5th fruit (dfs(5)), which will just be the cost of the 5th fruit itself.

 \circ The cost will be 3 + 1 = 4 coins (because buying the 3rd fruit yields all remaining fruits for free).

So, the minimum number of coins required is 4, which is the result of $dfs(1) \rightarrow dfs(3)$.

To summarize, the solution uses a combination of DFS for traversing different combinations, memoization for optimization, and

The time complexity of this algorithm is better than the naive recursive approach due to memoization, which reduces the number

Example Walkthrough Let's consider a small example where the array prices is [3, 5, 1, 4, 2]. This array represents the cost of each fruit, with the

We want to find the minimum number of coins required to purchase all the fruits, following the market offer that allows us to get

We use the DFS approach along with memoization to solve this. We start by calling the function dfs(1). Here's how it will work step by step:

index starting from 1.

minimum cost from acquiring fruits from these starting points. Calling dfs(2): The cost of the second fruit is 5 coins.

• Then, we consider the minimum cost of buying fruits starting from indices 2 through 3 (dfs(2) and dfs(3)). We need to calculate the

 Purchasing this fruit, we get the next 3 fruits for free, and since it only leaves the 5th fruit, the cost here is just 1 coin. Calling dfs (5): Directly returns the cost of the 5th fruit (2 coins), since there are no more fruits after this.

With memoization, the results of dfs(2) and dfs(5) are stored in the cache. If there were other recursive calls that needed

Calling dfs(3): The cost of the third fruit is 1 coin.

Now let's calculate the total costs for each path:

• If we take the path of dfs(1) -> dfs(2) -> dfs(5):

 \circ The cost will be 3 + 5 + 2 = 10 coins.

• If we take the path of dfs(1) -> dfs(3):

Solution Implementation

from functools import lru_cache

def dfs(index: int) -> int:

if index * 2 >= len(prices):

return prices[index - 1] + min(

dfs(child index)

and add the price at the current index to it.

Python

class Solution:

these results, the algorithm would not recompute them but retrieve them from the cache, thereby saving time and making the algorithm more efficient. This is how the DFS and memoization work together to find the minimum cost in this scenario.

Import the 'functools' library to use the 'lru_cache' decorator for memoization

def minimumCoins(self, prices: List[int]) -> int: # Decorator to provide memoization to the dfs function # This will store results of subproblems to avoid recomputation @lru cache(maxsize=None)

If the current index * 2 is greater than or equal to the length of 'prices',

then we're at a leaf node and should return the price at this index.

return prices[index - 1] # -1 because the prices are 1-indexed

Otherwise, recursively calculate the minimum cost to reach each child,

// Method to calculate the minimum number of coins needed, starting with a given price array

// Start the recursive depth-first search from the first index (1-based index)

// If this state has not been calculated before, calculate it

memoizationArray[i] = Integer.MAX_VALUE;

for (int j = i + 1; $j \le i * 2 + 1$; ++j) {

// Return the minimum coins found for this state

// Explore all possible next states

// Initialize to a very high value to ensure proper min comparison

// Minimize the coins by making the best choice at each step

// Cache array to store minimum cost for calculated states to avoid repeated calculations

return prices[currentIndex - 1]; // Return the item's price at the current index

// Recursive call to find the minimum cost and update the memo arrav

If the current index * 2 is greater than or equal to the length of 'prices',

then we're at a leaf node and should return the price at this index.

return prices[index - 1] # -1 because the prices are 1-indexed

Otherwise, recursively calculate the minimum cost to reach each child,

// For every next index within the allowed range (not more than double the current index)

for (let nextIndex = currentIndex + 1; nextIndex <= currentIndex * 2 + 1; ++nextIndex) {</pre>

memo[currentIndex] = Math.min(memo[currentIndex], prices[currentIndex - 1] + dfs(nextIndex));

// Helper function for Depth-First Search (DFS) to find minimum total cost

// Check whether the current state already has computed minimum cost

// Initialize memo[currentIndex] with a large number

memo[currentIndex] = Number.MAX_SAFE_INTEGER;

// Return the minimum cost for the current index

// Base case: when the double of the current index is out of array bounds

const memo: number[] = Array(num0fPrices + 1).fill(0);

const dfs = (currentIndex: number): number => {

if (currentIndex * 2 >= numOfPrices) {

if (memo[currentIndex] === 0) {

return memo[currentIndex];

from functools import lru_cache

@lru cache(maxsize=None)

def dfs(index: int) -> int:

// Call DFS starting from the first element

def minimumCoins(self, prices: List[int]) -> int:

if index * 2 >= len(prices):

return prices[index - 1] + min(

dfs(child index)

Decorator to provide memoization to the dfs function

and add the price at the current index to it.

This will store results of subproblems to avoid recomputation

Find the minimal total cost among all possible next steps.

for child_index in range(index + 1, index * 2 + 2)

Start the recursive function from the first index (1-indexed)

memoizationArray[i] = Math.min(memoizationArray[i], prices[i - 1] + dfs(j));

Find the minimal total cost among all possible next steps.

for child_index in range(index + 1, index * 2 + 2)

Start the recursive function from the first index (1-indexed) return dfs(1) Java class Solution { // Member variables with more conventional naming

```
// Each call calculates the minimum price from the given index 'i'
private int dfs(int i) {
   // If the current index is in the last half of the array,
   // a single coin of the corresponding value is used
   if (i * 2 >= numOfPrices) {
        return prices[i - 1]; // Return the face value as we have reached the end
```

return dfs(1);

private int[] prices;

private int numOfPrices;

private int[] memoizationArray;

this.prices = prices;

public int minimumCoins(int[] prices) {

// Helper method for depth-first search

if (memoizationArray[i] == 0) {

return memoizationArray[i];

memoizationArray = new int[numOfPrices + 1];

numOfPrices = prices.length;

```
C++
class Solution {
public:
    // Function to compute the minimum number of coins needed
    // It takes a vector of prices as an argument
    int minimumCoins(vector<int>& prices) {
        int n = prices.size(); // Get the number of elements in prices
        vector<int> minCoins(n + 1, INT_MAX); // Initialize the minCoins DP array with a high value (INT_MAX)
        // Recursive function to find the minimum coins required.
        // It uses memoization to store the results of subproblems.
        function<int(int)> computeMinCoins = [&](int currentIndex) {
            // Base case: if the current index is more than half the size
            // we are at a leaf node and return its price
            if (currentIndex * 2 >= n) {
                return prices[currentIndex - 1];
            // If the current index has not been previously computed
            if (minCoins[currentIndex] == INT MAX) {
                // Iterate over the range of possible next indices within the problem constraints
                for (int nextIndex = currentIndex + 1; nextIndex <= currentIndex * 2 + 1; ++nextIndex) {</pre>
                    // Update the minimum coins for the current index by comparing the previously stored value
                    // with the sum of the current index's price and the computed minimum coins required
                    // for the next index (computed recursively)
                    minCoins[currentIndex] = min(minCoins[currentIndex], prices[currentIndex - 1] + computeMinCoins(nextIndex));
            // Return the minimum coins required for the current index
            return minCoins[currentIndex];
        };
        // Call the recursive function with the starting index of 1
        return computeMinCoins(1);
TypeScript
// Function to determine the minimum total cost to buy all the products
function minimumCoins(prices: number[]): number {
    // The number of prices given in the input array
    const numOfPrices = prices.length;
```

```
return dfs(1);
# Import the 'functools' library to use the 'lru_cache' decorator for memoization
```

class Solution:

};

return dfs(1) Time and Space Complexity The given Python code takes advantage of recursion with memoization to find the minimum number of coins needed from a list of prices. The function dfs is decorated with @cache, which stores the results of previous computations to avoid redundant calculations. **Time Complexity:**

• The total number of states is bound by O(n), where n is the length of the prices list. This is because the depth of recursion is at most log2(n) due to the condition i * 2 >= len(prices) that halts deeper recursive calls. • Therefore, the time complexity is 0(n * log(n)), as each state can result in a linear number of calls based on the current index, and the depth

a unique starting index i.

is logarithmic relative to n.

• For each call to dfs(i), the function iterates from i + 1 to i * 2 + 1, which results in a maximum of i recursive calls in the worst case.

Space Complexity: • Memoization will require additional space for caching the results. At most, there will be 0(n) cache entries because each entry corresponds to

• The maximum depth of the recursive call stack will also be O(log(n)) due to the nature of the problem, as the recursion depth is determined by

how many times the index can be doubled before reaching len(prices). • Combining the space needed for memoization and the maximum recursion depth, we have O(n) space complexity for caching and O(log(n)) space complexity for the call stack, leading to a total space complexity of O(n).

In conclusion, the time complexity of the code is O(n * log(n)), and the space complexity is O(n).

• Since memoization is used, each state dfs(i) for i = 1 to len(prices) / 2 is computed only once.